

# Cryptography for **Privacy-Preserving Machine Learning**

Morten Dahl

*The tf-encrypted Project and Dropout Labs*

*Applied ML Days, AI & Trust track, EPFL, January 2019*

**Why?**

# Machine Learning Process

IMAGENET



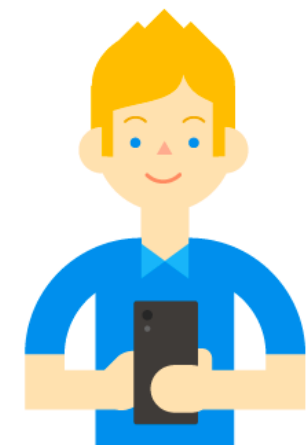
data set



training



prediction service



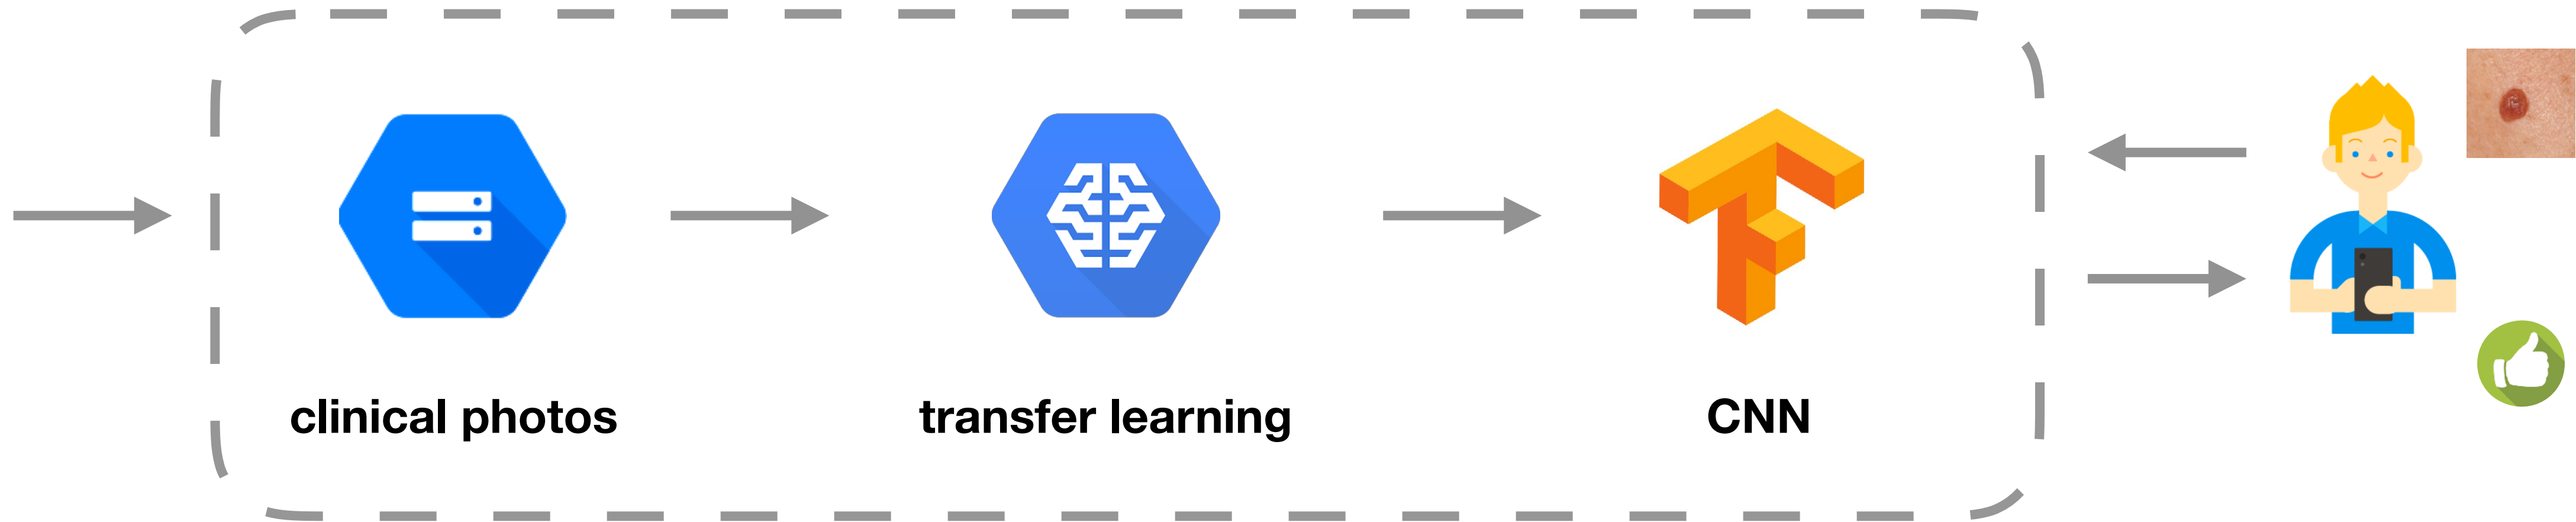
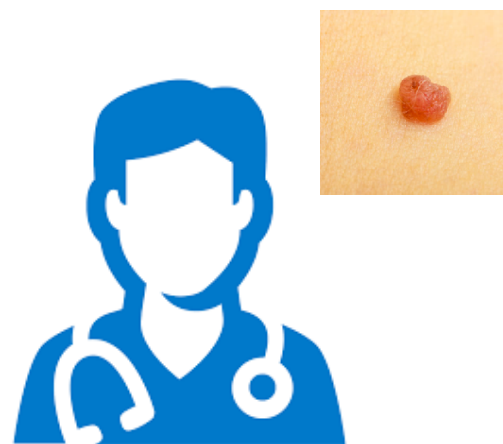


# Skin Cancer Image Classification

Brett Kuprel

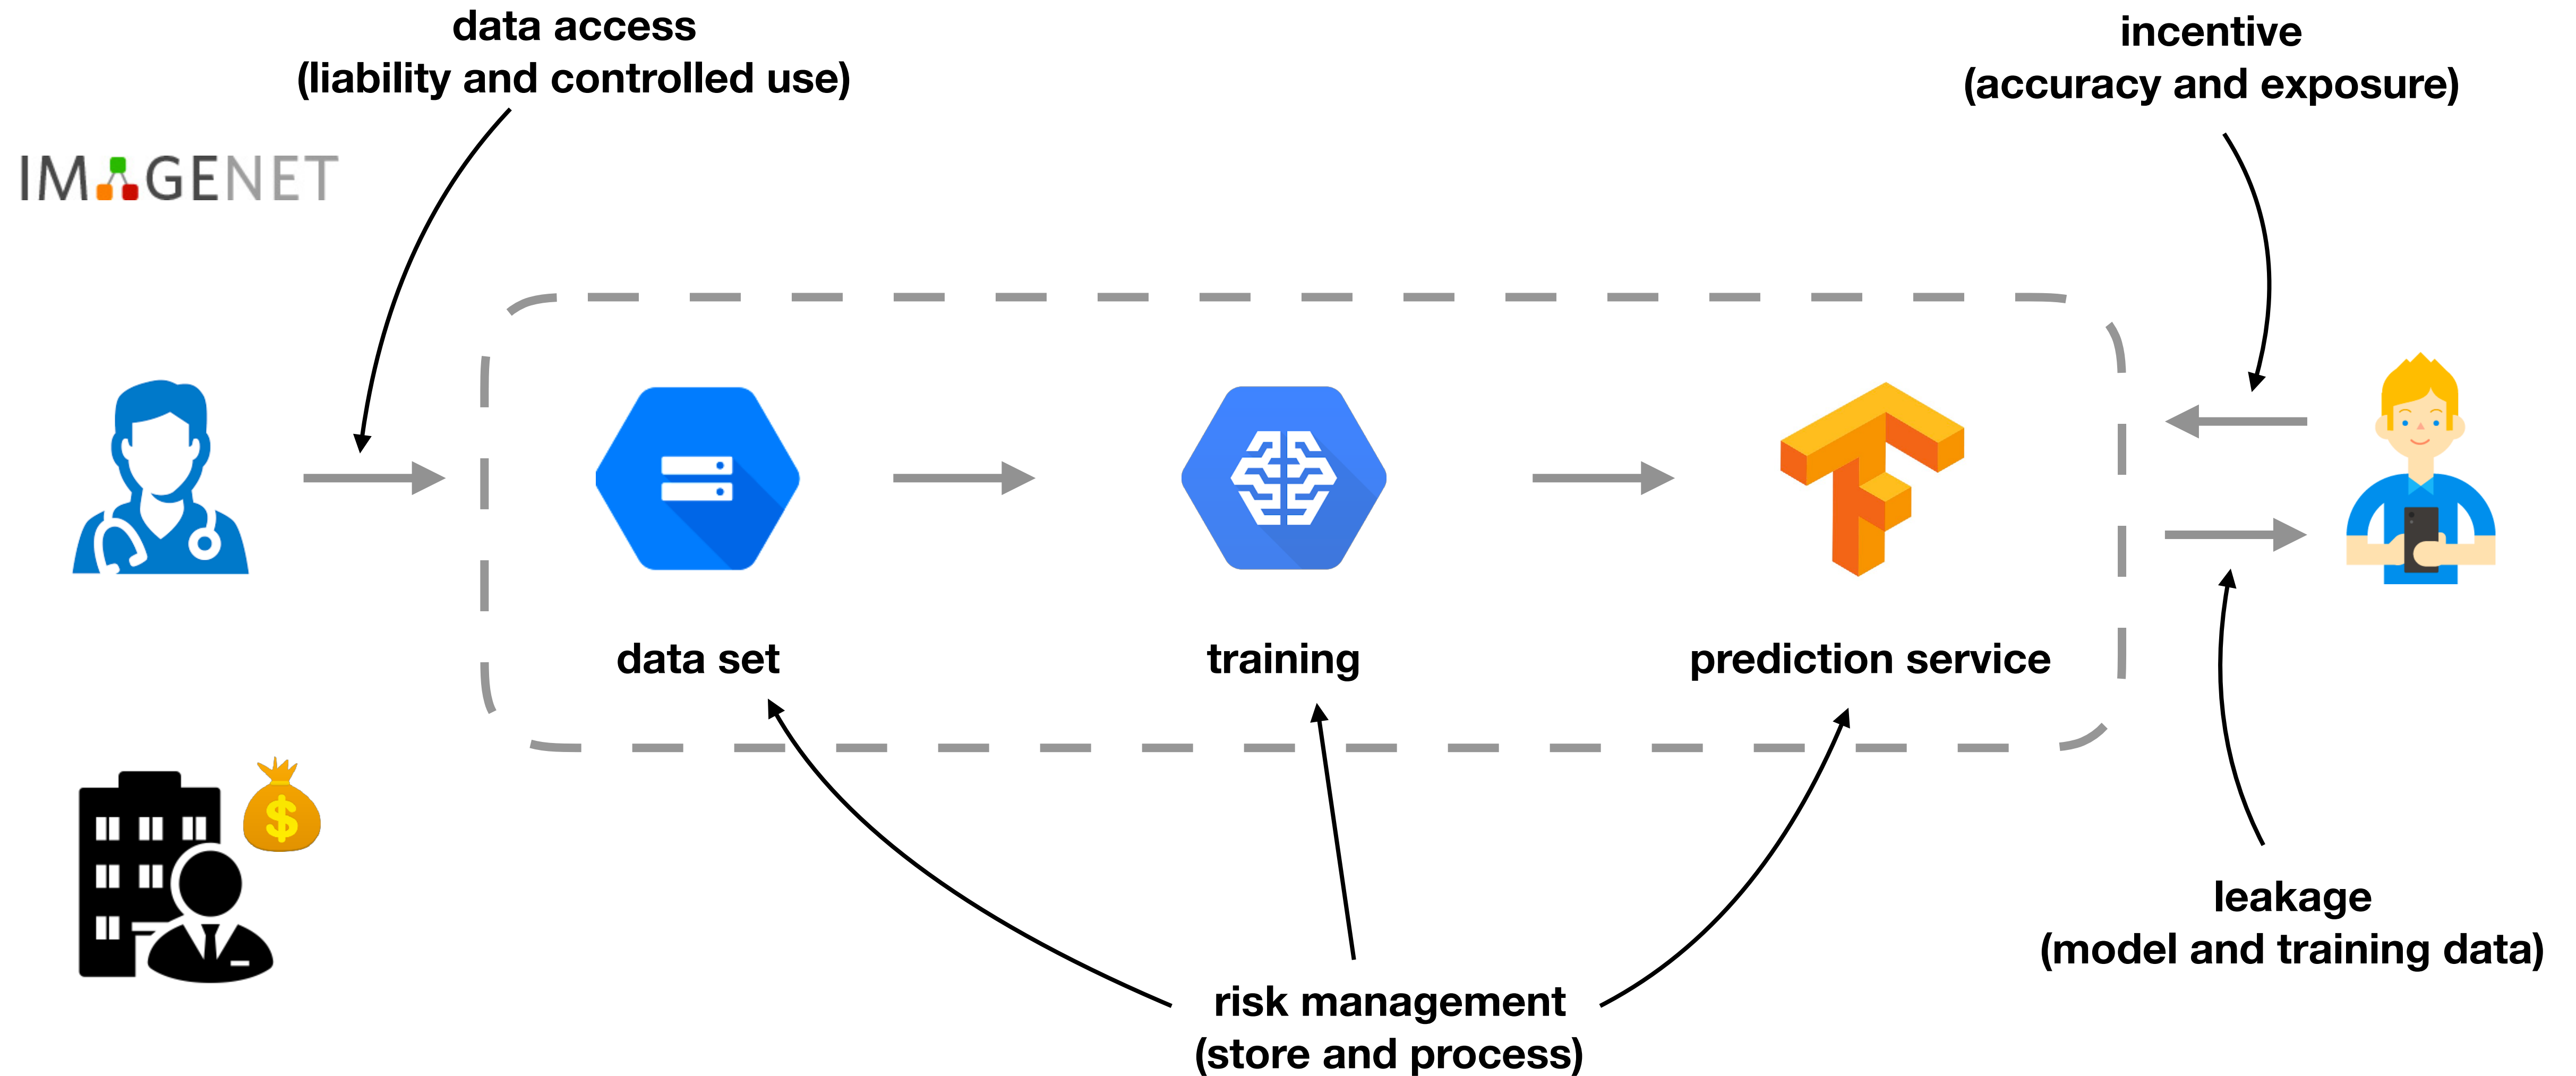
12:30-12:40pm

Join Brett Kuprel, and see how TensorFlow was used by the artificial intelligence lab and medical school of Stanford to classify skin cancer images. He'll describe the project steps: from acquiring a dataset, training a deep network, and evaluating of the results. To wrap up, Brett will give his take on the future of skin cancer image classification.

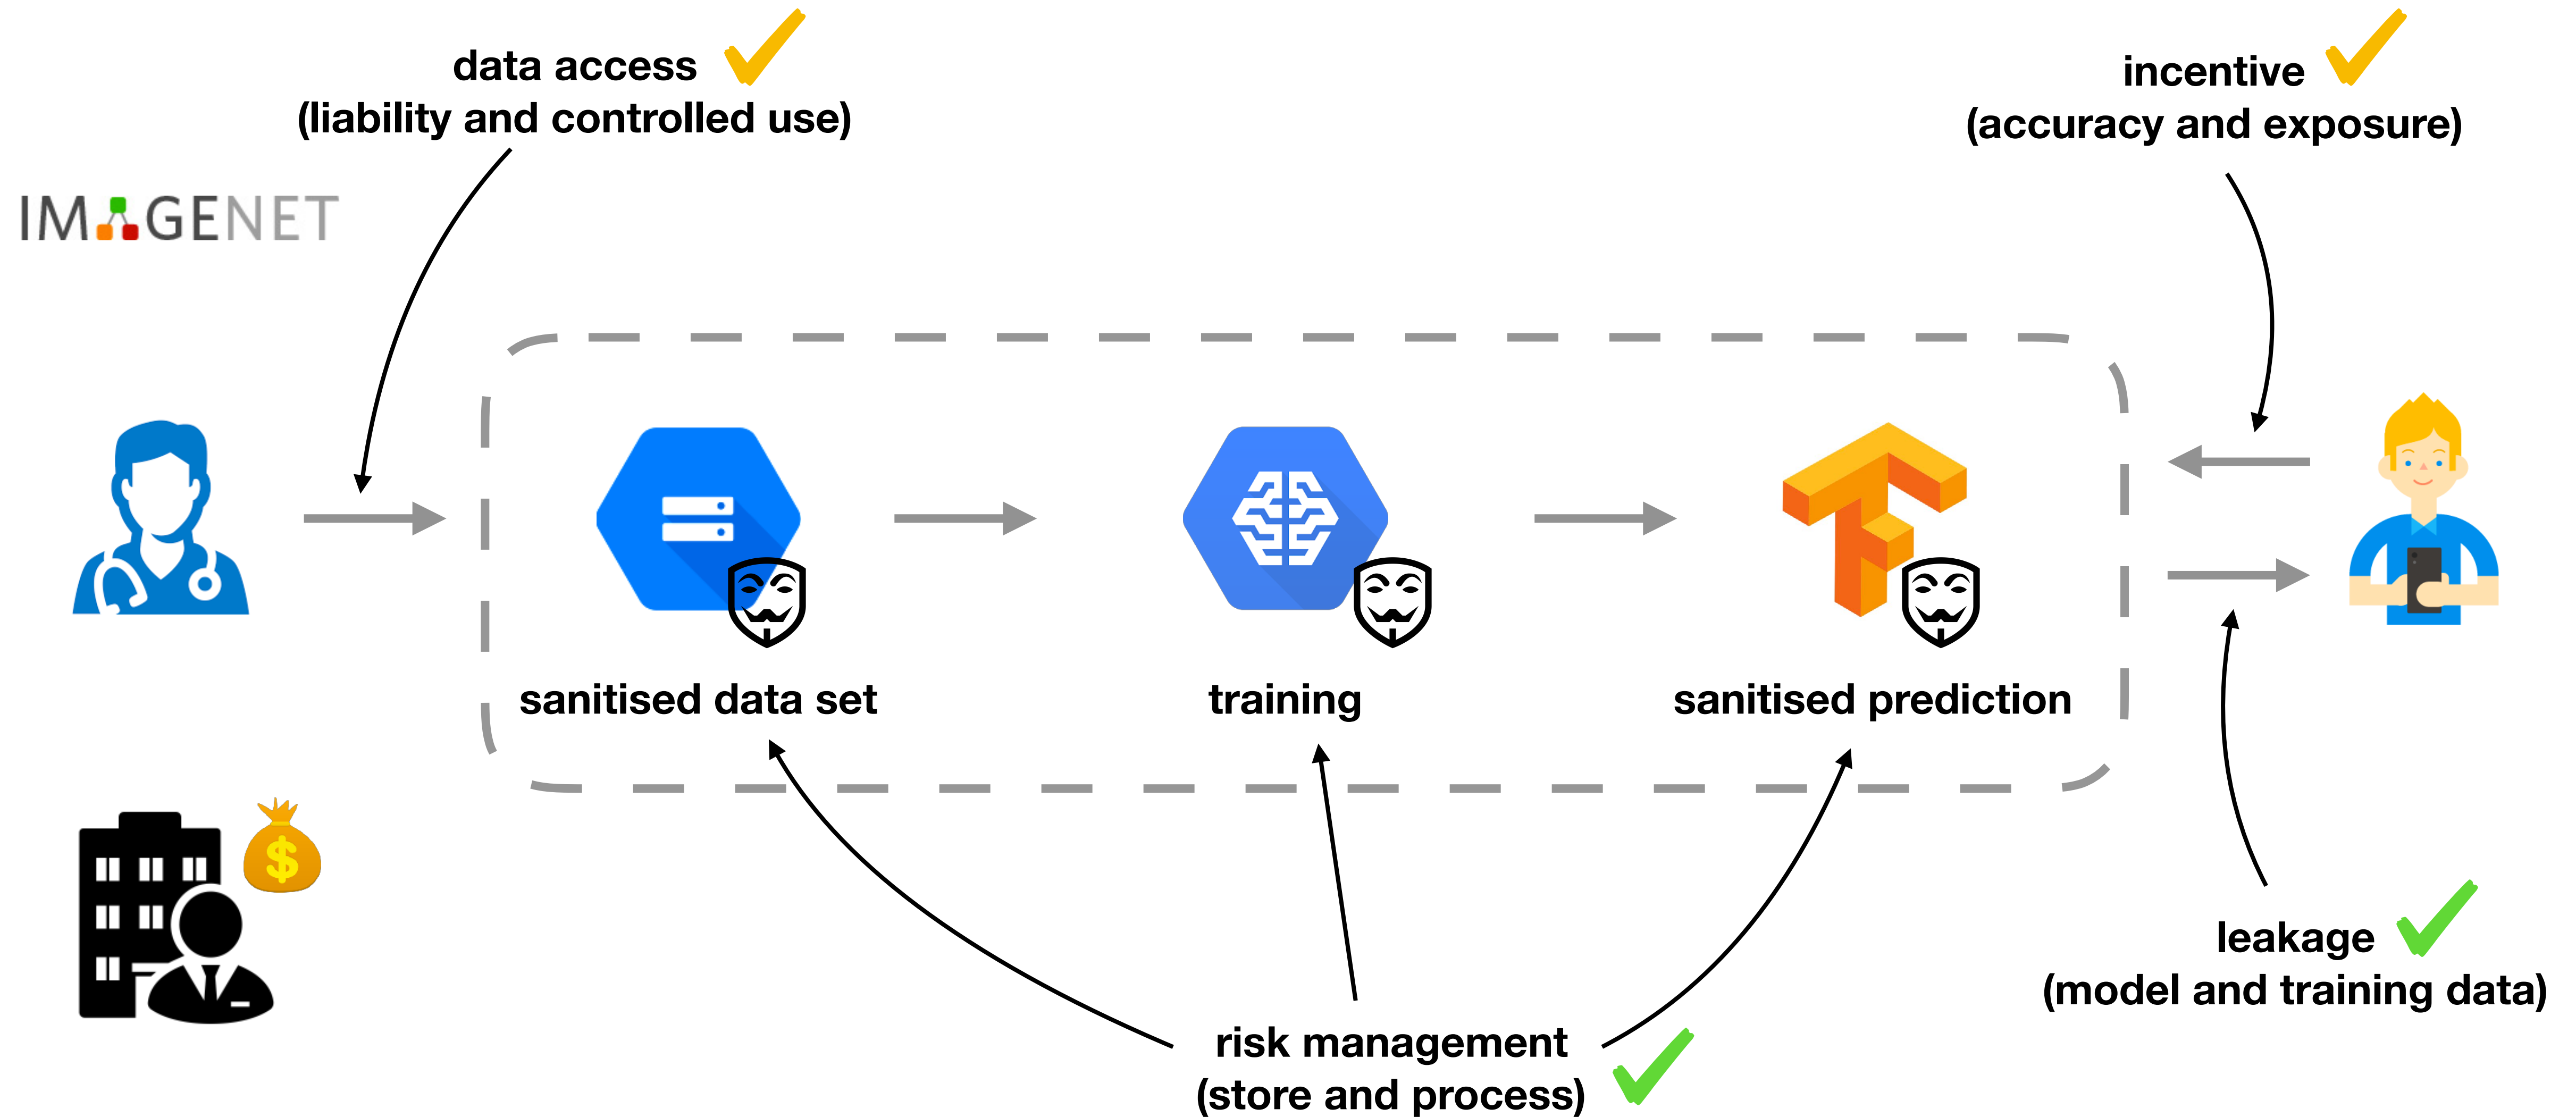


machine learning positioned to have huge impact on health care

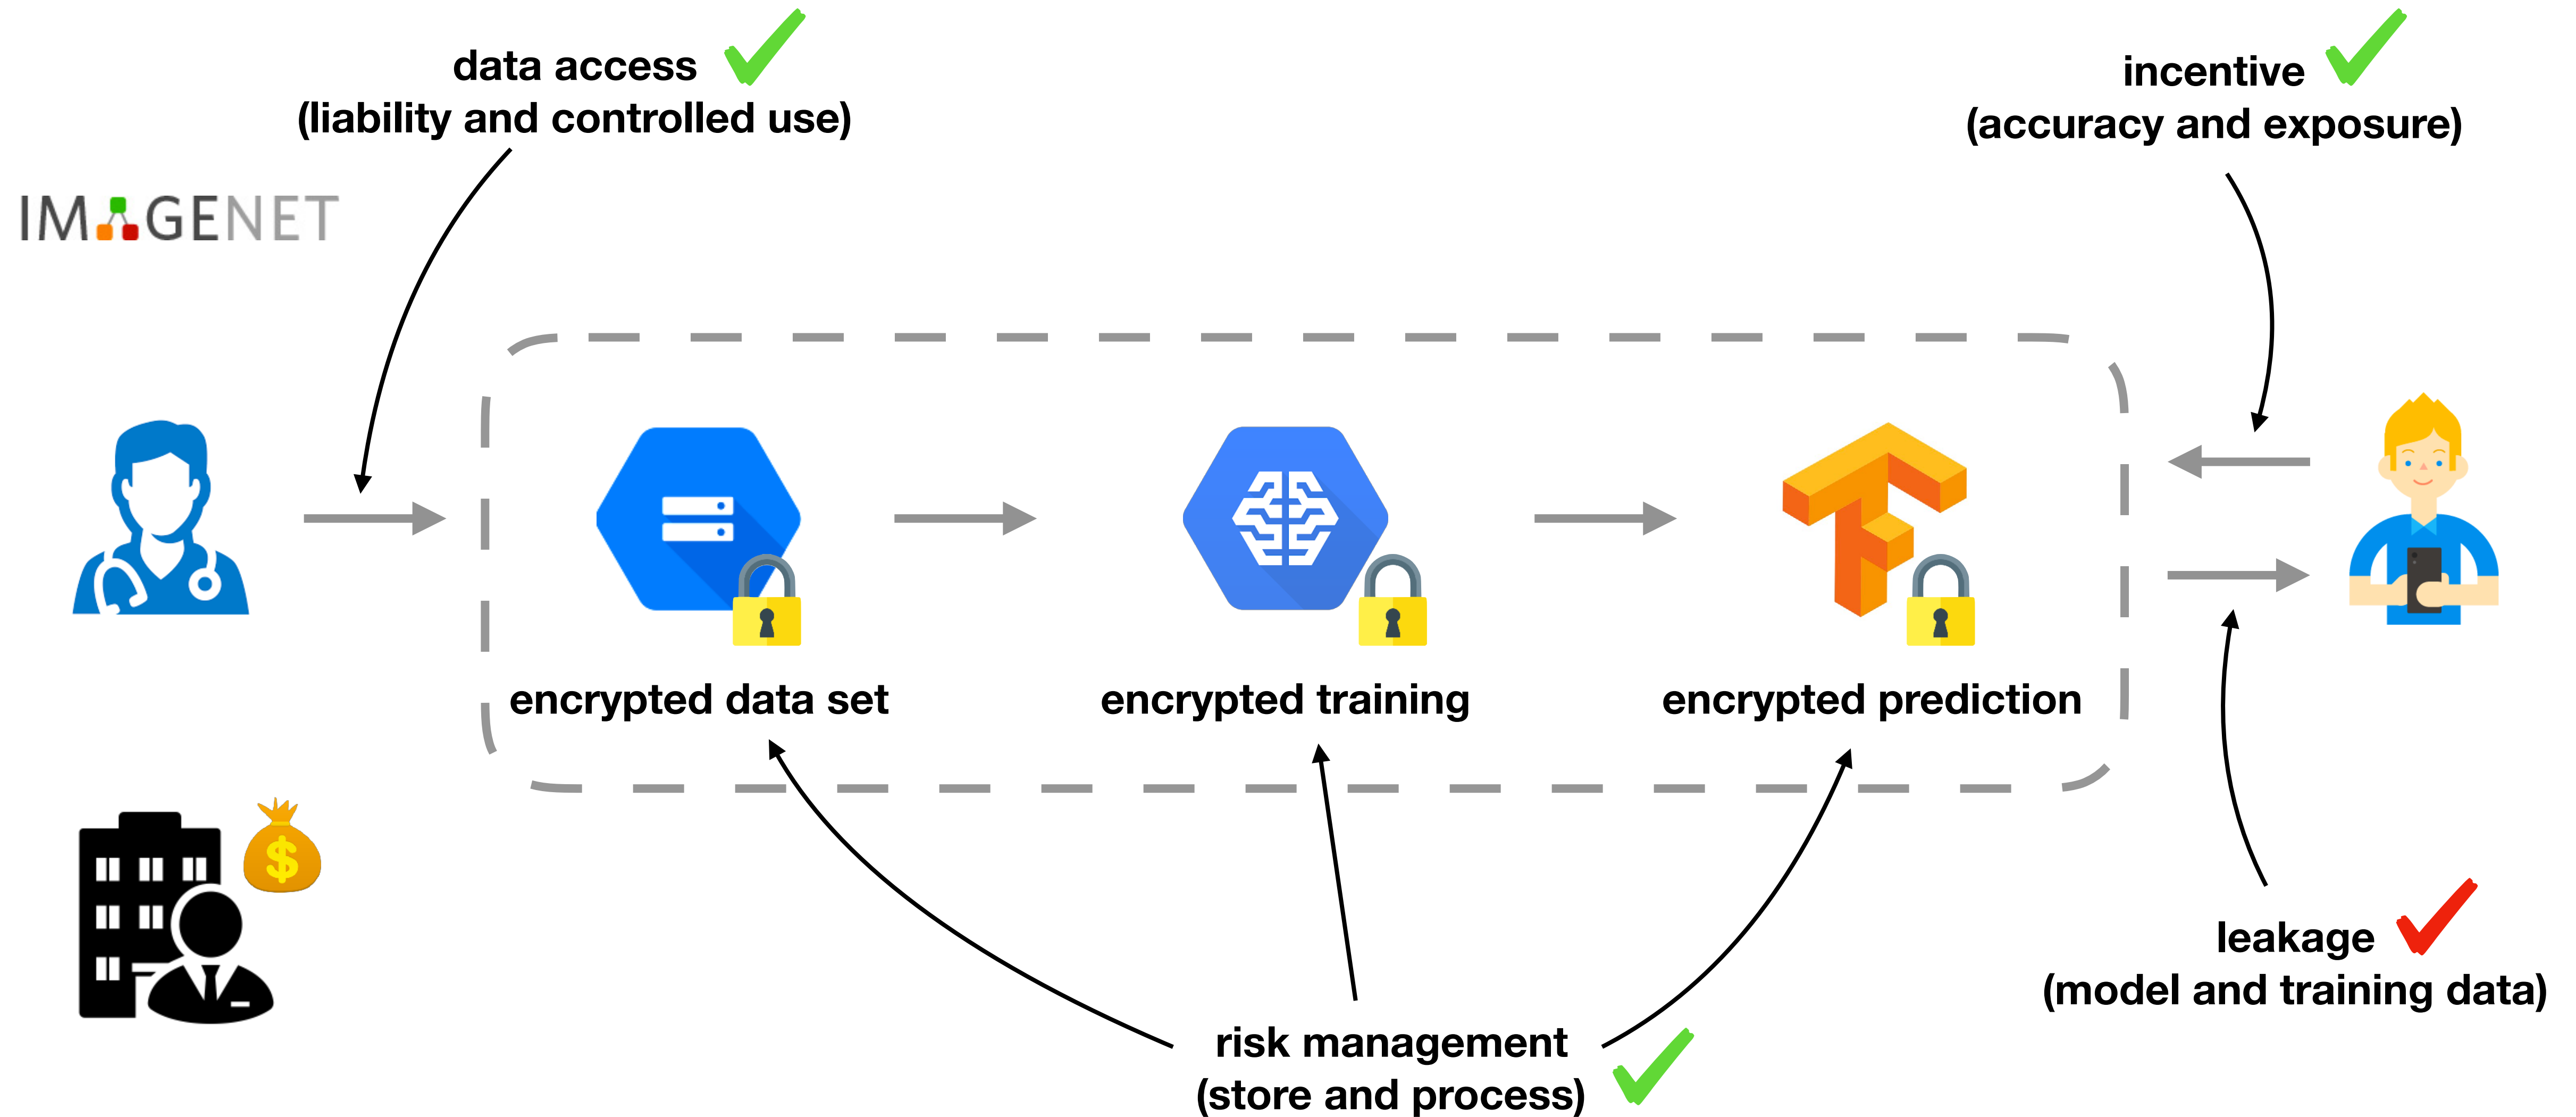
# Potential Bottlenecks



# Differential Privacy

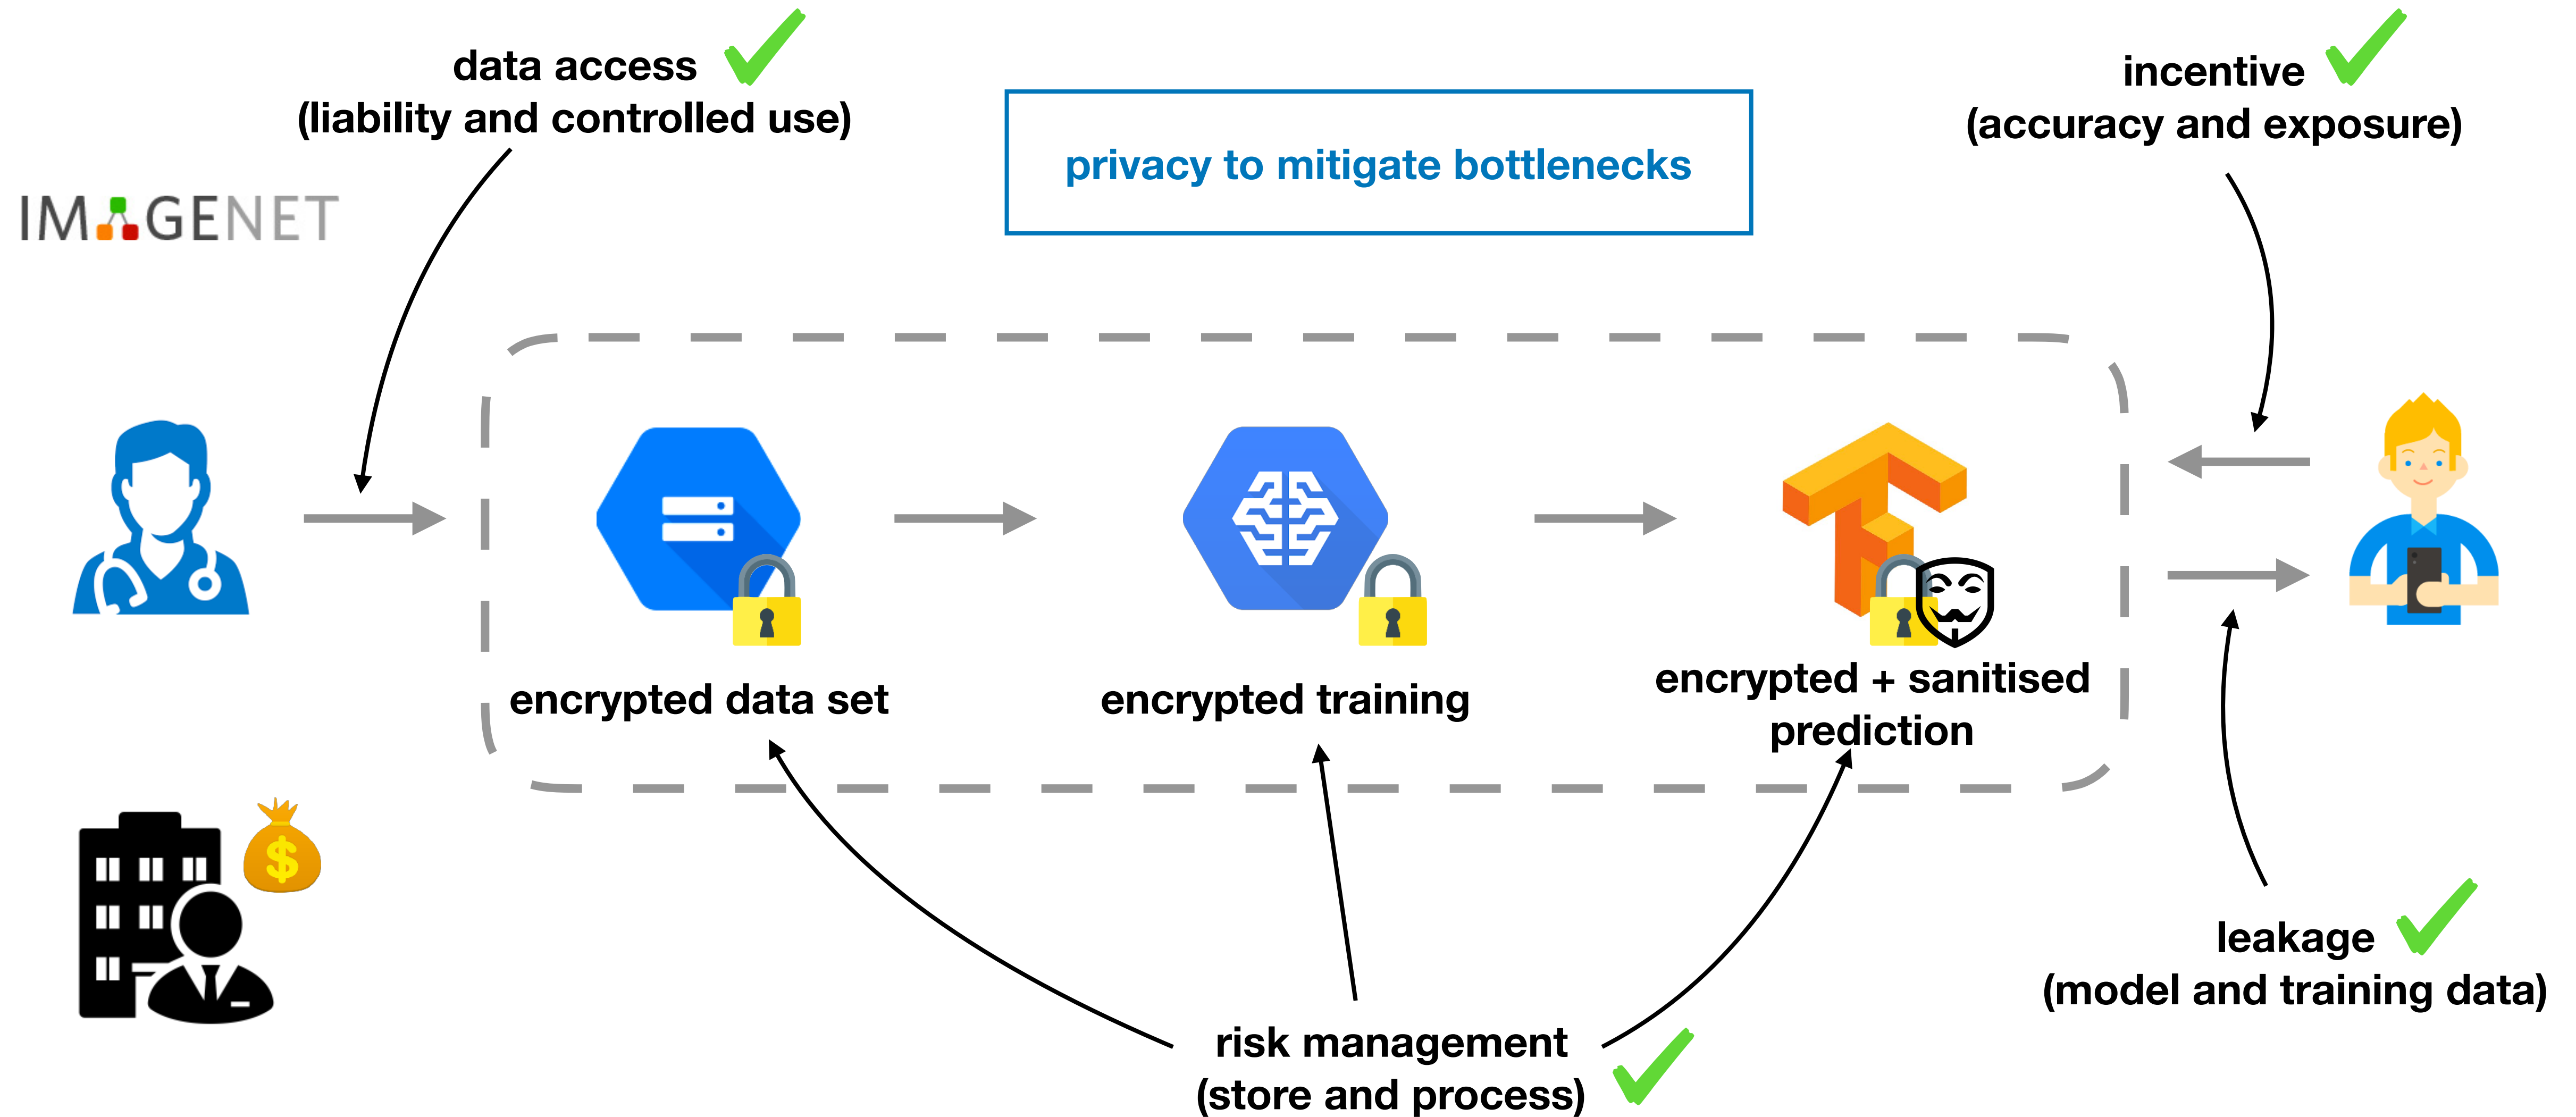


# Secure Computation





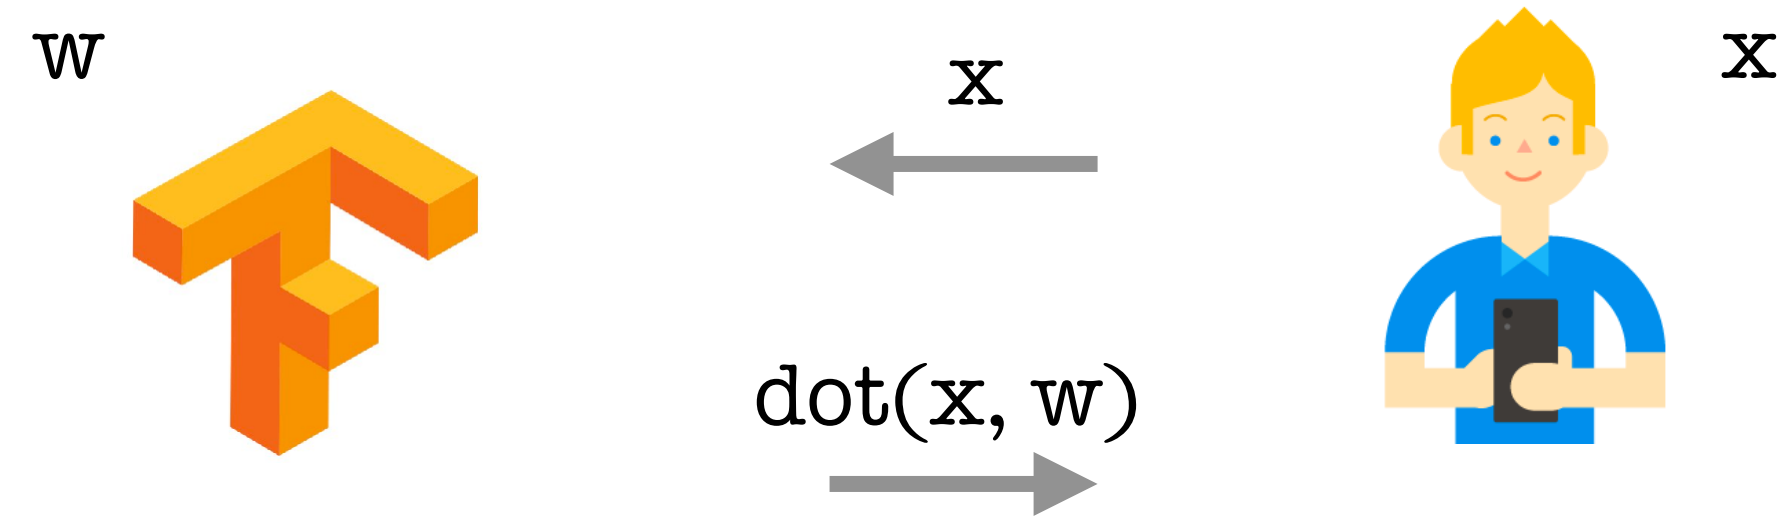
# Hybrid





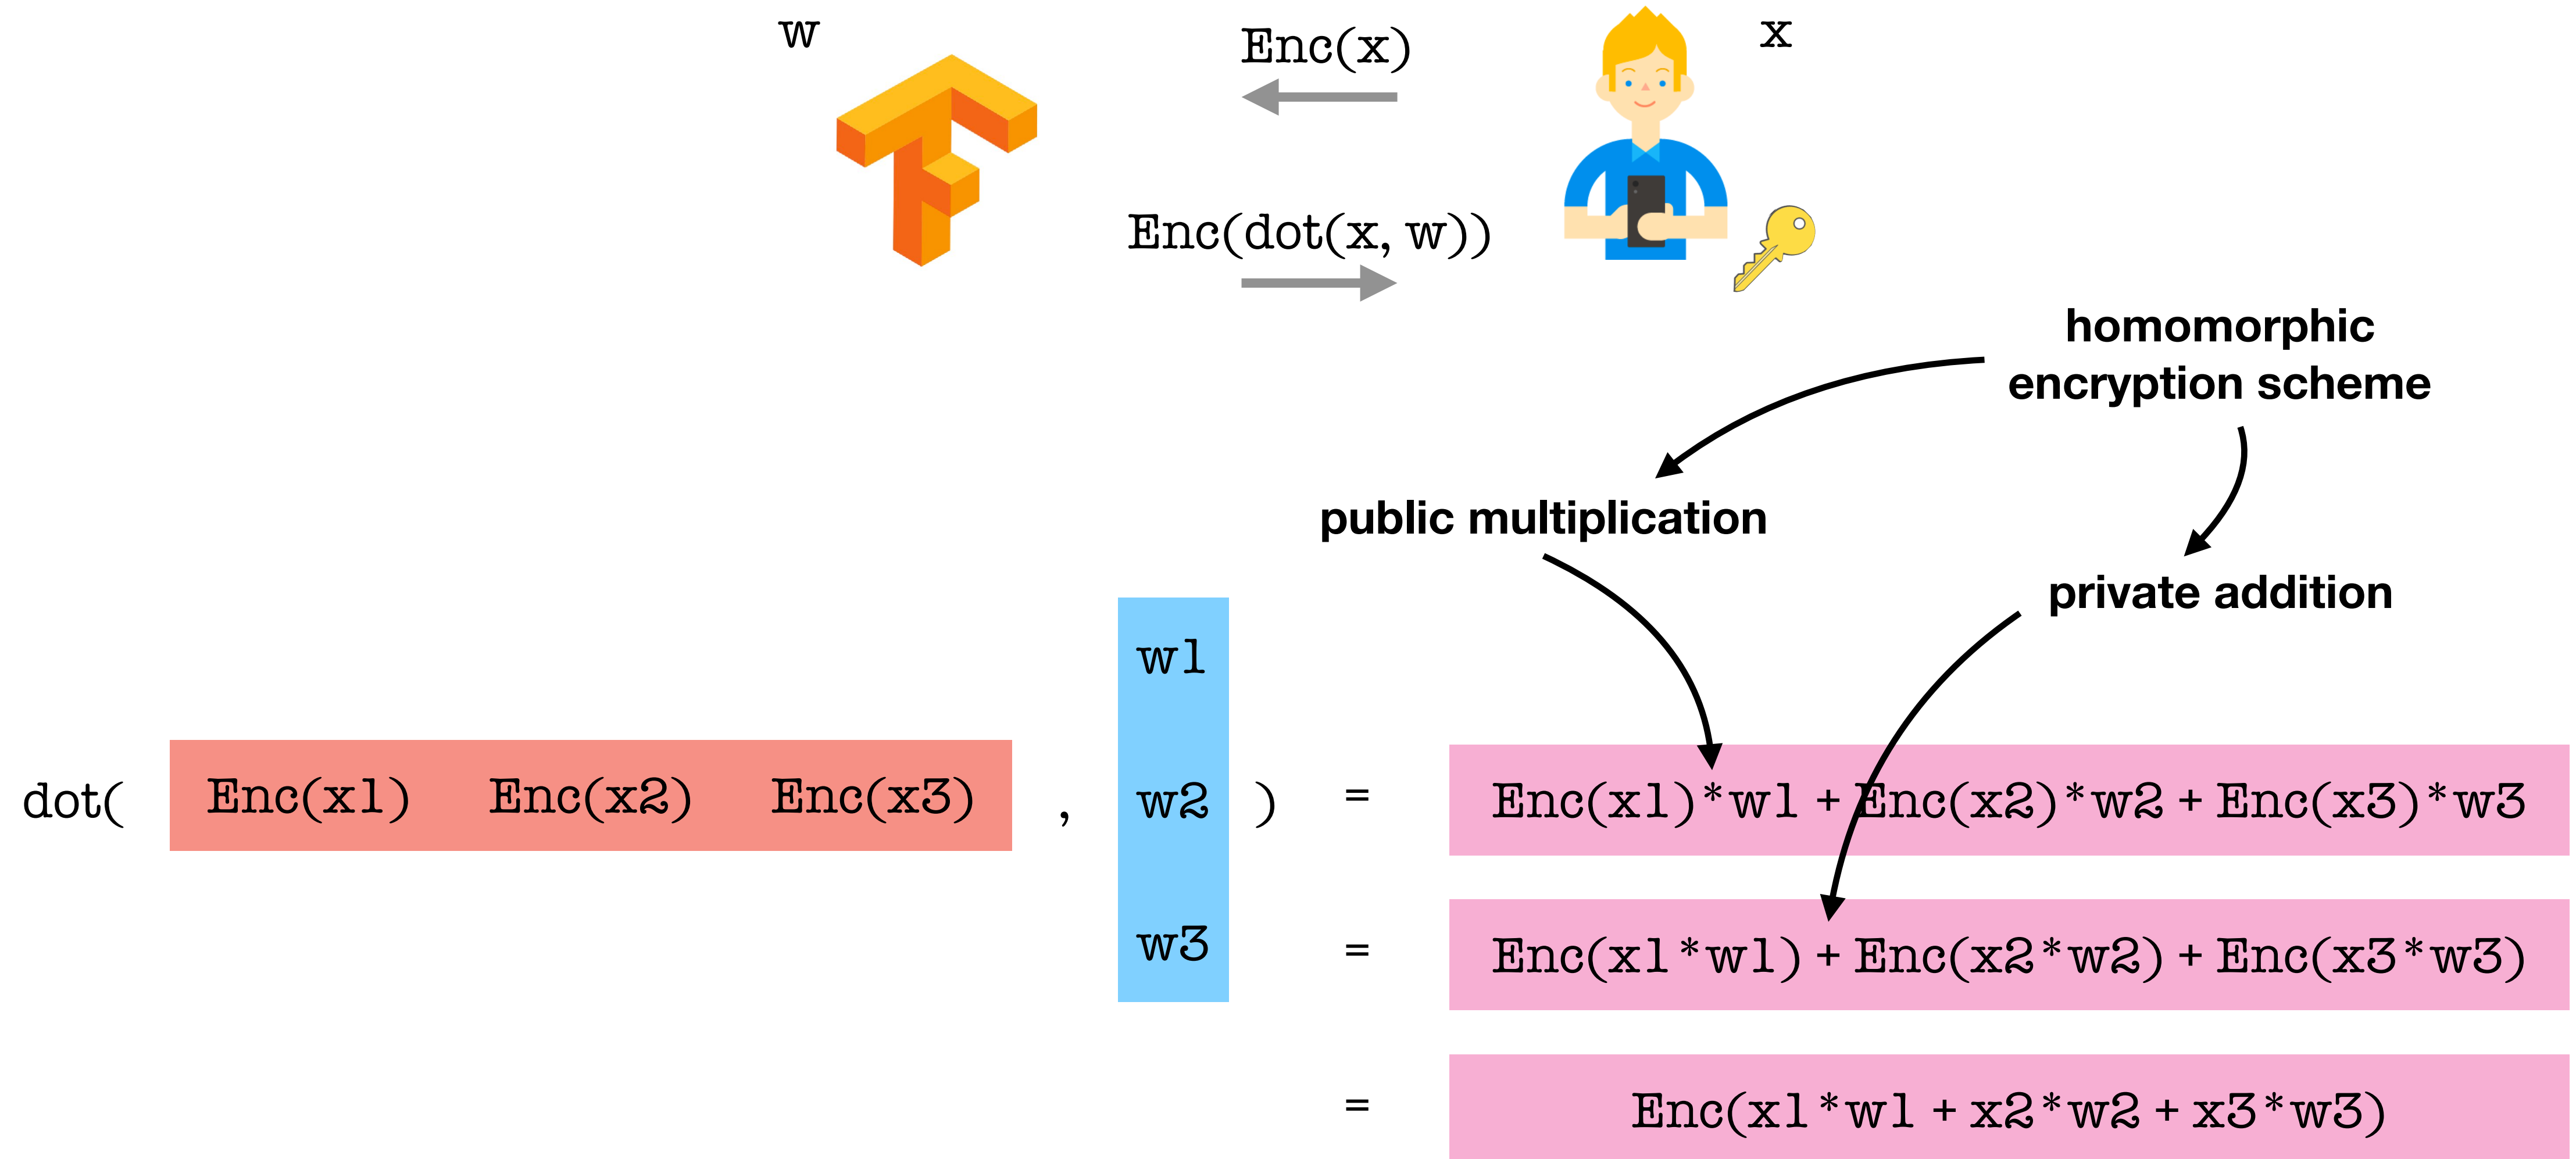
# Prediction

# Prediction with Linear Model

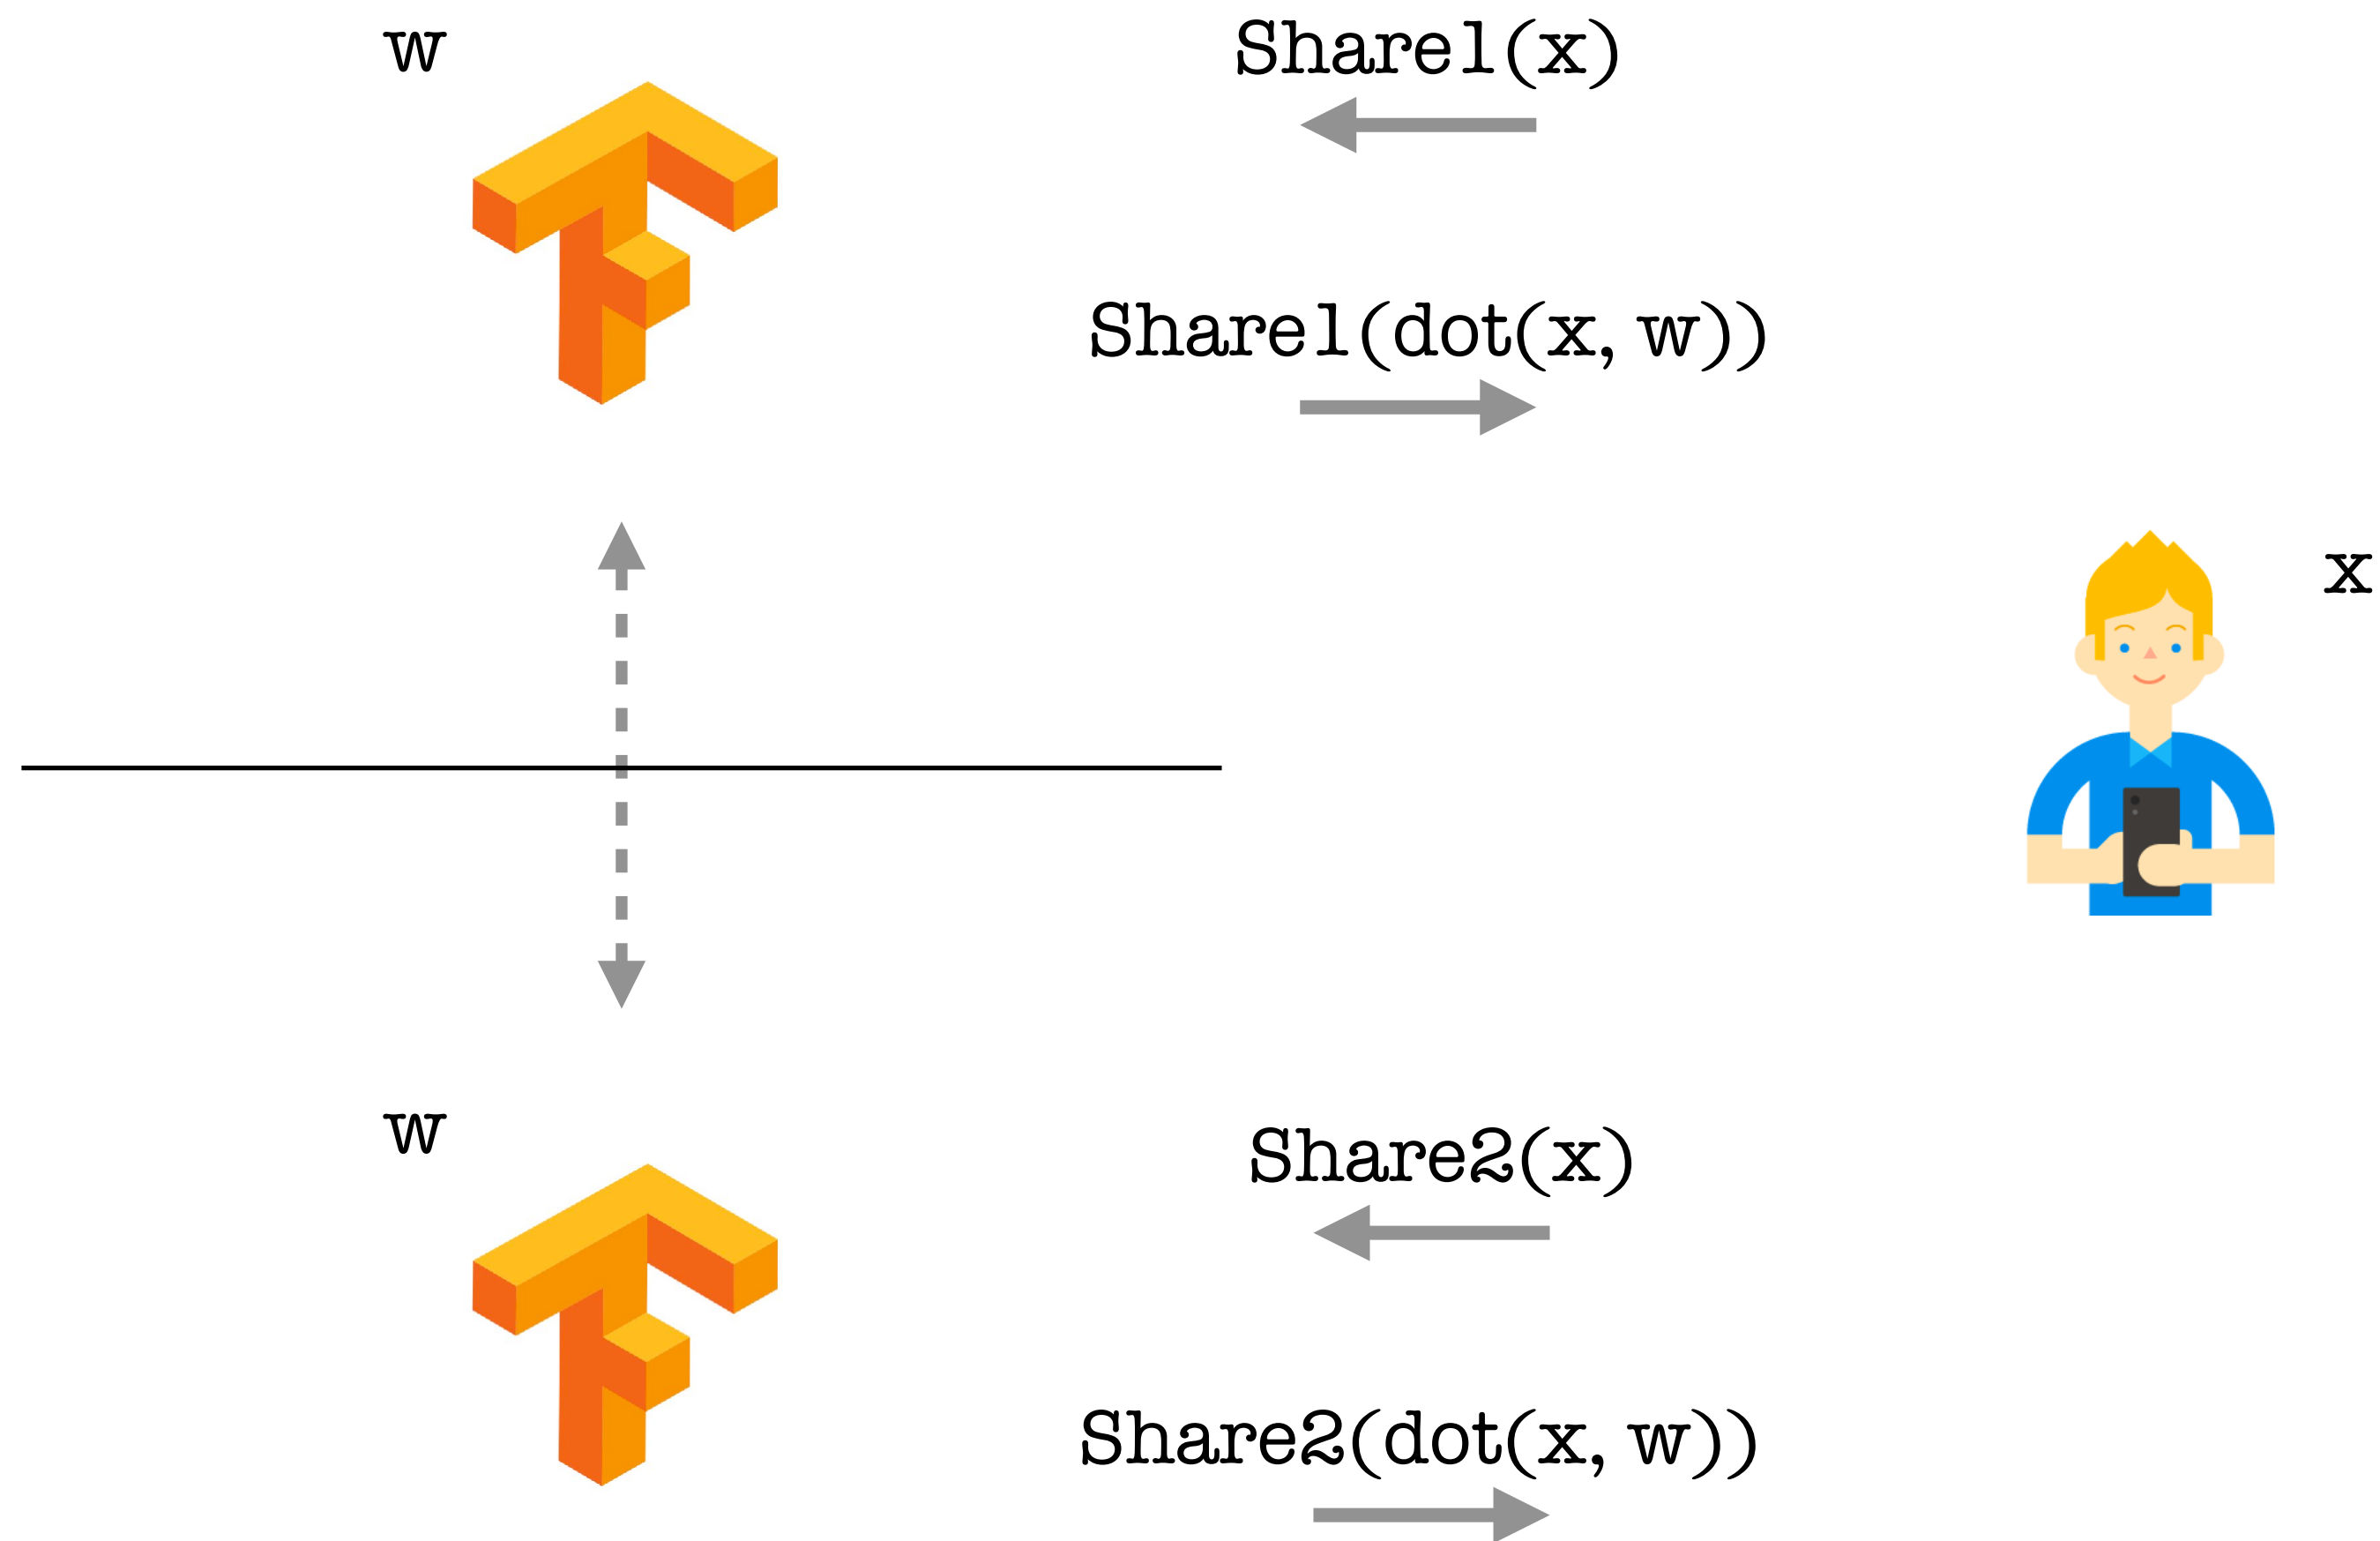


$$\text{dot}\left( \begin{array}{|c|c|c|} \hline x1 & x2 & x3 \\ \hline \end{array}, \begin{array}{|c|} \hline w1 \\ w2 \\ w3 \\ \hline \end{array} \right) = \begin{array}{|c|} \hline x1 * w1 + x2 * w2 + x3 * w3 \\ \hline \end{array}$$

# ... using Homomorphic Encryption

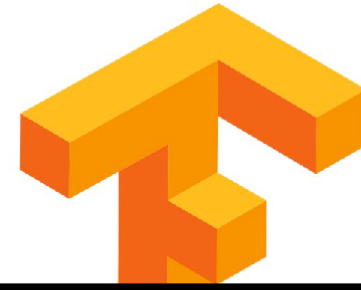


# ... using Secret Sharing



# ... using Secret Sharing

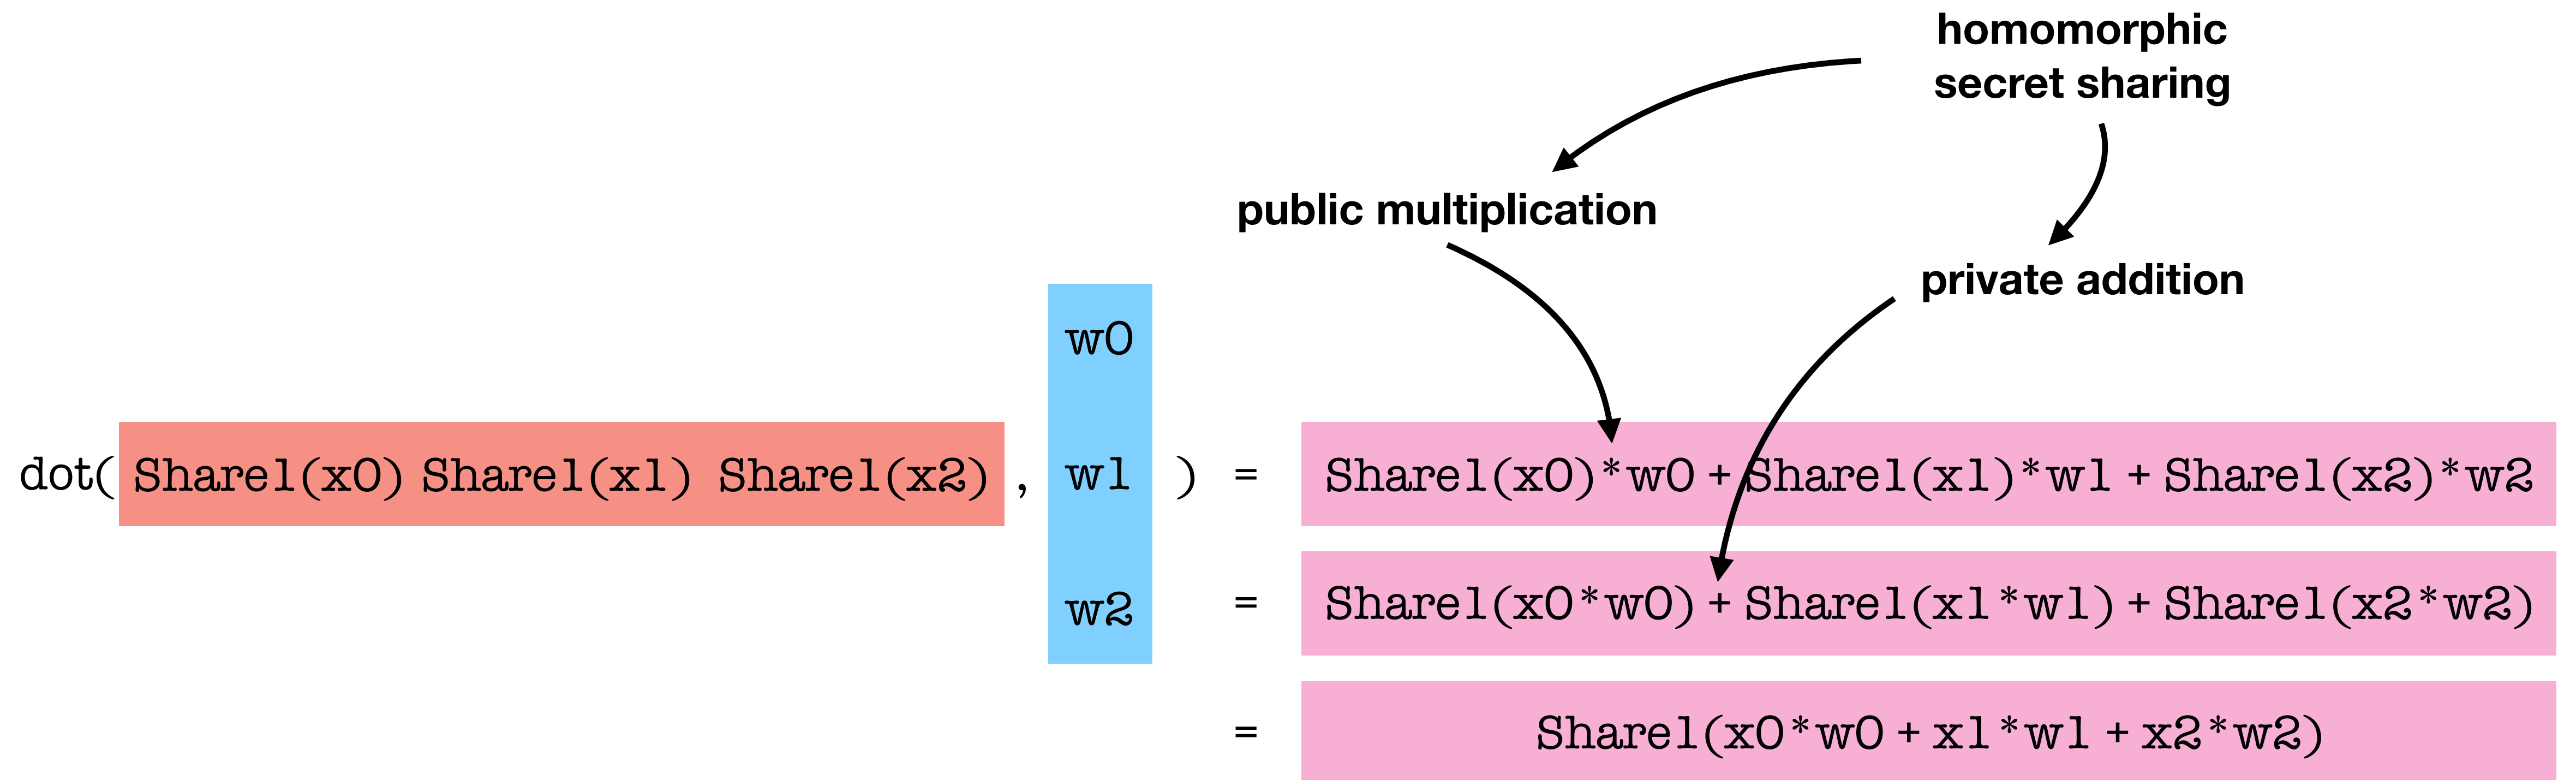
$w$



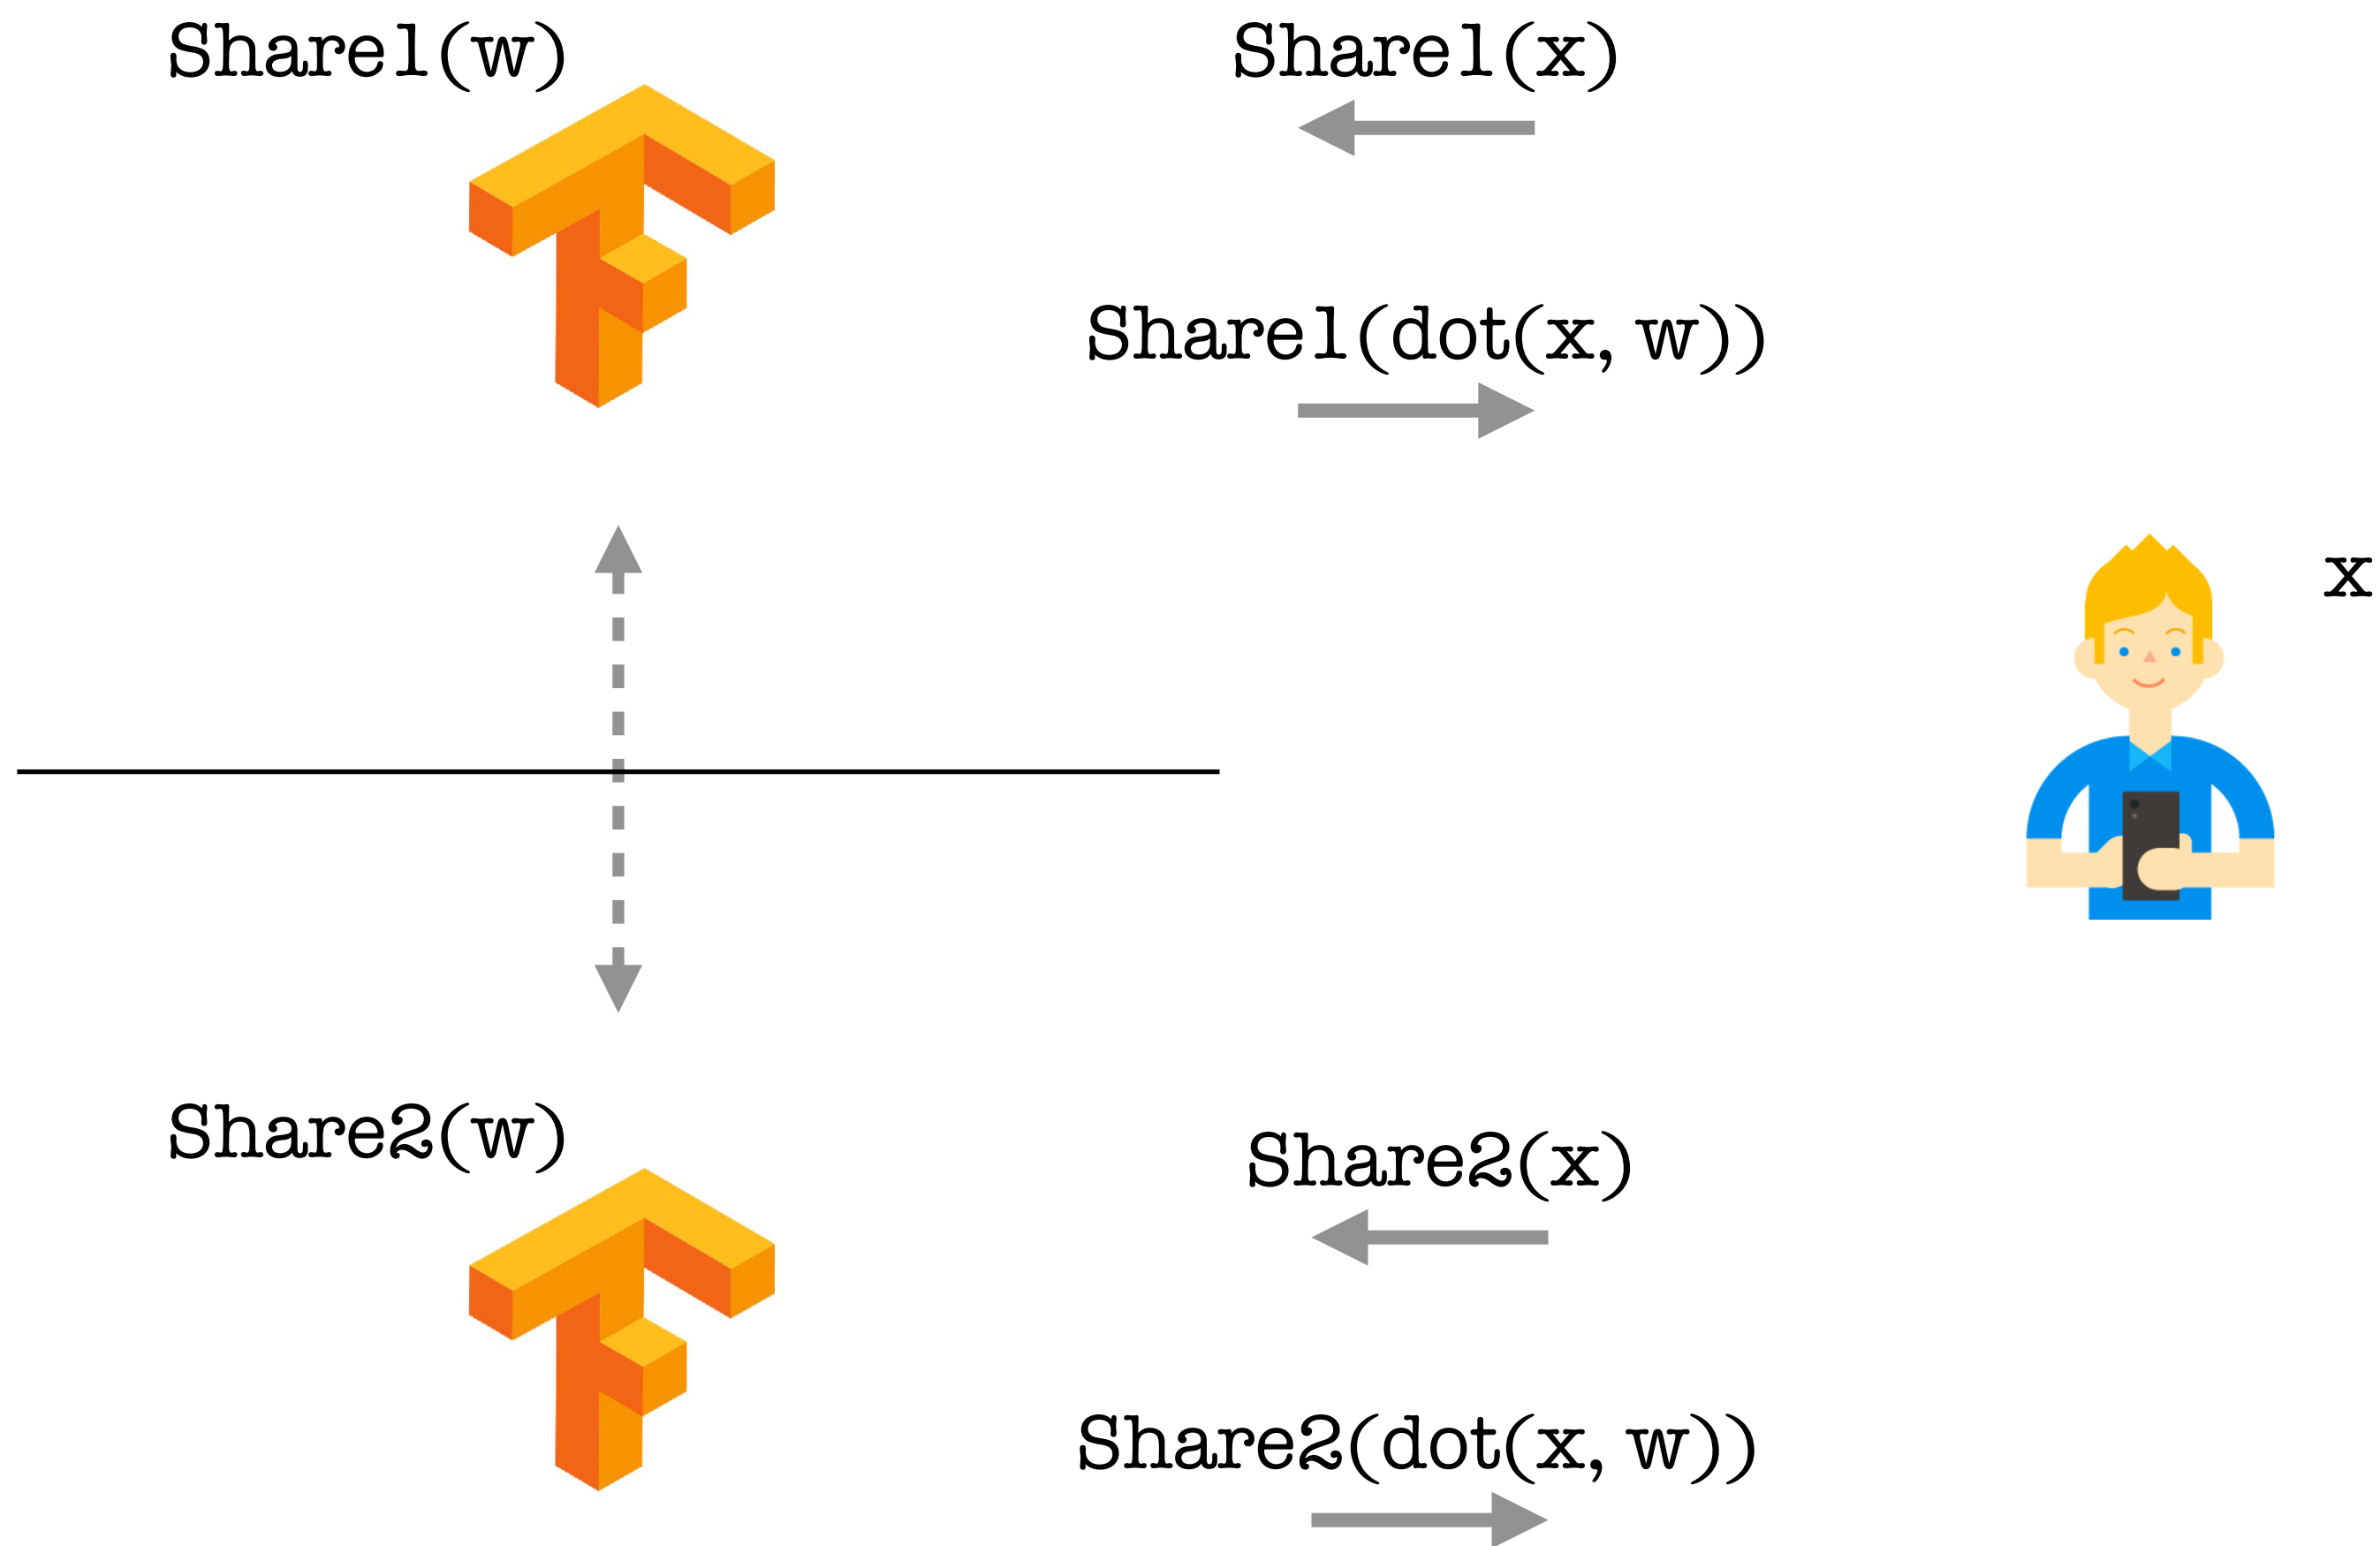
$\text{Share1}(x)$



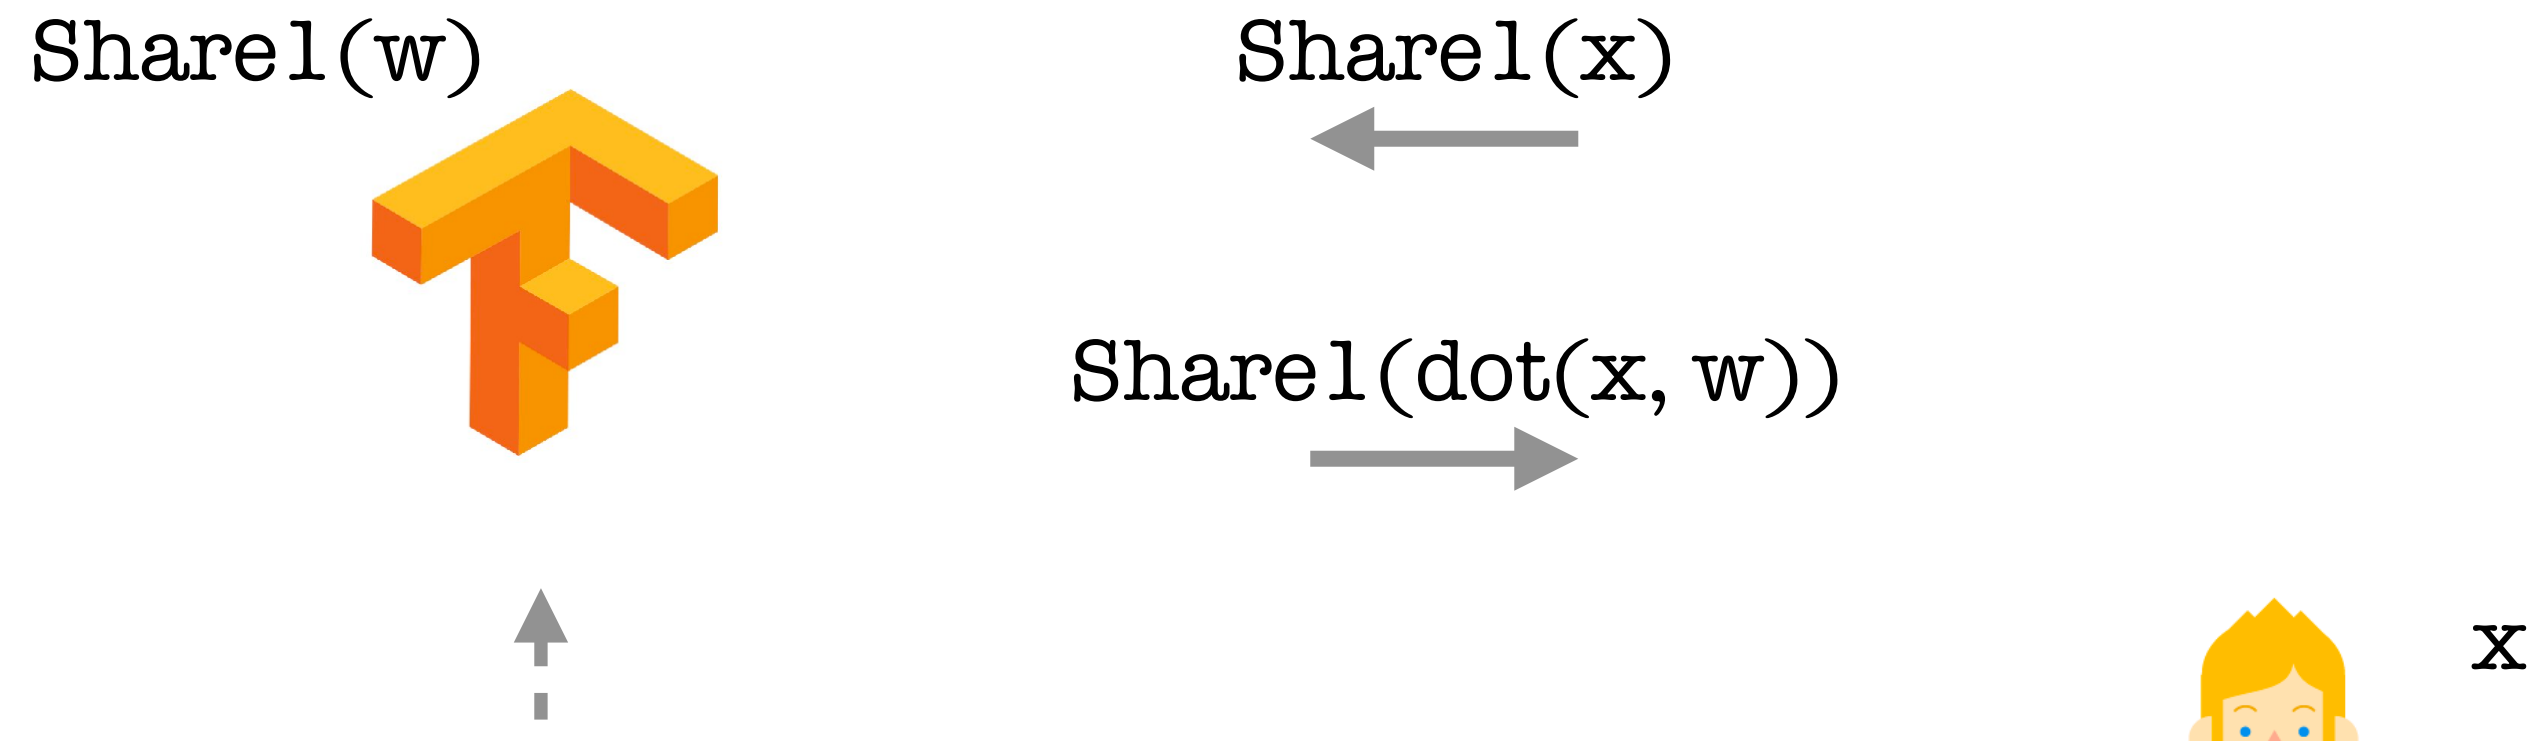
$\text{Share1}(\text{dot}(x, w))$



# ... using Secret Sharing



# ... using Secret Sharing



**private multiplication**

$$\text{dot}(\text{Share1}(x_0) \text{ Share1}(x_1) \text{ Share1}(x_2), \begin{matrix} \text{Share}(w_0) \\ \text{Share}(w_1) \\ \text{Share}(w_2) \end{matrix}) = \begin{matrix} \text{Share1}(x_0) * \text{Share}(w_0) + \dots \\ \text{Share1}(x_0 * w_0) + \dots \\ \text{Share1}(x_0 * w_0 + x_1 * w_1 + x_2 * w_2) \end{matrix}$$



# Training

# Server-Aided



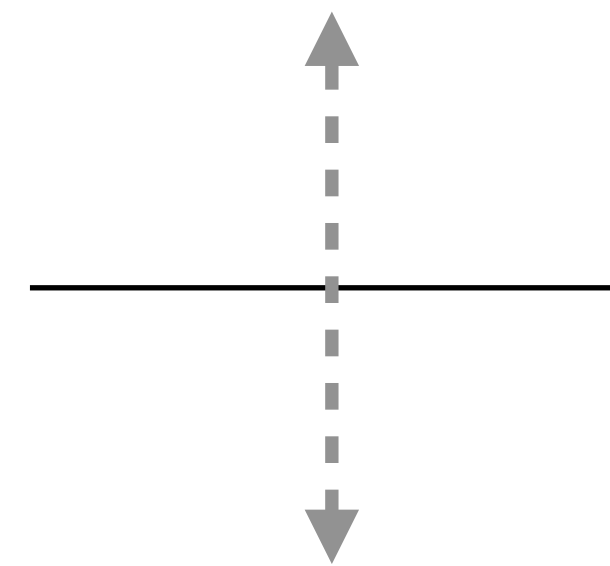
$\text{Share1}(x2, y2)$



$\text{Share2}(x2, y2)$



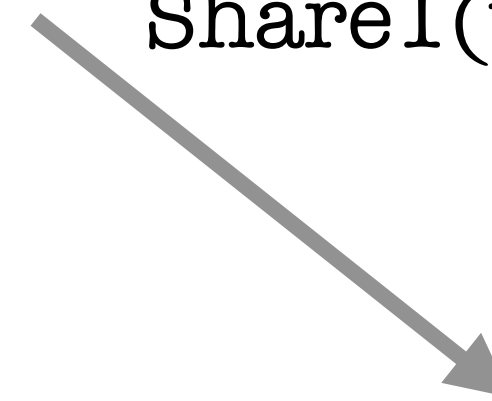
# Server-Aided



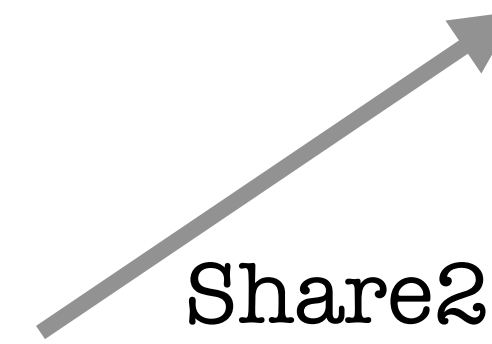
# Server-Aided



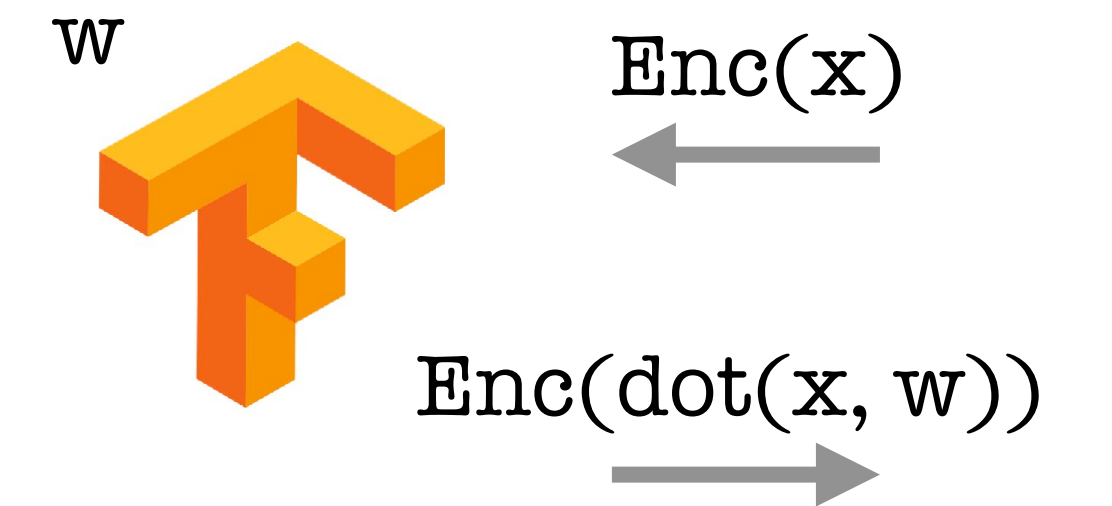
Share1(w)



Share2(w)



# Server-Aided



# Server-Aided



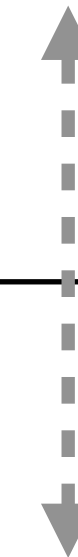
Share1(w)



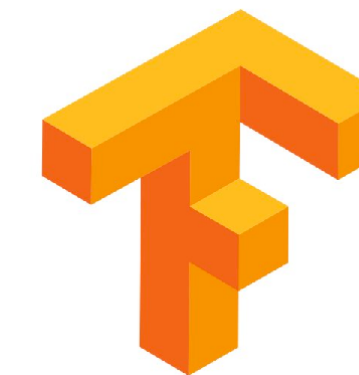
Share1(x)



Share1(pred)



Share2(w)



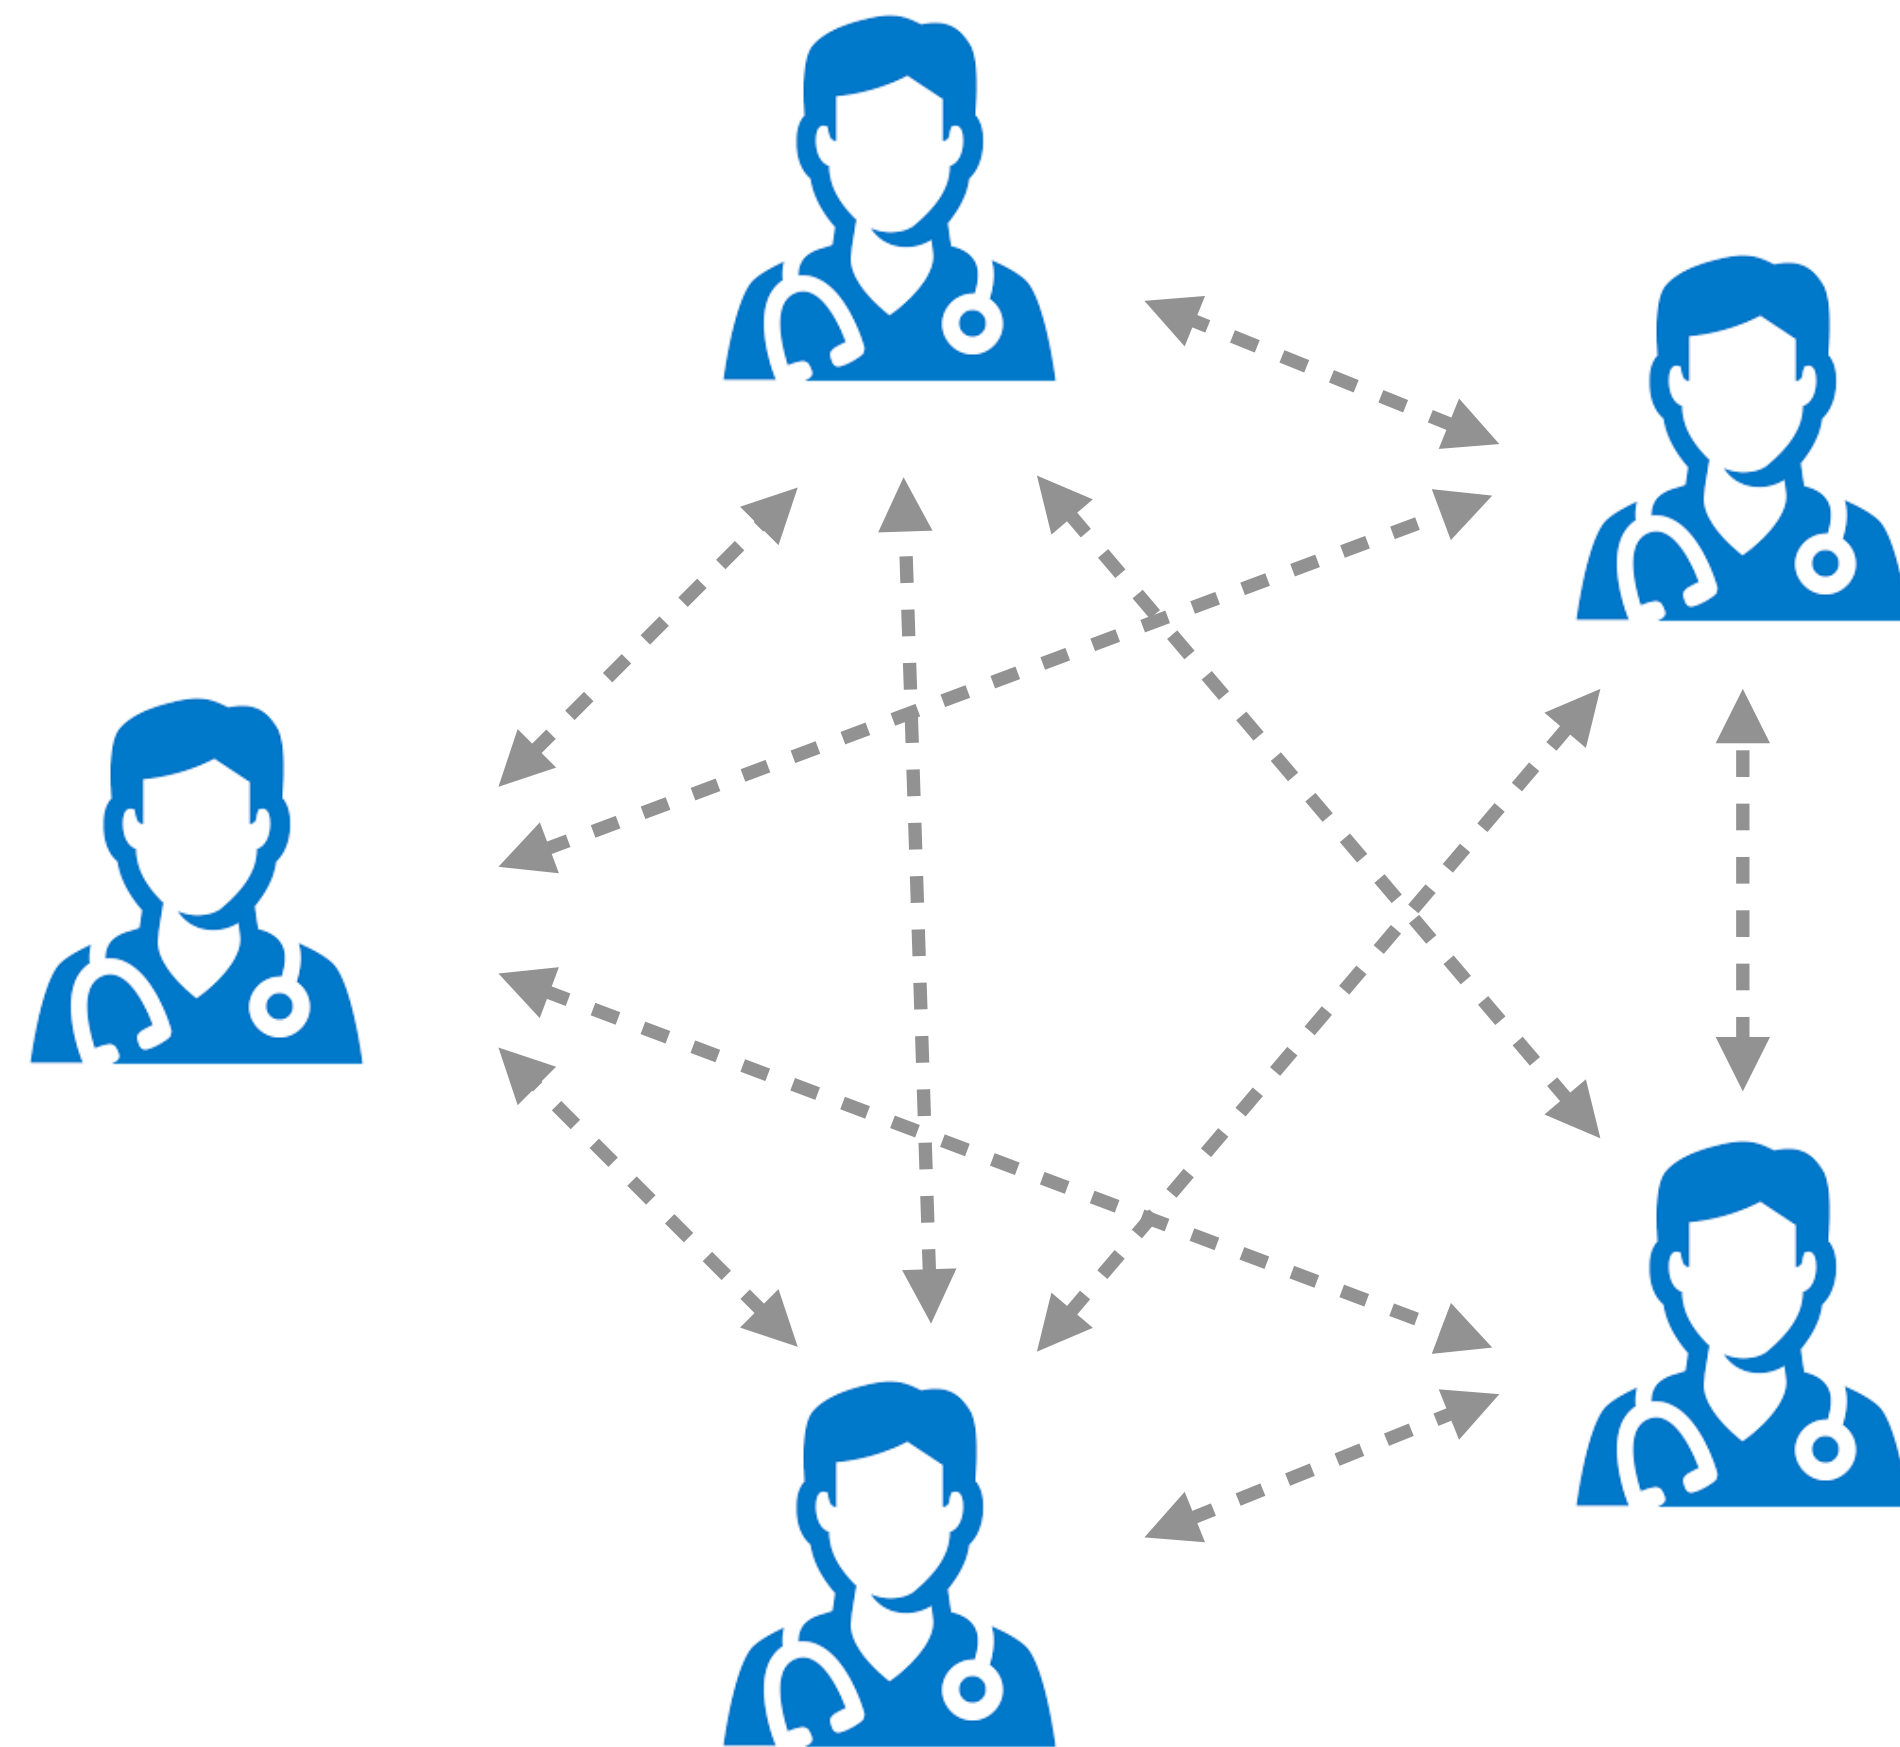
Share2(x)



Share2(pred)



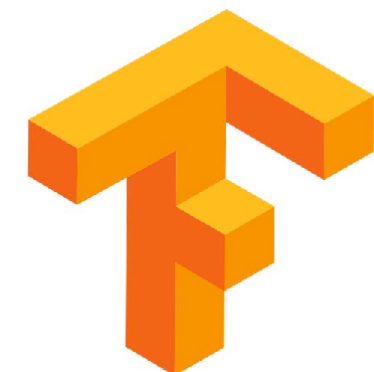
# Multi-Party



Share1(w)



Share2(w)





**Making It Accessible**

# Projects and Literature

## Recent research papers using secure computation

**CryptoNets**: *Applying Neural Networks to Encrypted Data with High Throughput and Accuracy*, Dowlin et al.

**SecureML**: *A System for Scalable Privacy-Preserving Machine Learning*, Mohassel and Zhang

**DeepSecure**: *Scalable Provably-Secure Deep Learning*, Rouhani et al.

**Gazelle**: *A Low Latency Framework for Secure Neural Network Inference*, Juvekar et al.

**ABY3**: *A Mixed Protocol Framework for Machine Learning*, Mohassel and Rindal

**SecureNN**: *Efficient and Private Neural Network Training*, Wagh et al.

**Blind Justice**: *Fairness with Encrypted Sensitive Attributes*, Kilbertus et al.

(also great summary in <https://eprint.iacr.org/2017/1190>)

## Specialised projects

**tf-encrypted** (<https://github.com/mortendahl/tf-encrypted>)

**PySyft** (<https://github.com/OpenMined/PySyft>)

## Secure computation frameworks

**SCALE-MAMBA** (<https://homes.esat.kuleuven.be/~nsmart/SCALE/>)

**MP-SPDZ** (<https://github.com/n1analytics/MP-SPDZ>)

**ABY** (<https://github.com/encryptogroup/ABY>)

**OblivC** (<http://oblivc.org/>)

(much more at <https://github.com/rdragos/awesome-mpc>)

# Multidisciplinary Challenge

Data science  
(use-cases, workflow, monitoring)

Cryptography  
(techniques, protocols, trust)

Machine learning  
(models, approx, precision)

Engineering  
(distributed, multi-core, readability)

**need for common language**

# tf-encrypted

open source project for exploring and experimenting with  
privacy-preserving machine learning in TensorFlow

separate concerns, take expertise out of equation,  
and provide tight integration with ecosystem

# Private Prediction with tf-encrypted

```
1  import tensorflow as tf
2
3  def provide_weights():""" Load from disk """
4  def provide_input(): """ Pre-process """
5  def receive_output(logits): return tf.Print([], [tf.argmax(logits)])
6
7  # get model weights
8  w0, b0, w1, b1, w2, b2 = provide_weights()
9
10 # get prediction input
11 x = provide_input()
12
13 # compute prediction
14 layer0 = tf.nn.relu((tf.matmul(x, w0) + b0))
15 layer1 = tf.nn.relu((tf.matmul(layer0, w1) + b1))
16 logits = tf.matmul(layer1, w2) + b2
17
18 # process result of prediction
19 prediction_op = receive_output(logits)
20
21 # run graph execution in a tf.Session
22 with tf.Session() as sess:
23     sess.run(tf.global_variables_initializer())
24     sess.run(prediction_op)
```

```
1  import tensorflow as tf
2  import tf_encrypted as tfe
3
4  def provide_weights():""" Load from disk """
5  def provide_input(): """ Pre-process """
6  def receive_output(logits): return tf.Print([], [tf.argmax(logits)])
7
8  # get model weights as private tensors from owner
9  w0, b0, w1, b1, w2, b2 = tfe.define_private_input("model-owner", provide_weights)
10
11 # get prediction input as private tensors from client
12 x = tfe.define_private_input("prediction-client", provide_input)
13
14 # compute private prediction on servers
15 layer0 = tfe.relu((tfe.matmul(x, w0) + b0))
16 layer1 = tfe.relu((tfe.matmul(layer0, w1) + b1))
17 logits = tfe.matmul(layer1, w2) + b2
18
19 # process result of prediction on client
20 prediction_op = tfe.define_output("prediction-client", logits, receive_output)
21
22 # run secure graph execution in a tf.Session
23 with tfe.Session() as sess:
24     sess.run(tf.global_variables_initializer())
25     sess.run(prediction_op)
```

# Private Prediction with tf-encrypted

```
1 import tensorflow as tf
2
3 def provide_weights(): """ Load from disk """
4 def provide_input(): """ Pre-process """
5 def receive_output(logits): return tf.Print([], [tf.argmax(logits)])
6
7 # get model weights
8 w0, b0, w1, b1, w2, b2 = provide_weights()
9
10 # get prediction input
11 x = provide_input()
12
13 # compute prediction
14 layer0 = tf.nn.relu((tf.matmul(x, w0) + b0))
15 layer1 = tf.nn.relu((tf.matmul(layer0, w1) + b1))
16 logits = tf.matmul(layer1, w2) + b2
17
18 # process result of prediction
19 prediction_op = receive_output(logits)
20
21 # run graph execution in a tf.Session
22 with tf.Session() as sess:
23     sess.run(tf.global_variables_initializer())
24     sess.run(prediction_op)
```

```
1 import tensorflow as tf
2 import tf_encrypted as tfe
3
4 def provide_weights(): """ Load from disk """
5 def provide_input(): """ Pre-process """
6 def receive_output(logits): return tf.Print([], [tf.argmax(logits)])
7
8 # get model weights as private tensors from owner
9 w0, b0, w1, b1, w2, b2 = tfe.define_private_input("model-owner", provide_weights)
10
11 # get prediction input as private tensors from client
12 x = tfe.define_private_input("prediction-client", provide_input)
13
14 # compute private prediction on servers
15 layer0 = tfe.relu((tfe.matmul(x, w0) + b0))
16 layer1 = tfe.relu((tfe.matmul(layer0, w1) + b1))
17 logits = tfe.matmul(layer1, w2) + b2
18
19 # process result of prediction on client
20 prediction_op = tfe.define_output("prediction-client", logits, receive_output)
21
22 # run secure graph execution in a tf.Session
23 with tfe.Session() as sess:
24     sess.run(tf.global_variables_initializer())
25     sess.run(prediction_op)
```

# Wrap-Up

You can **compute on encrypted data**,  
without the ability to decrypt

Privacy-preserving ML mitigate **bottlenecks** and  
**enable access** to sensitive information

Secure computation **distributes trust and control**,  
and is complementary to e.g. differential privacy

Privacy-preserving ML is a multidisciplinary field  
benefitting from **adaptations** on both sides

Focus on **usability** and **integration**

# Thank you!