

Cryptography for Privacy-Preserving Machine Learning

Morten Dahl

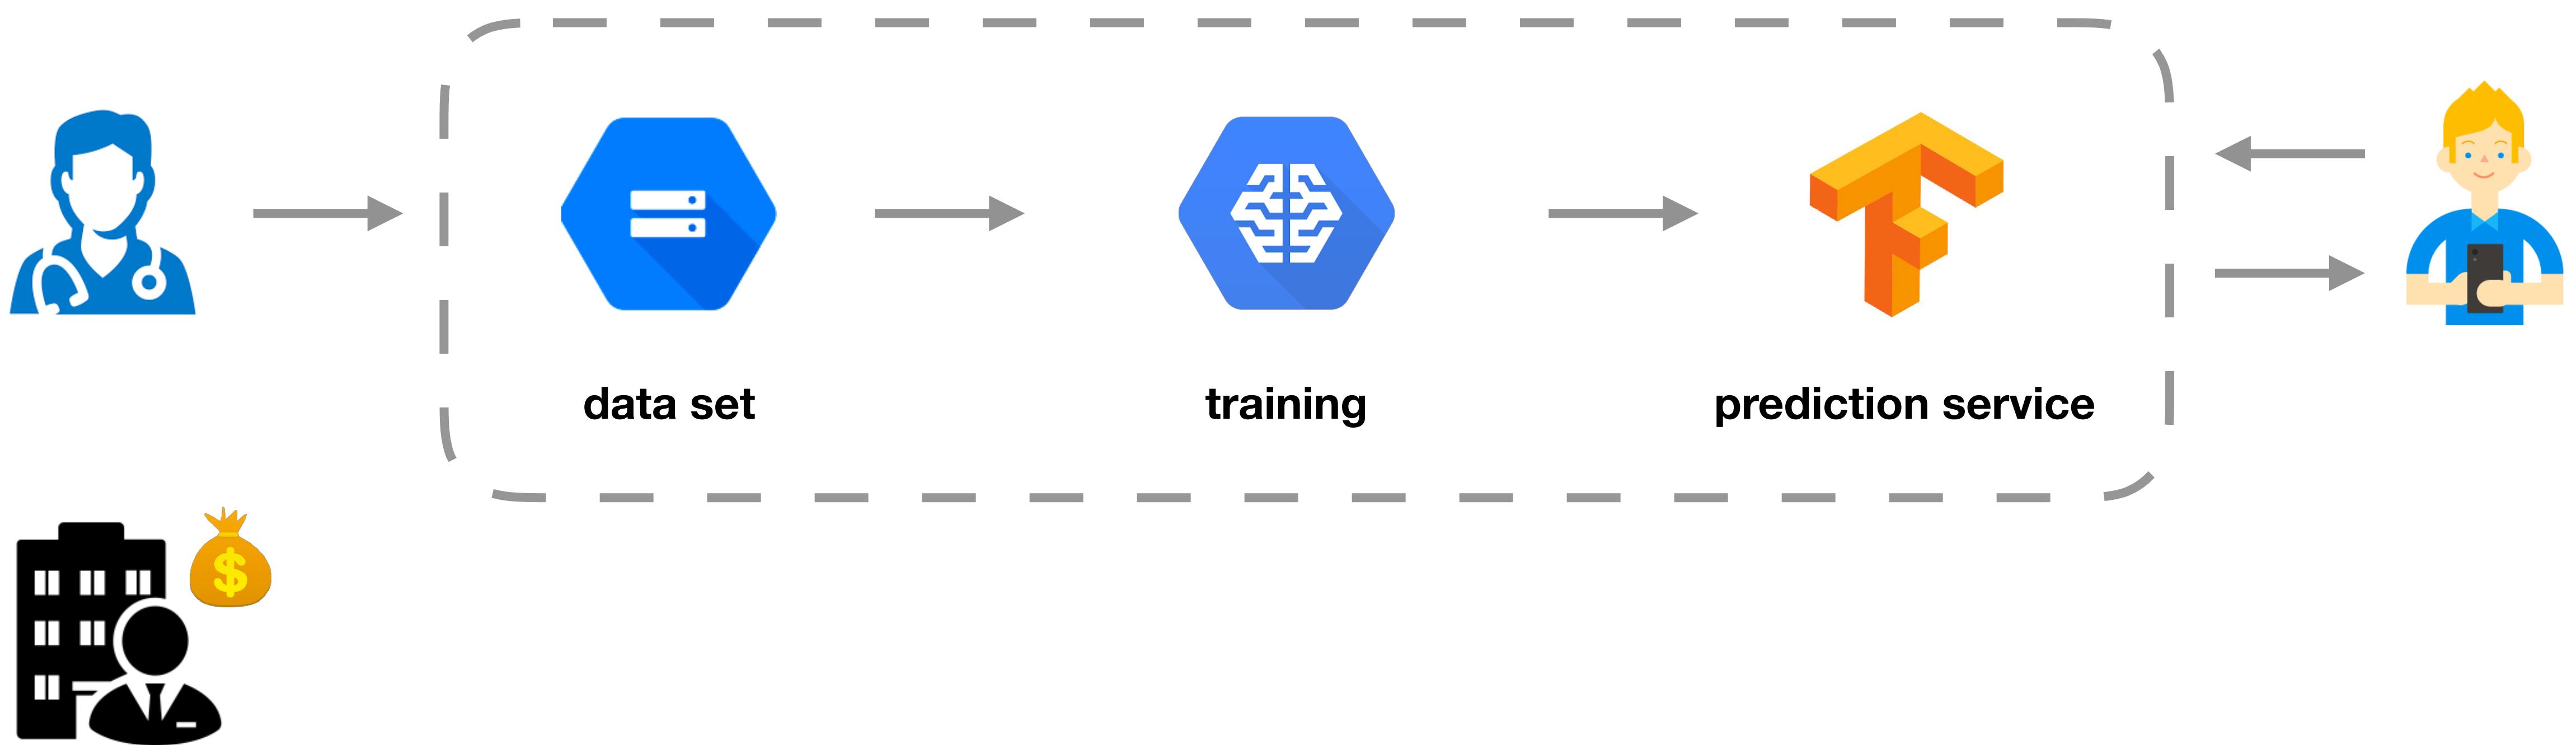
The tf-encrypted Project and Dropout Labs

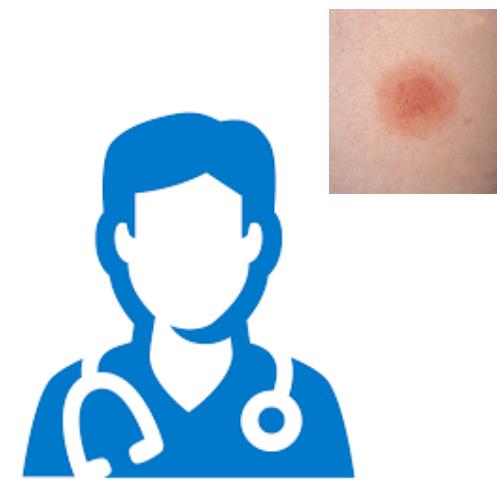
Hacking Deep Learning, Bar-Ilan University, January 2019

Why?

Machine Learning Process

IMAGENET



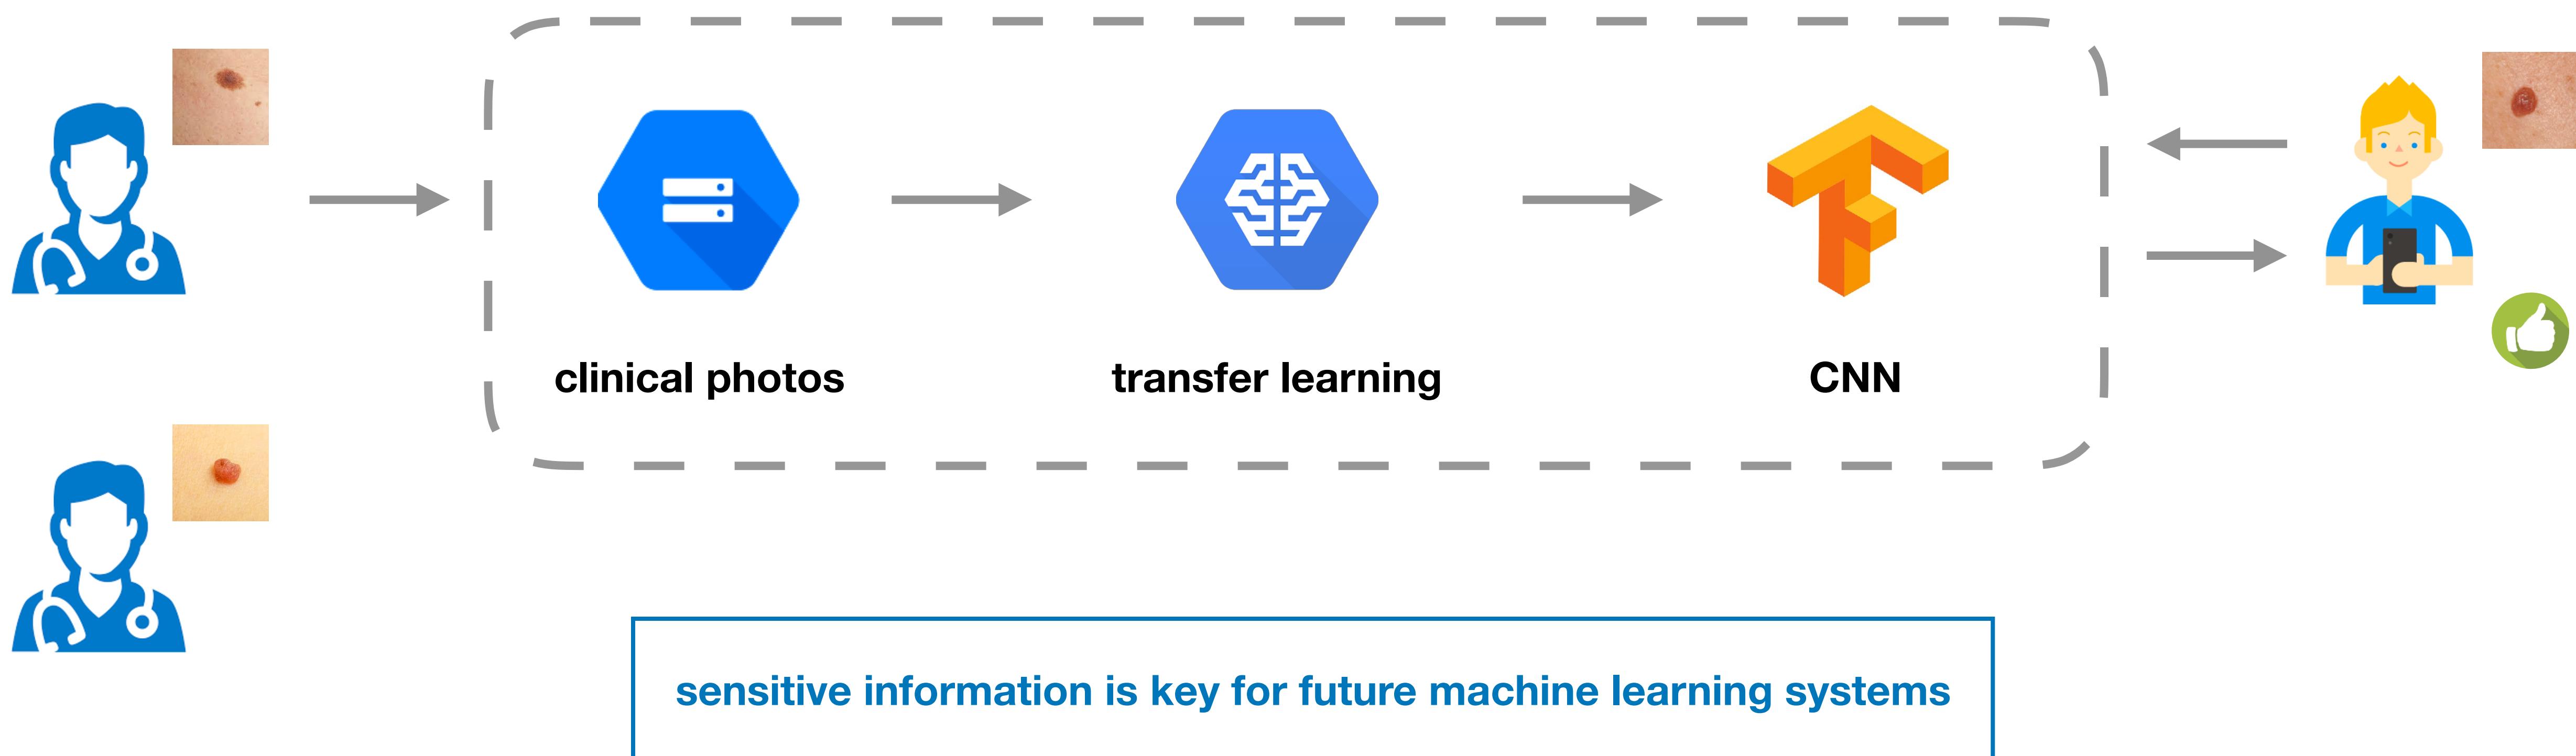


Skin Cancer Image Classification

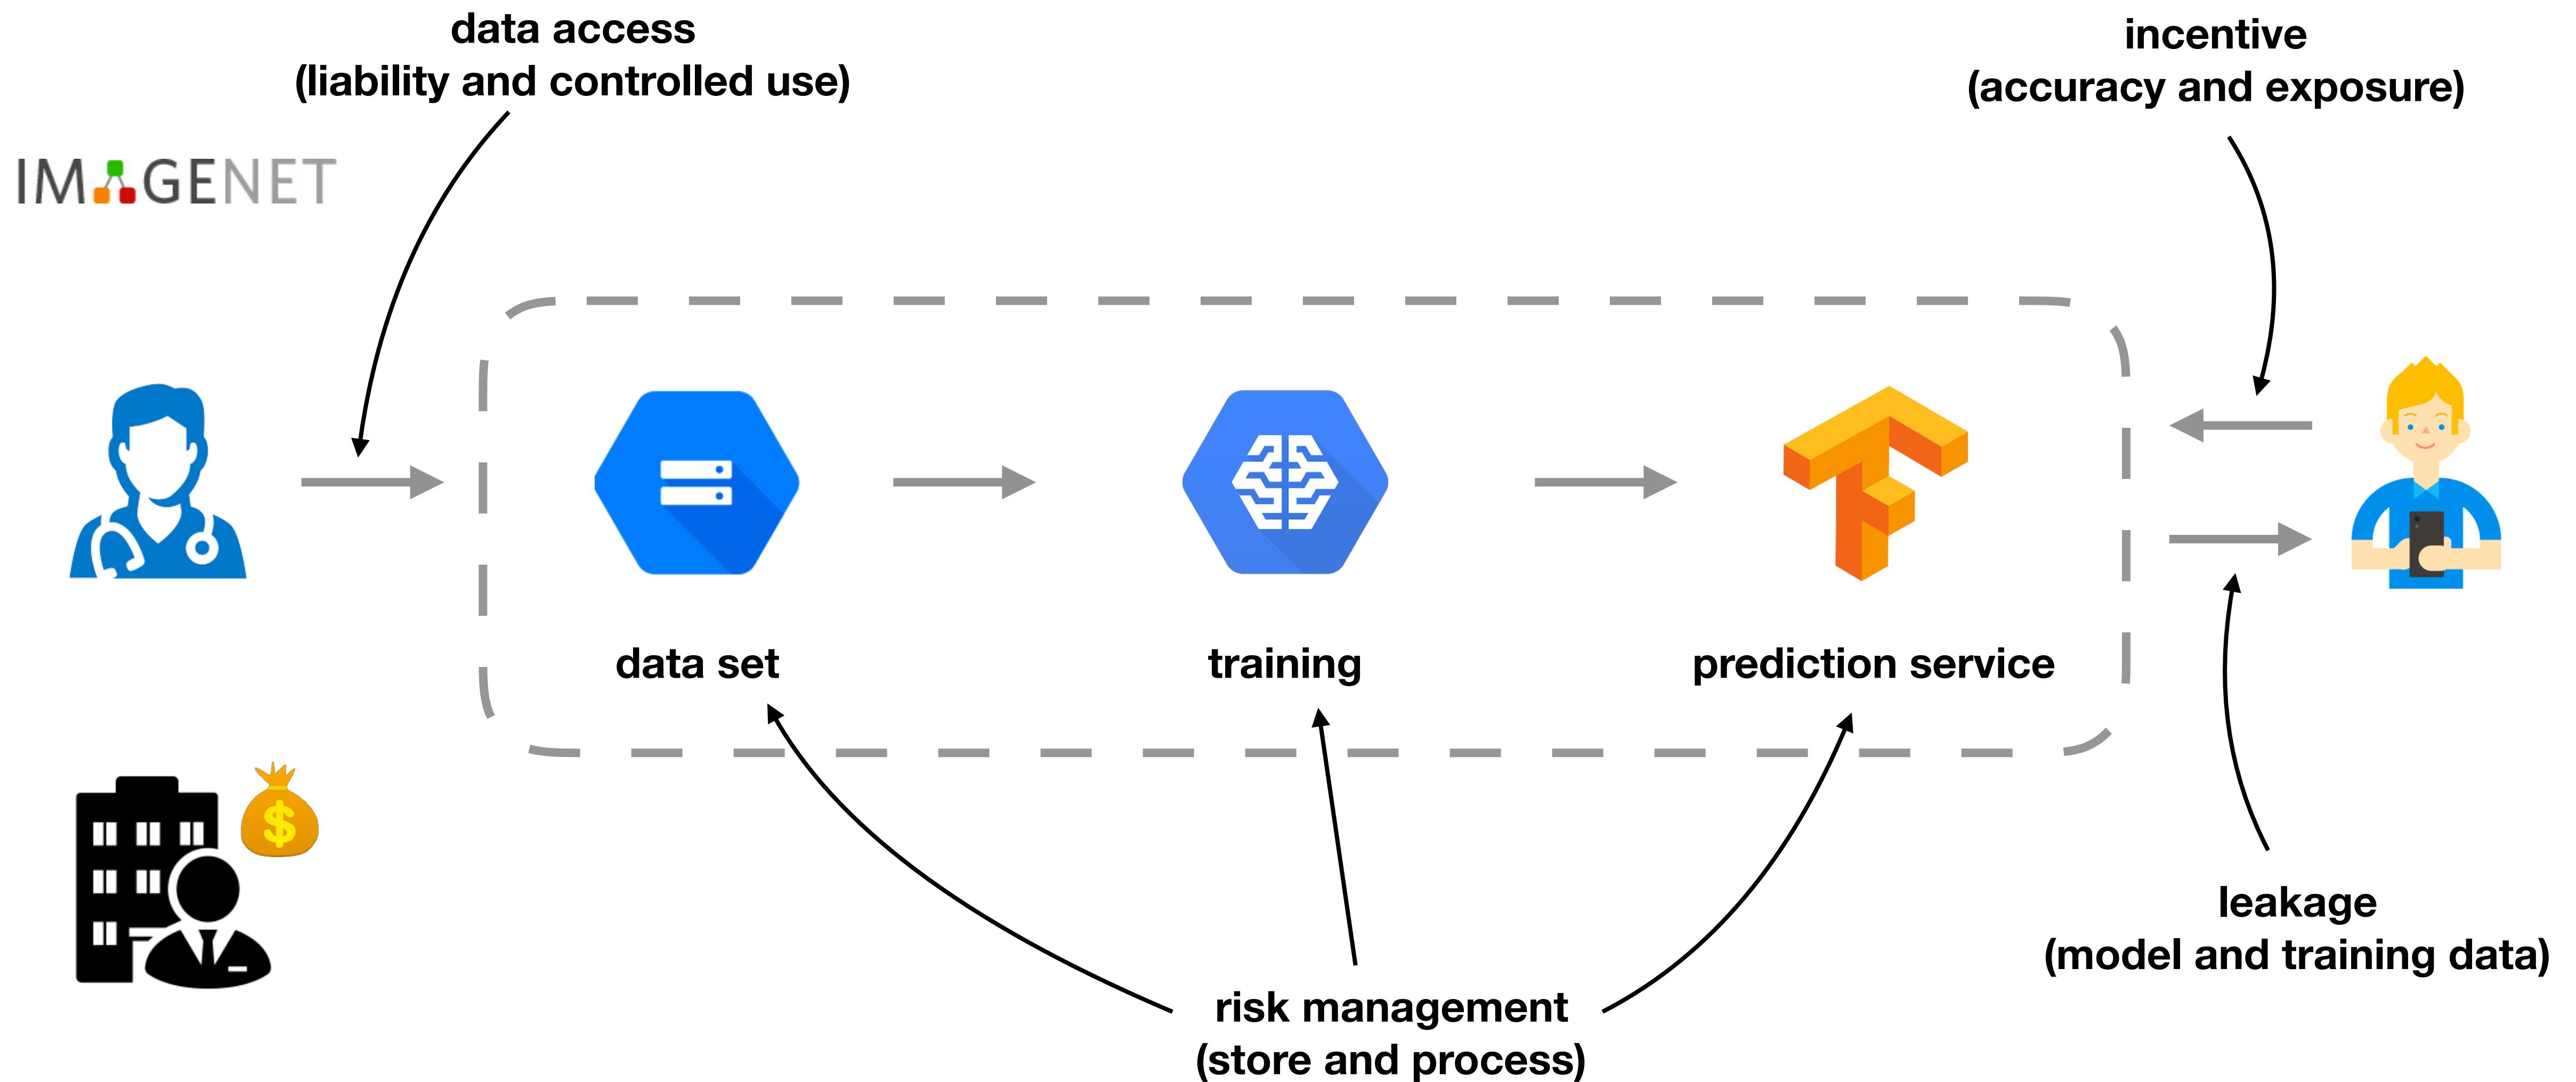
Brett Kuprel

12:30-12:40pm

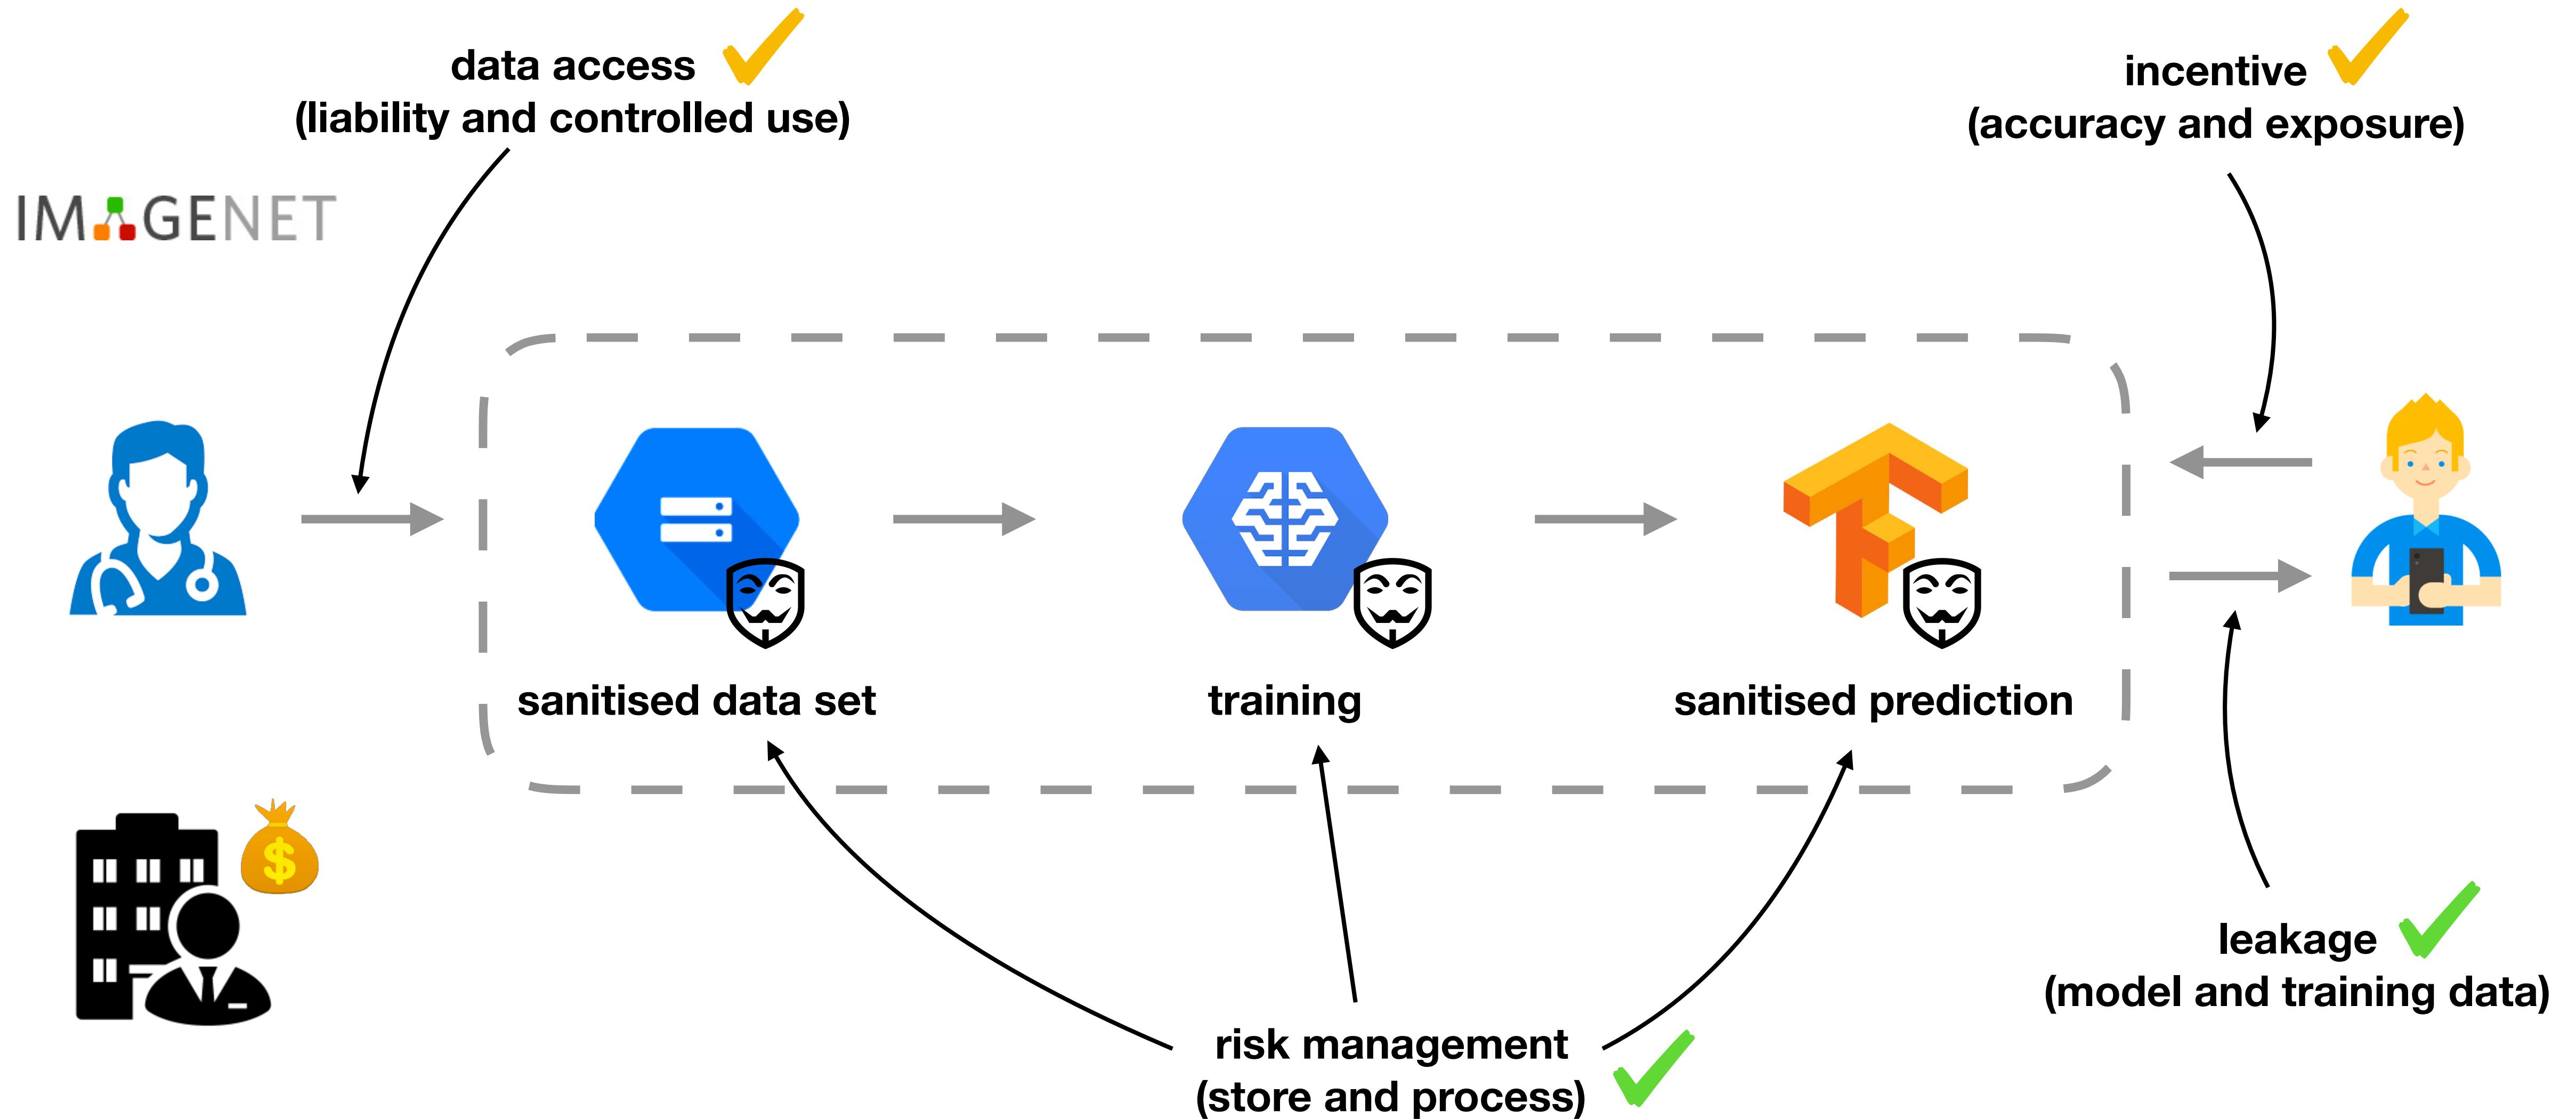
Join Brett Kuprel, and see how TensorFlow was used by the artificial intelligence lab and medical school of Stanford to classify skin cancer images. He'll describe the project steps: from acquiring a dataset, training a deep network, and evaluating of the results. To wrap up, Brett will give his take on the future of skin cancer image classification.



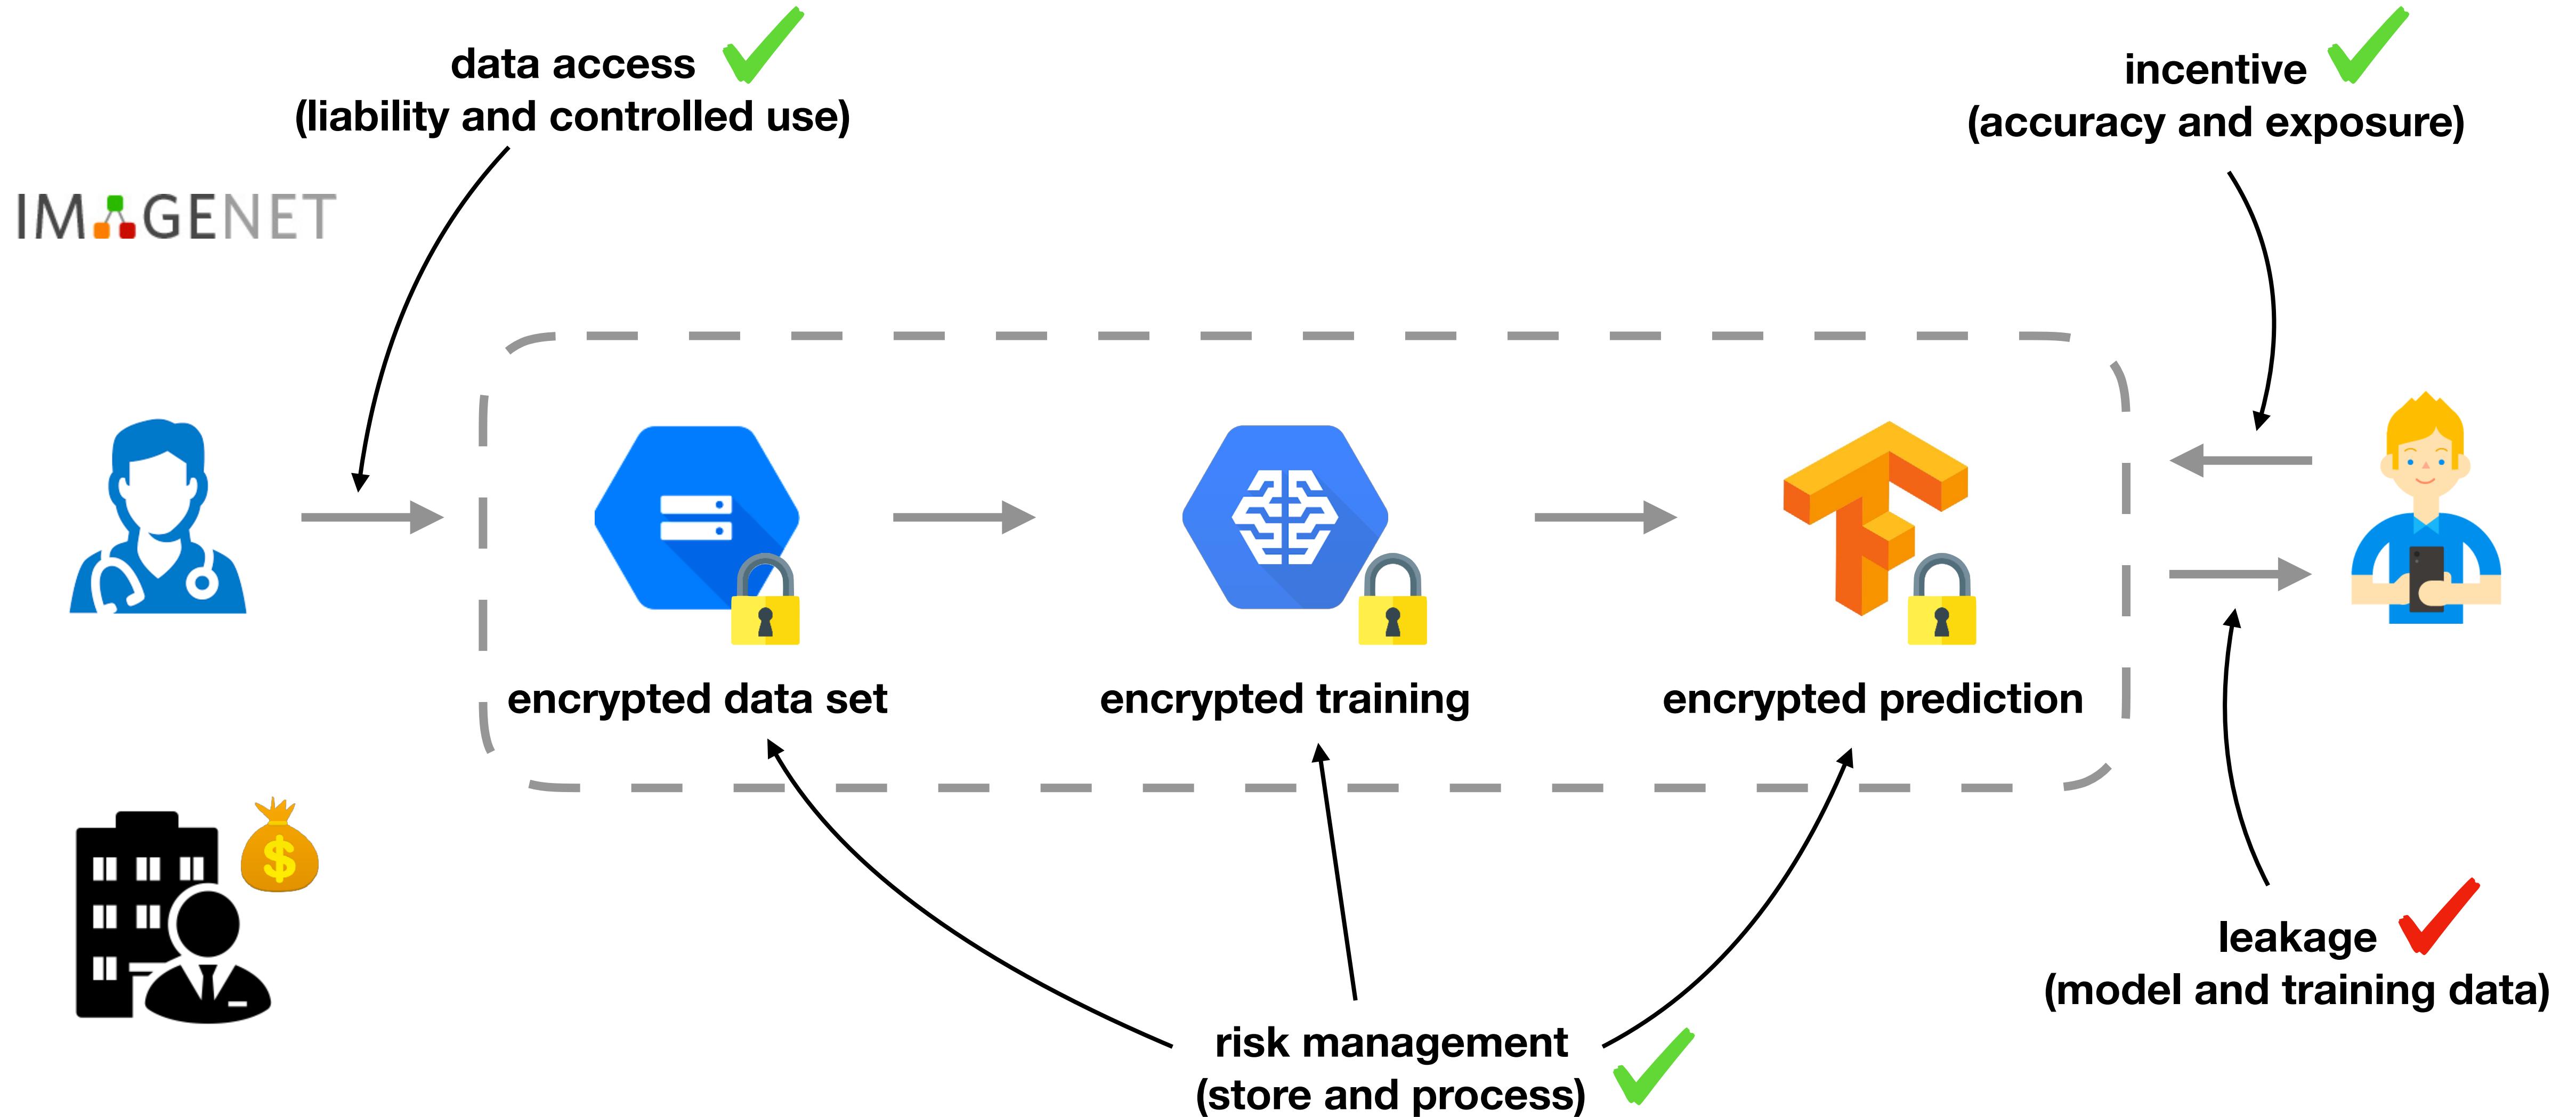
Potential Bottlenecks



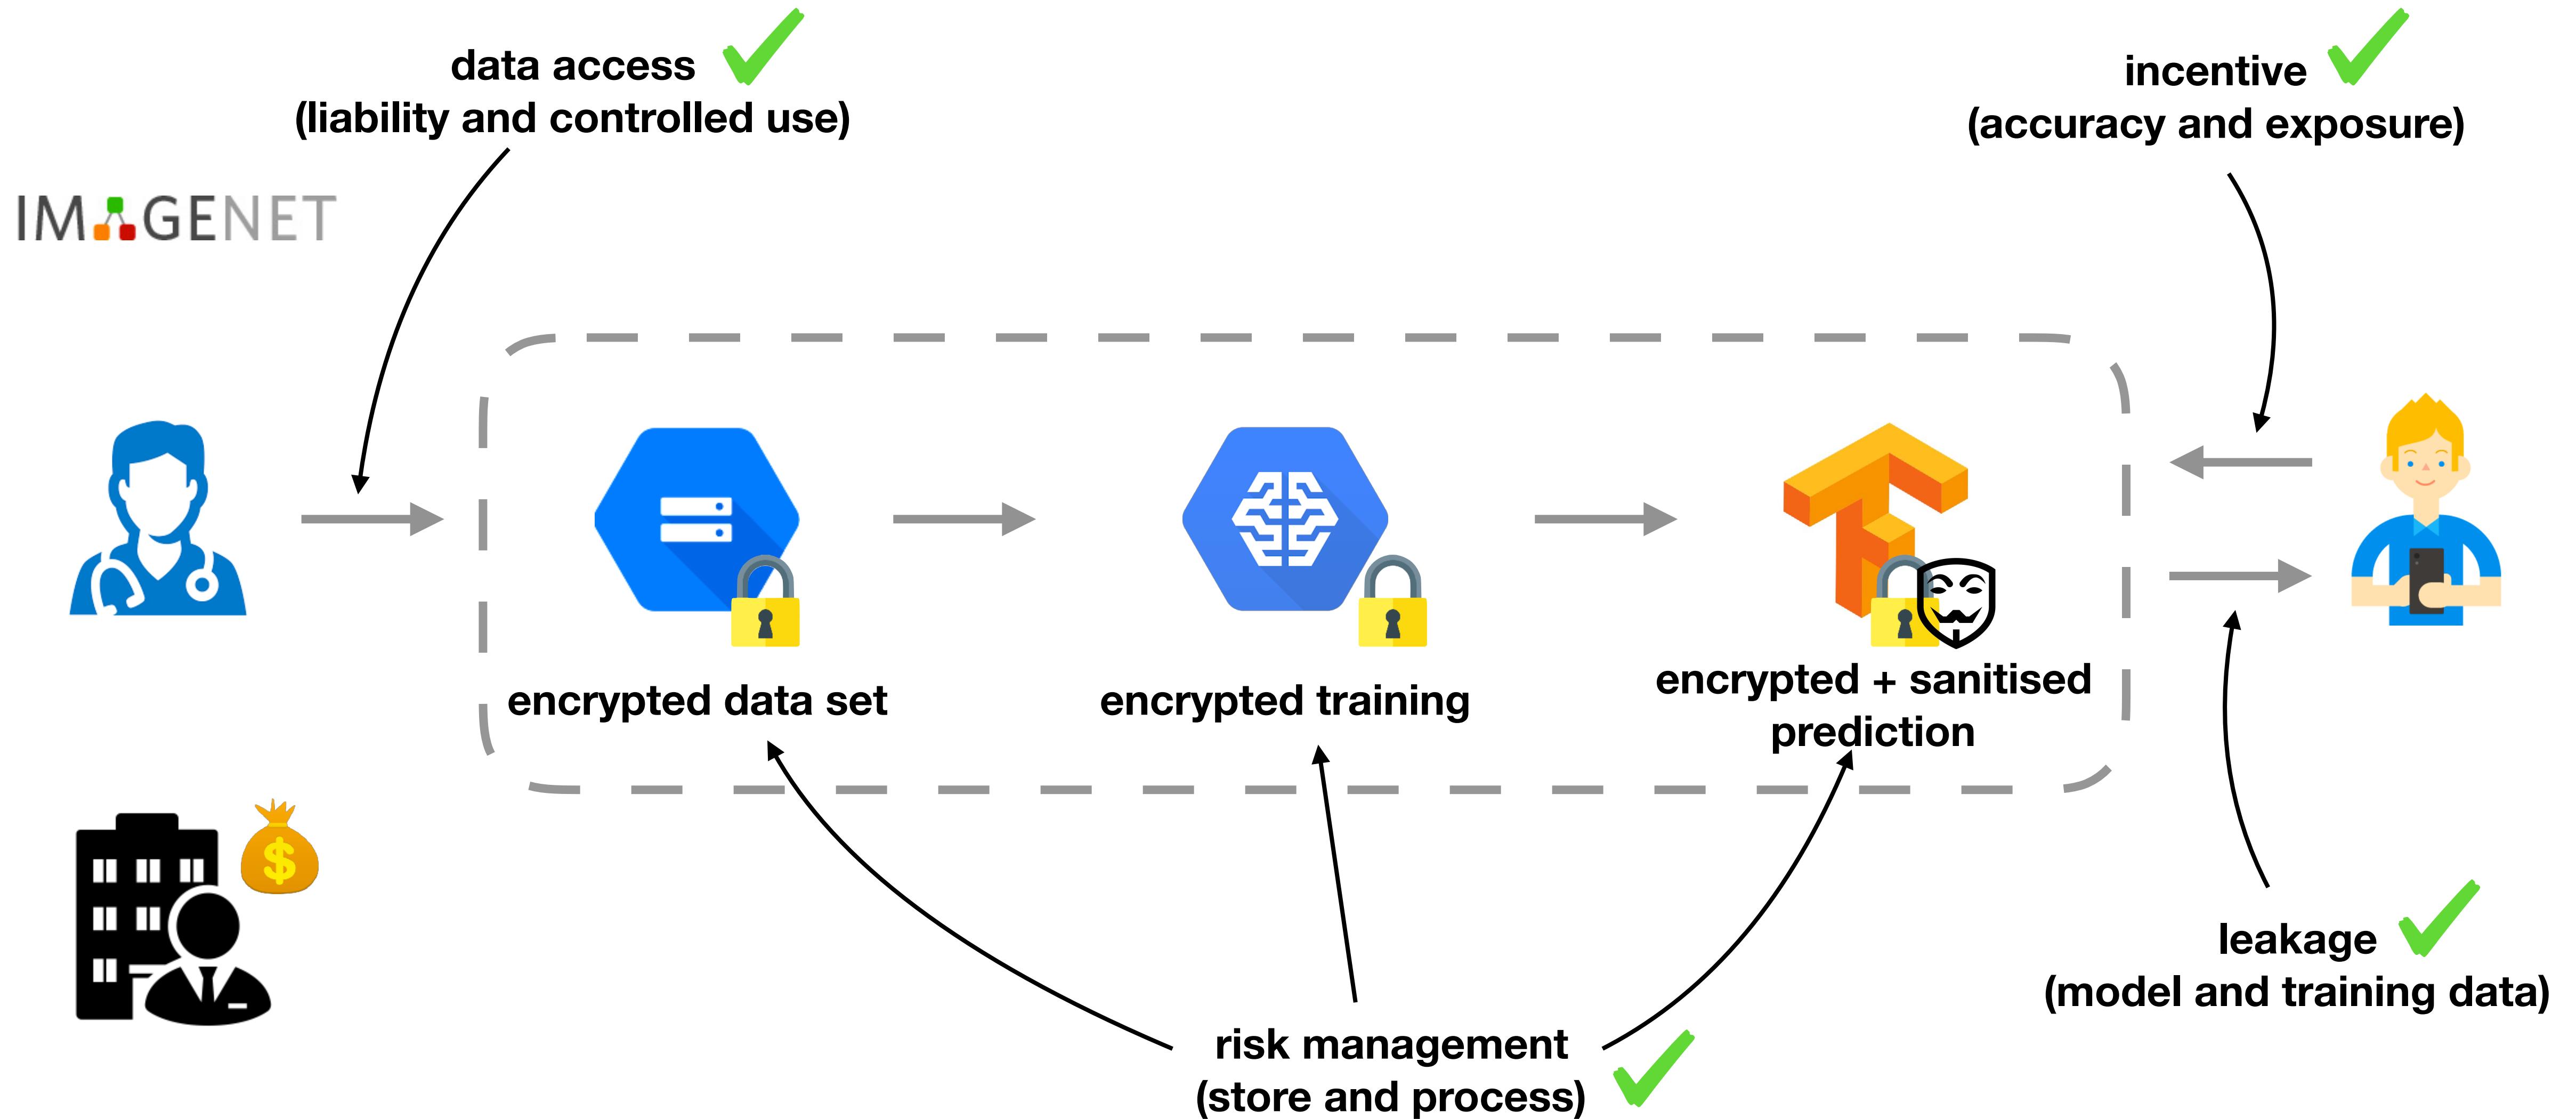
Differential Privacy



Secure Computation



Hybrid



Secure Computation

fixed-point encoding

$$\lfloor \pi * 10^4 \rfloor = 31415$$

Homomorphic Encryption

Computation

heavy

Communication

light

Data types

integers + bits

Platform

software

Secret Sharing

light

heavy

integers + bits

software

Garbled Circuits

medium

medium

bits

software

Secure Enclaves

light

light

any

hardware

CryptoNets'16



SecureML'17



DeepSecure'17



Gazelle'18



ABY3'18



SecureNN'18



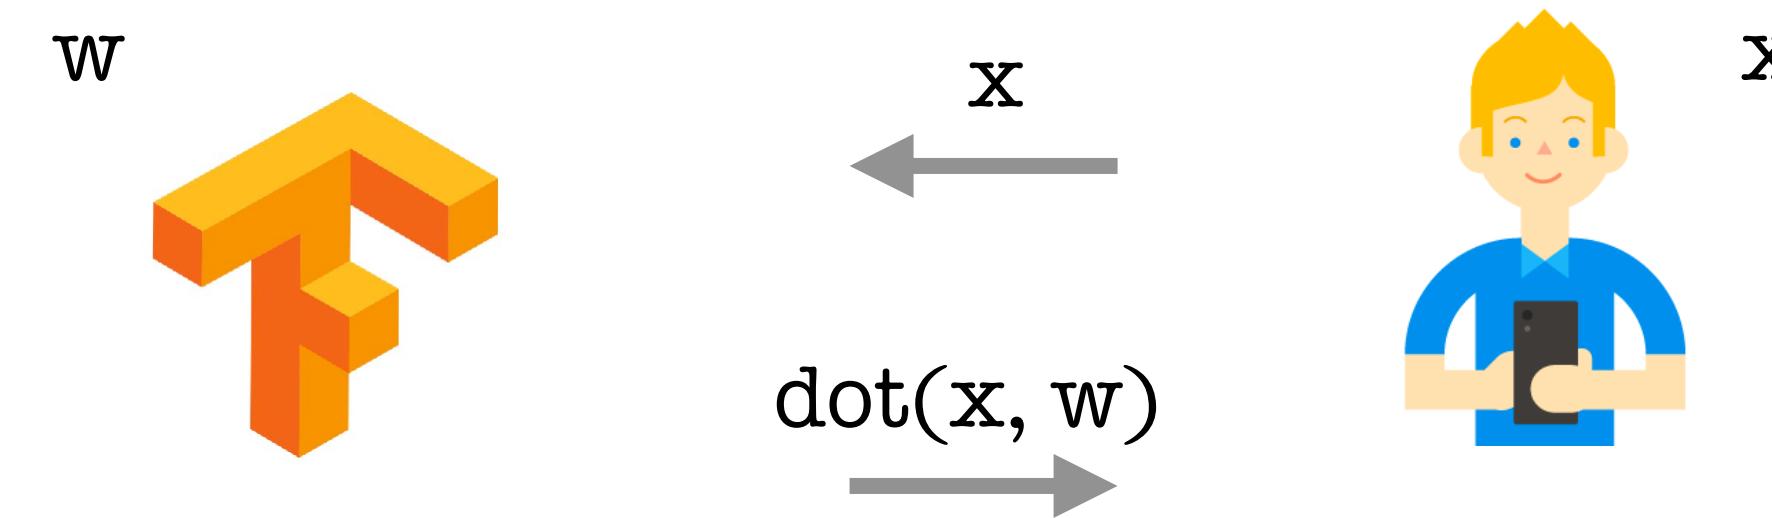
Slalom'18



tendency to mix techniques

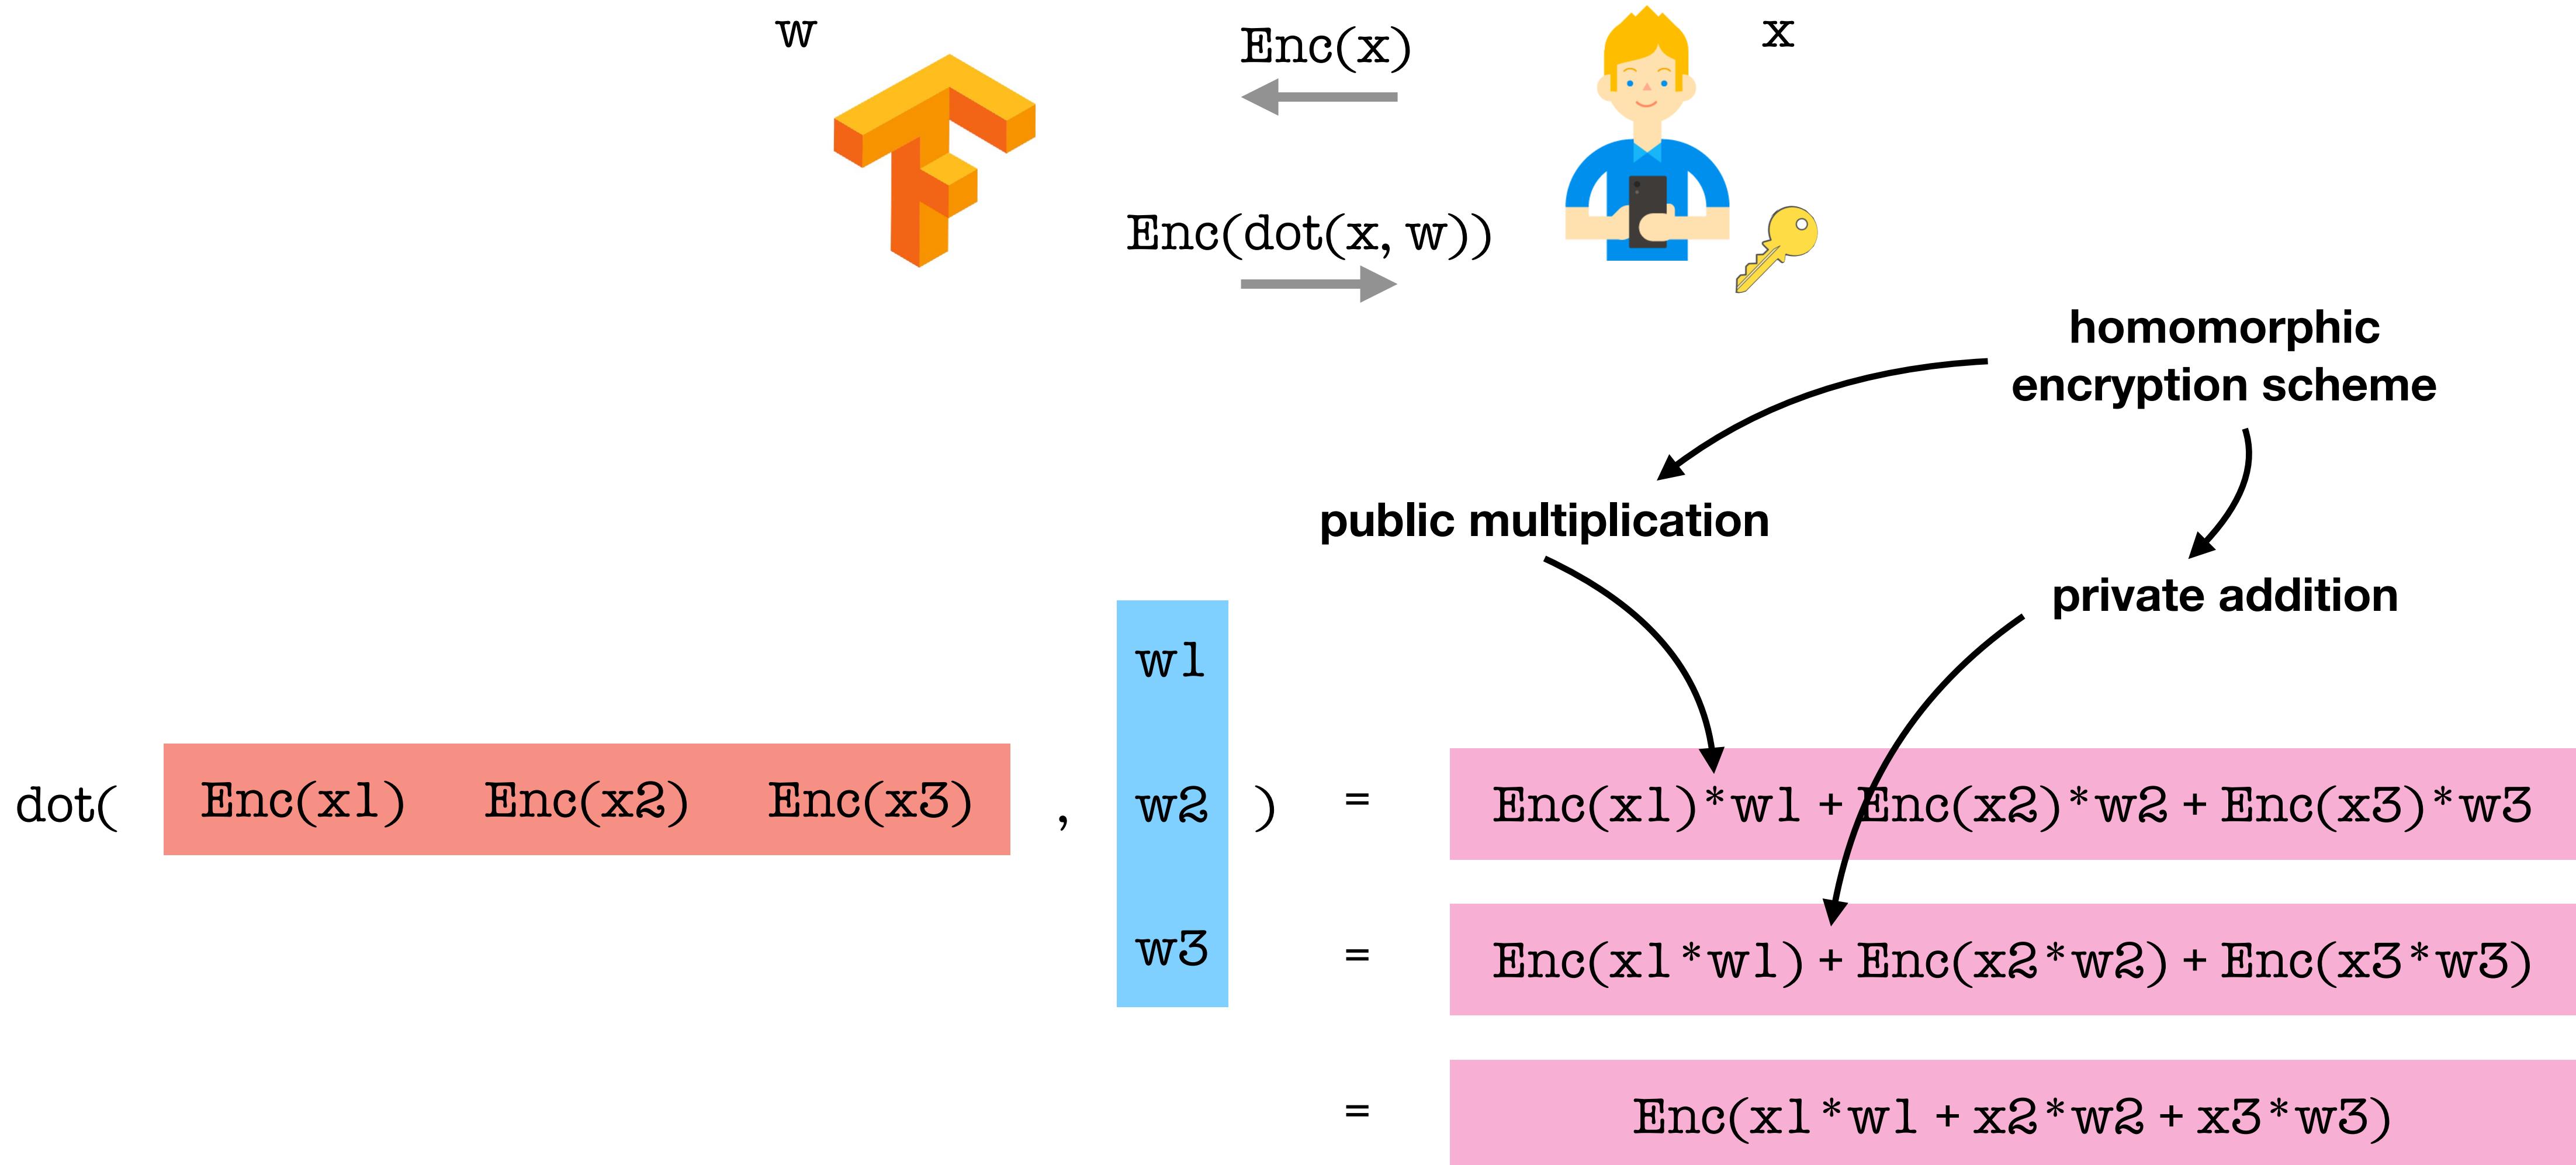
Linear Regression

Prediction with Linear Regression Model



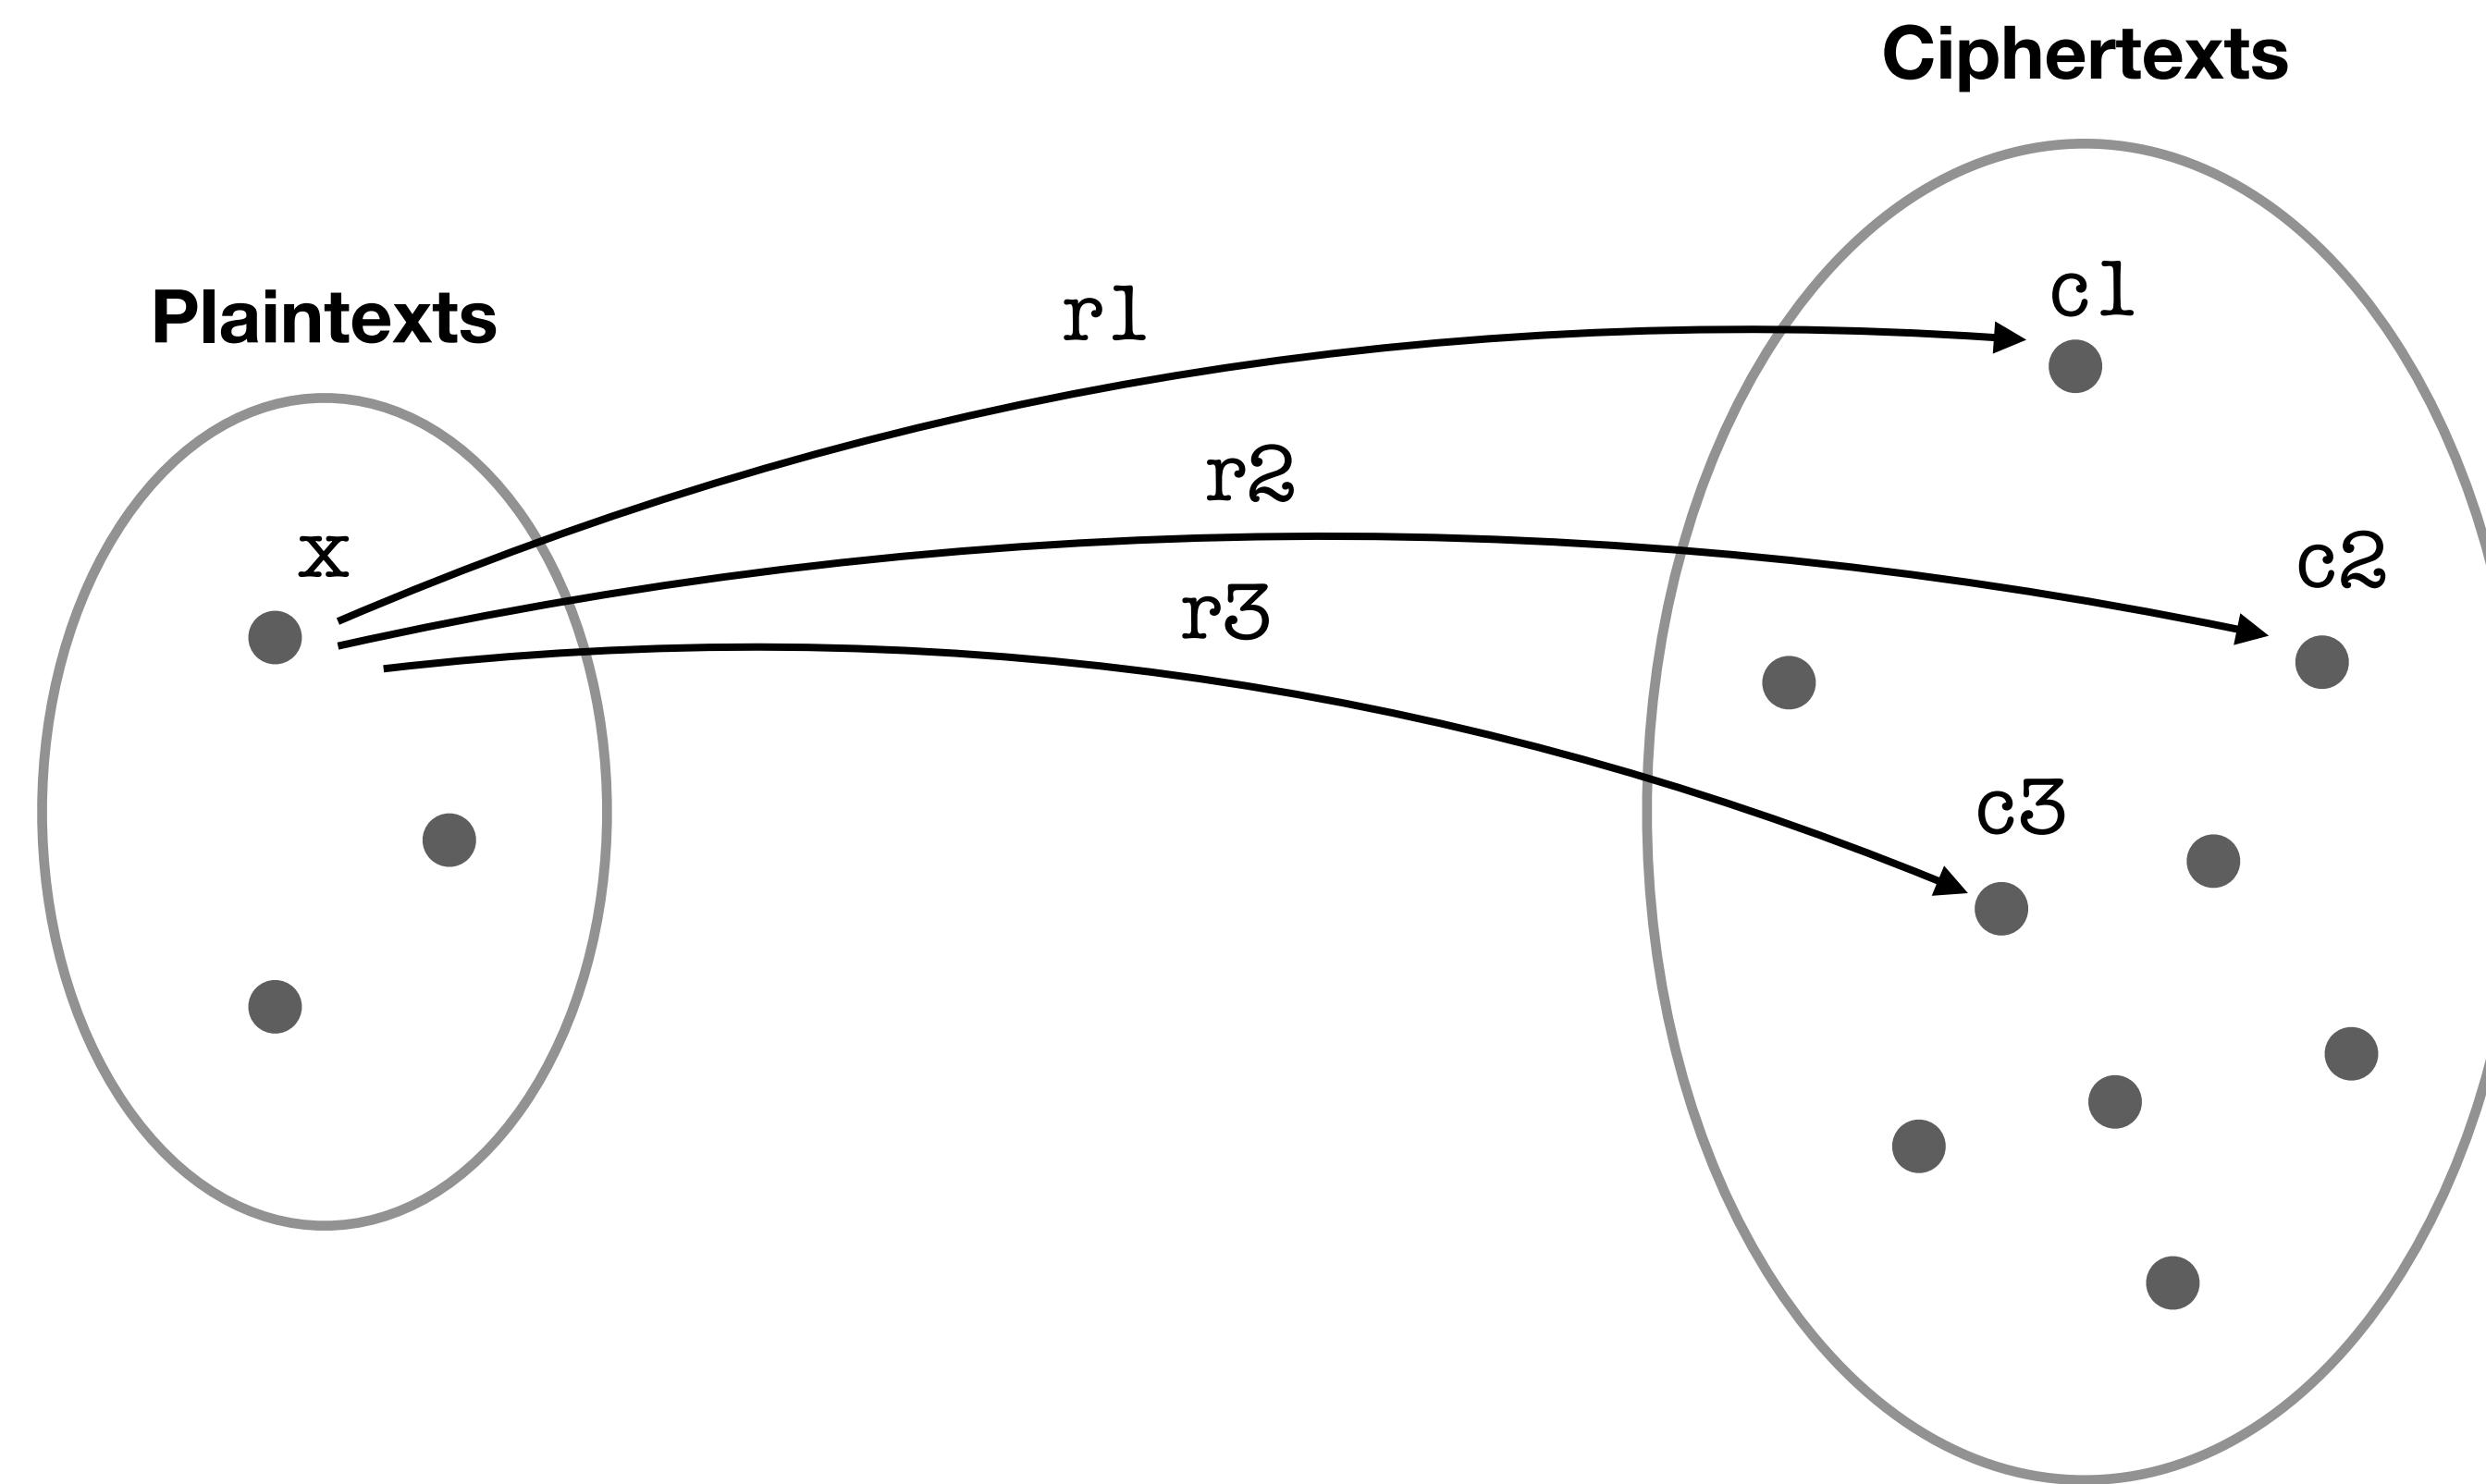
$$\text{dot}(\begin{matrix} x_1 & x_2 & x_3 \end{matrix}, \begin{matrix} w_1 \\ w_2 \\ w_3 \end{matrix}) = x_1 * w_1 + x_2 * w_2 + x_3 * w_3$$

Using Homomorphic Encryption



Probabilistic Encryption

$$c = \text{Enc}(x, r)$$



Paillier Homomorphic Encryption

typically ~4000 bits:
computation is significantly
more expensive

$$c = \text{Enc}(x, r) = g^x * r^n \bmod n^2$$

public encryption key

public encryption key

$$\begin{aligned} g &= 36 \\ n &= 35 \\ n^2 &= 1225 \end{aligned}$$

$$\begin{aligned} \text{Enc}(5, 2) &= 36^5 * 2^{35} \bmod 1225 = 718 \\ \text{Enc}(5, 4) &= 36^5 * 4^{35} \bmod 1225 = 674 \end{aligned}$$

Private Addition in Paillier

$$\begin{aligned} & \text{Enc}(\mathbf{x}, \mathbf{r}) * \text{Enc}(\mathbf{y}, \mathbf{s}) \\ &= (\mathbf{g}^{\mathbf{x}} * \mathbf{r}^n \bmod n^2) * (\mathbf{g}^{\mathbf{y}} * \mathbf{s}^n \bmod n^2) \\ &= \mathbf{g}^{\mathbf{x+y}} * (\mathbf{r} * \mathbf{s})^n \bmod n^2 \\ &= \text{Enc}(\mathbf{x+y}, \mathbf{r*s}) \end{aligned}$$

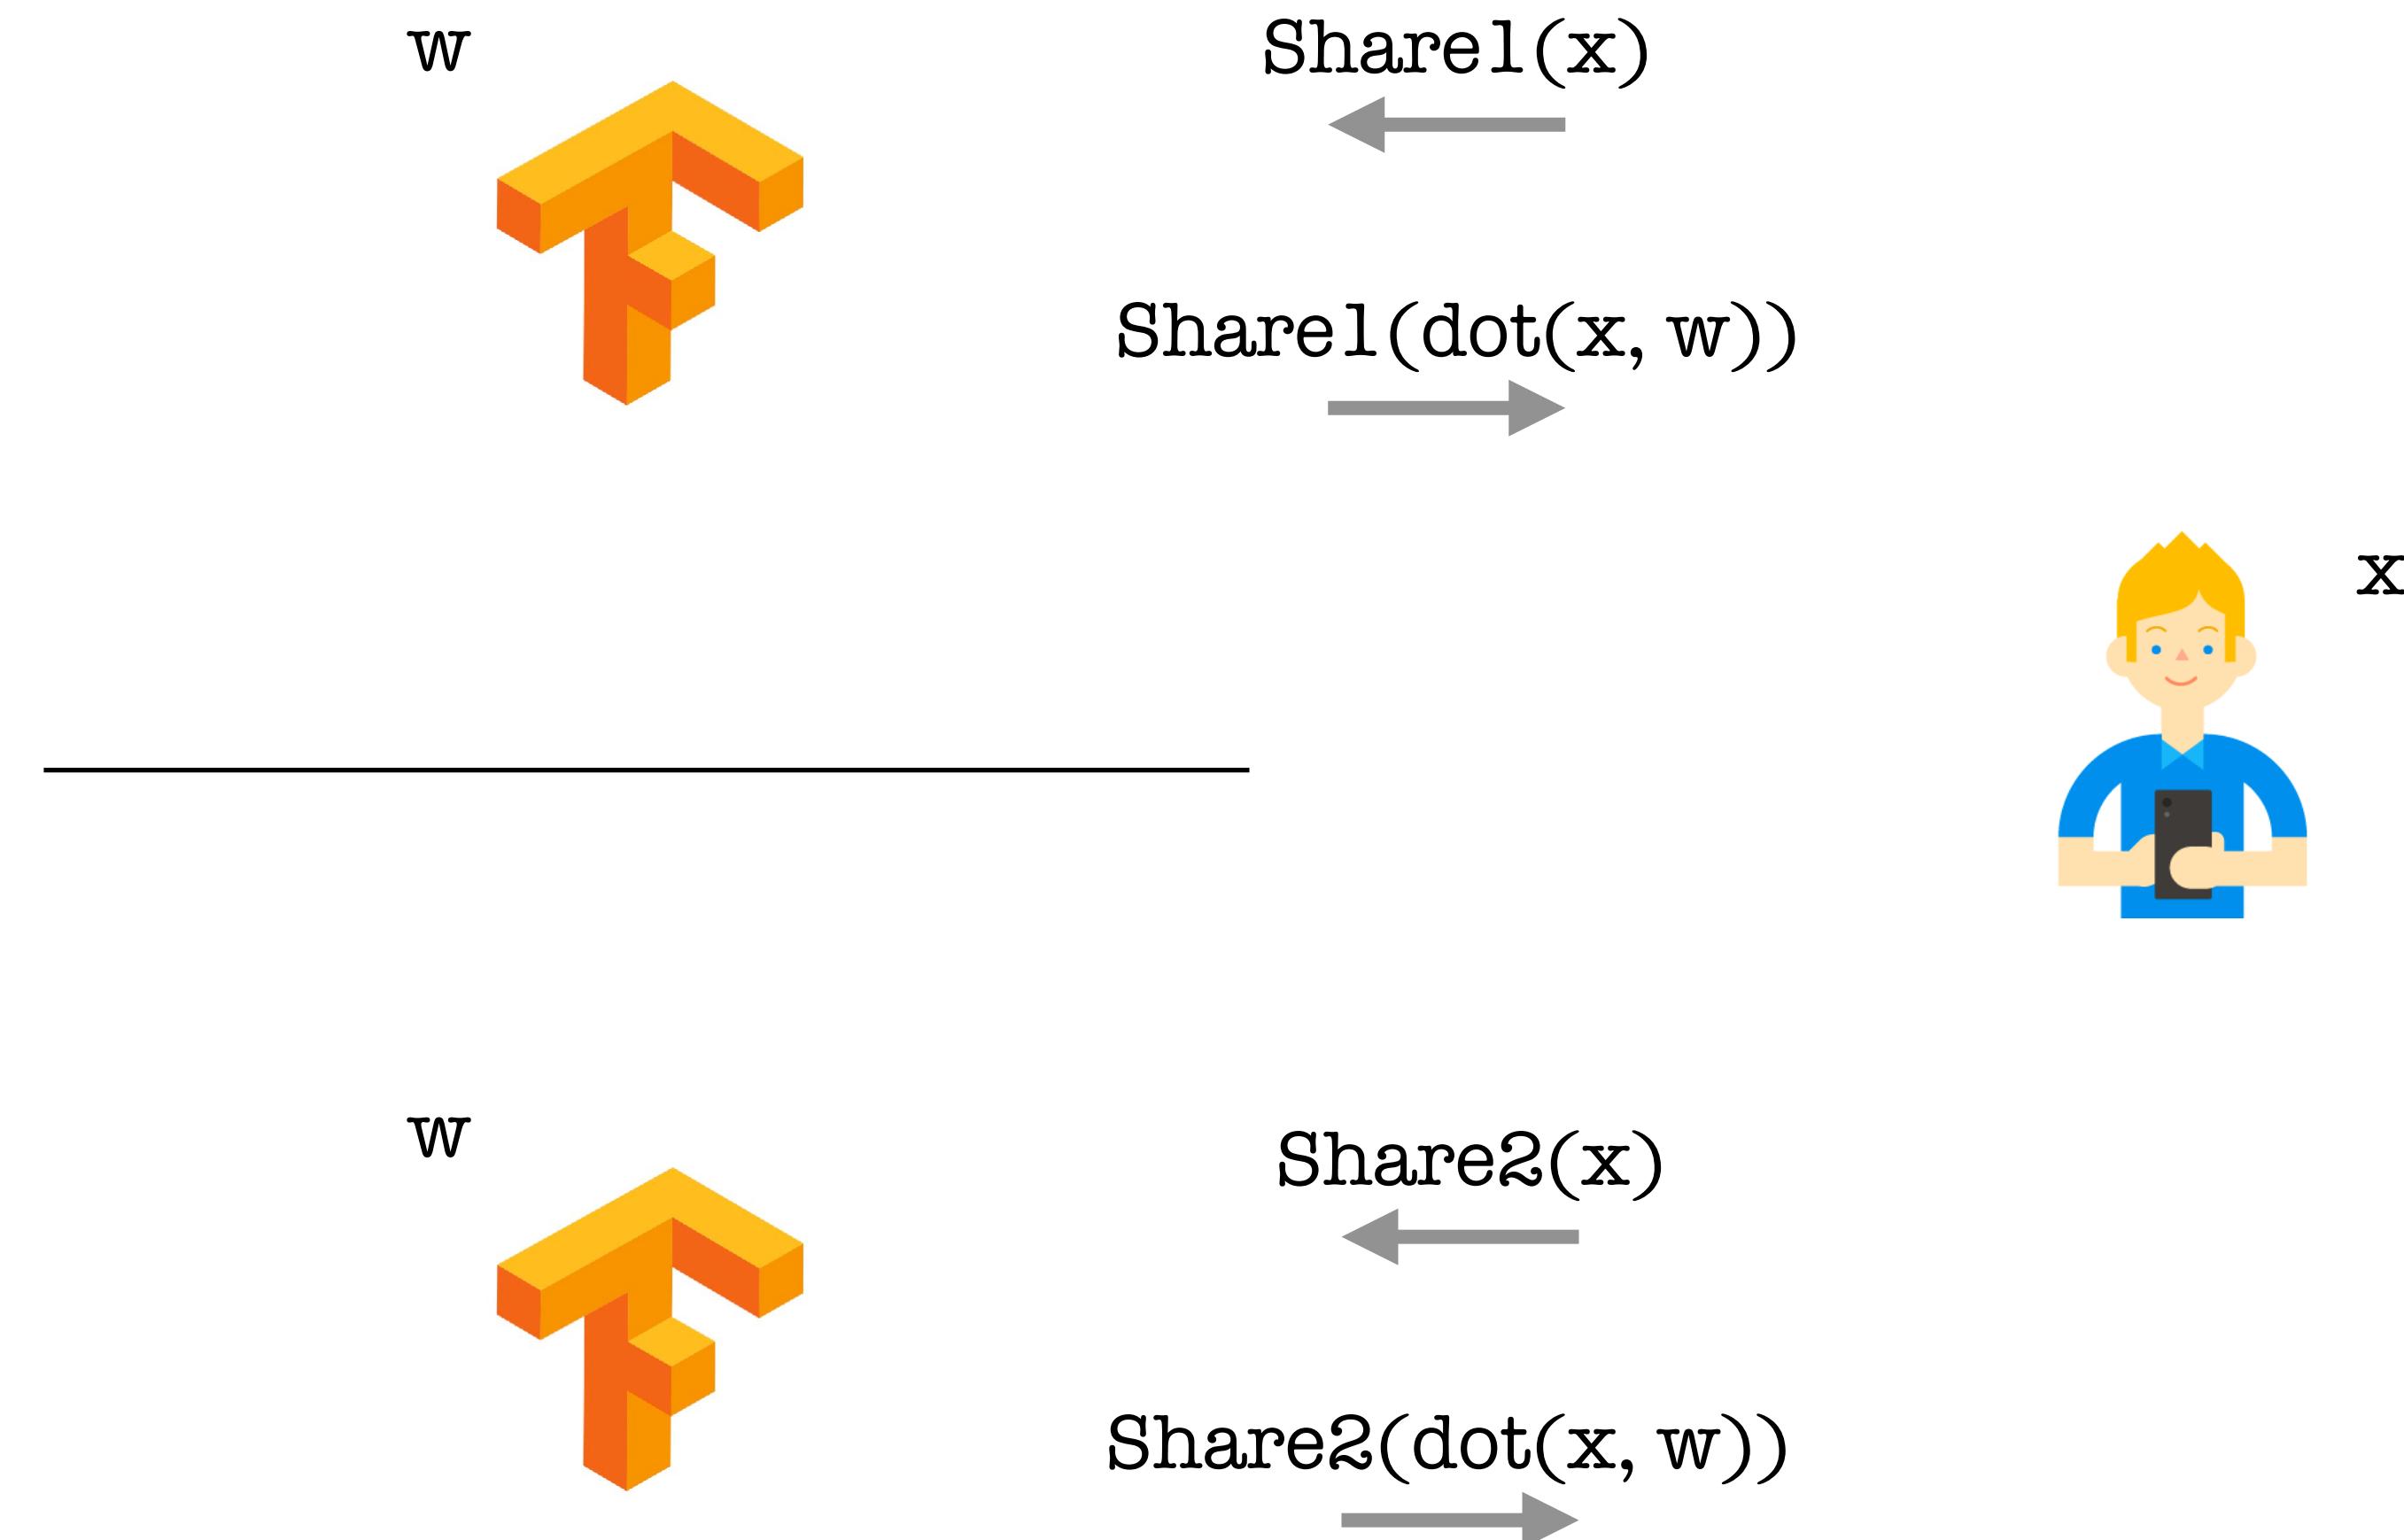
$$\begin{aligned} & \text{Enc}(5, 2) * \text{Enc}(5, 4) \\ &= 718 * 674 \\ &= 57 \\ &= 36^{10} * 8^{35} \\ &= \text{Enc}(10, 8) \end{aligned}$$

Public Multiplication in Paillier

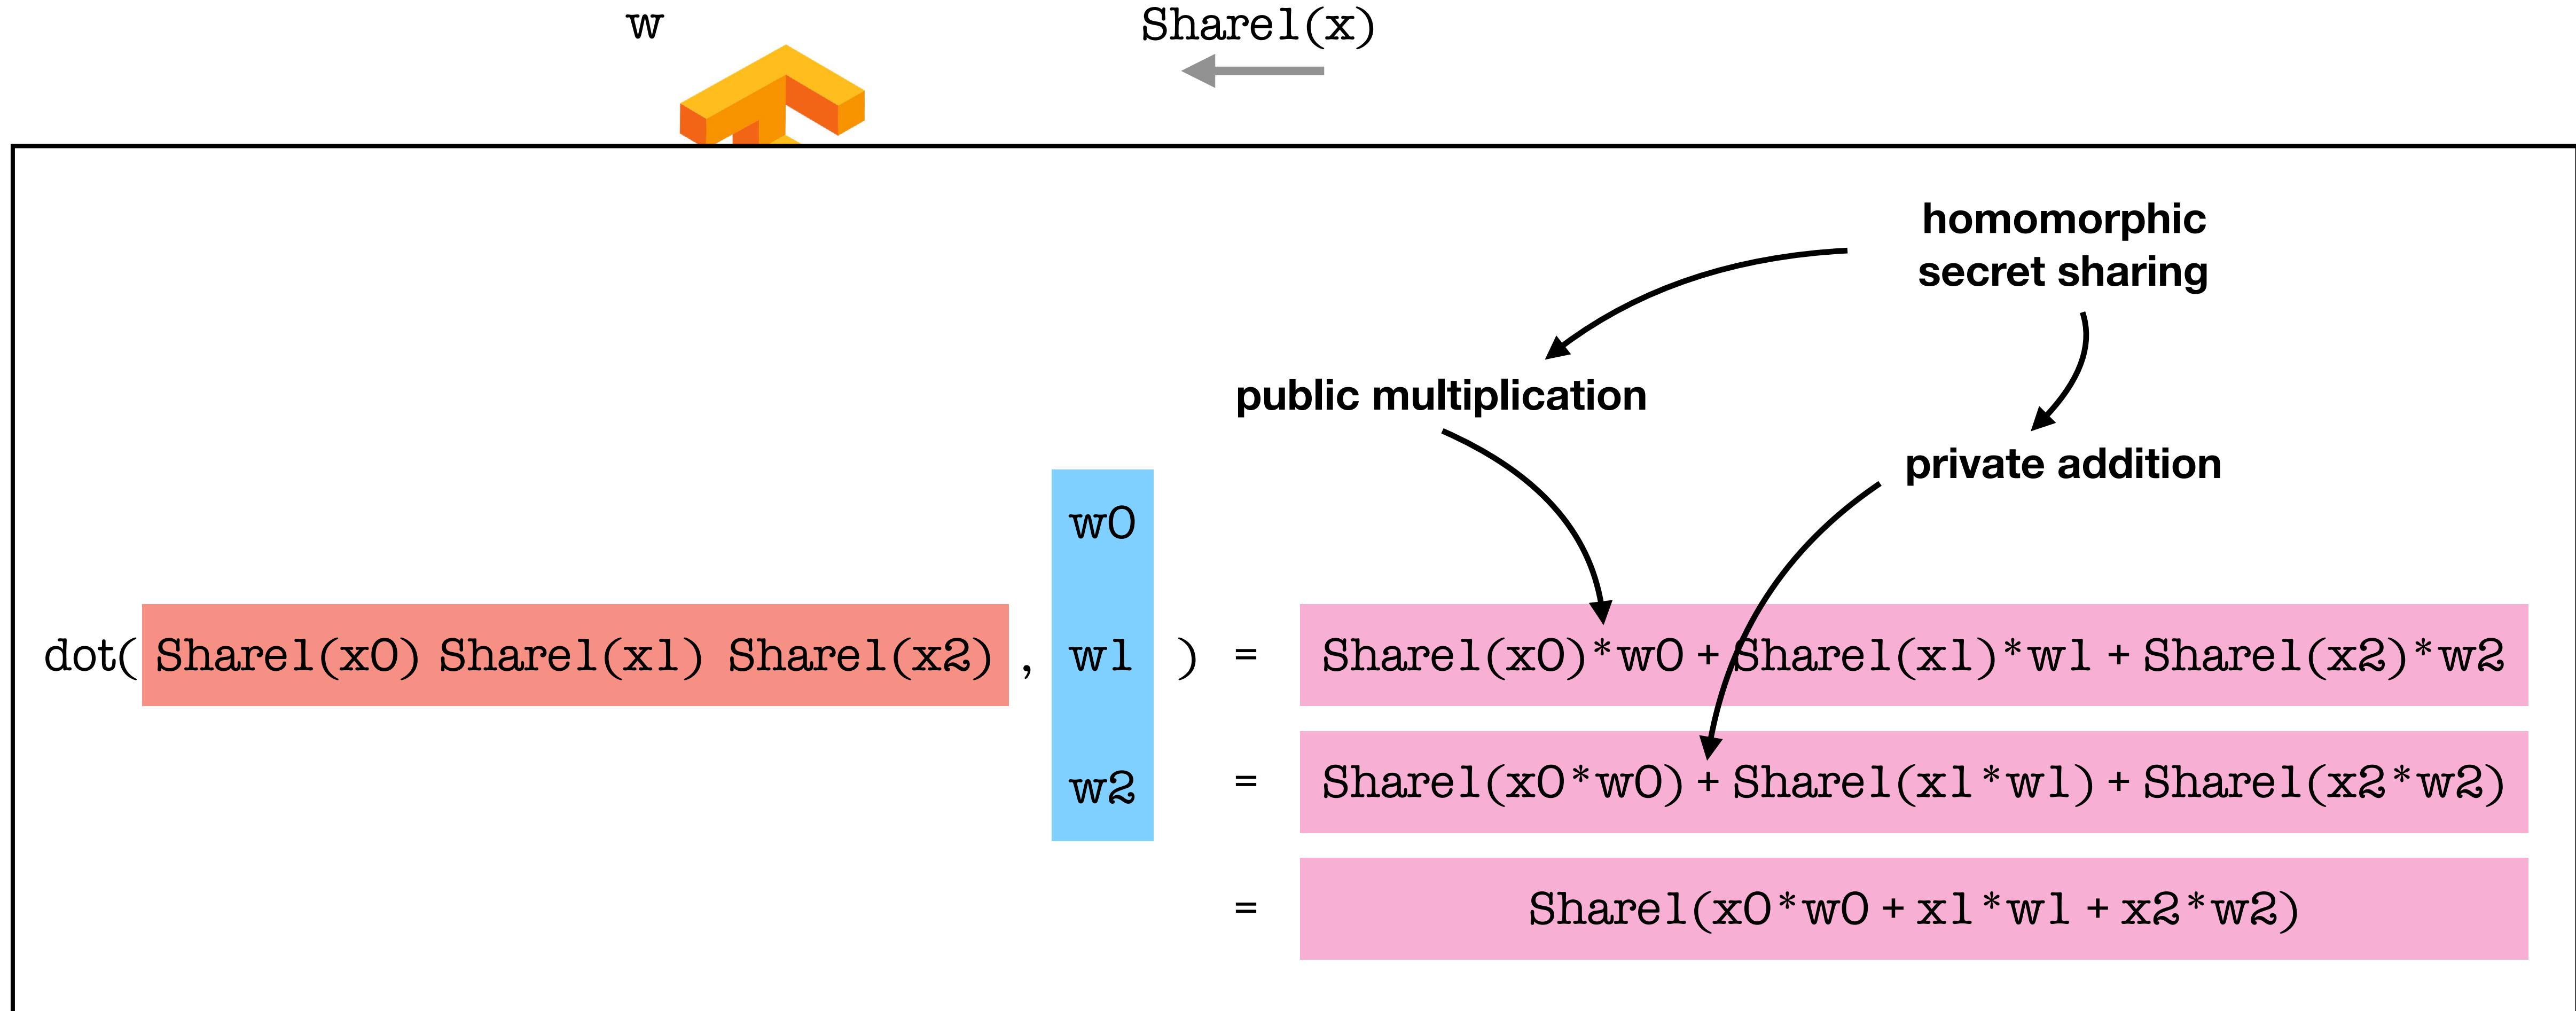
$$\begin{aligned} & \text{Enc}(\text{x}, \text{r})^{\wedge w} \\ &= (\text{g}^{\wedge x} * \text{r}^{\wedge n} \bmod \text{n}^2)^{\wedge w} \\ &= \text{g}^{\wedge(x * w)} * (\text{r}^{\wedge w})^{\wedge n} \bmod \text{n}^2 \\ &= \text{Enc}(\text{x} * w, \text{r}^{\wedge w}) \end{aligned}$$

$$\begin{aligned} & \text{Enc}(5, 2)^{\wedge 2} \\ &= 718 * 718 \\ &= 1024 \\ &= 36^{\wedge 10} * 4^{\wedge 35} \\ &= \text{Enc}(10, 4) \end{aligned}$$

Using Secret Sharing



Using Secret Sharing



SPDZ Secret Sharing

$$\text{Share1}(x, r) = r \bmod m$$

public parameter

$$\text{Share2}(x, r) = x - r \bmod m$$

$$x = \text{Share1}(x, r) + \text{Share2}(x, r) \bmod m$$

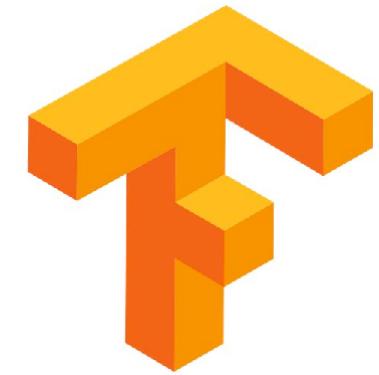
$$m = 10$$

$$\text{Share1}(5, 7) = 7 \bmod 10 = 7$$

$$\text{Share2}(5, 7) = 5 - 7 \bmod 10 = 8$$

$$7 + 8 = 15 = 5 \bmod 10$$

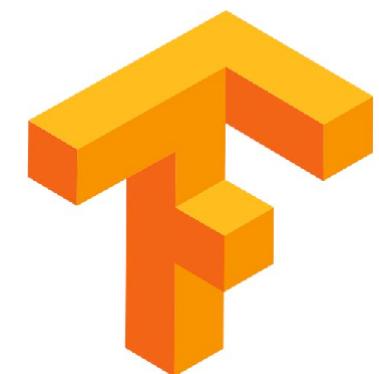
Private Addition in SPDZ



x₁

y₁

z₁ = x₁ + y₁



x₂

y₂

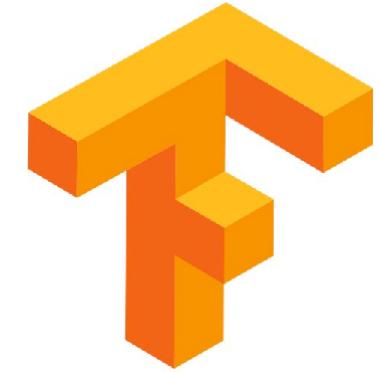
z₂ = x₂ + y₂

$$x = x_1 + x_2$$

$$y = y_1 + y_2$$

$$\begin{aligned}x + y &= (x_1 + x_2) + (y_1 + y_2) \\&= (x_1 + y_1) + (x_2 + y_2) \\&= z_1 + z_2\end{aligned}$$

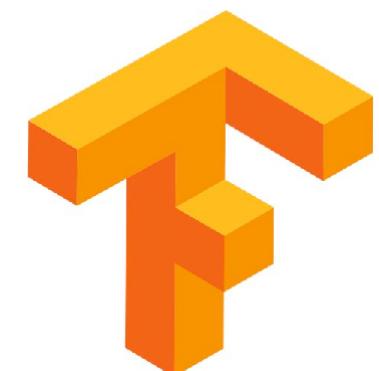
Public Multiplication in SPDZ



x1

w

z1 = x1 * w



x2

w

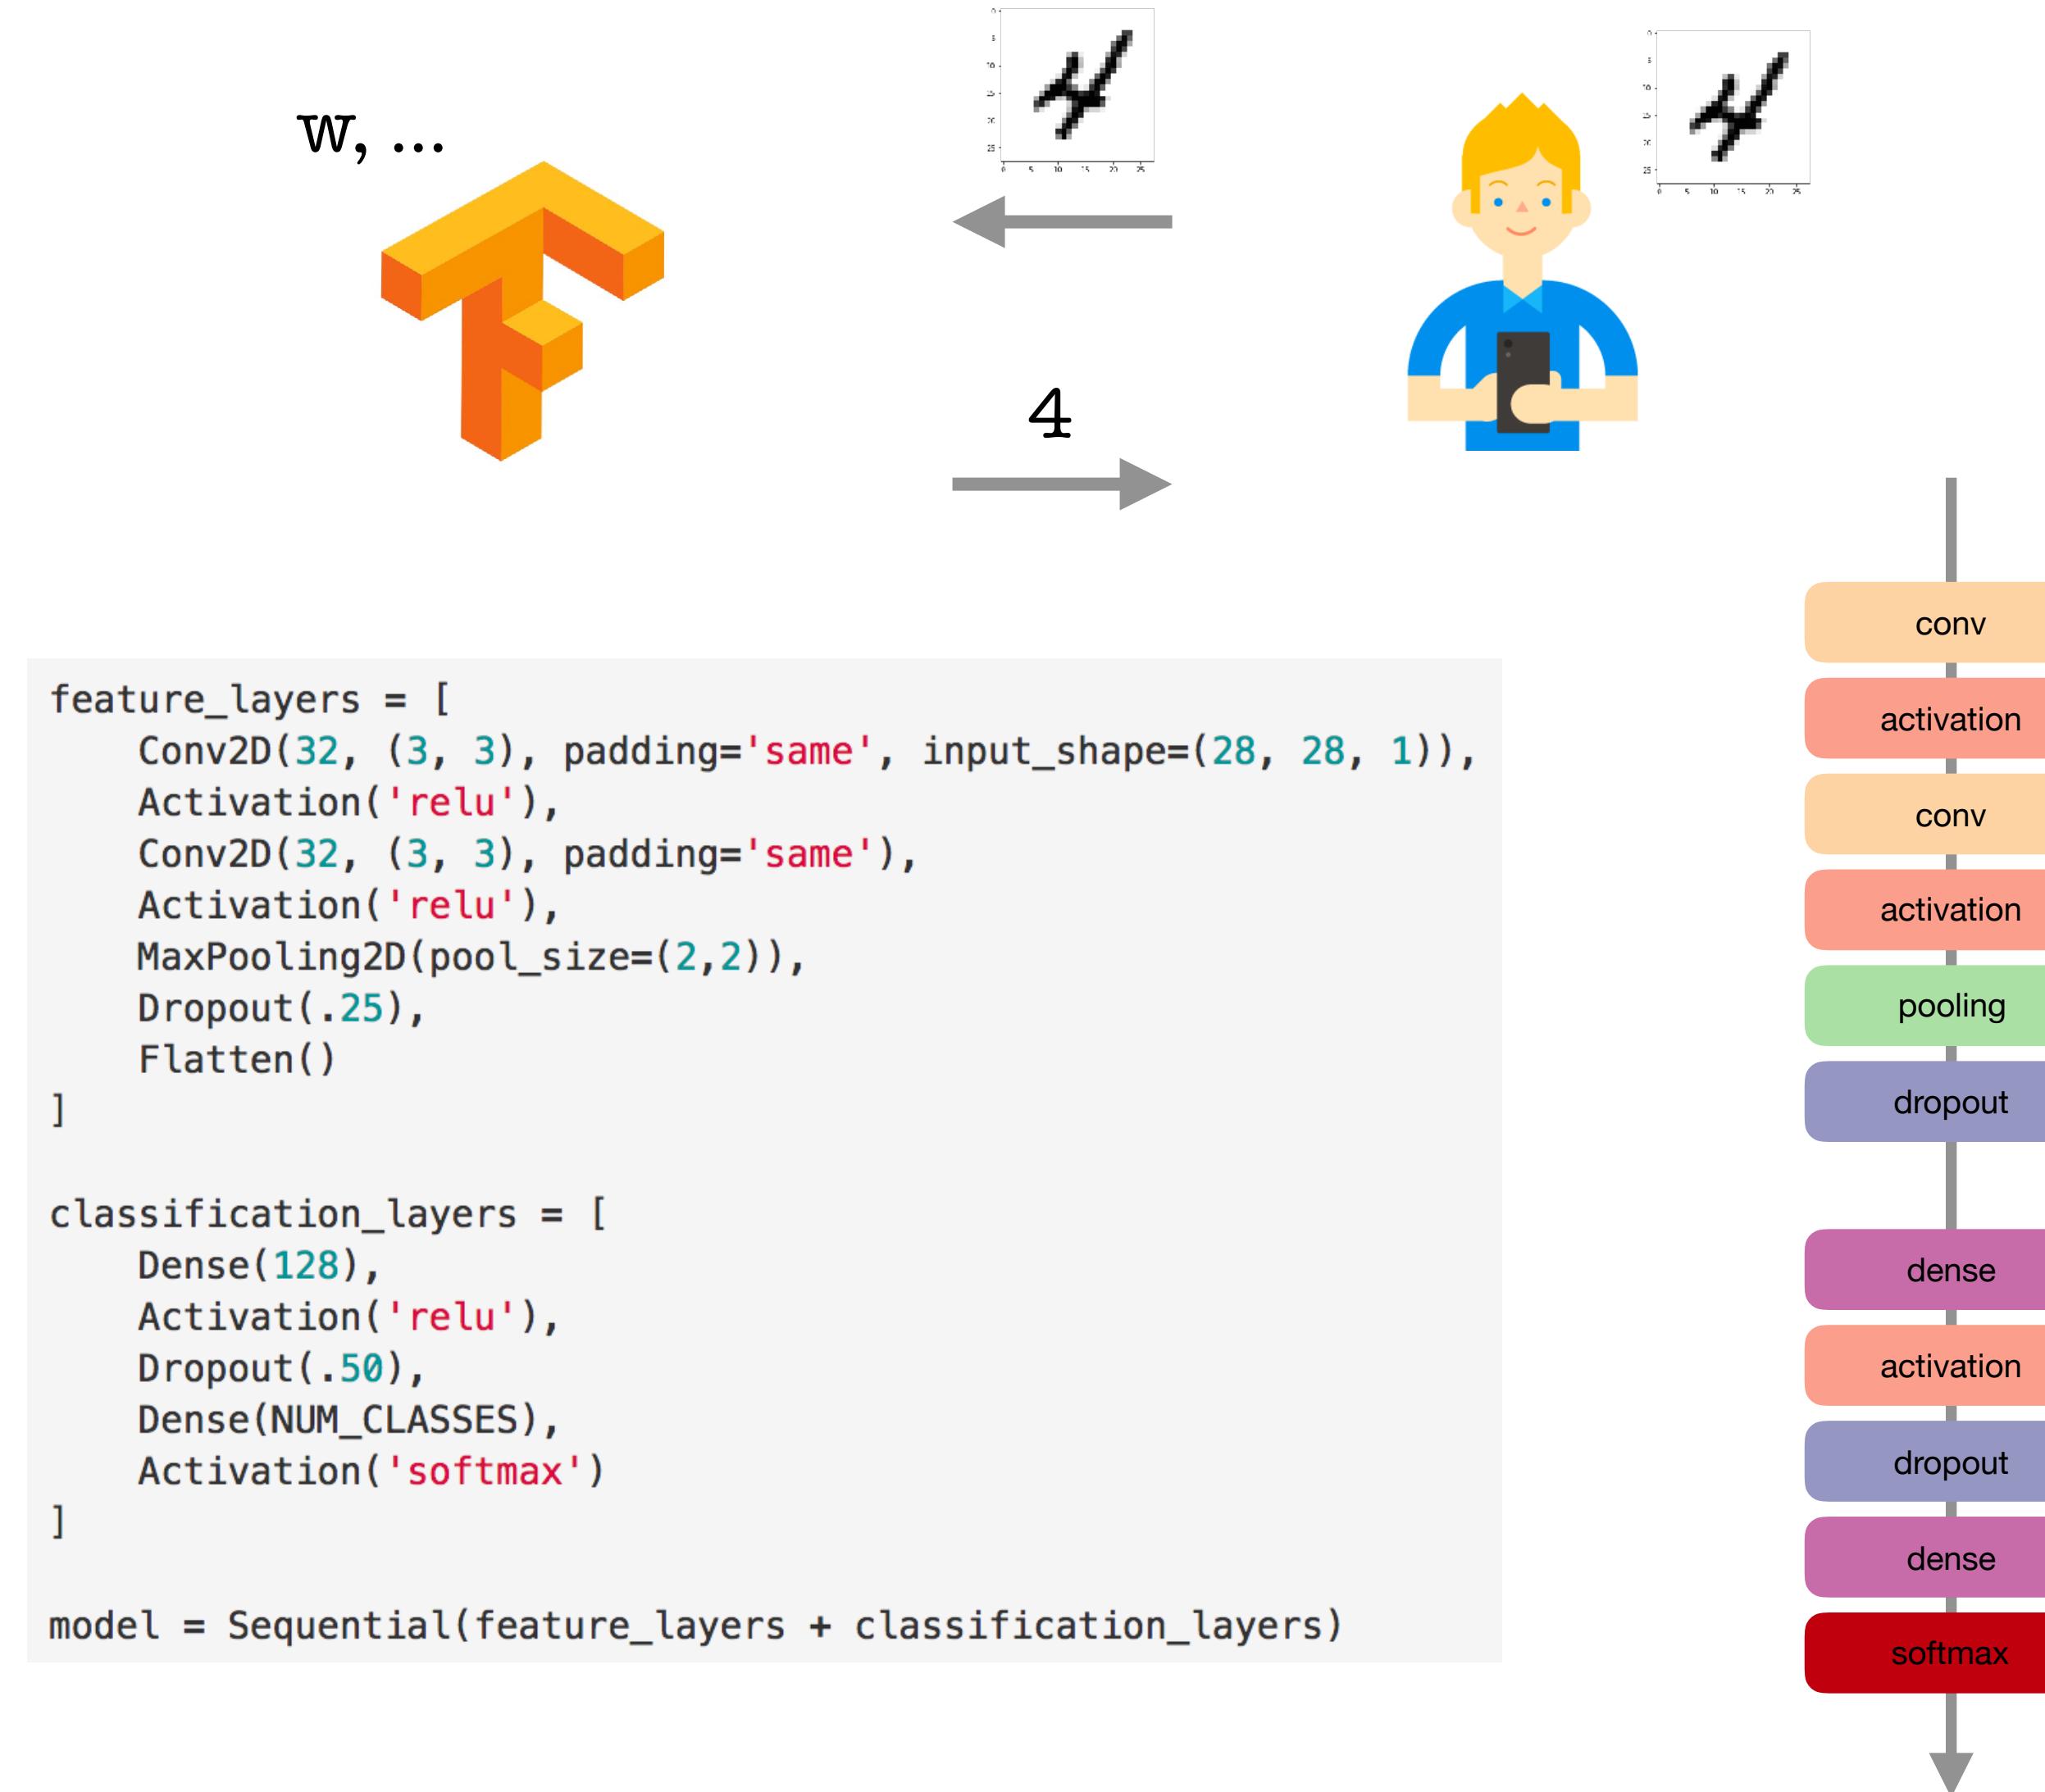
z2 = x2 * w

$$x = x1 + x2$$

$$\begin{aligned} x * w &= (x1 + x2) * w \\ &= (x1 * w) + (x2 * w) \\ &= z1 + z2 \end{aligned}$$

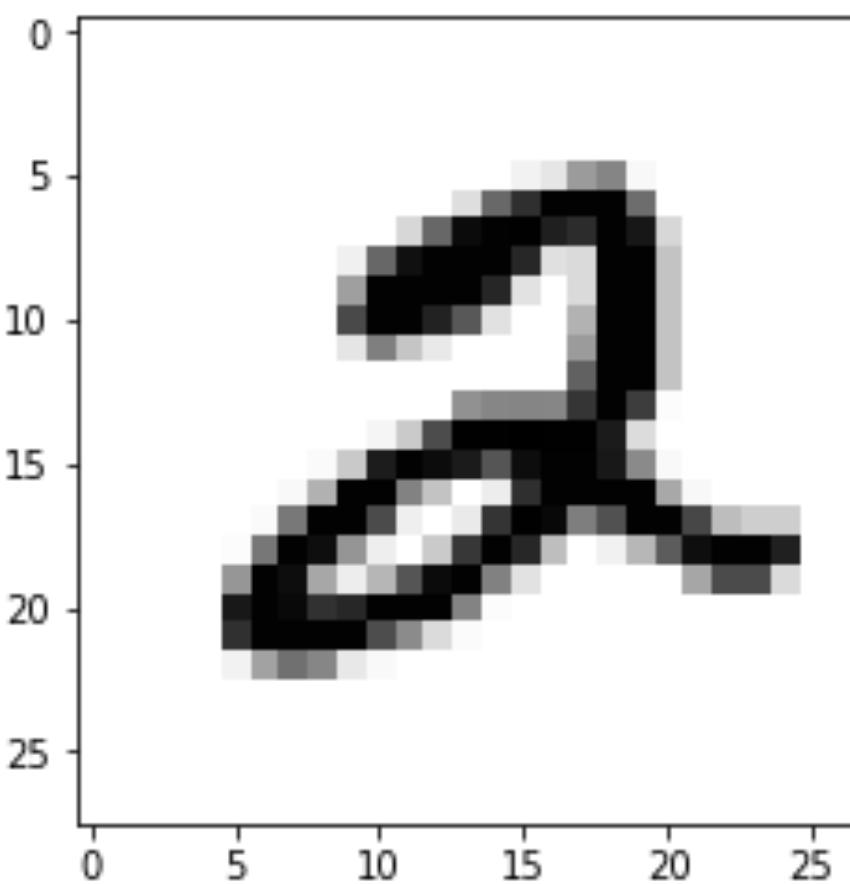
Convolutional Neural Networks

Digit Classification with CNNs

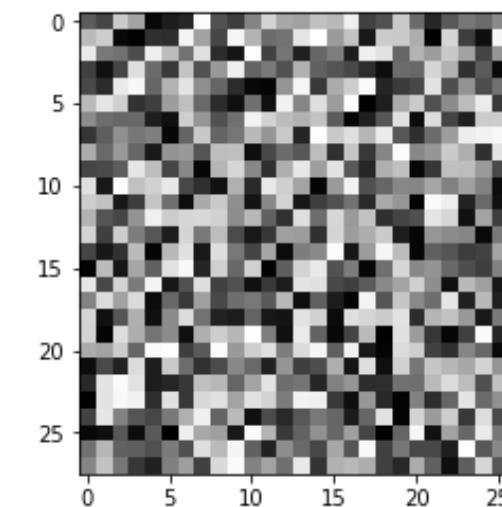


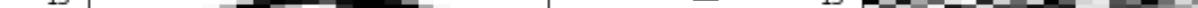
Secret Sharing Images

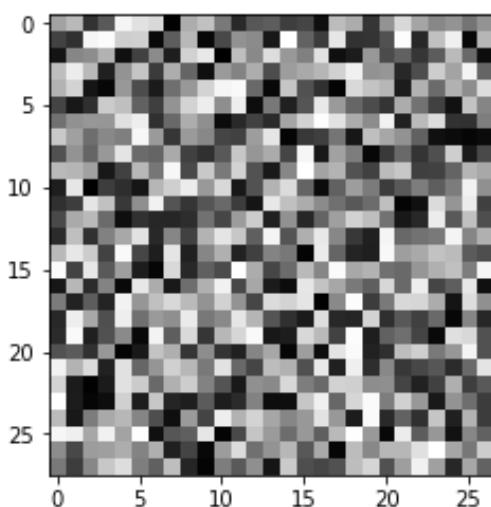
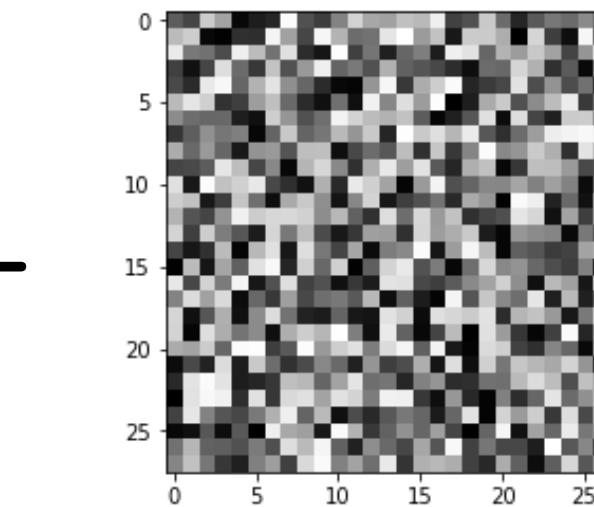
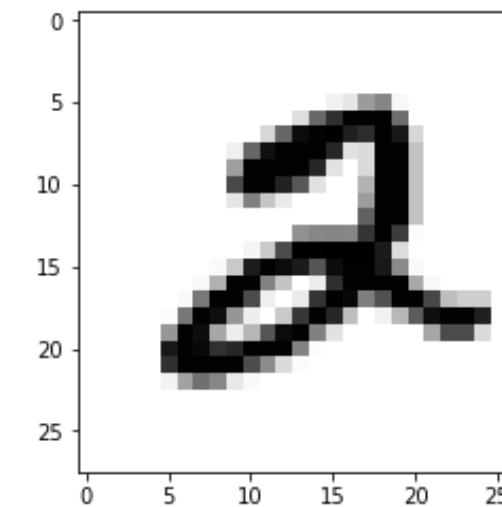
Secret Sharing Images



$$\text{Share}_1(x, r) = r \quad =$$

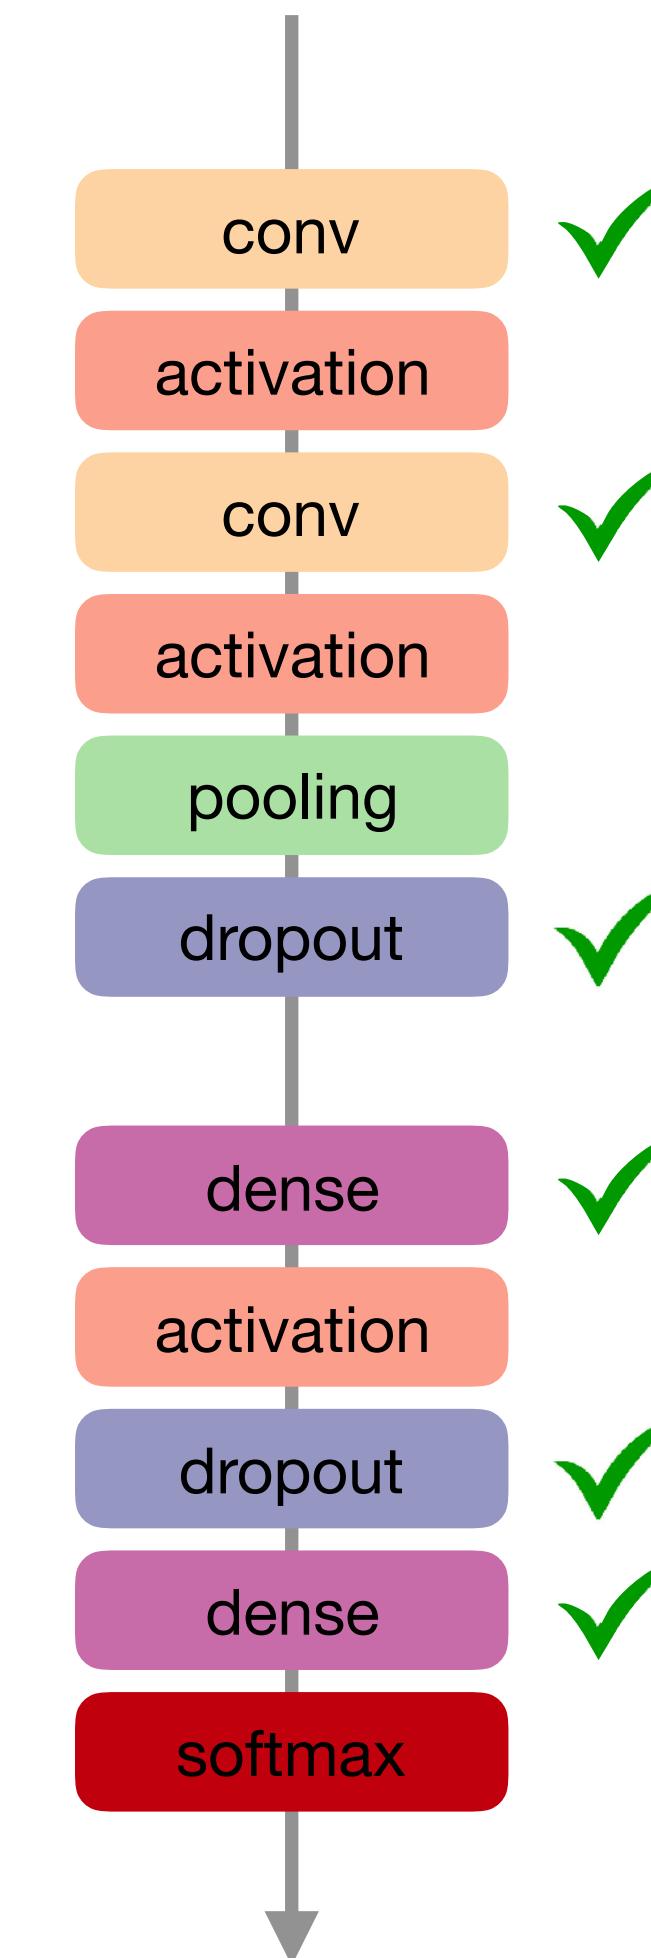


$$\text{Share2}(x, r) = x - r =$$




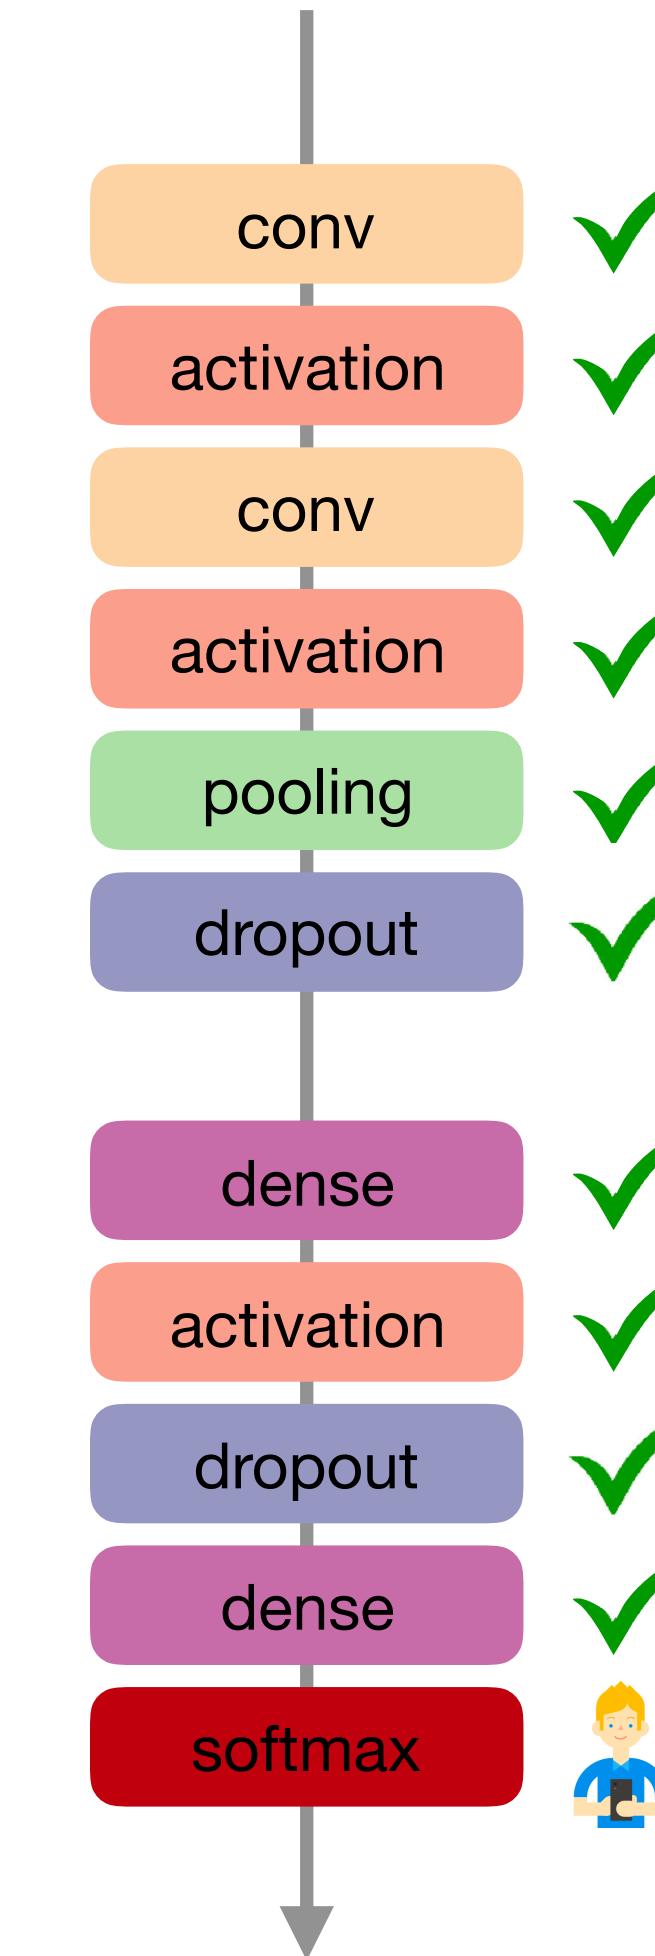
Digit Classification with Secret Sharing

```
feature_layers = [  
    Conv2D(32, (3, 3), padding='same', input_shape=(28, 28, 1)),  
    Activation('relu'),  
    Conv2D(32, (3, 3), padding='same'),  
    Activation('relu'),  
    MaxPooling2D(pool_size=(2,2)),  
    Dropout(.25),  
    Flatten()  
]  
  
classification_layers = [  
    Dense(128),  
    Activation('relu'),  
    Dropout(.50),  
    Dense(NUM_CLASSES),  
    Activation('softmax')  
]  
  
model = Sequential(feature_layers + classification_layers)
```



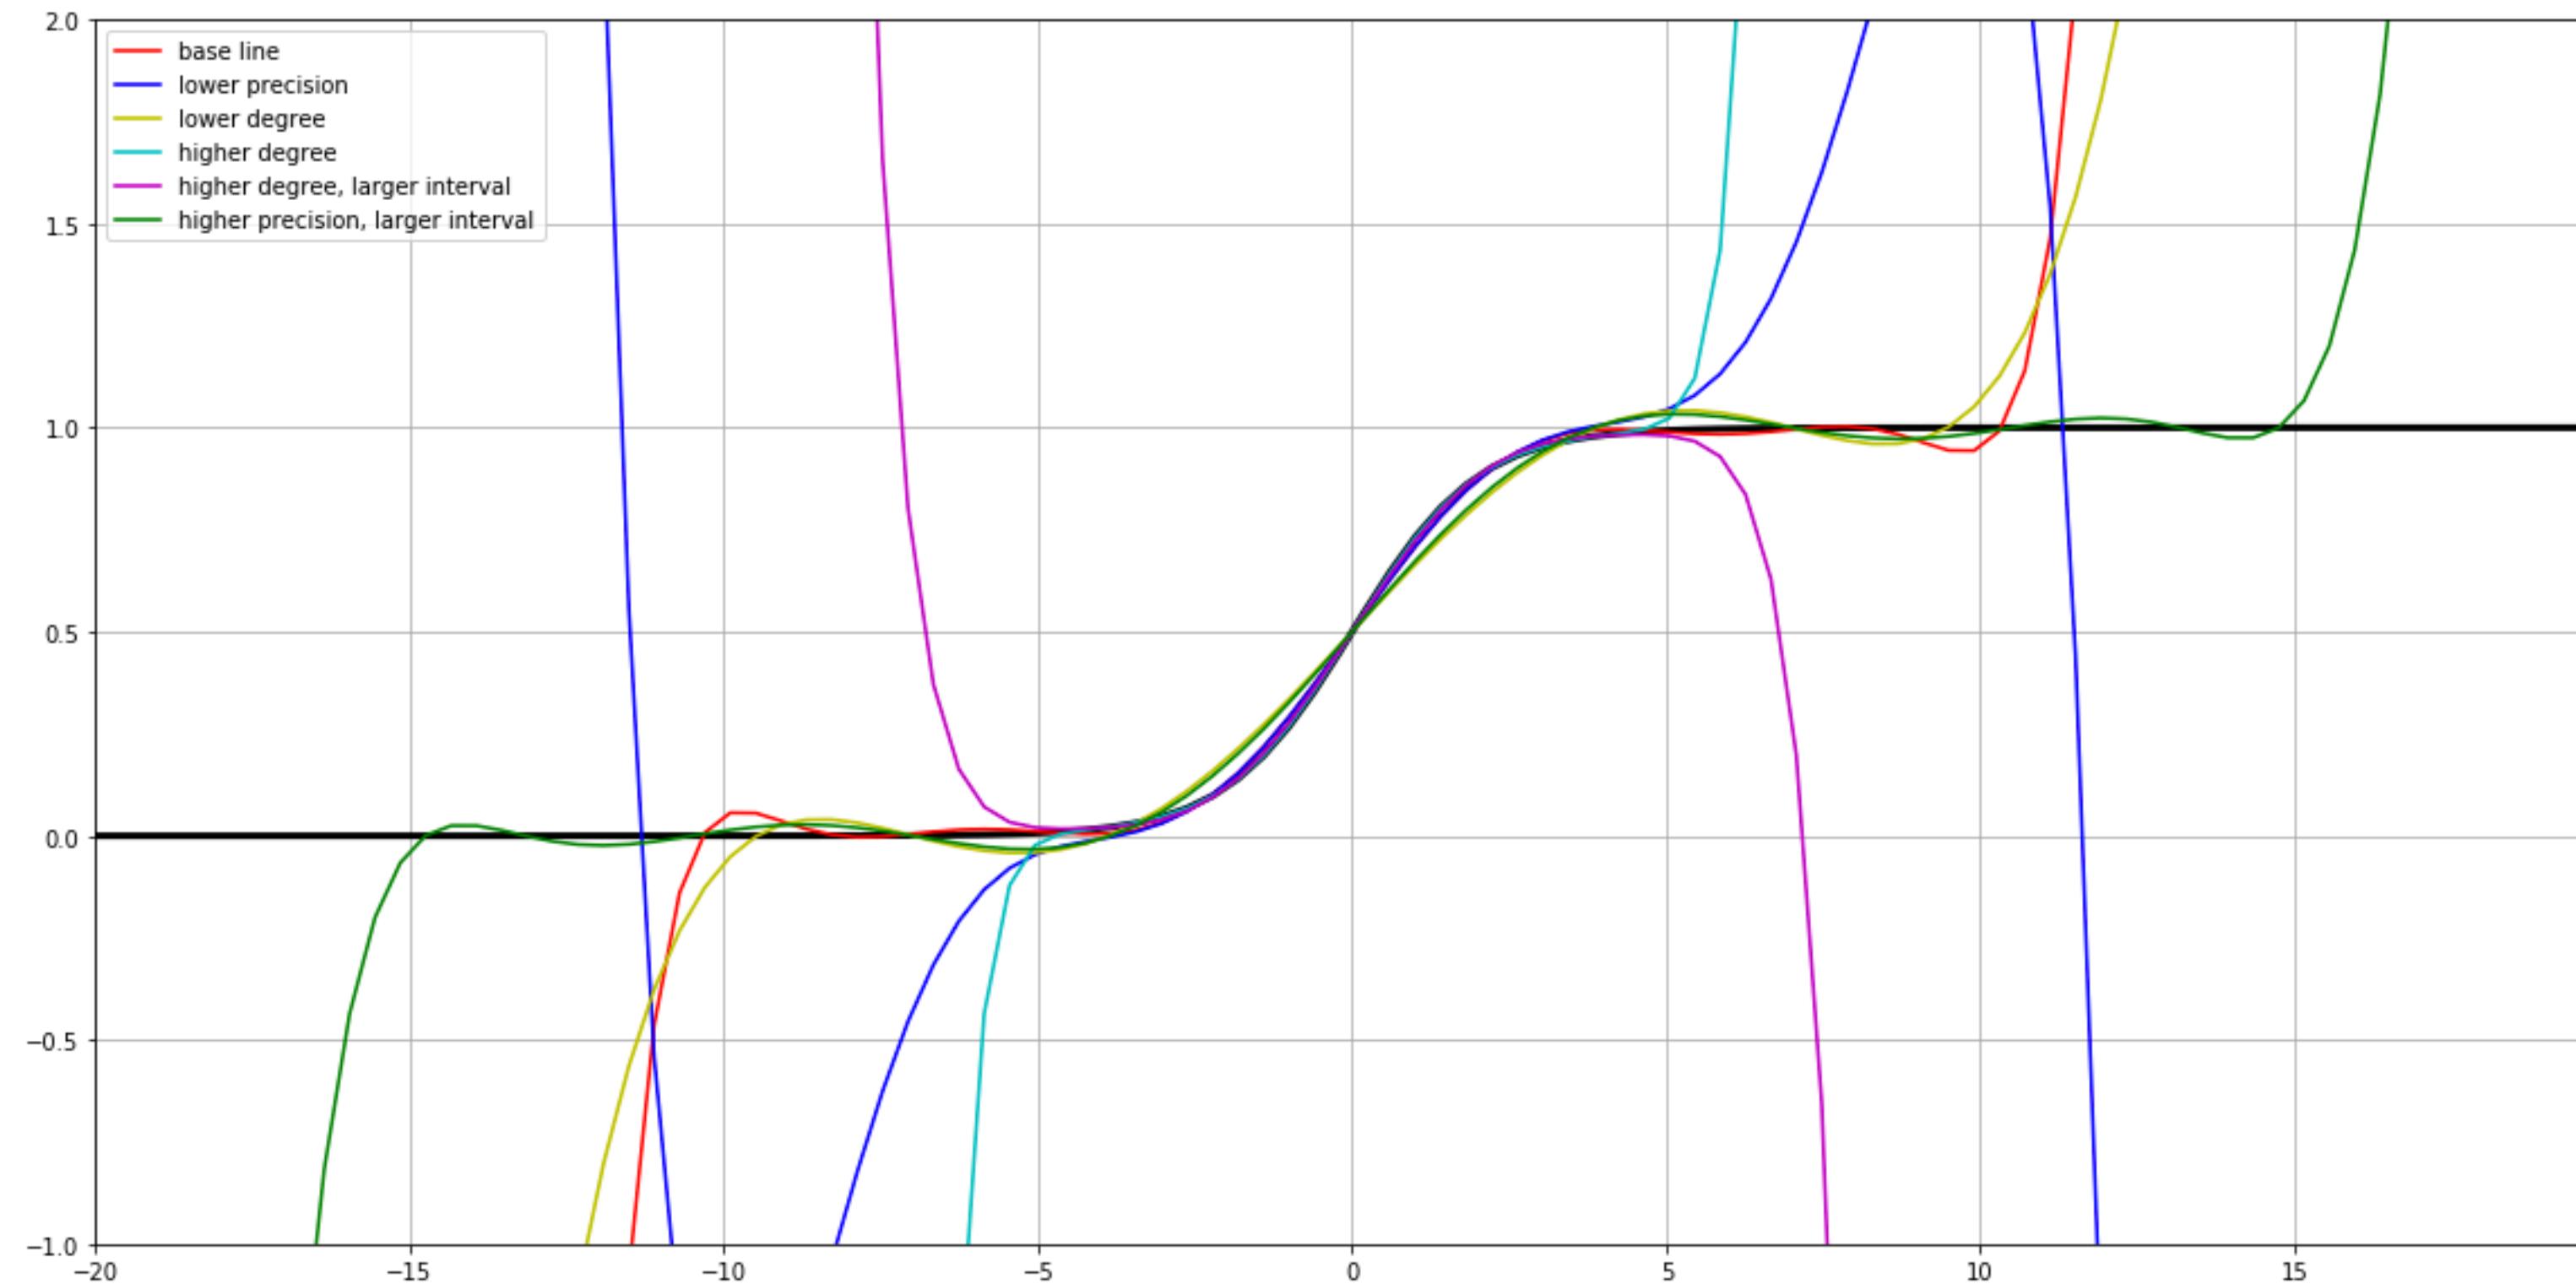
Digit Classification with Secret Sharing

```
feature_layers = [
    feature_layers = [
        Conv2D(32, (3, 3), padding='same', input_shape=(28, 28, 1)),
        Activation('sigmoid'),
        Conv2D(32, (3, 3), padding='same'),
        Activation('sigmoid'),
        AveragePooling2D(pool_size=(2,2)),
        Dropout(.25),
    ] Flatten()
]
c1
classification_layers = [
    Dense(128),
    Activation('sigmoid'),
    Dropout(.50),
    Dense(5),
] Activation('softmax')
]
mc
model = Sequential(feature_layers + classification_layers)
```



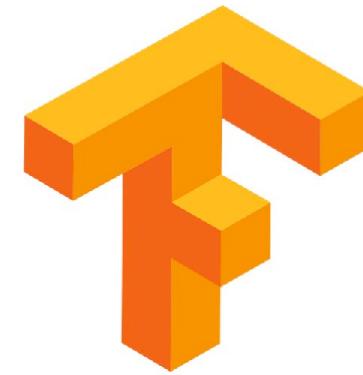
model adaptation can improve performance significantly

Sigmoid via Polynomial Approximation



$$f(x) = c_7 * x^7 + c_5 * x^5 + \dots + c_1 * x + c_0$$

Private Multiplication in SPDZ



random triple
(a_1 , b_1 , c_1)

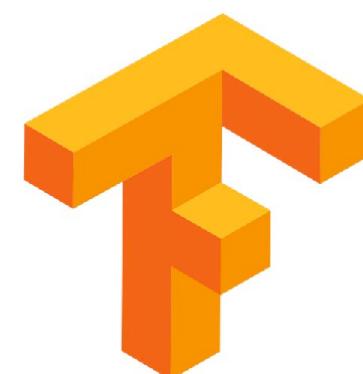
x_1

y_1

alpha

beta

$$z_1 = \text{alpha} * \text{beta} + \text{alpha} * b_1 + \text{beta} * a_1 + c_1$$



(a_2 , b_2 , c_2)

x_2

y_2

alpha

beta

$$z_2 = \text{alpha} * b_2 + \text{beta} * a_2 + c_2$$

$$\begin{aligned} a &= a_1 + a_2 \\ b &= b_1 + b_2 \\ c &= a * b = c_1 + c_2 \end{aligned}$$

$$x = x_1 + x_2 \quad y = y_1 + y_2$$

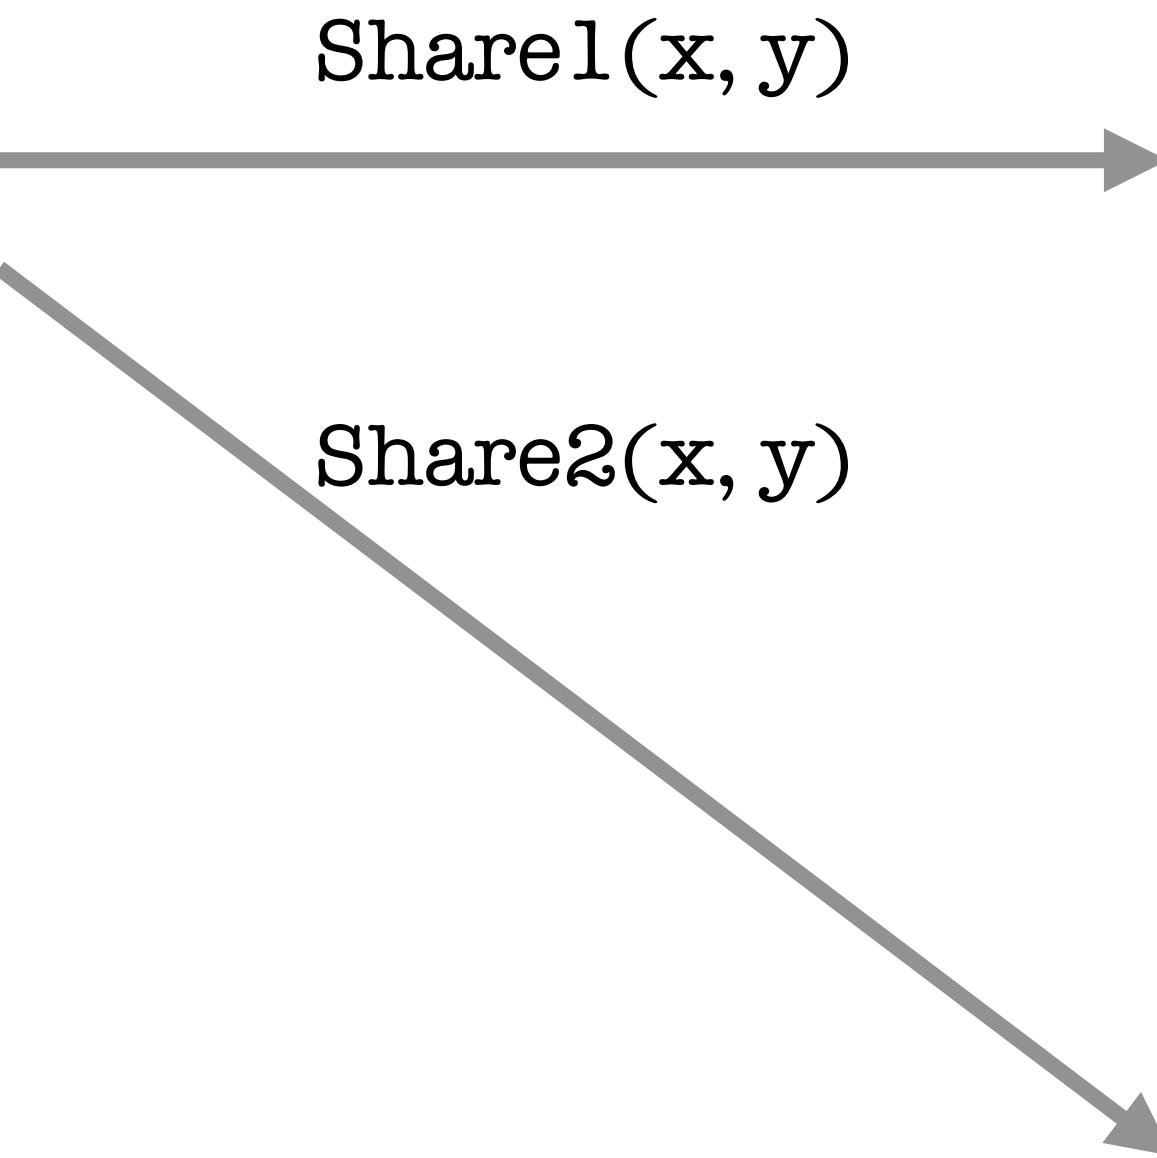
$$\text{alpha} = x - a$$

$$\text{beta} = y - b$$

$$\begin{aligned} x * y &= \dots \\ &= z_1 + z_2 \end{aligned}$$

Training

Participants



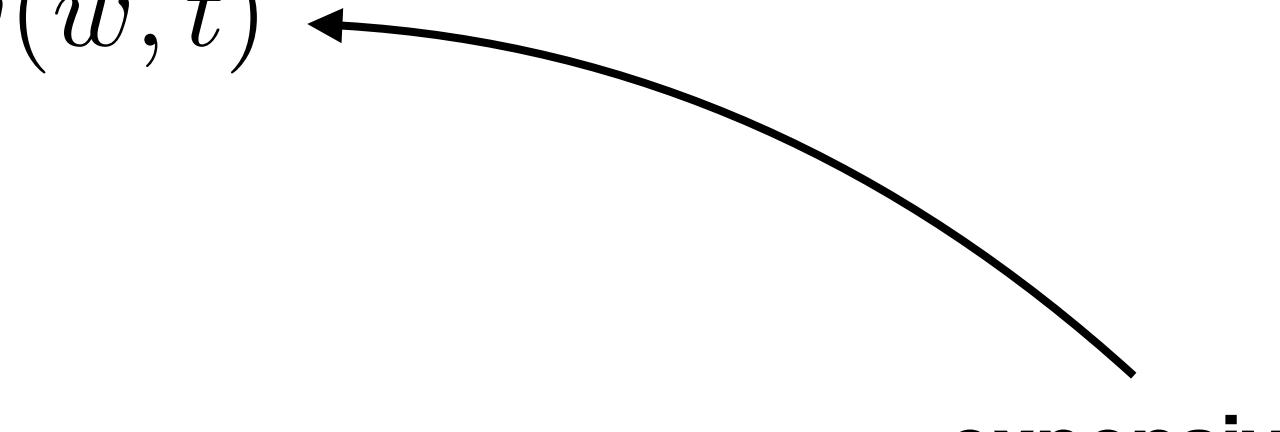
SoftMax Alternatives

$$f(x_i) = \frac{e^{-x_i}}{\sum_j e^{-x_j}}$$

$$g(x_i) = \frac{\text{ReLU}(x_i)}{\sum_j \text{ReLU}(x_j)}$$

Simpler Optimisation Strategies

RMSProp:

$$w = w - \frac{\eta}{\sqrt{v(w, t)}} \nabla Q(w)$$


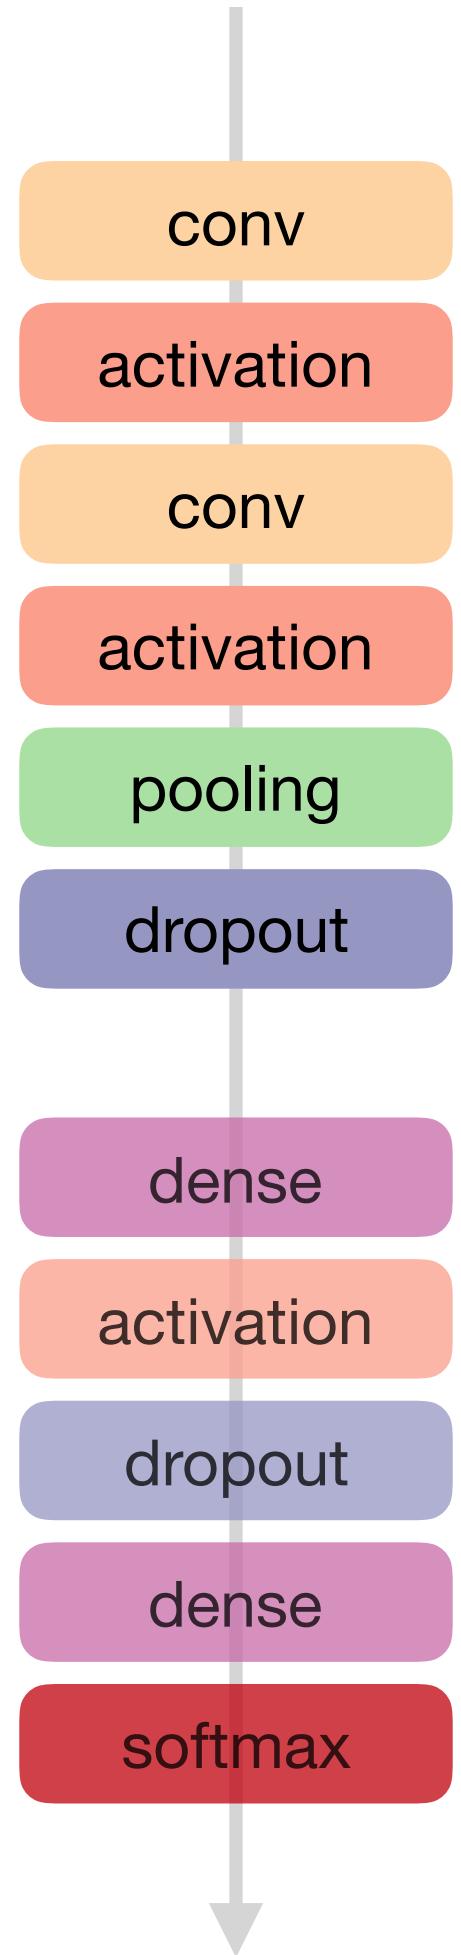
expensive

SGD:

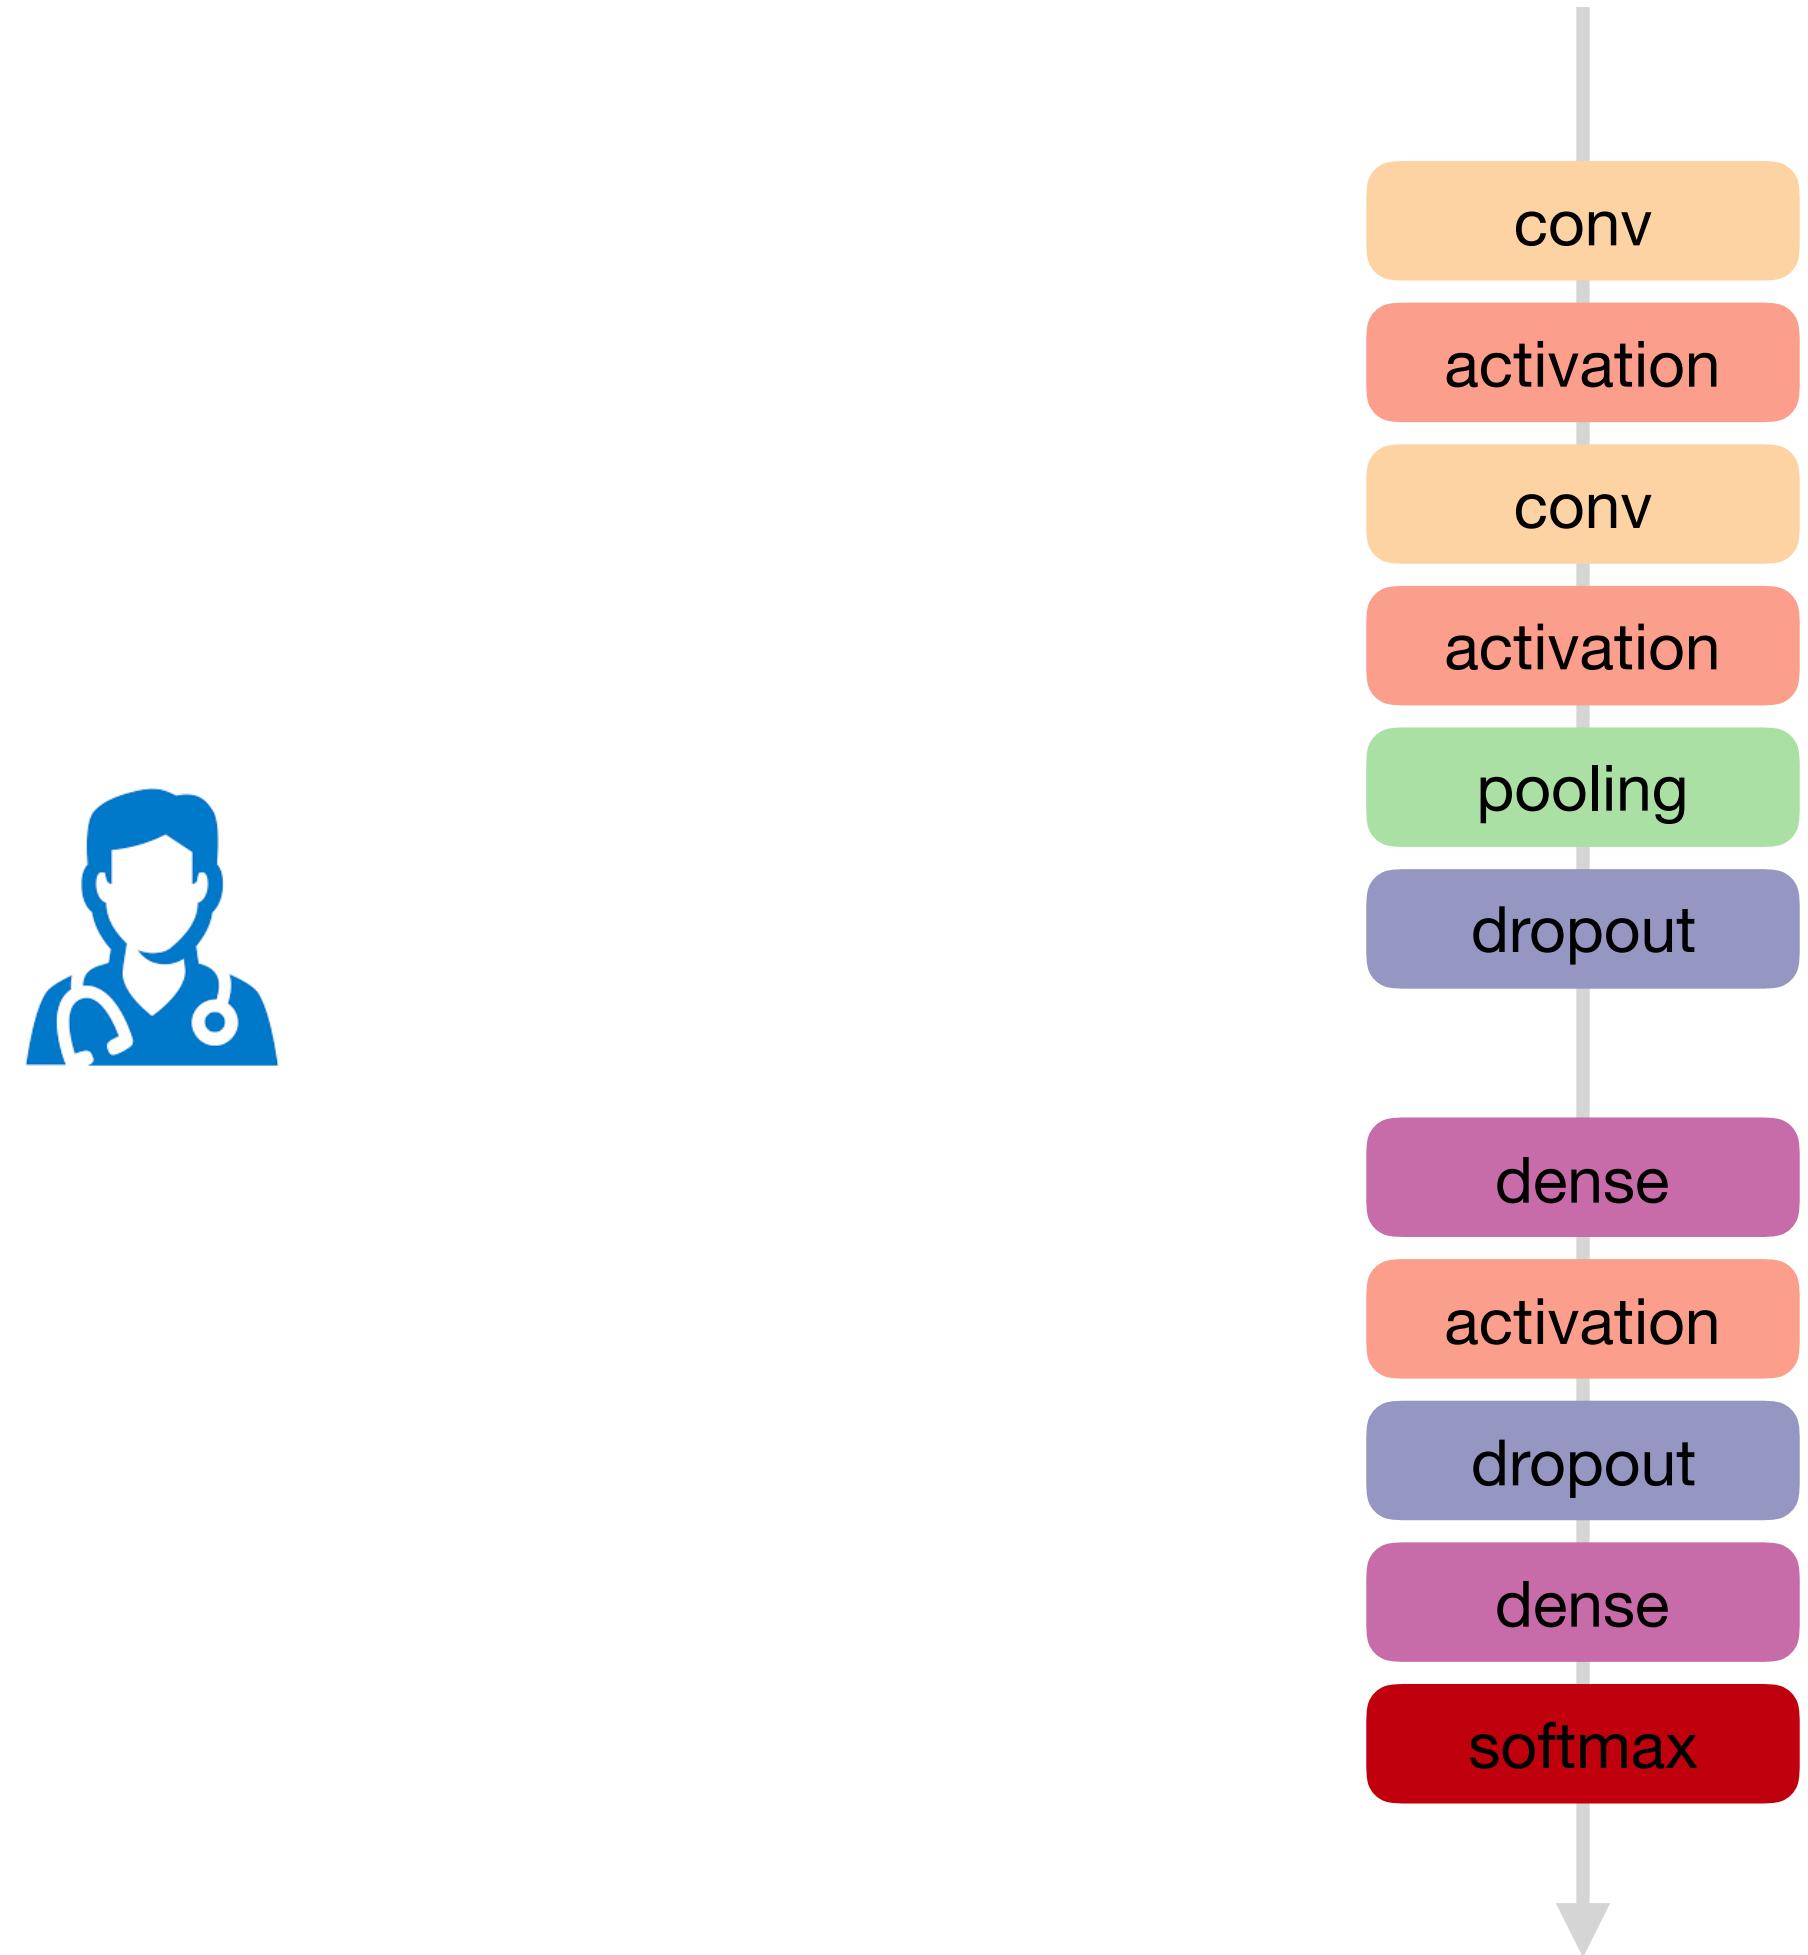
$$w = w - \eta \nabla Q(w)$$

Transfer Learning

IMAGENET



Transfer Learning



Making It Accessible

Projects and Literature

Recent research papers using secure computation

CryptoNets: Applying Neural Networks to Encrypted Data with High Throughput and Accuracy, Dowlin et al.

SecureML: A System for Scalable Privacy-Preserving Machine Learning, Mohassel and Zhang

DeepSecure: Scalable Provably-Secure Deep Learning, Rouhani et al.

Gazelle: A Low Latency Framework for Secure Neural Network Inference, Juvekar et al.

ABY3: A Mixed Protocol Framework for Machine Learning, Mohassel and Rindal

SecureNN: Efficient and Private Neural Network Training, Wagh et al.

(great summary in <https://eprint.iacr.org/2017/1190>)

Secure computation frameworks

SCALE-MAMBA (<https://homes.esat.kuleuven.be/~nsmart/SCALE/>)

ABY (<https://github.com/encryptogroup/ABY>)

OblivC (<http://oblivc.org/>)

Specialised projects

tf-encrypted (<https://github.com/mortendahl/tf-encrypted>)

PySyft (<https://github.com/OpenMined/PySyft>)

(much more at <https://github.com/rdragos/awesome-mpc>)

Multidisciplinary Challenge

Data science
(use-cases, workflow, monitoring)

Cryptography
(techniques, protocols, guarantees)

Machine learning
(models, approx, quantisation)

Engineering
(distributed, multi-core, readability)

useful to have common framework

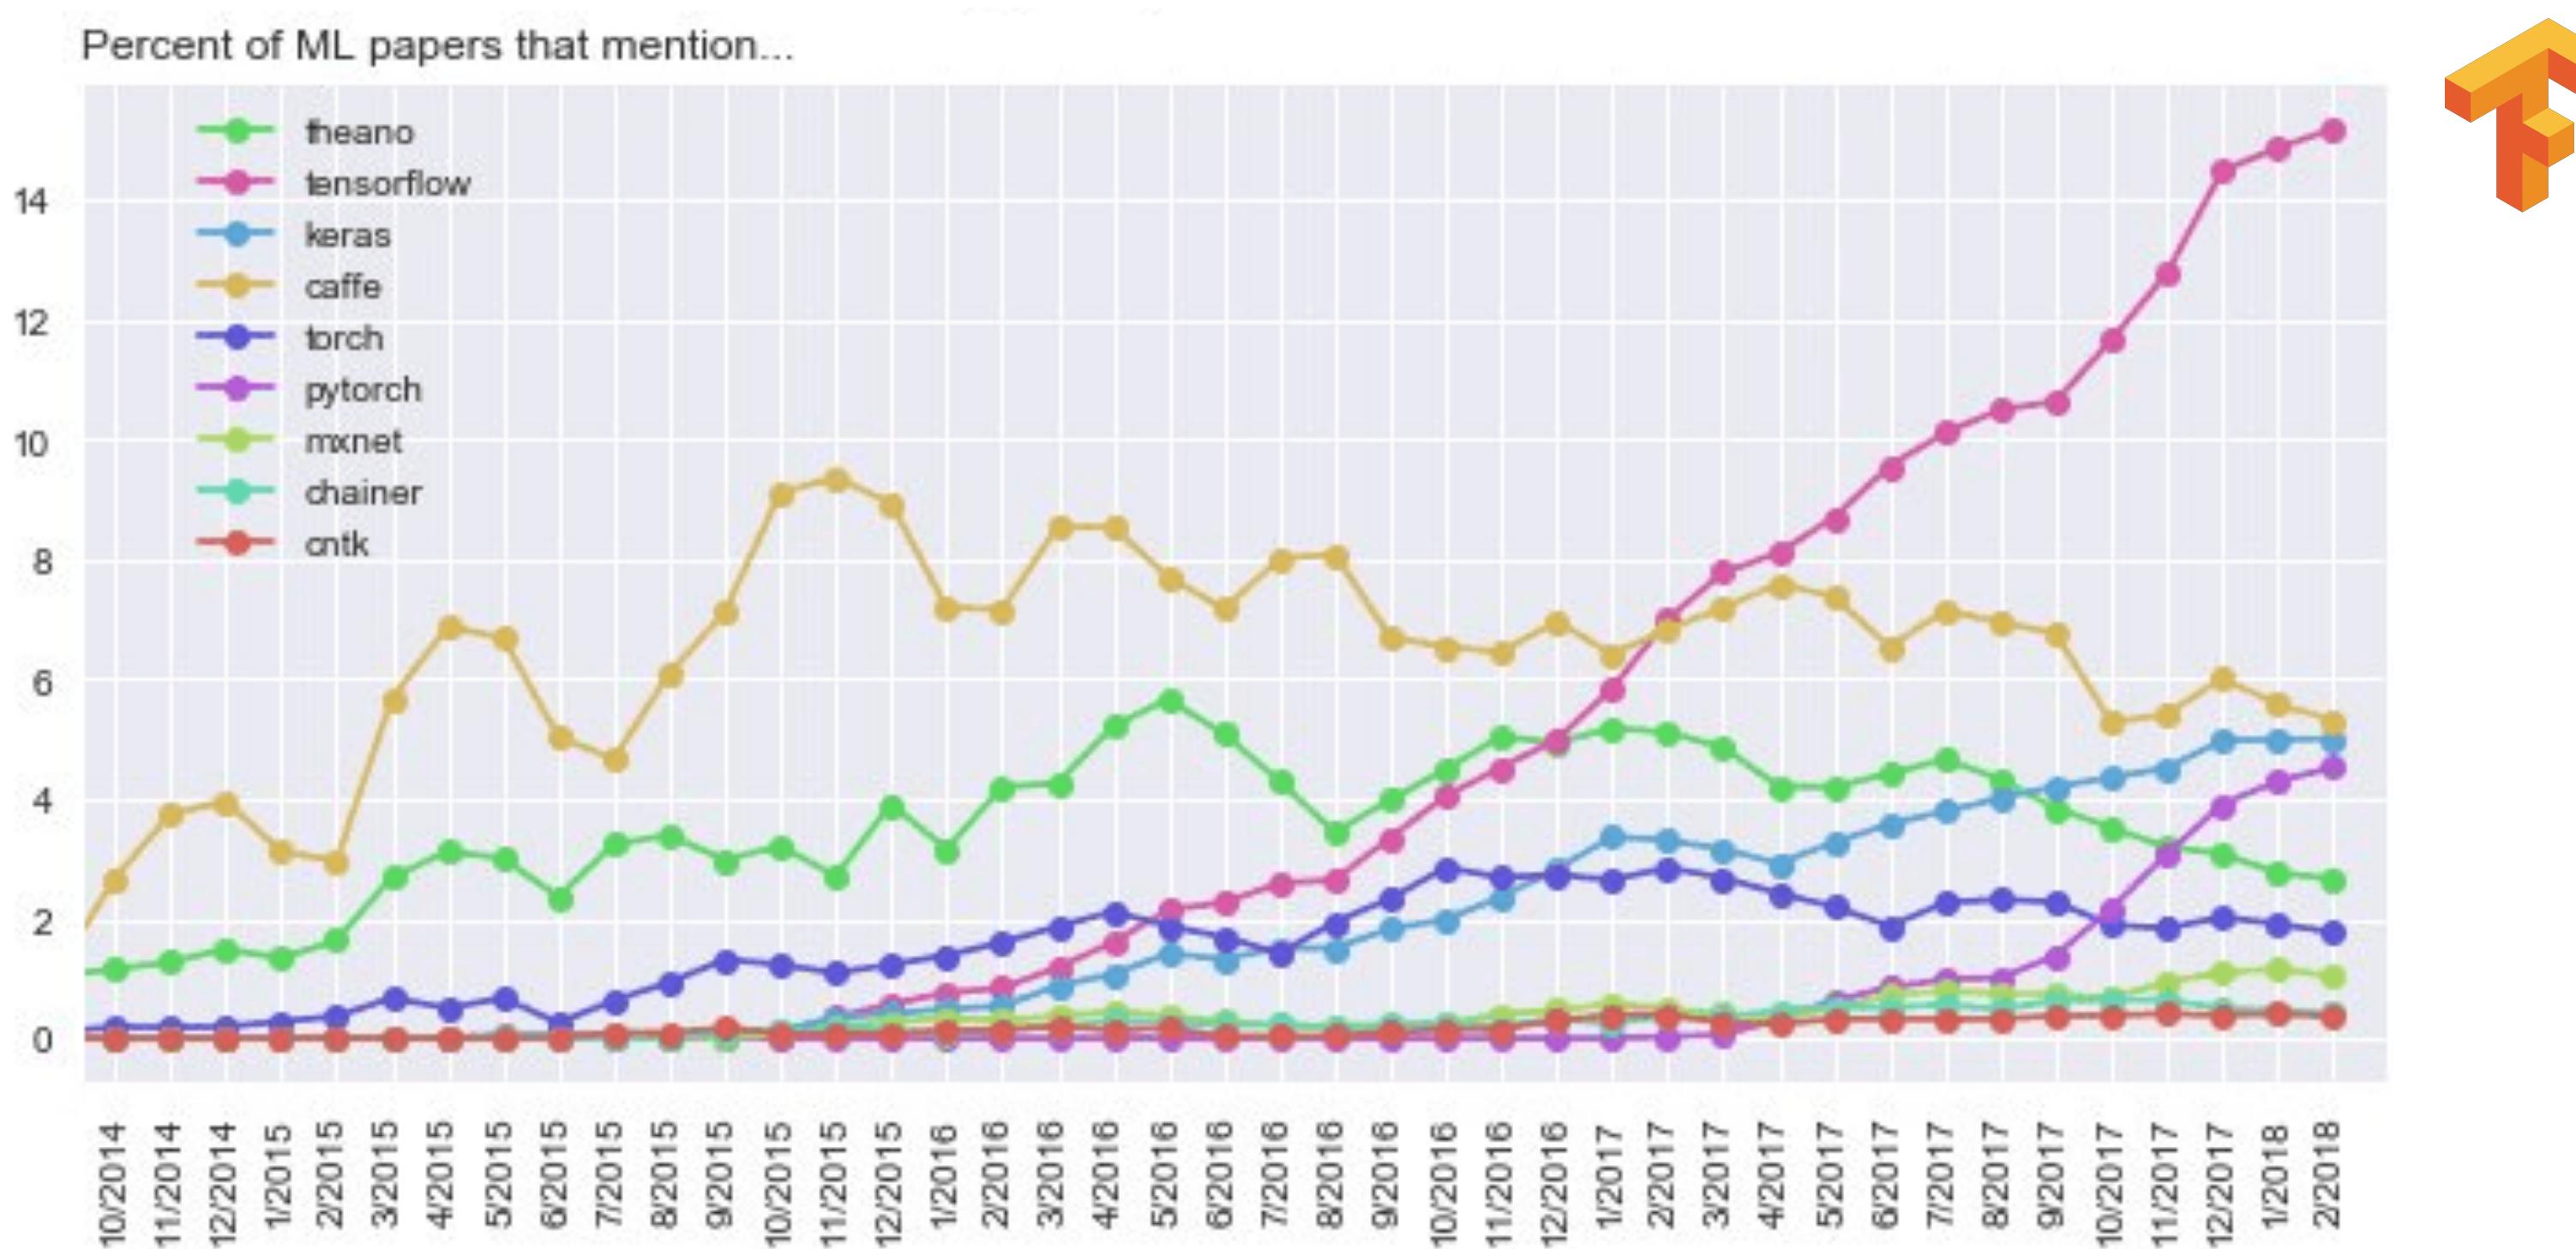
tf-encrypted

open source project for exploring and experimenting with
privacy-preserving machine learning in TensorFlow

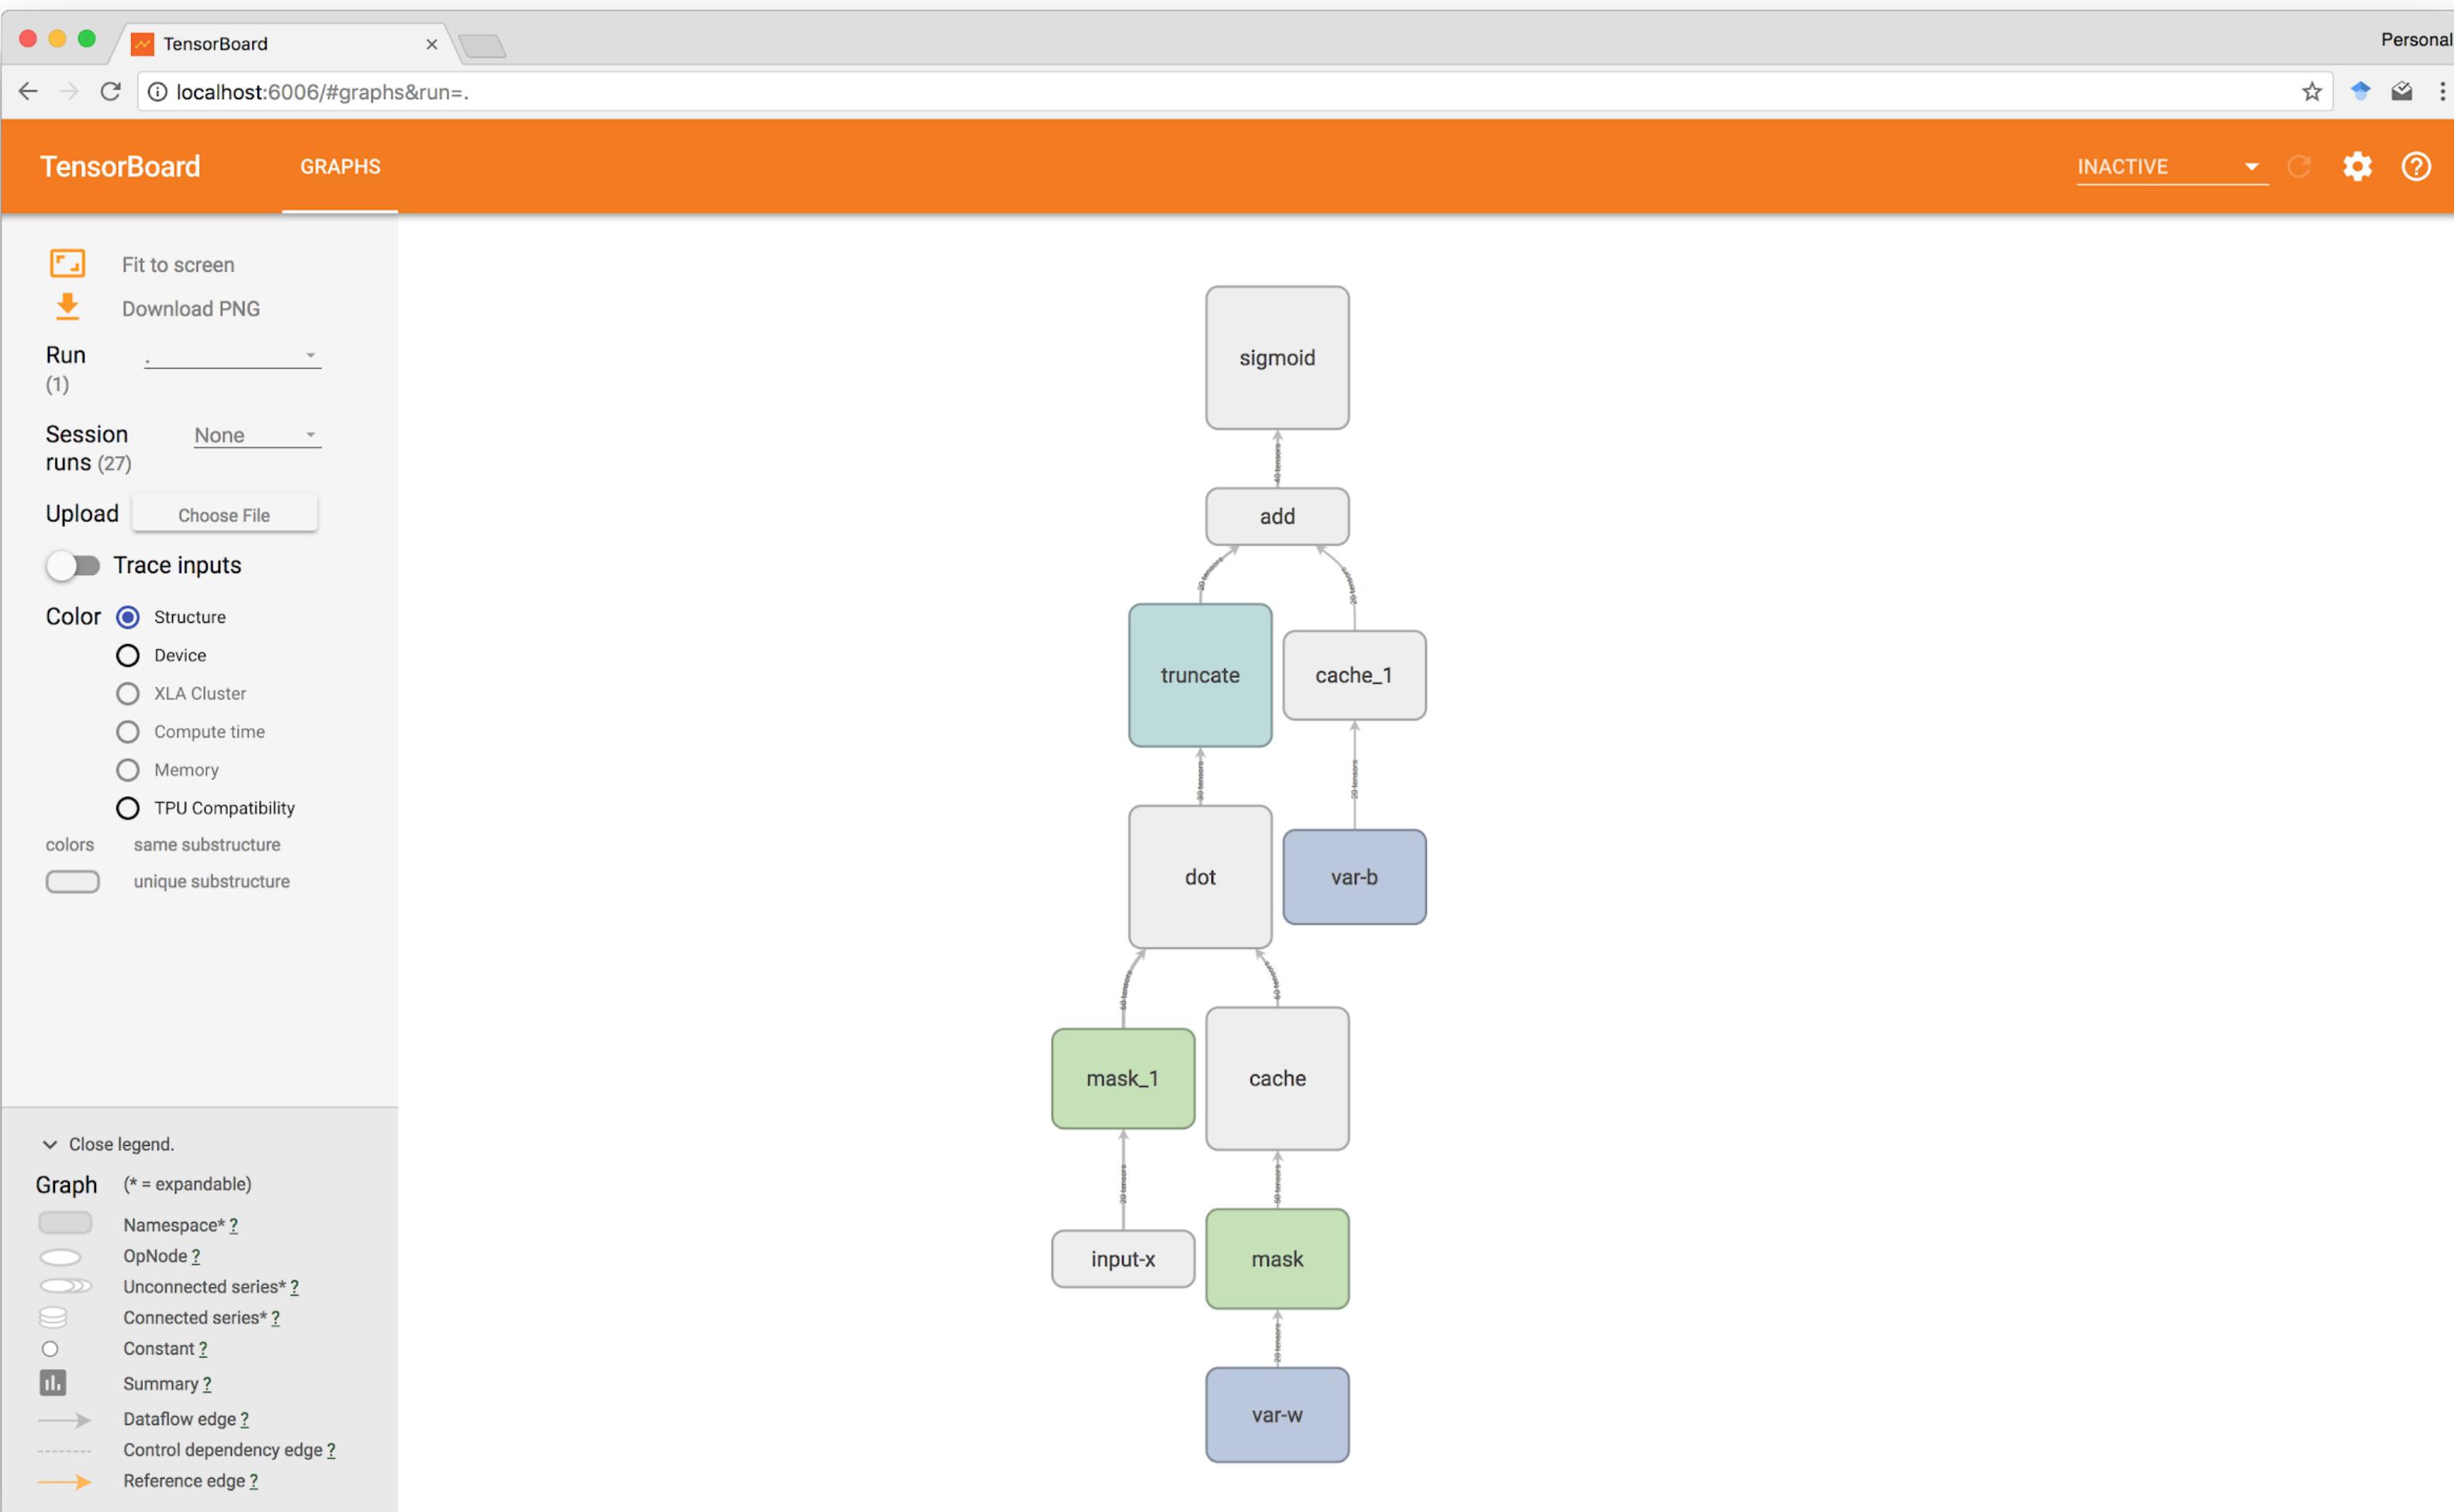
TensorFlow

platform for production-level training and deployment of models

popular and backed by Google



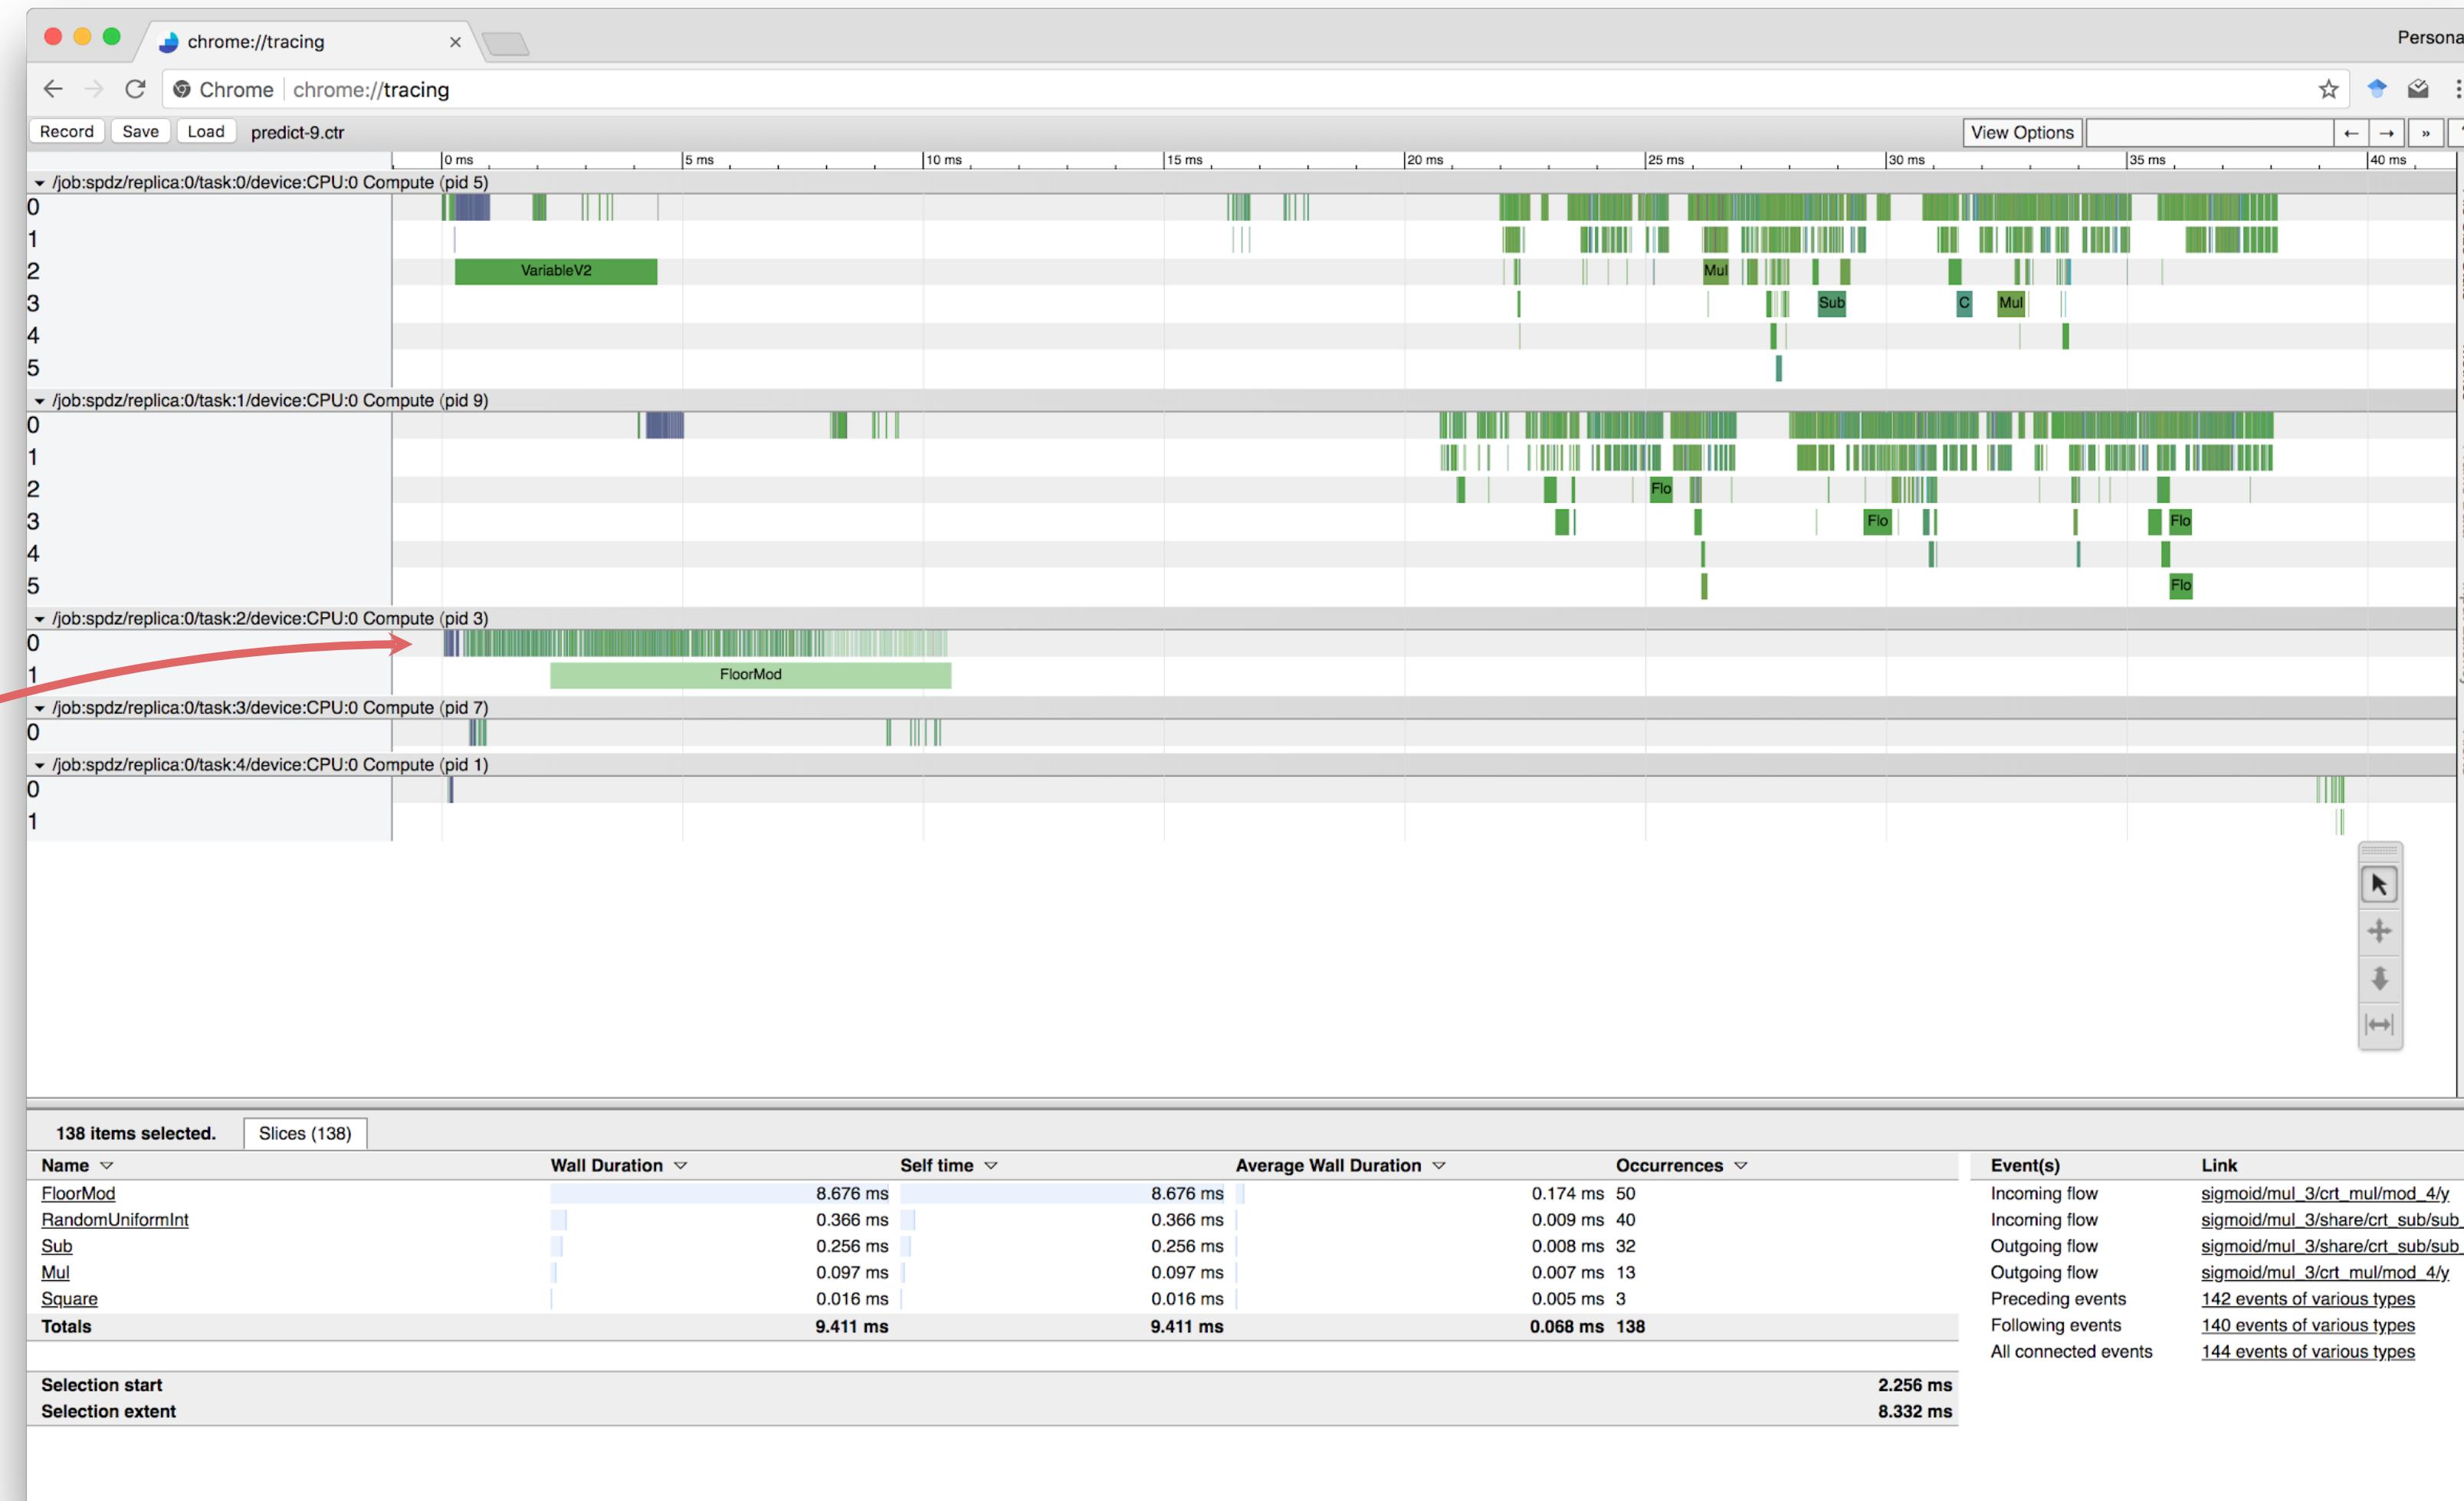
Dataflow Graphs



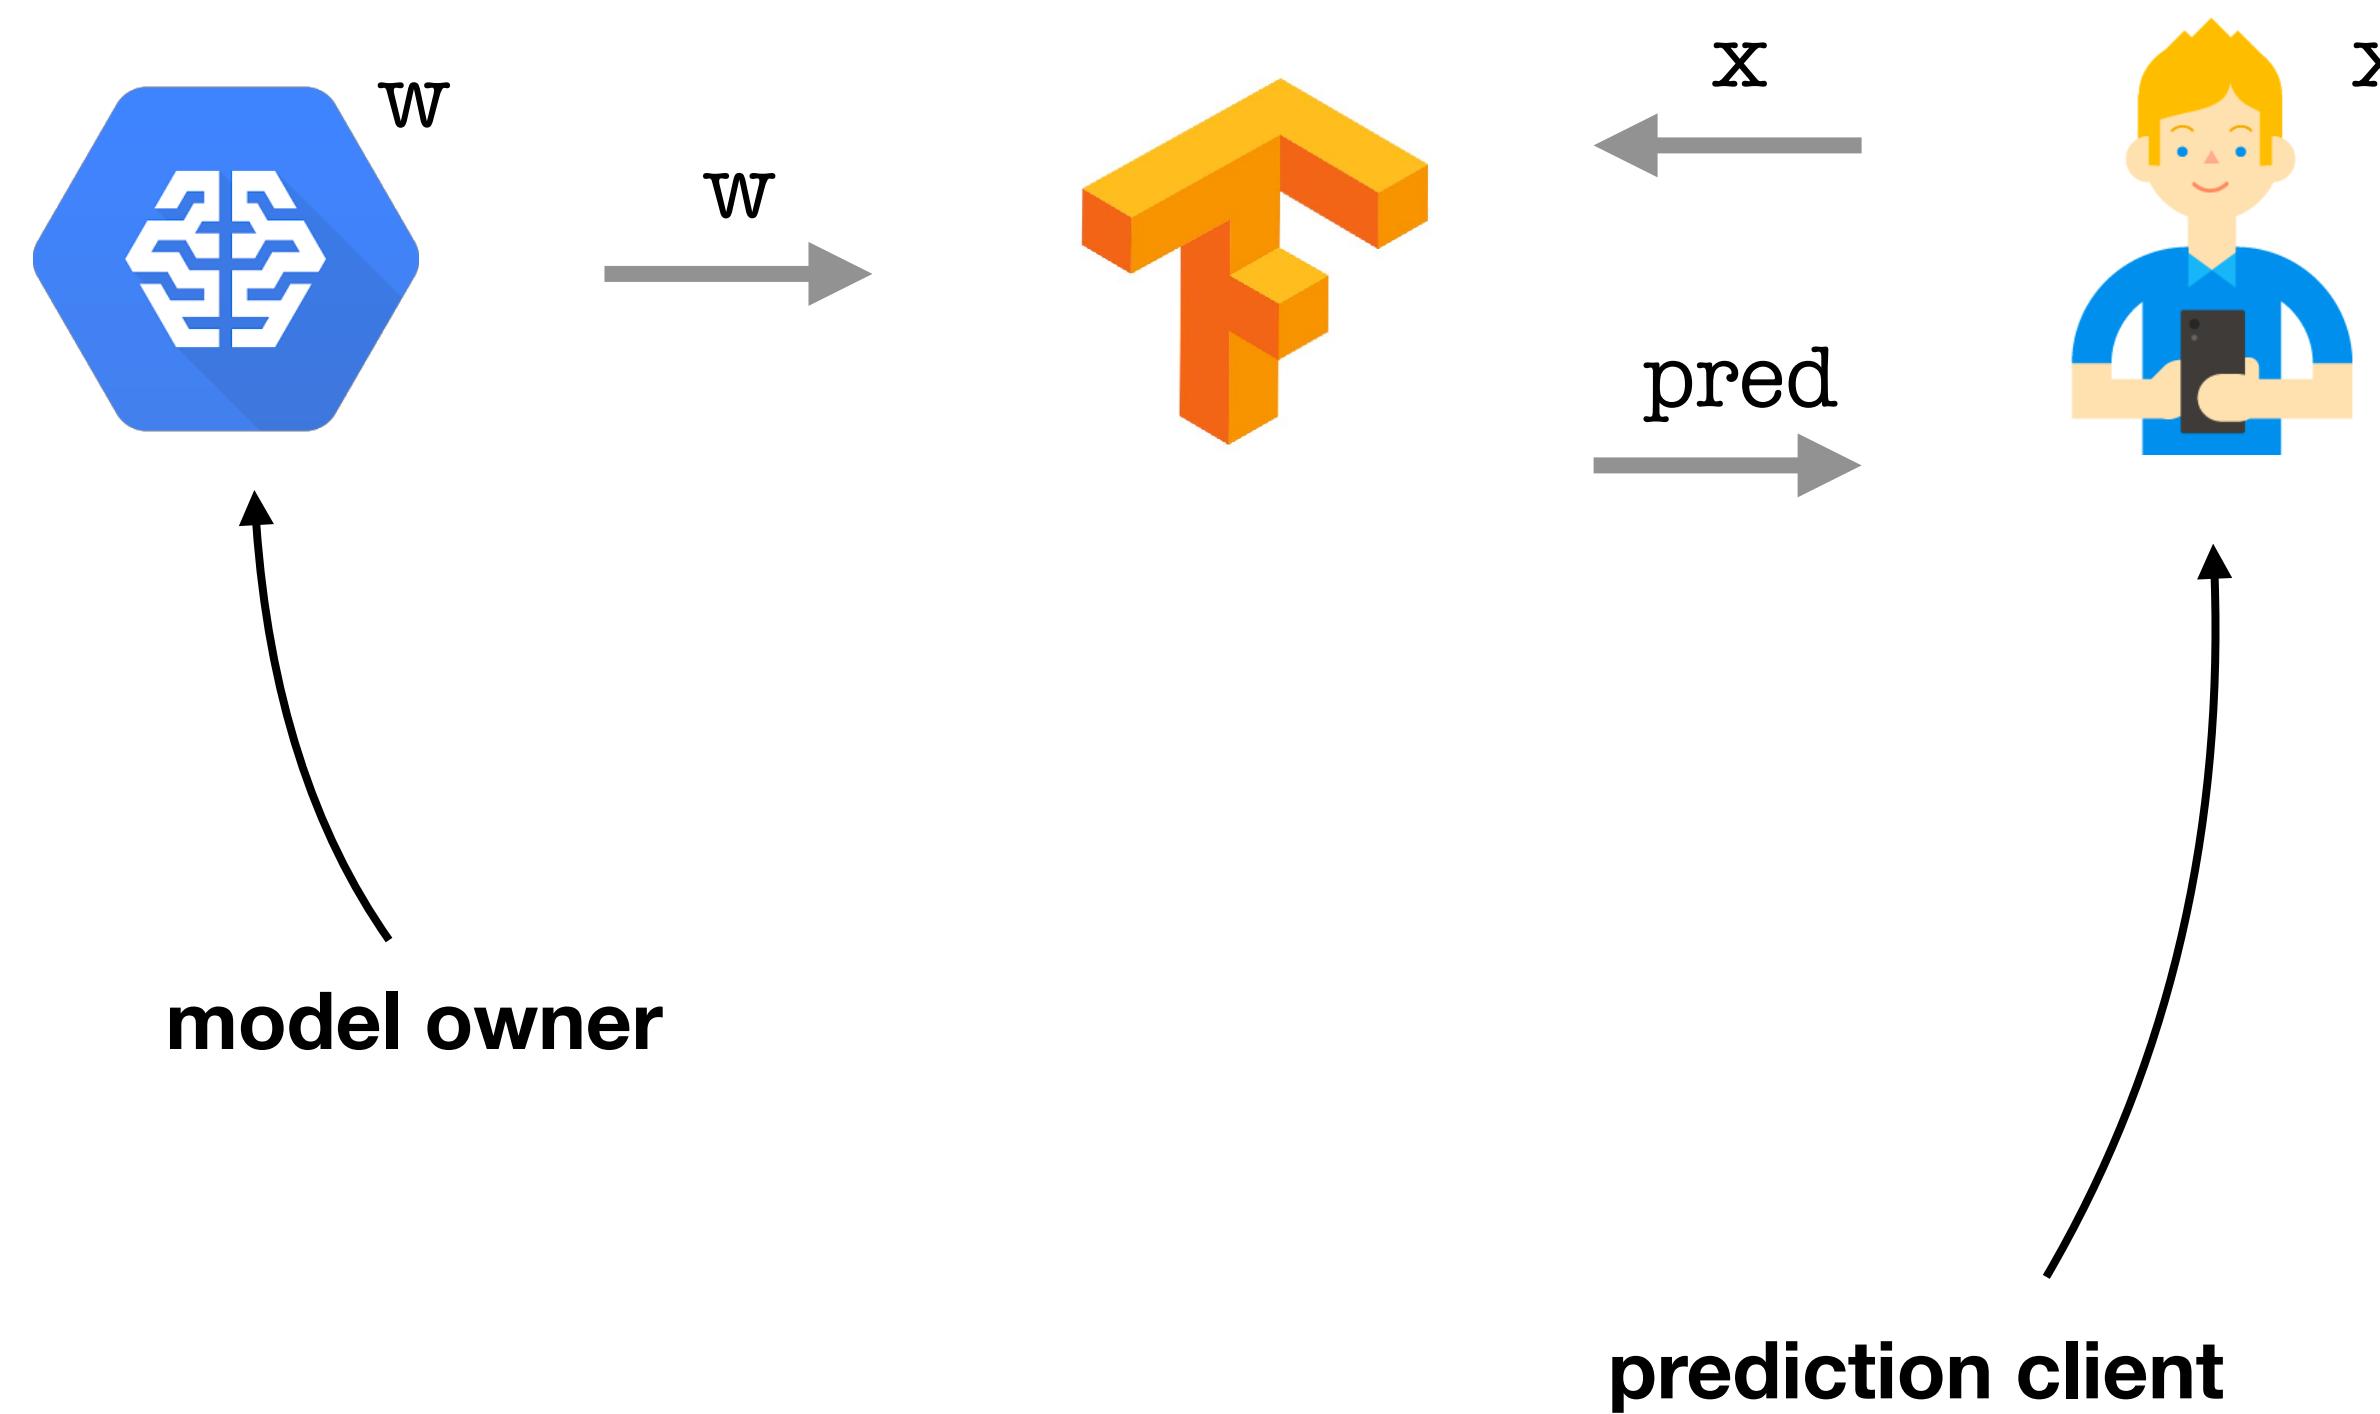
Optimized Distributed Execution

Concurrent

Batching

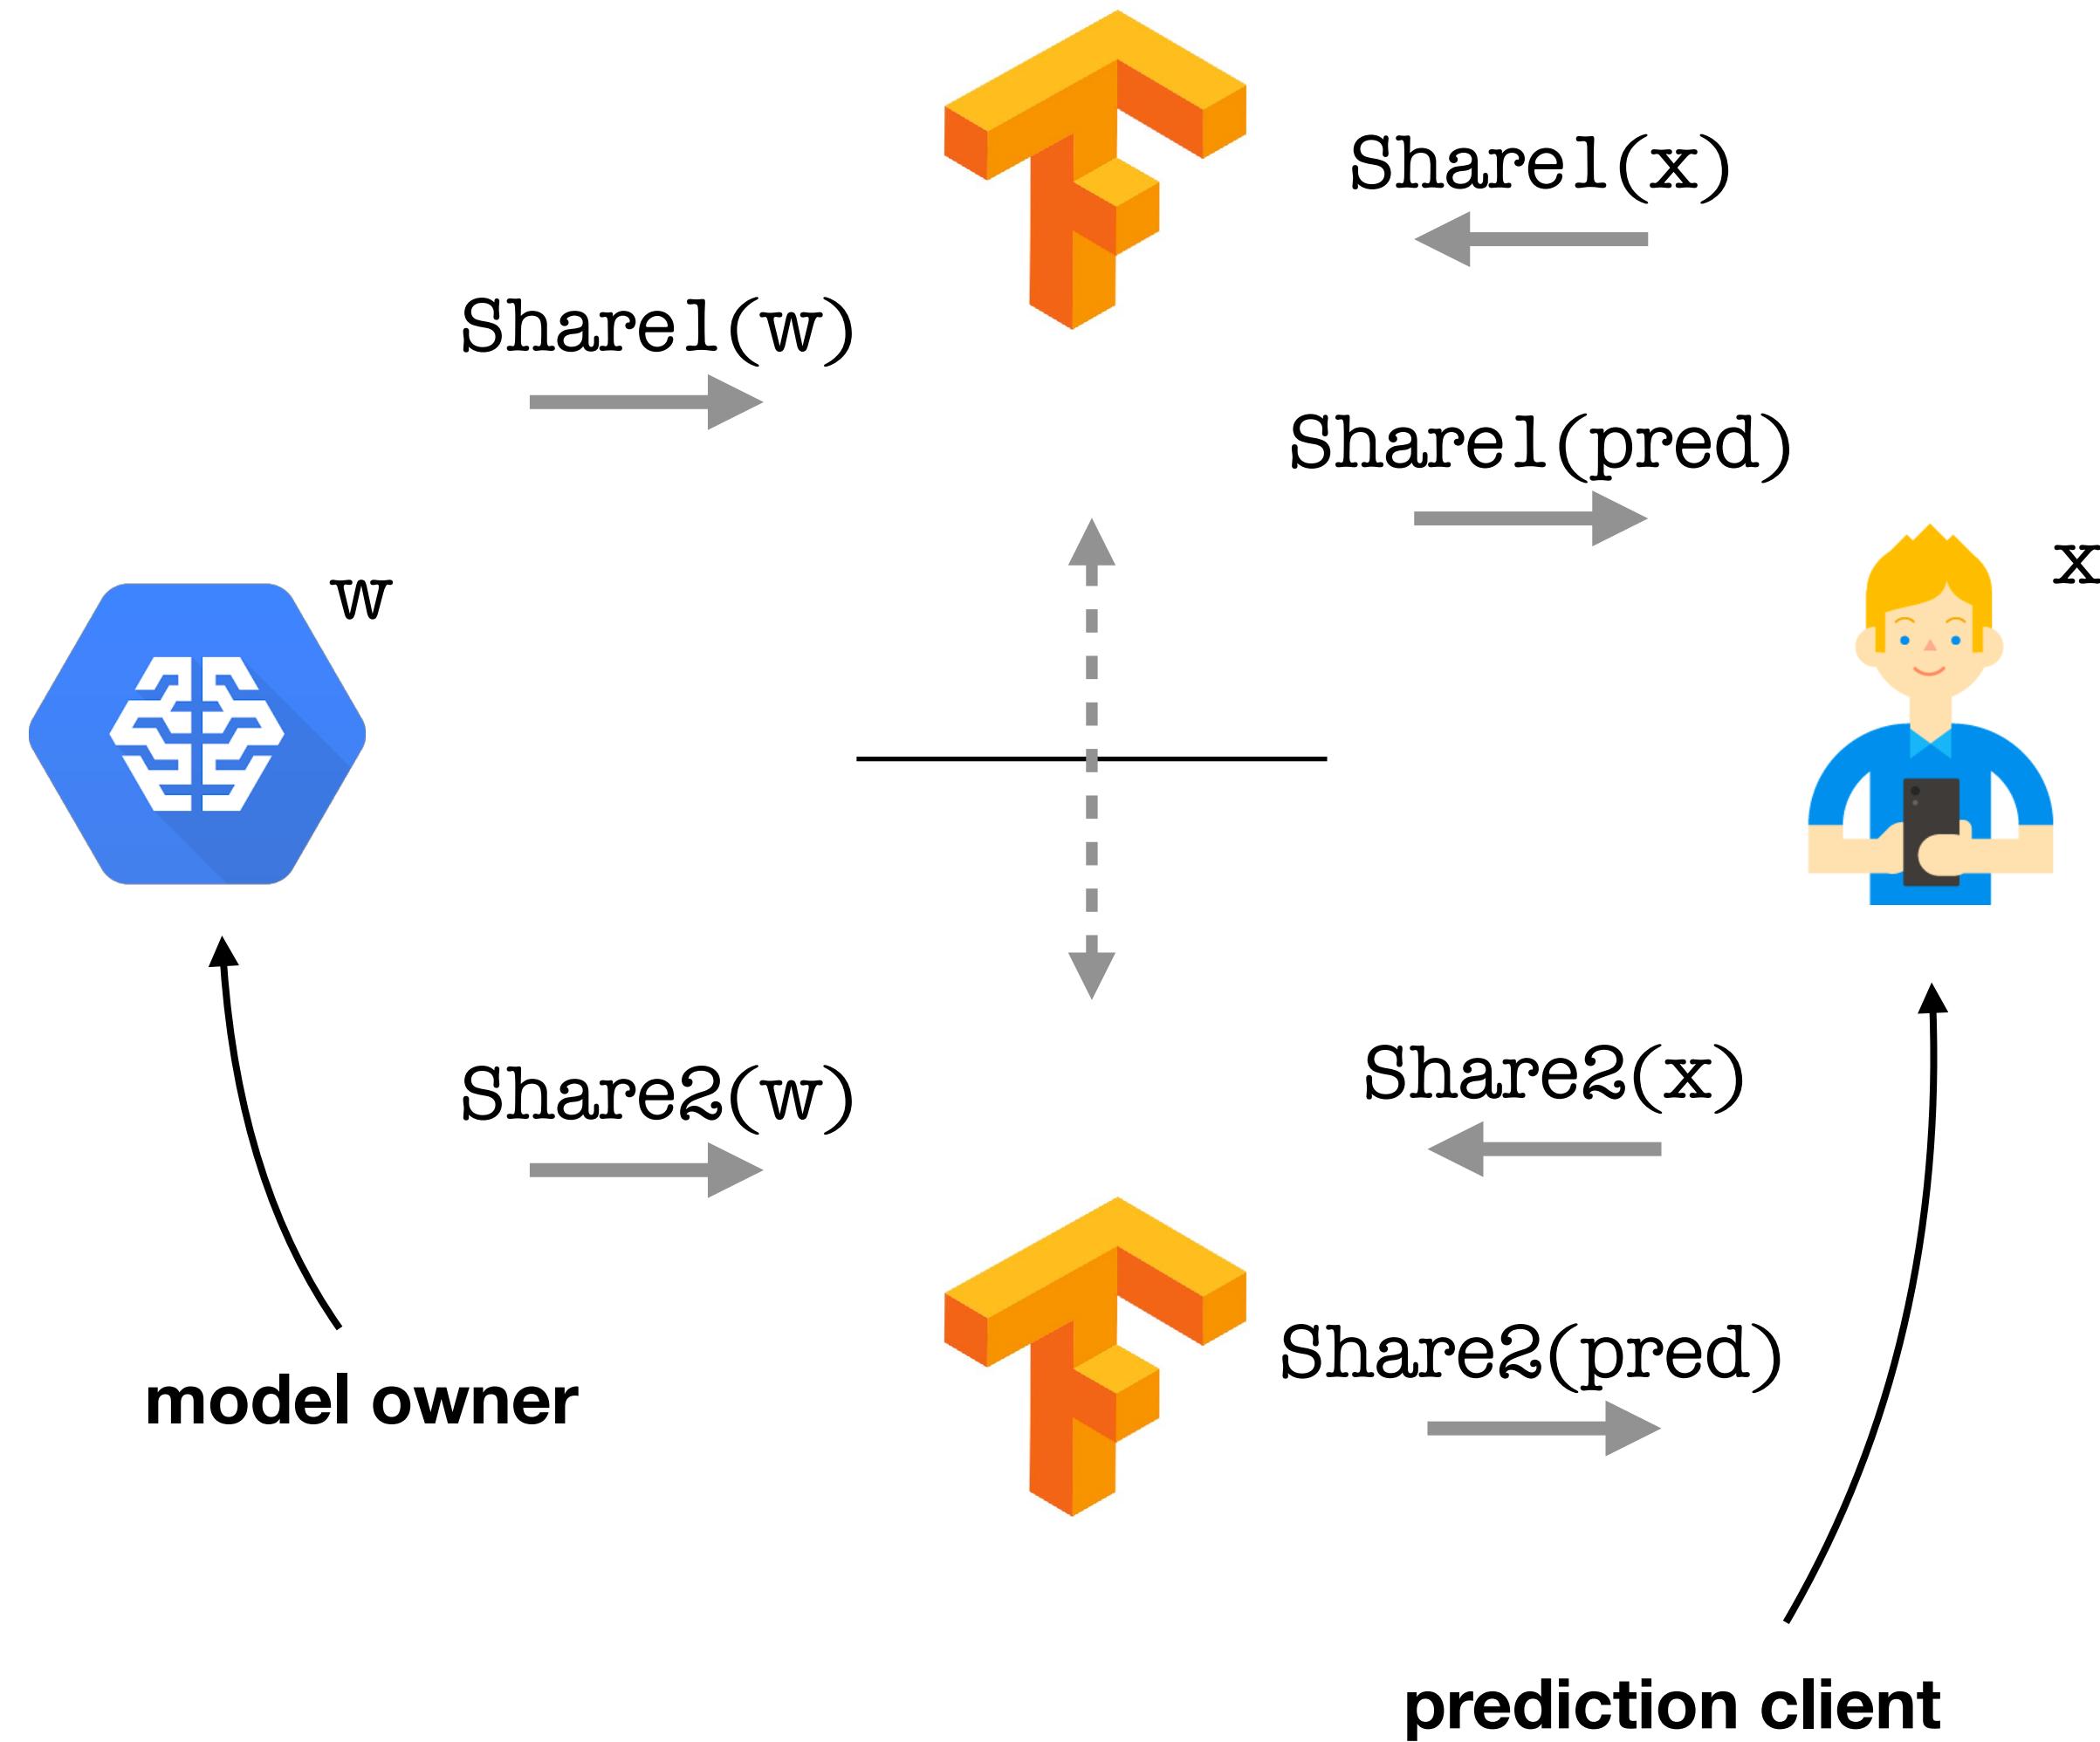


Public Prediction Example (TensorFlow)



```
1 import tensorflow as tf
2
3 def provide_weights(): """ Load from disk """
4 def provide_input(): """ Pre-process """
5 def receive_output(logits): return tf.Print([], [tf.argmax(logits)])
6
7 # get model weights
8 w0, b0, w1, b1, w2, b2 = provide_weights()
9
10 # get prediction input
11 x = provide_input()
12
13 # compute prediction
14 layer0 = tf.nn.relu(tf.matmul(x, w0) + b0)
15 layer1 = tf.nn.relu(tf.matmul(layer0, w1) + b1)
16 logits = tf.matmul(layer2, w2) + b2
17
18 # process result of prediction
19 prediction_op = receive_output(logits)
20
21 # run graph execution in a tf.Session
22 with tf.Session() as sess:
23     sess.run(tf.global_variables_initializer())
24     sess.run(prediction_op)
```

Private Prediction Example (tf-encrypted)



```
1 import tensorflow as tf
2 import tf_encrypted as tfe
3
4 def provide_weights(): """ Load from disk """
5 def provide_input(): """ Pre-process """
6 def receive_output(logits): return tf.Print([], [tf.argmax(logits)])
7
8 # get model weights as private tensors from owner
9 w0, b0, w1, b1, w2, b2 = tfe.define_private_input("model-owner", provide_weights)
10
11 # get prediction input as private tensors from client
12 x = tfe.define_private_input("prediction-client", provide_input)
13
14 # compute private prediction on servers
15 layer0 = tfe.relu((tfe.matmul(x, w0) + b0))
16 layer1 = tfe.relu((tfe.matmul(layer0, w1) + b1))
17 logits = tfe.matmul(layer1, w2) + b2
18
19 # process result of prediction on client
20 prediction_op = tfe.define_output("prediction-client", logits, receive_output)
21
22 # run secure graph execution in a tf.Session
23 with tfe.Session() as sess:
24     sess.run(tf.global_variables_initializer())
25     sess.run(prediction_op)
```

Wrap-Up

Privacy-preserving technologies are enablers for accessing sensitive information

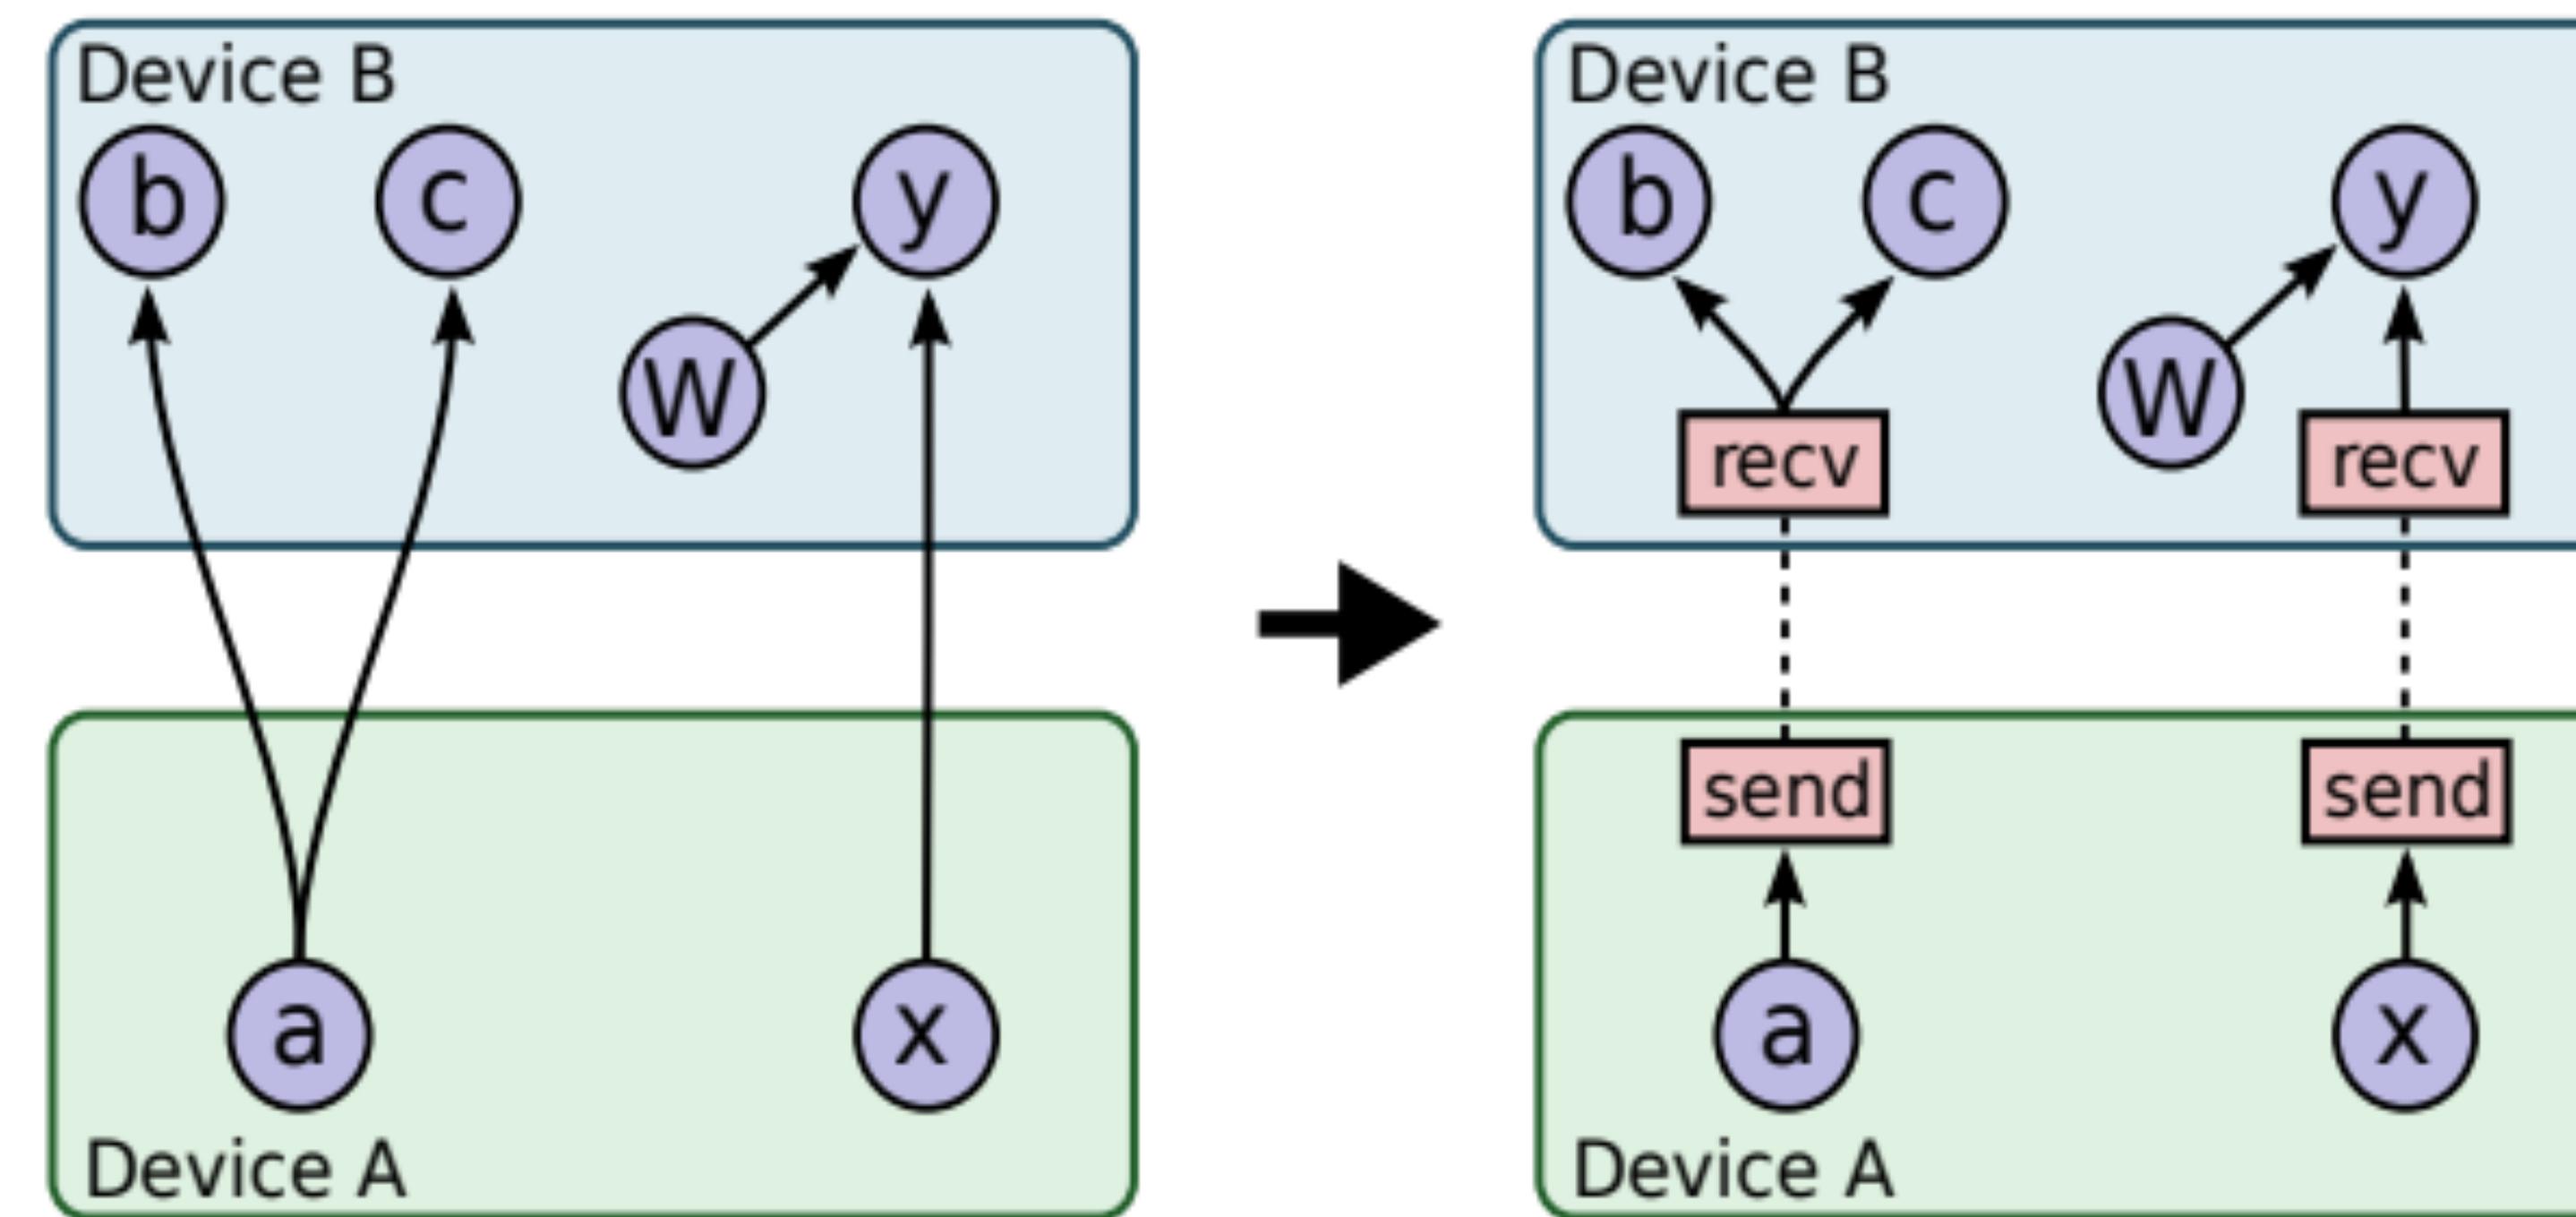
Secure computation is complementary, not a replacement

Multidisciplinary field benefitting from adaptations on both sides

Bonus:
some interest from safety (governance)
and fairness (validation)

Thank you!

Distributed TensorFlow



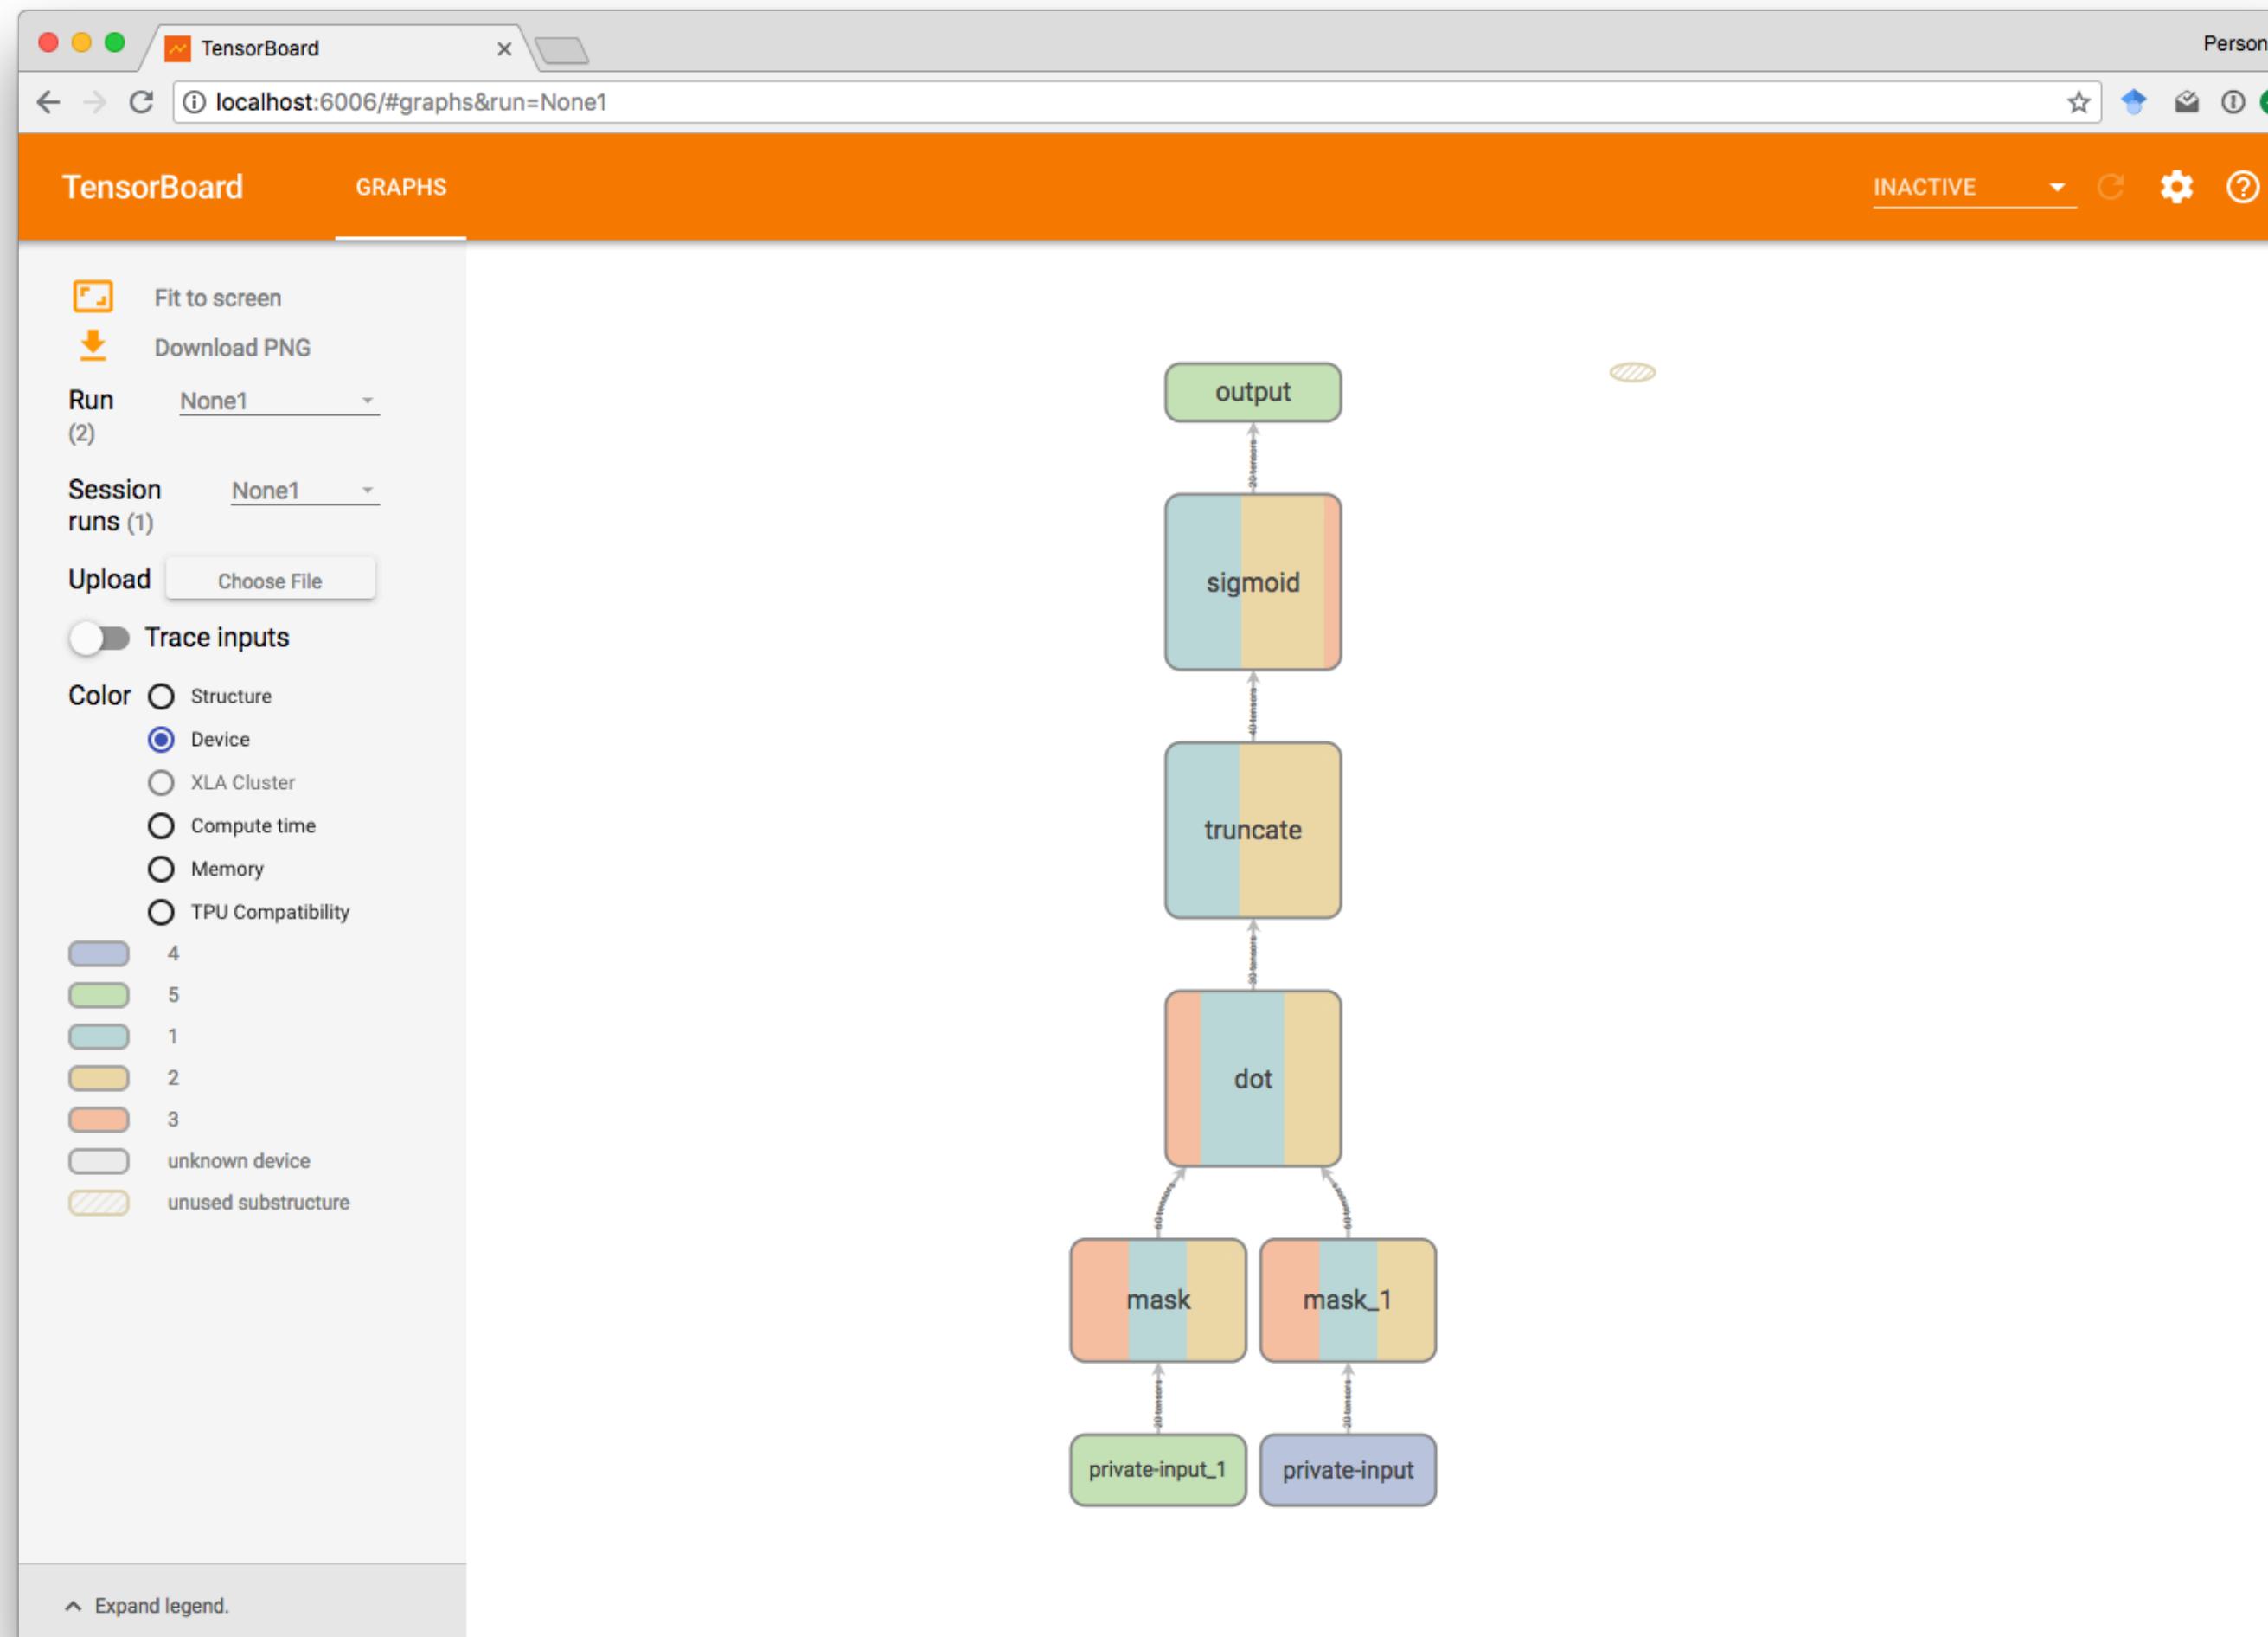
Protocols in TensorFlow

```
1  def _matmul_masked_prot(prot, x, y):
2      a, a0, a1, alpha_on_0, alpha_on_1 = x.unwrapped
3      b, b0, b1, beta_on_0, beta_on_1 = y.unwrapped
4
5      with tf.name_scope('matmul'):
6
7          with tf.device(prot.crypto_producer.device_name):
8              ab = a.matmul(b)
9              ab0, ab1 = prot._share(ab)
10
11         with tf.device(prot.server_0.device_name):
12             z0 = ab0 + a0.matmul(beta) + alpha.matmul(b0) + alpha.matmul(beta)
13
14         with tf.device(prot.server_1.device_name):
15             z1 = ab1 + a1.matmul(beta) + alpha.matmul(b1)
16
17     return PondPrivateTensor(prot, z0, z1)
```

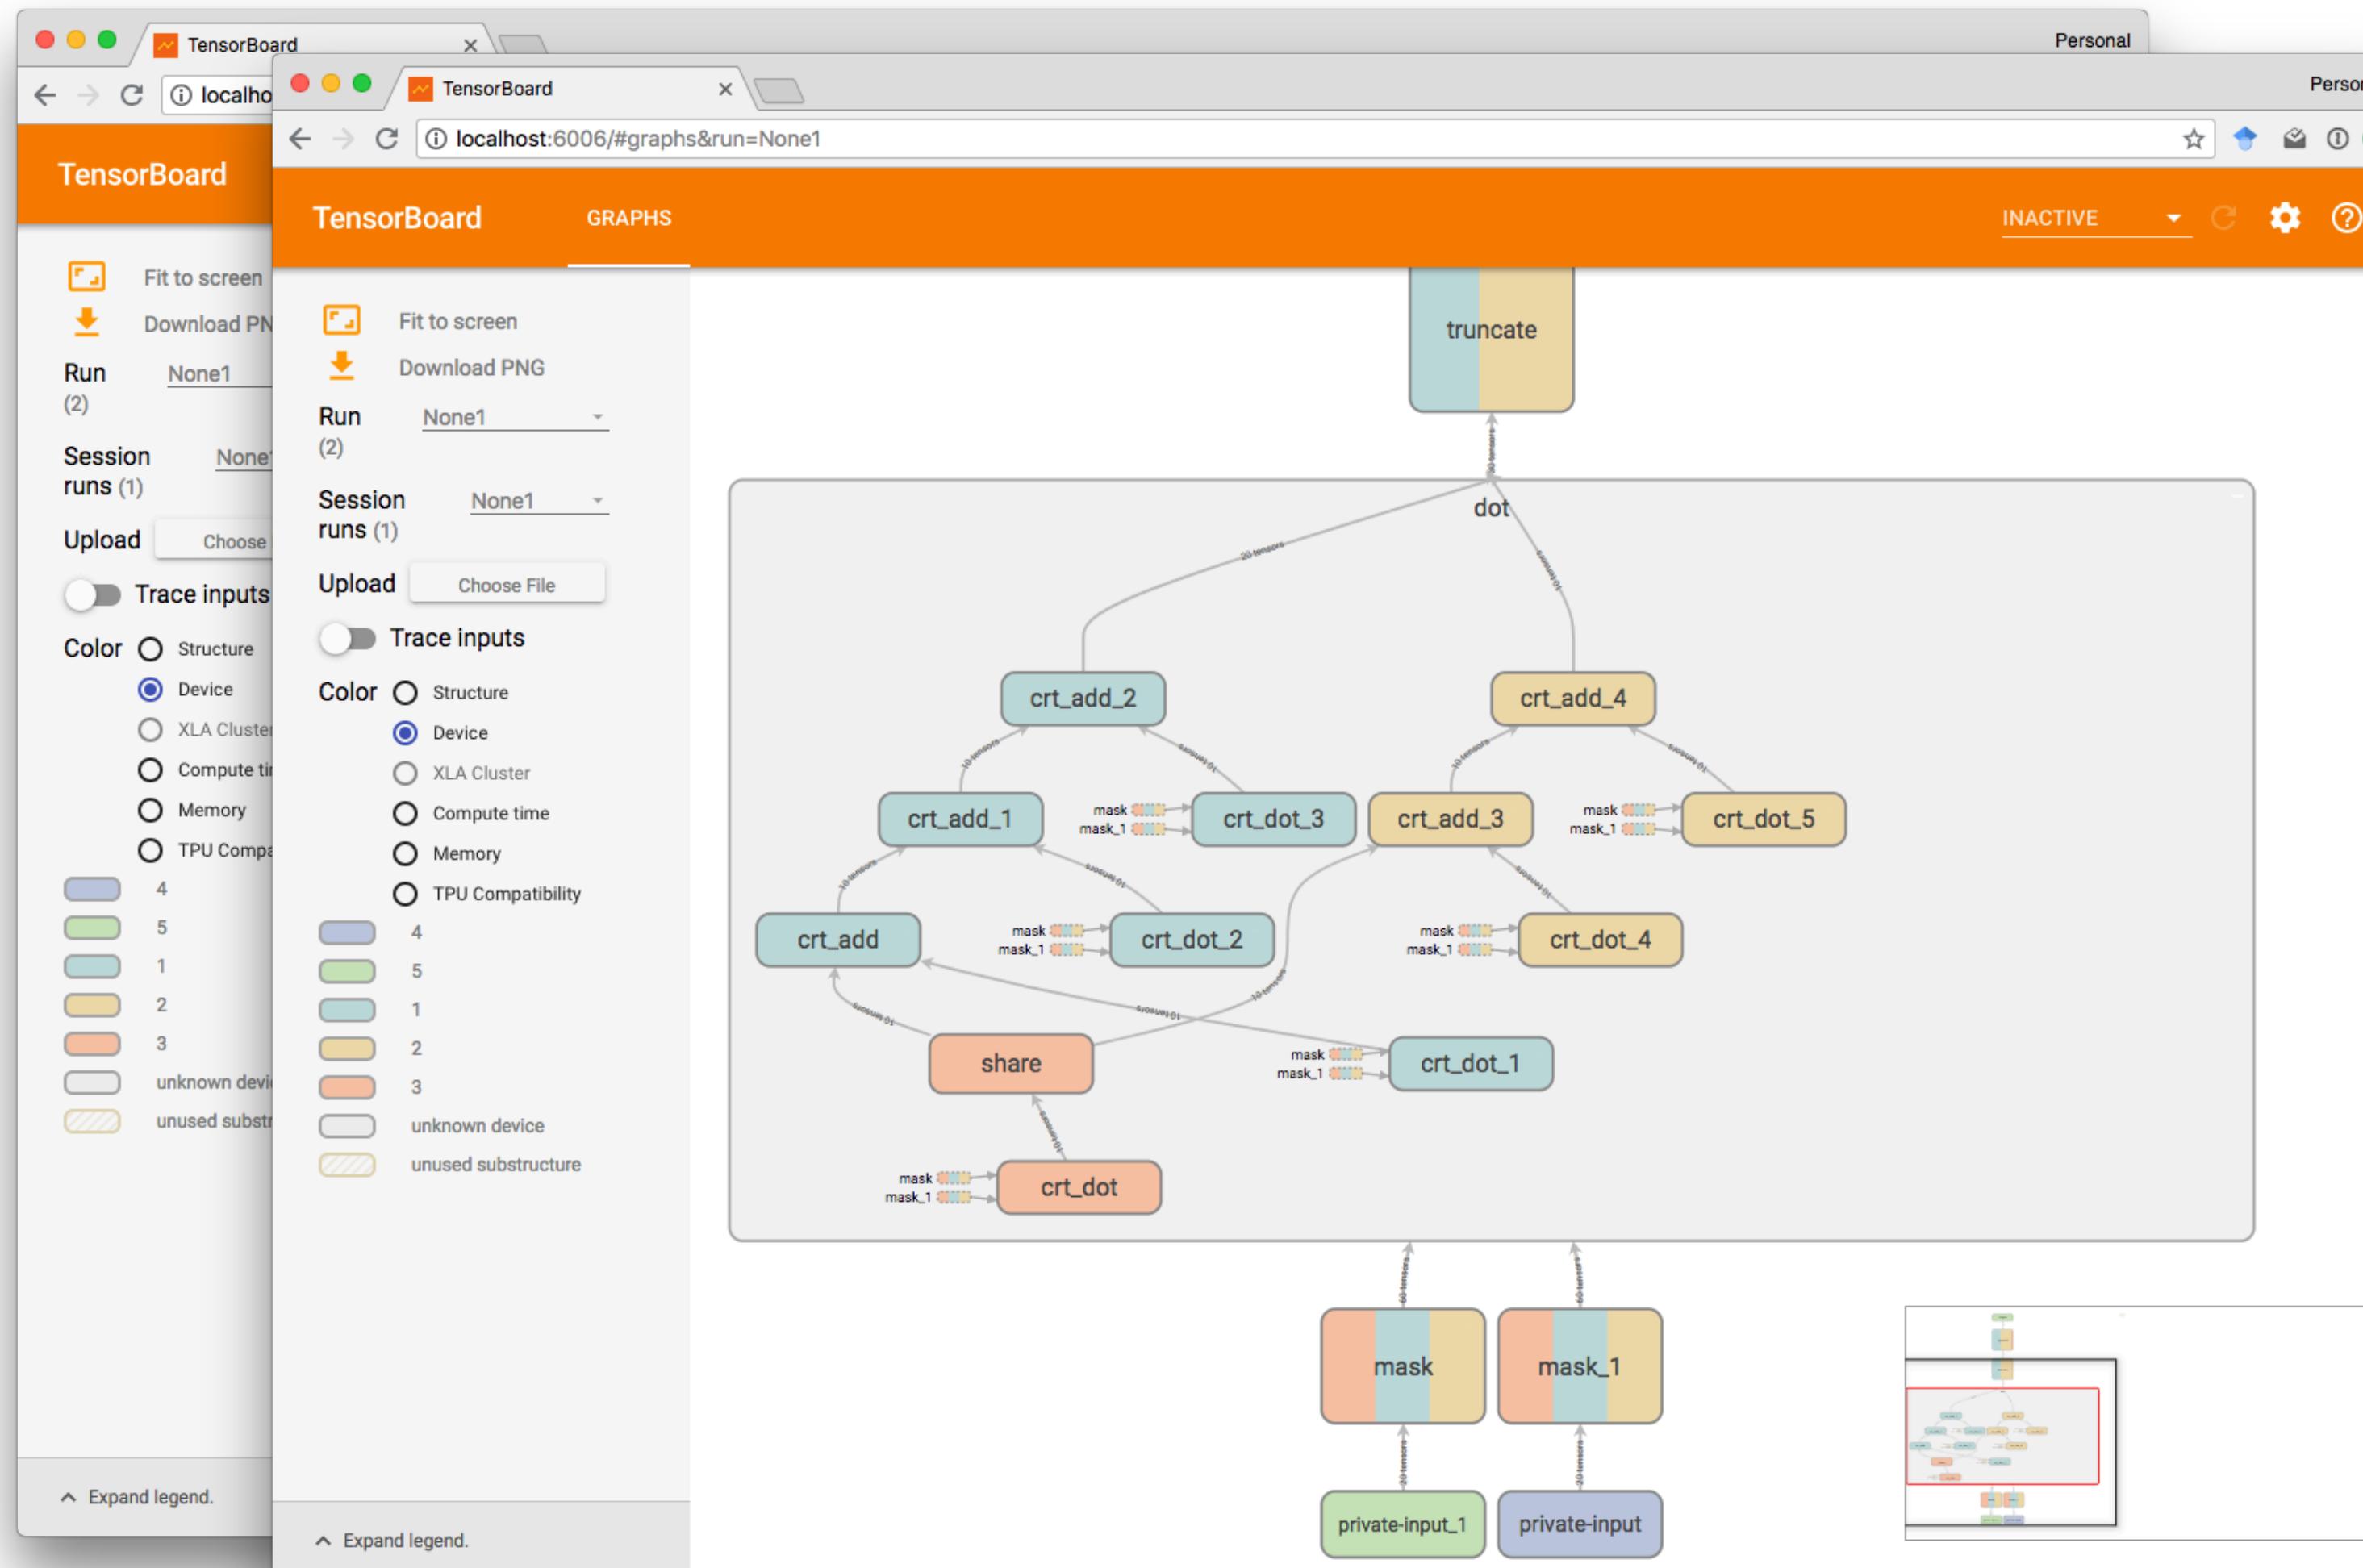
Protocols in TensorFlow

```
1  def _matmul_masked_prot(prot, x, y):
2      a, a0, a1, alpha_on_0, alpha_on_1 = x.unwrapped
3      b, b0, b1, beta_on_0, beta_on_1 = y.unwrapped
4
5      with tf.name_scope('matmul'):
6
7          with tf.device(prot.crypto_producer.device_name):
8              ab = a.matmul(b)
9              ab0, ab1 = prot._share(ab)
10
11         with tf.device(prot.server_0.device_name):
12             z0 = ab0 + a0.matmul(beta) + alpha.matmul(b0) + alpha.matmul(beta)
13
14         with tf.device(prot.server_1.device_name):
15             z1 = ab1 + a1.matmul(beta) + alpha.matmul(b1)
16
17     return PondPrivateTensor(prot, z0, z1)
```

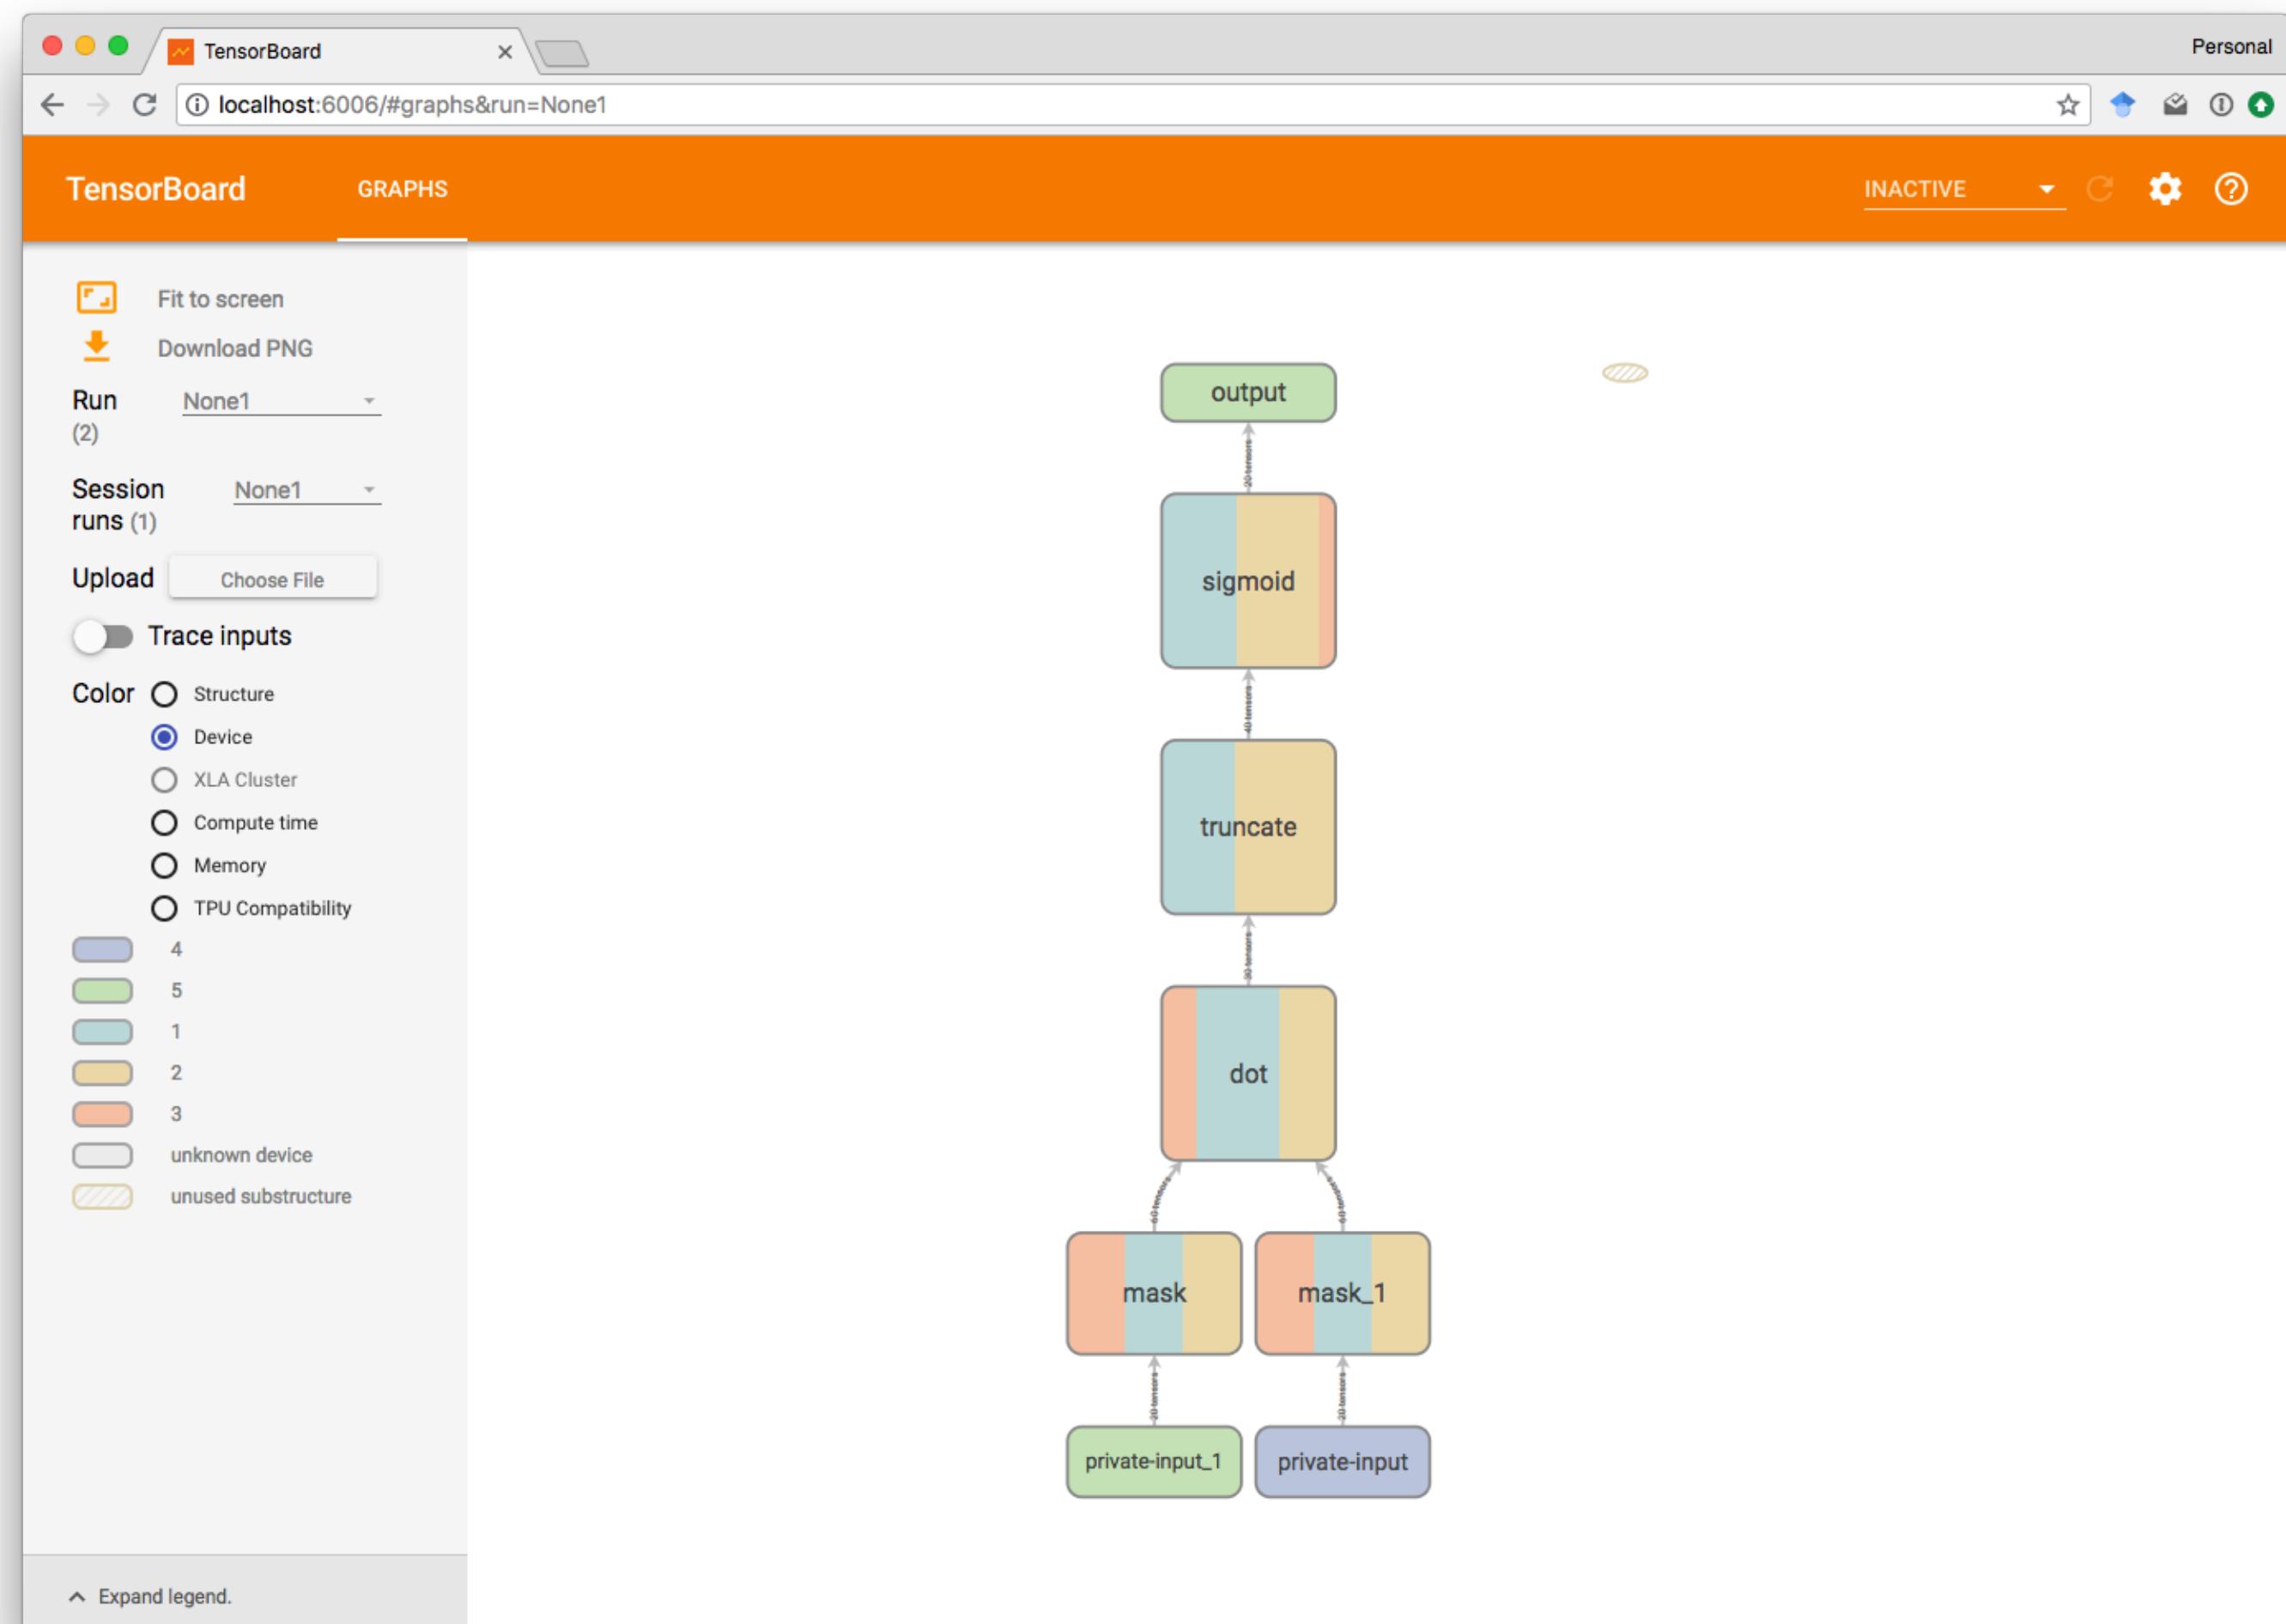
Compilation to Graphs



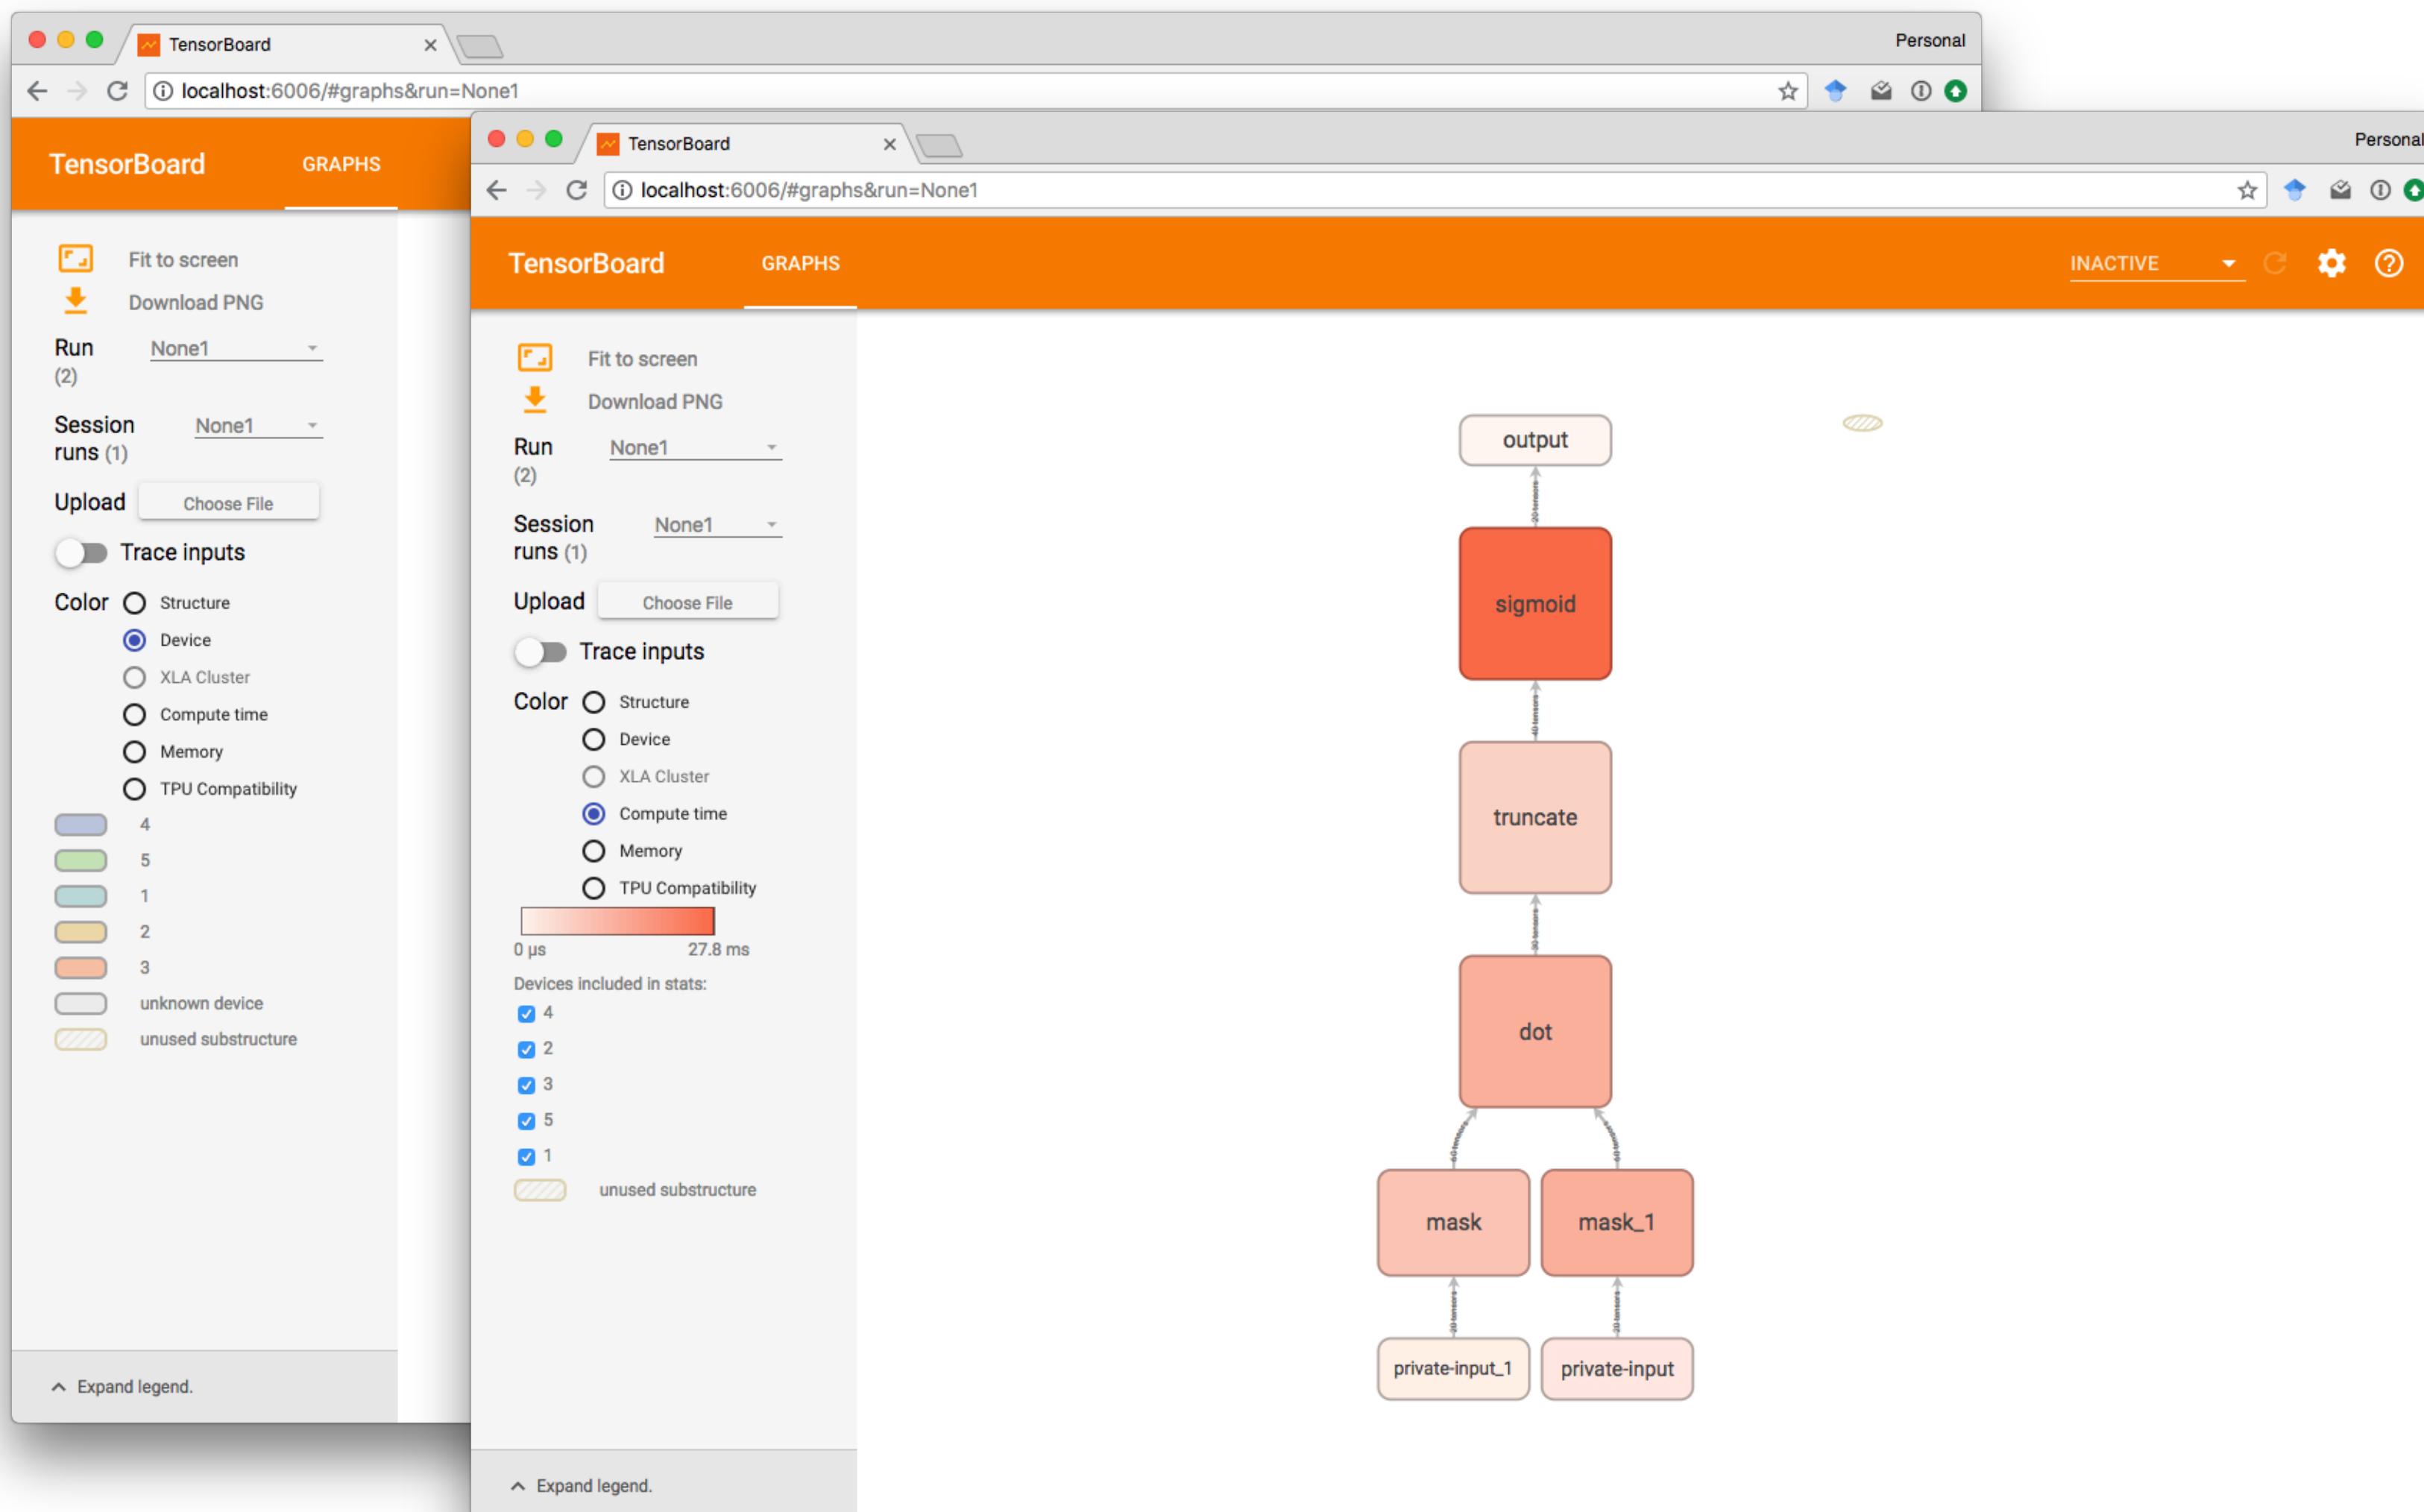
Compilation to Graphs



Profiling

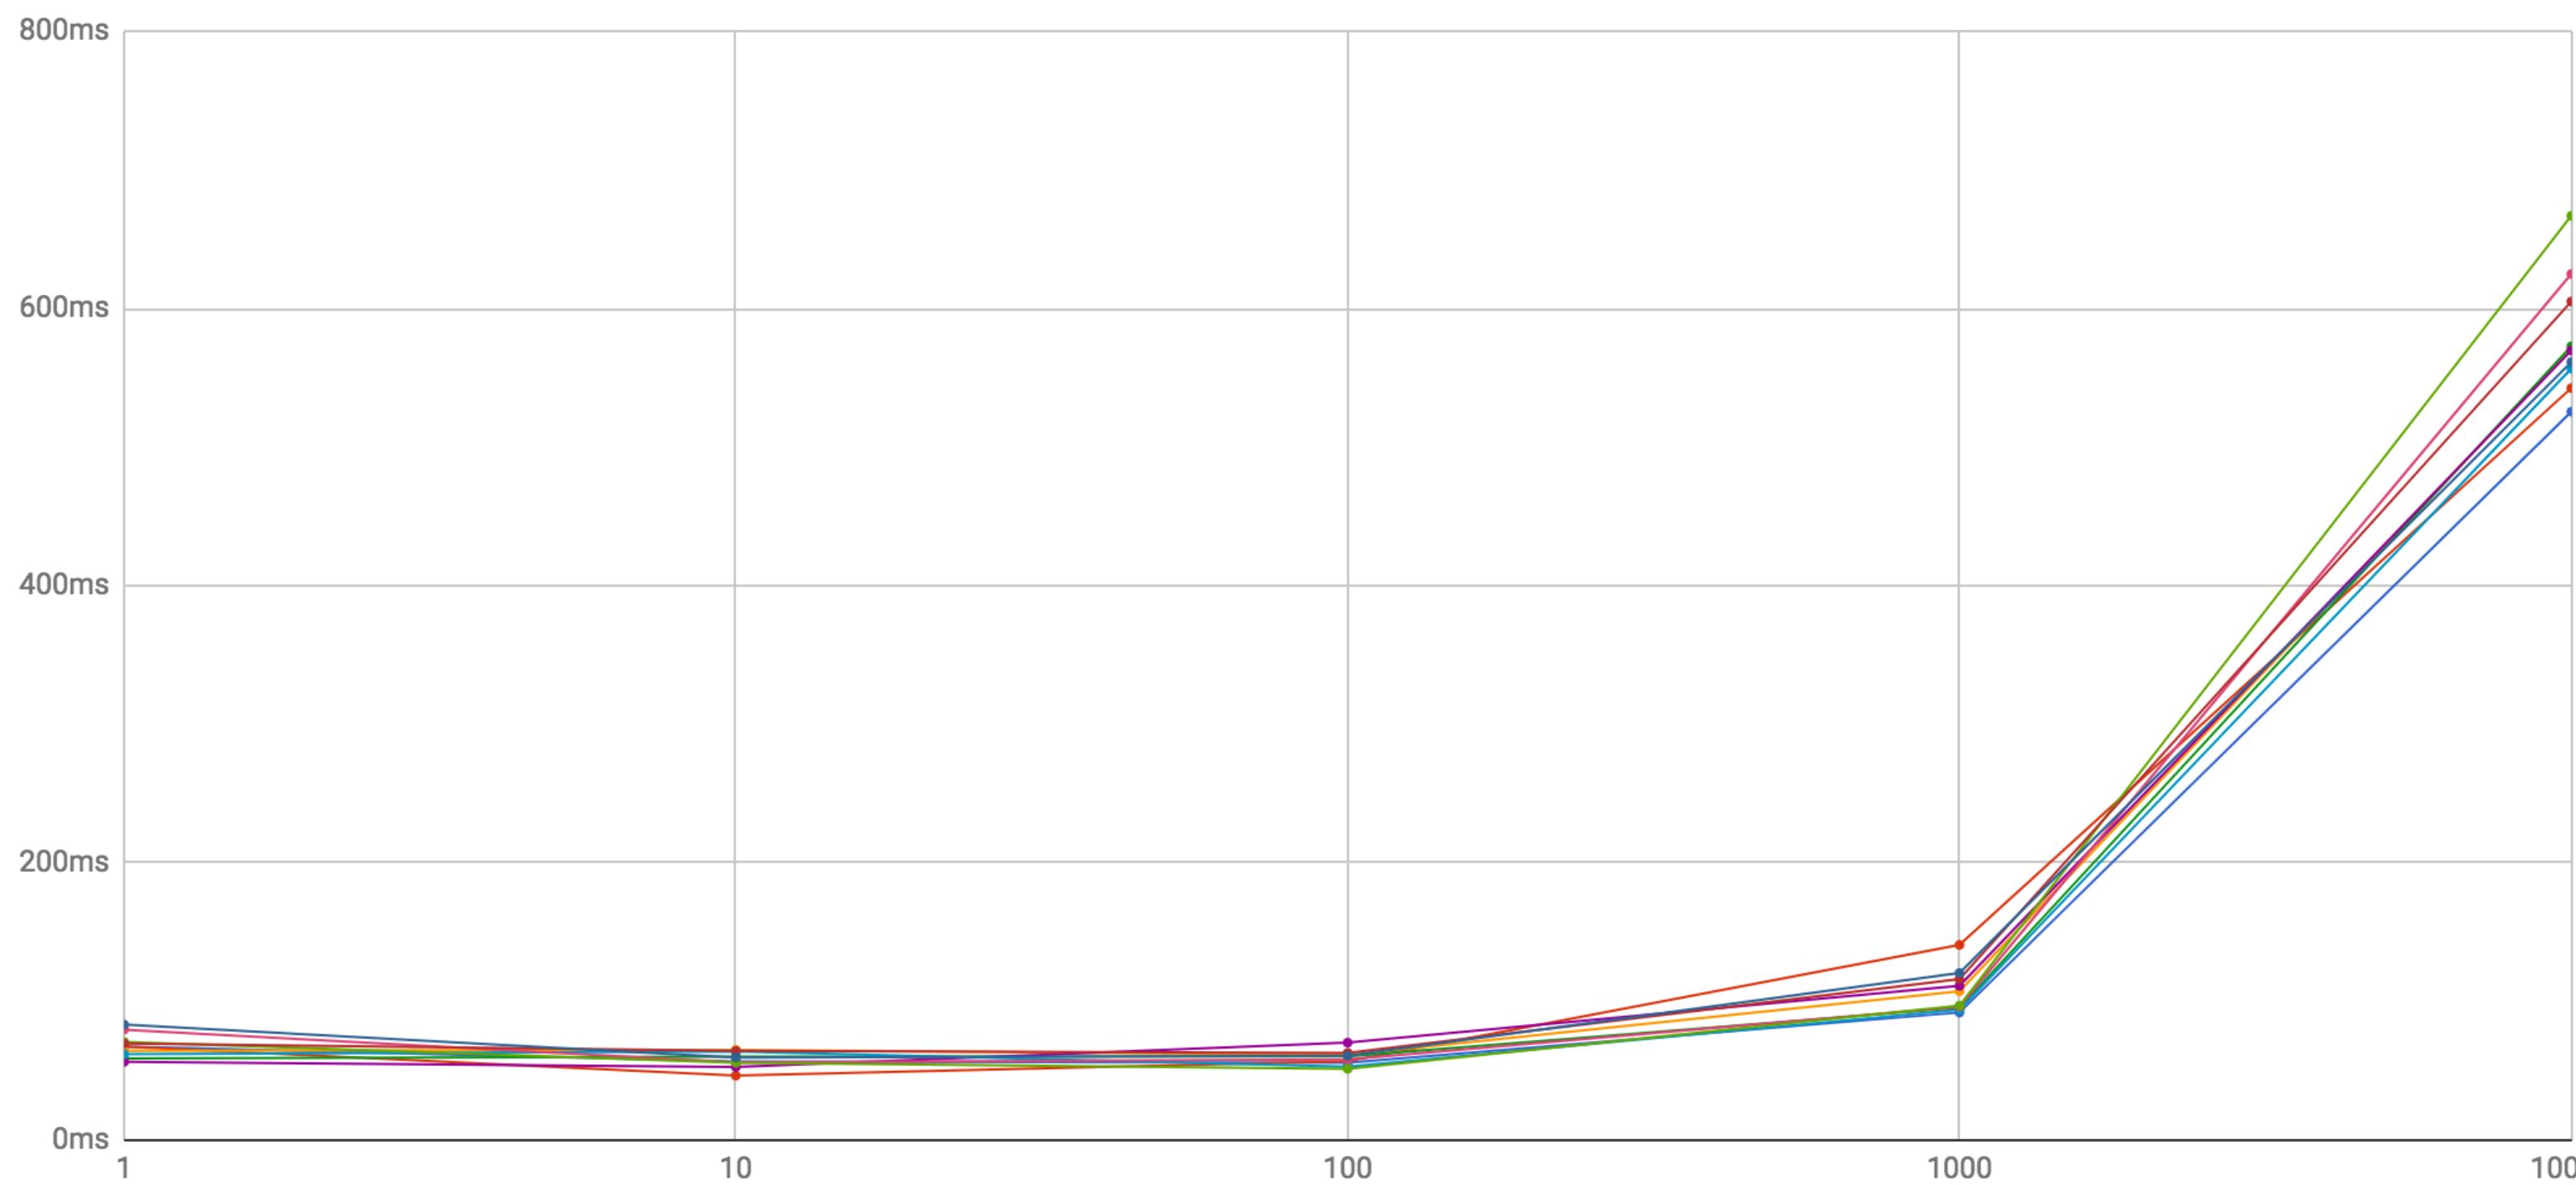


Profiling



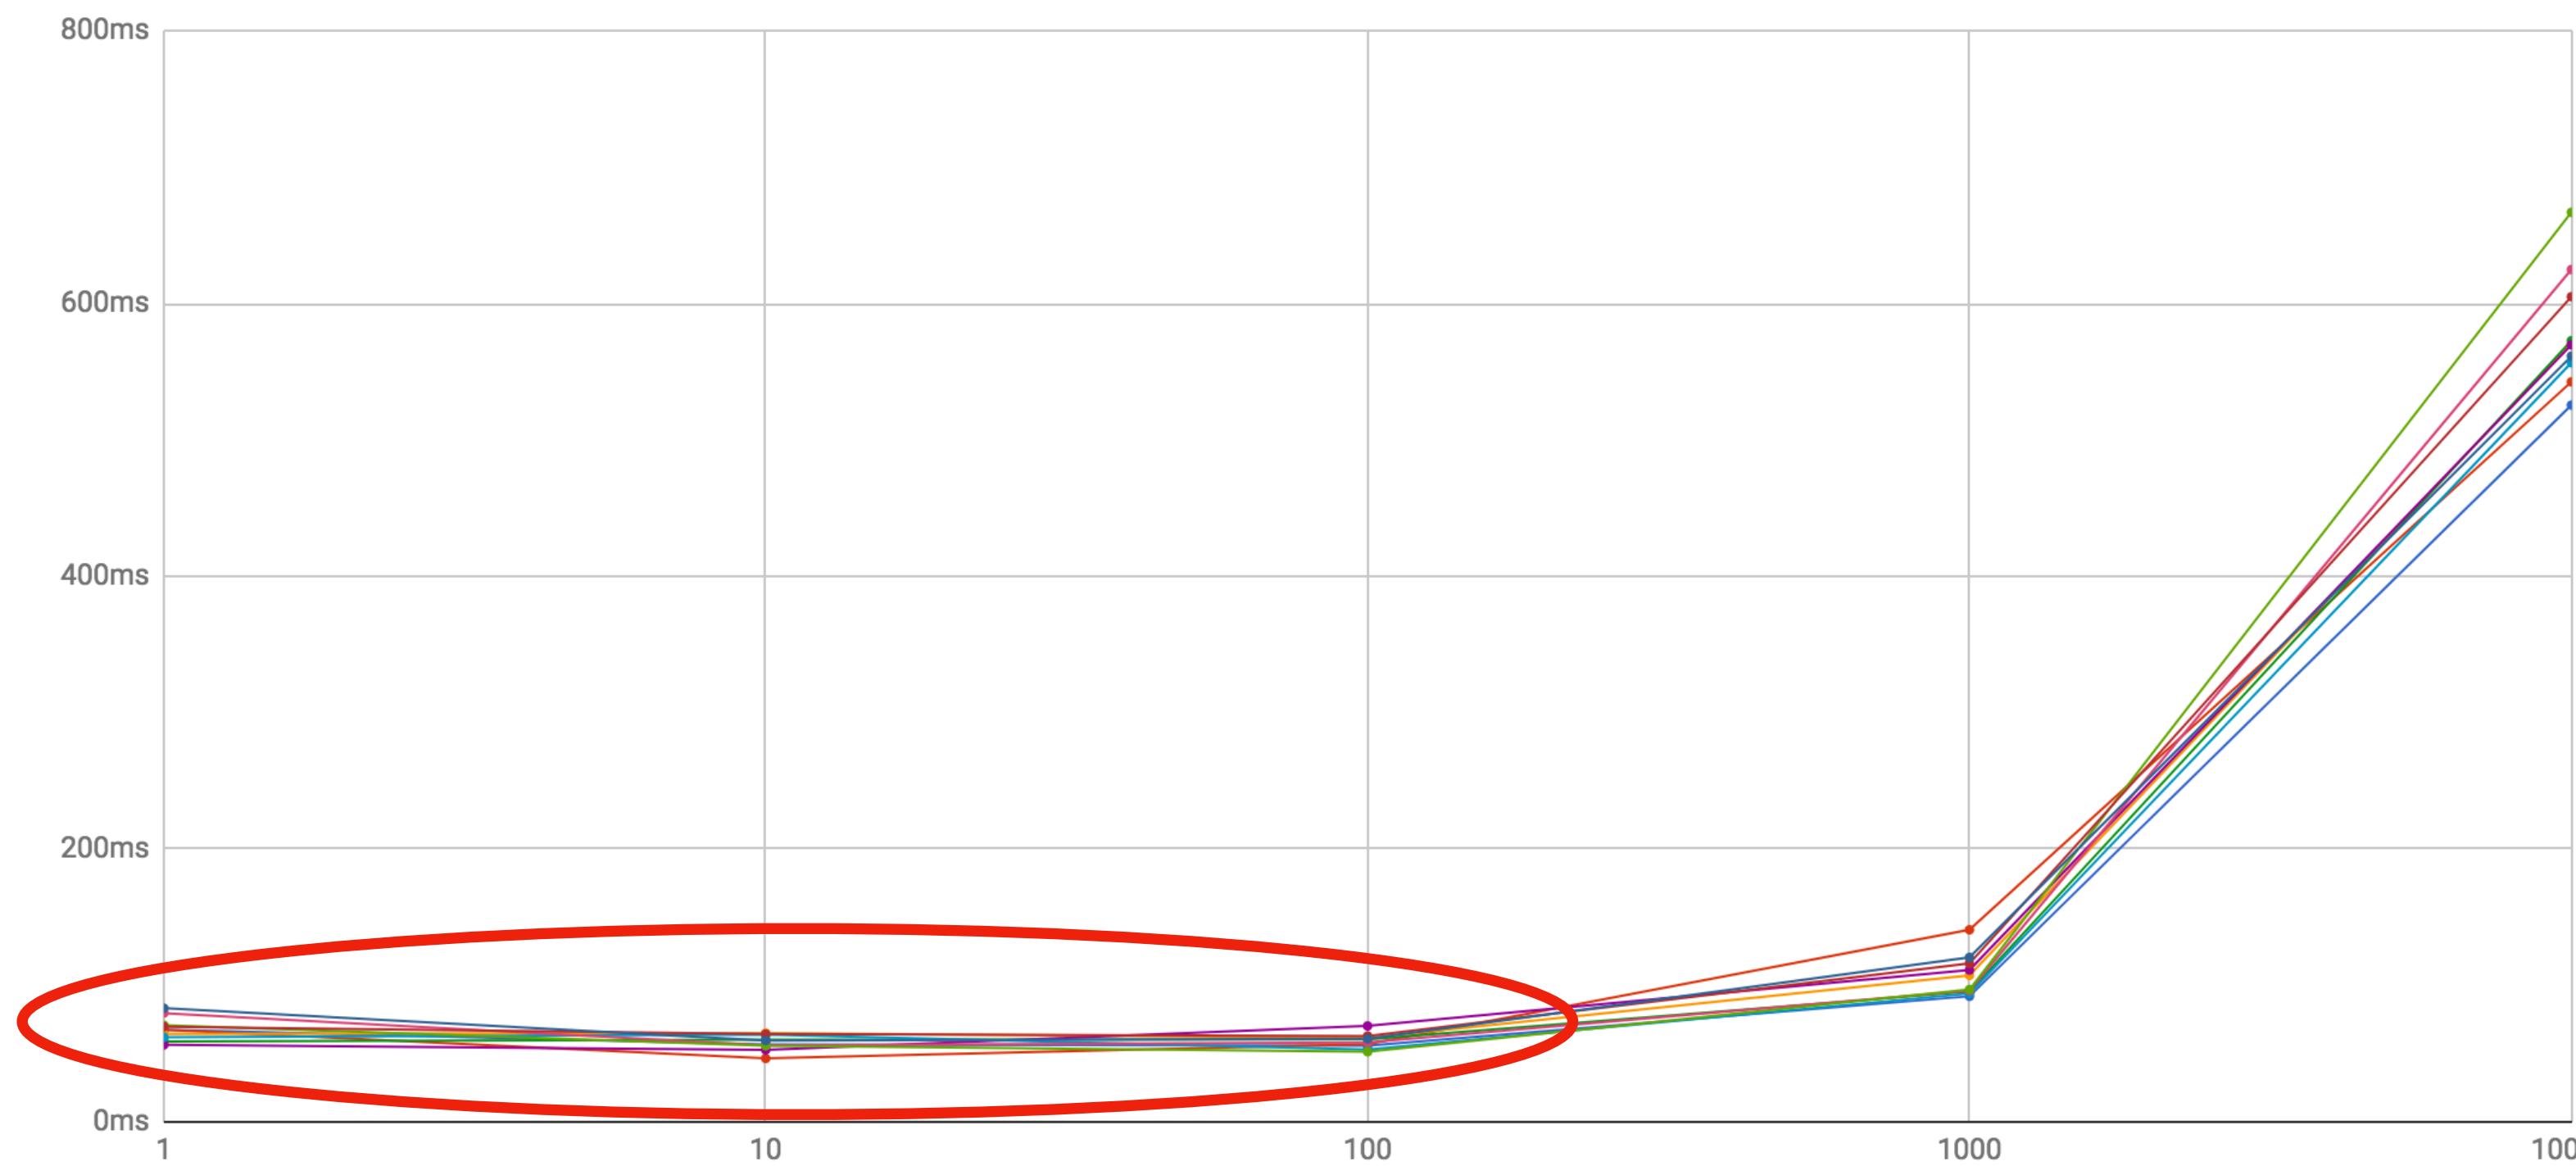
Performance

100 features, servers on Google cloud (2 vCPU, 10 GB)



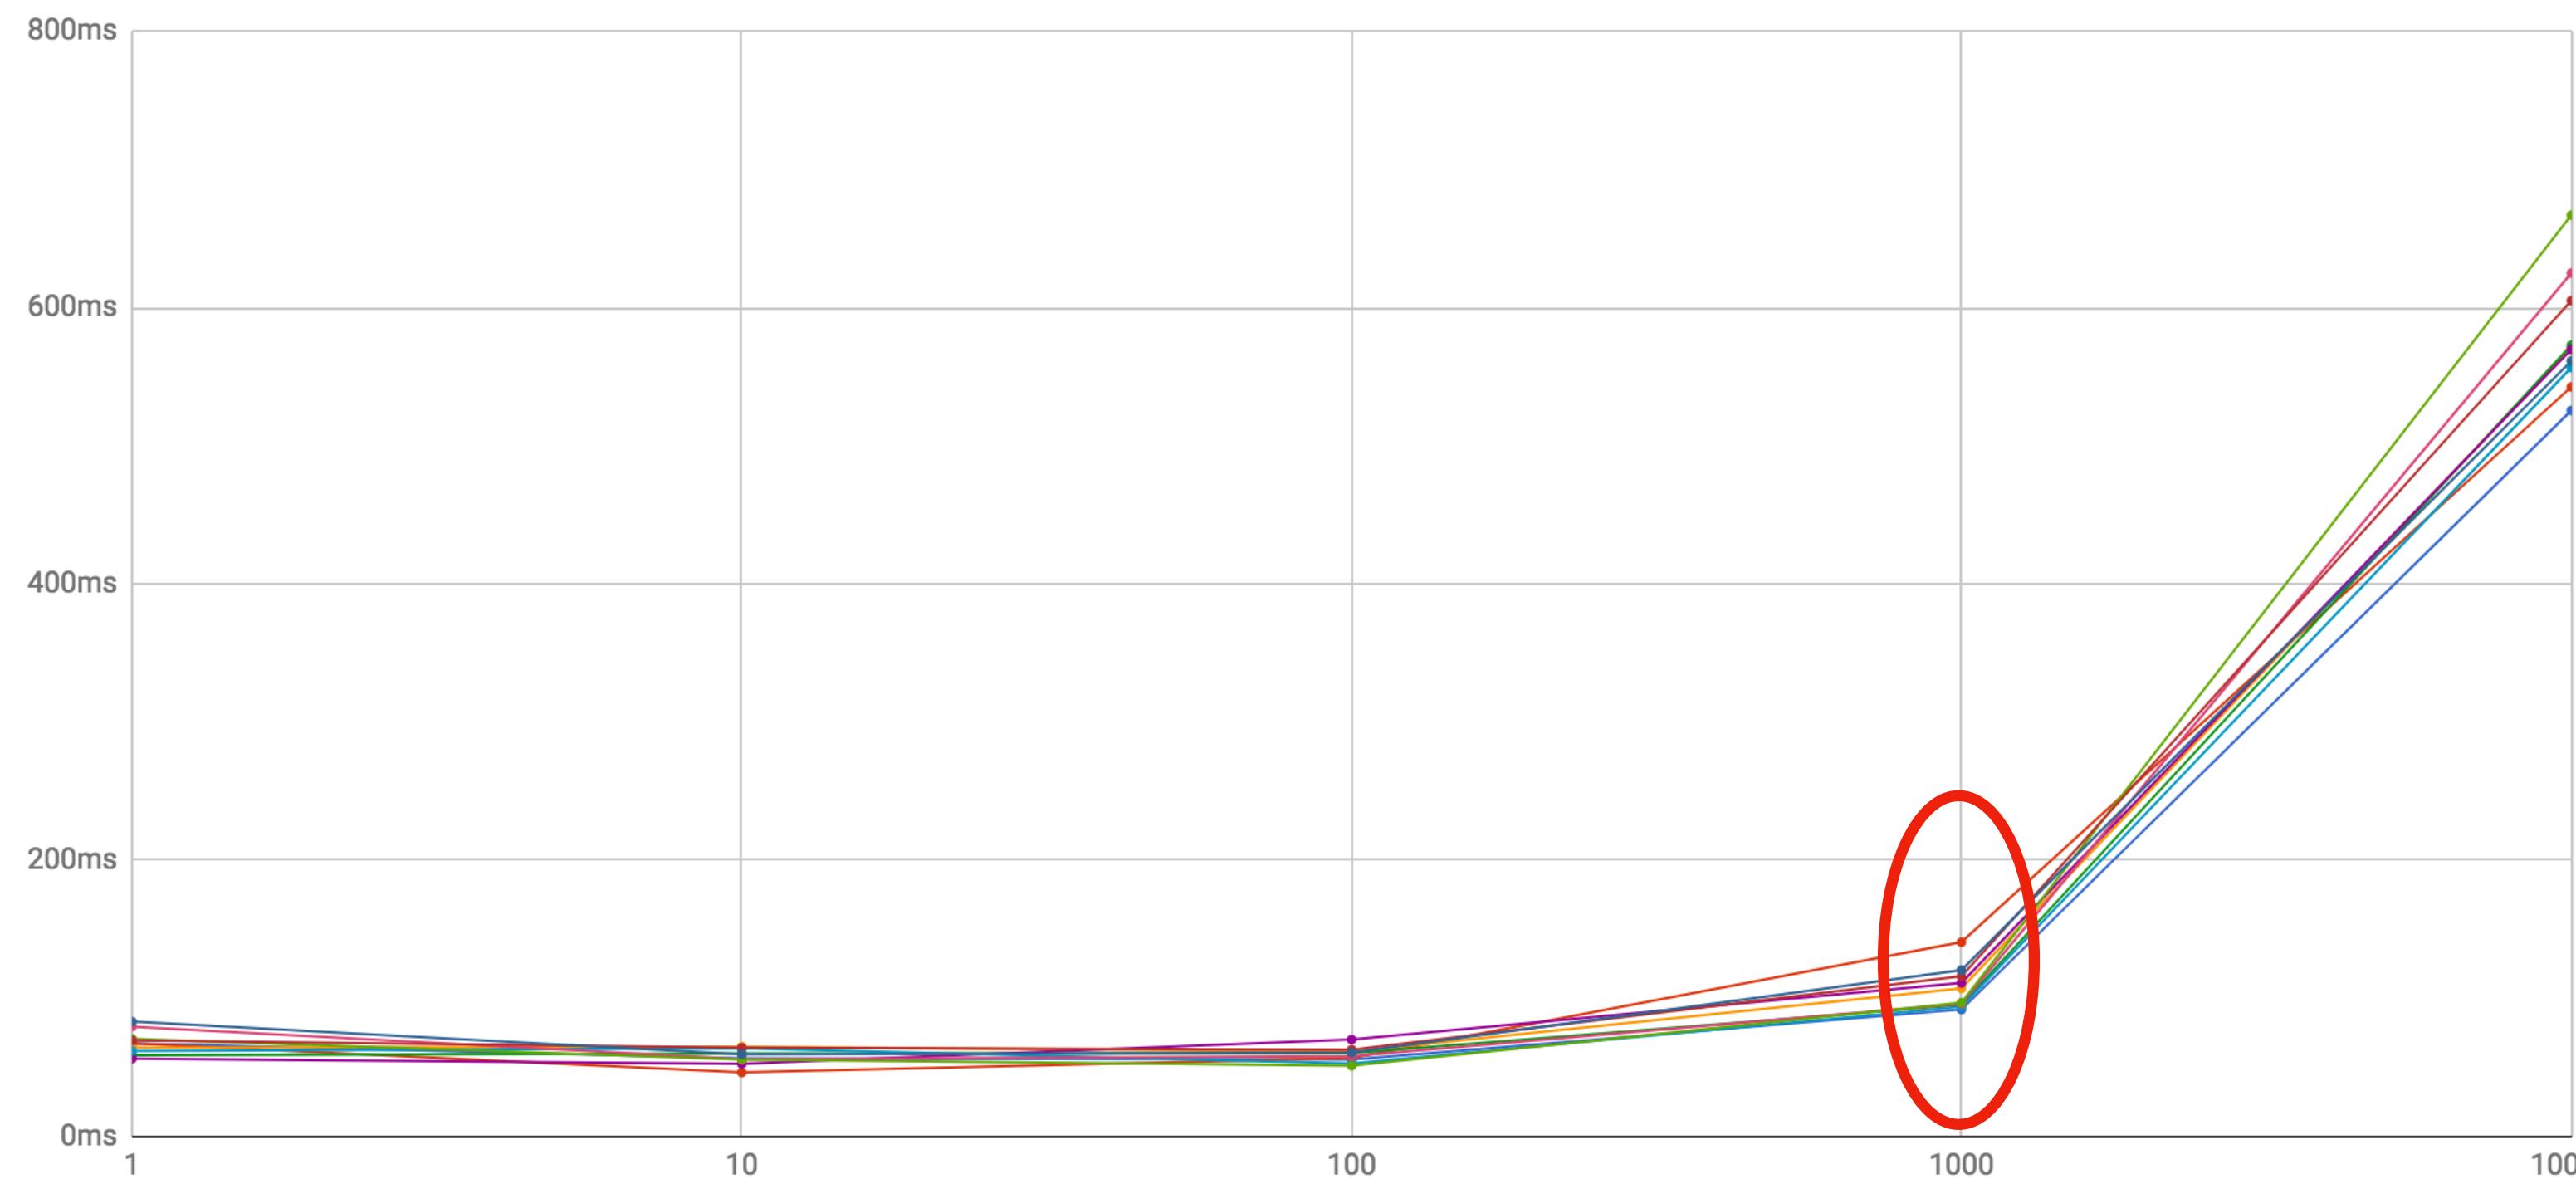
Performance

100 features, servers on Google cloud (2 vCPU, 10 GB)



Performance

100 features, servers on Google cloud (2 vCPU, 10 GB)



Performance

100 features, servers on Google cloud (2 vCPU, 10 GB)

