

Privacy-Preserving Machine Learning in TensorFlow with TF Encrypted

Morten Dahl

DropoutLabs

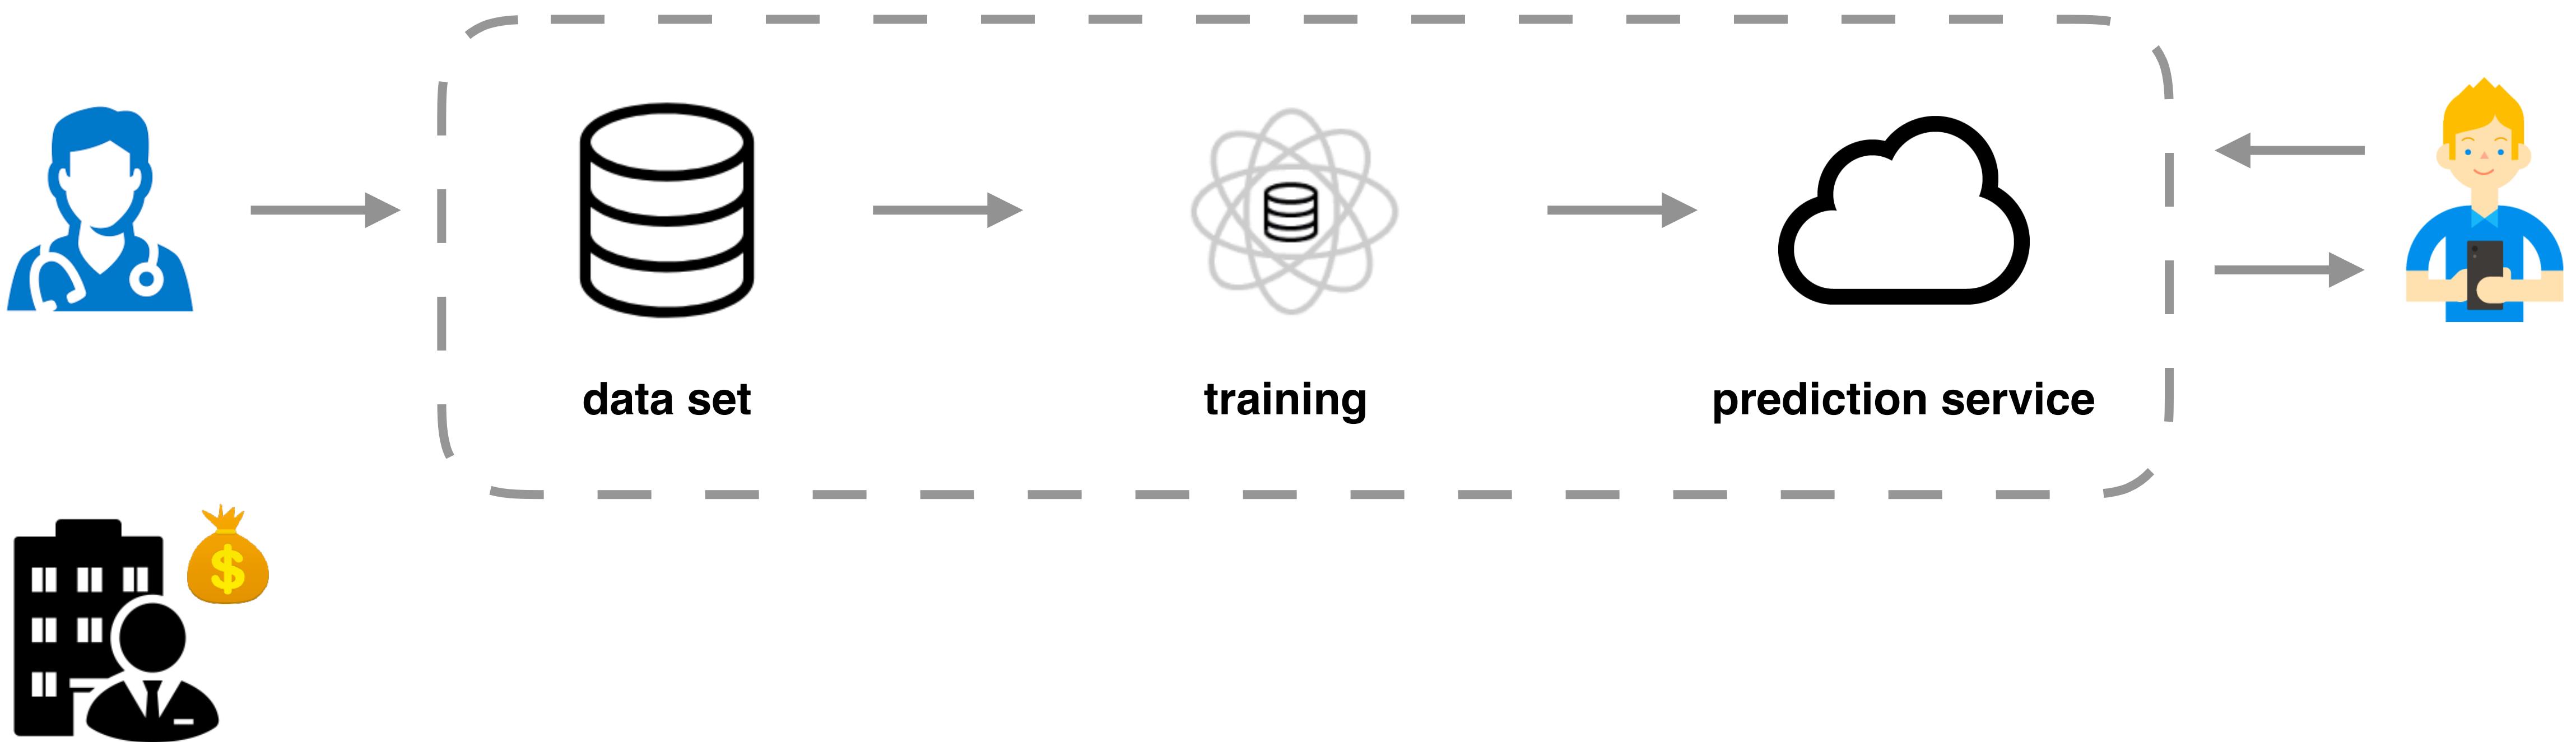
O'Reilly AI Conference, New York, April 2019

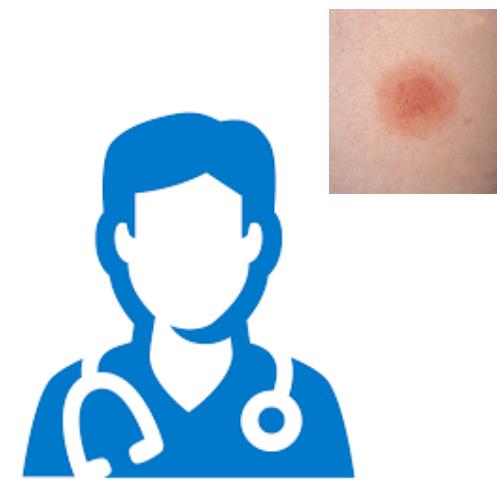
Why?

Privacy in machine learning

Machine Learning Process

IMAGENET



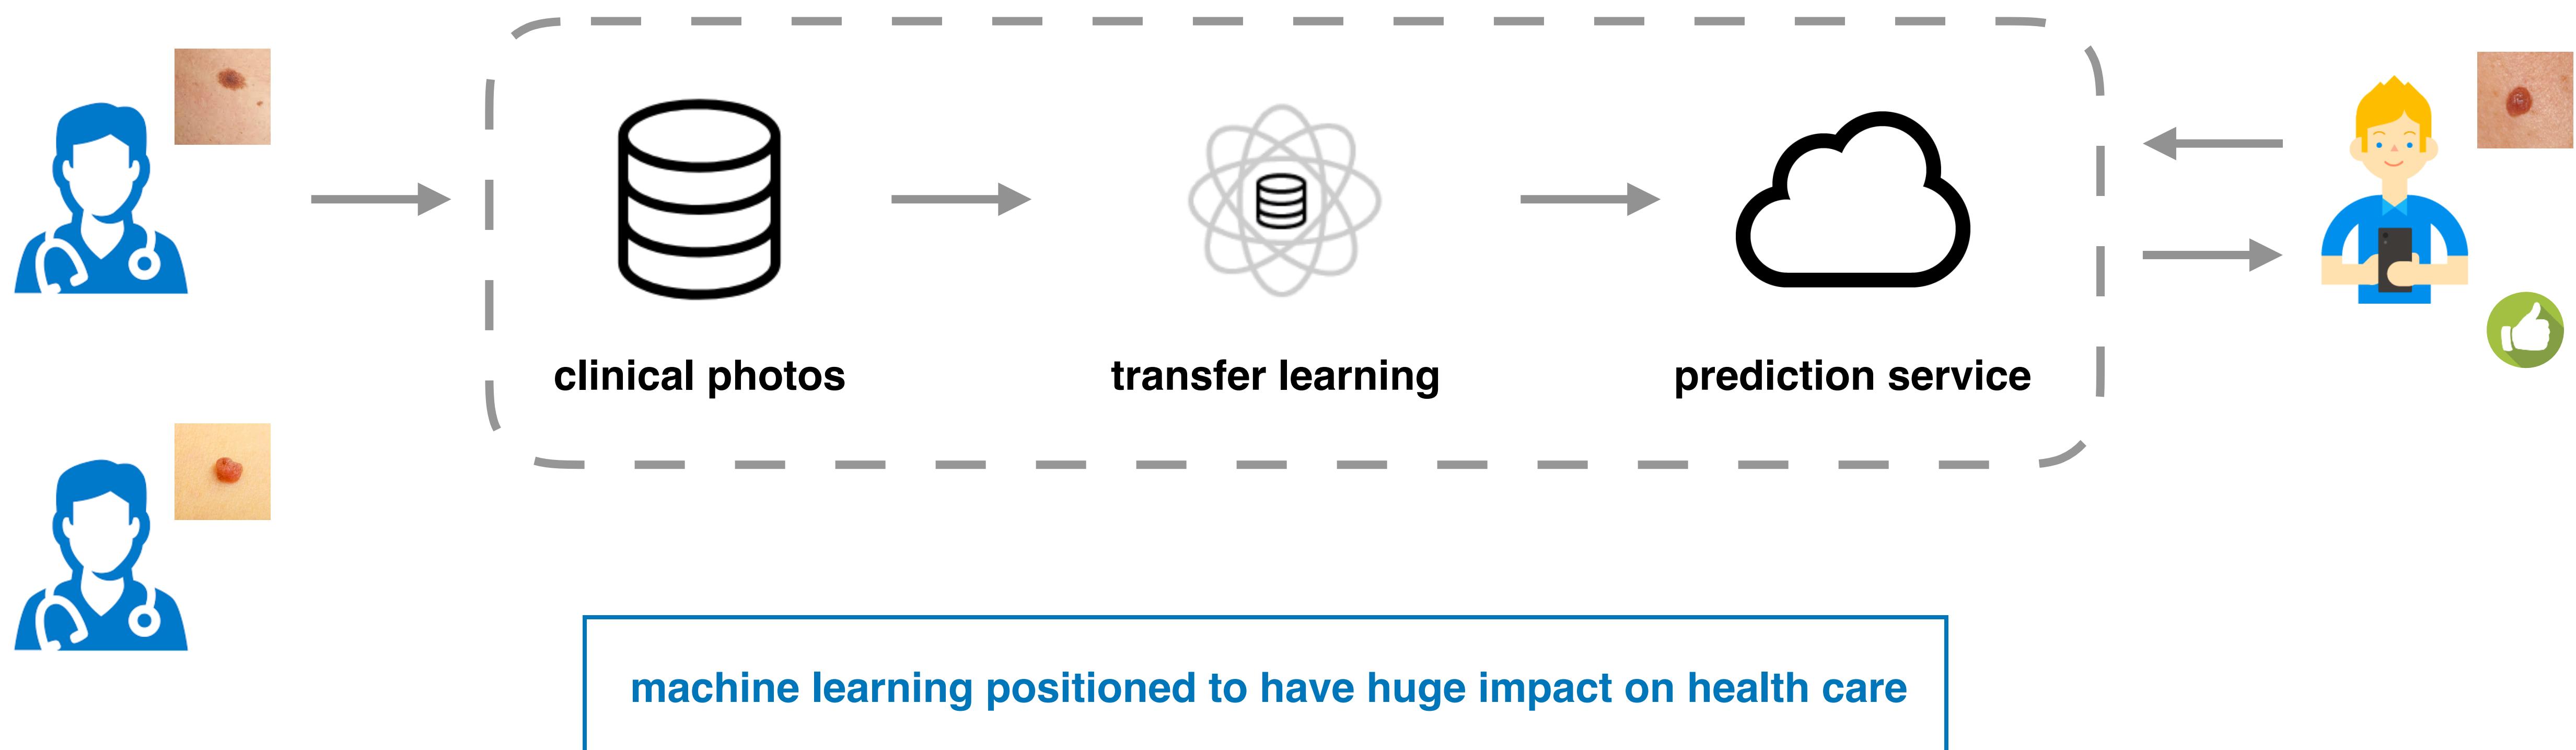


Skin Cancer Image Classification

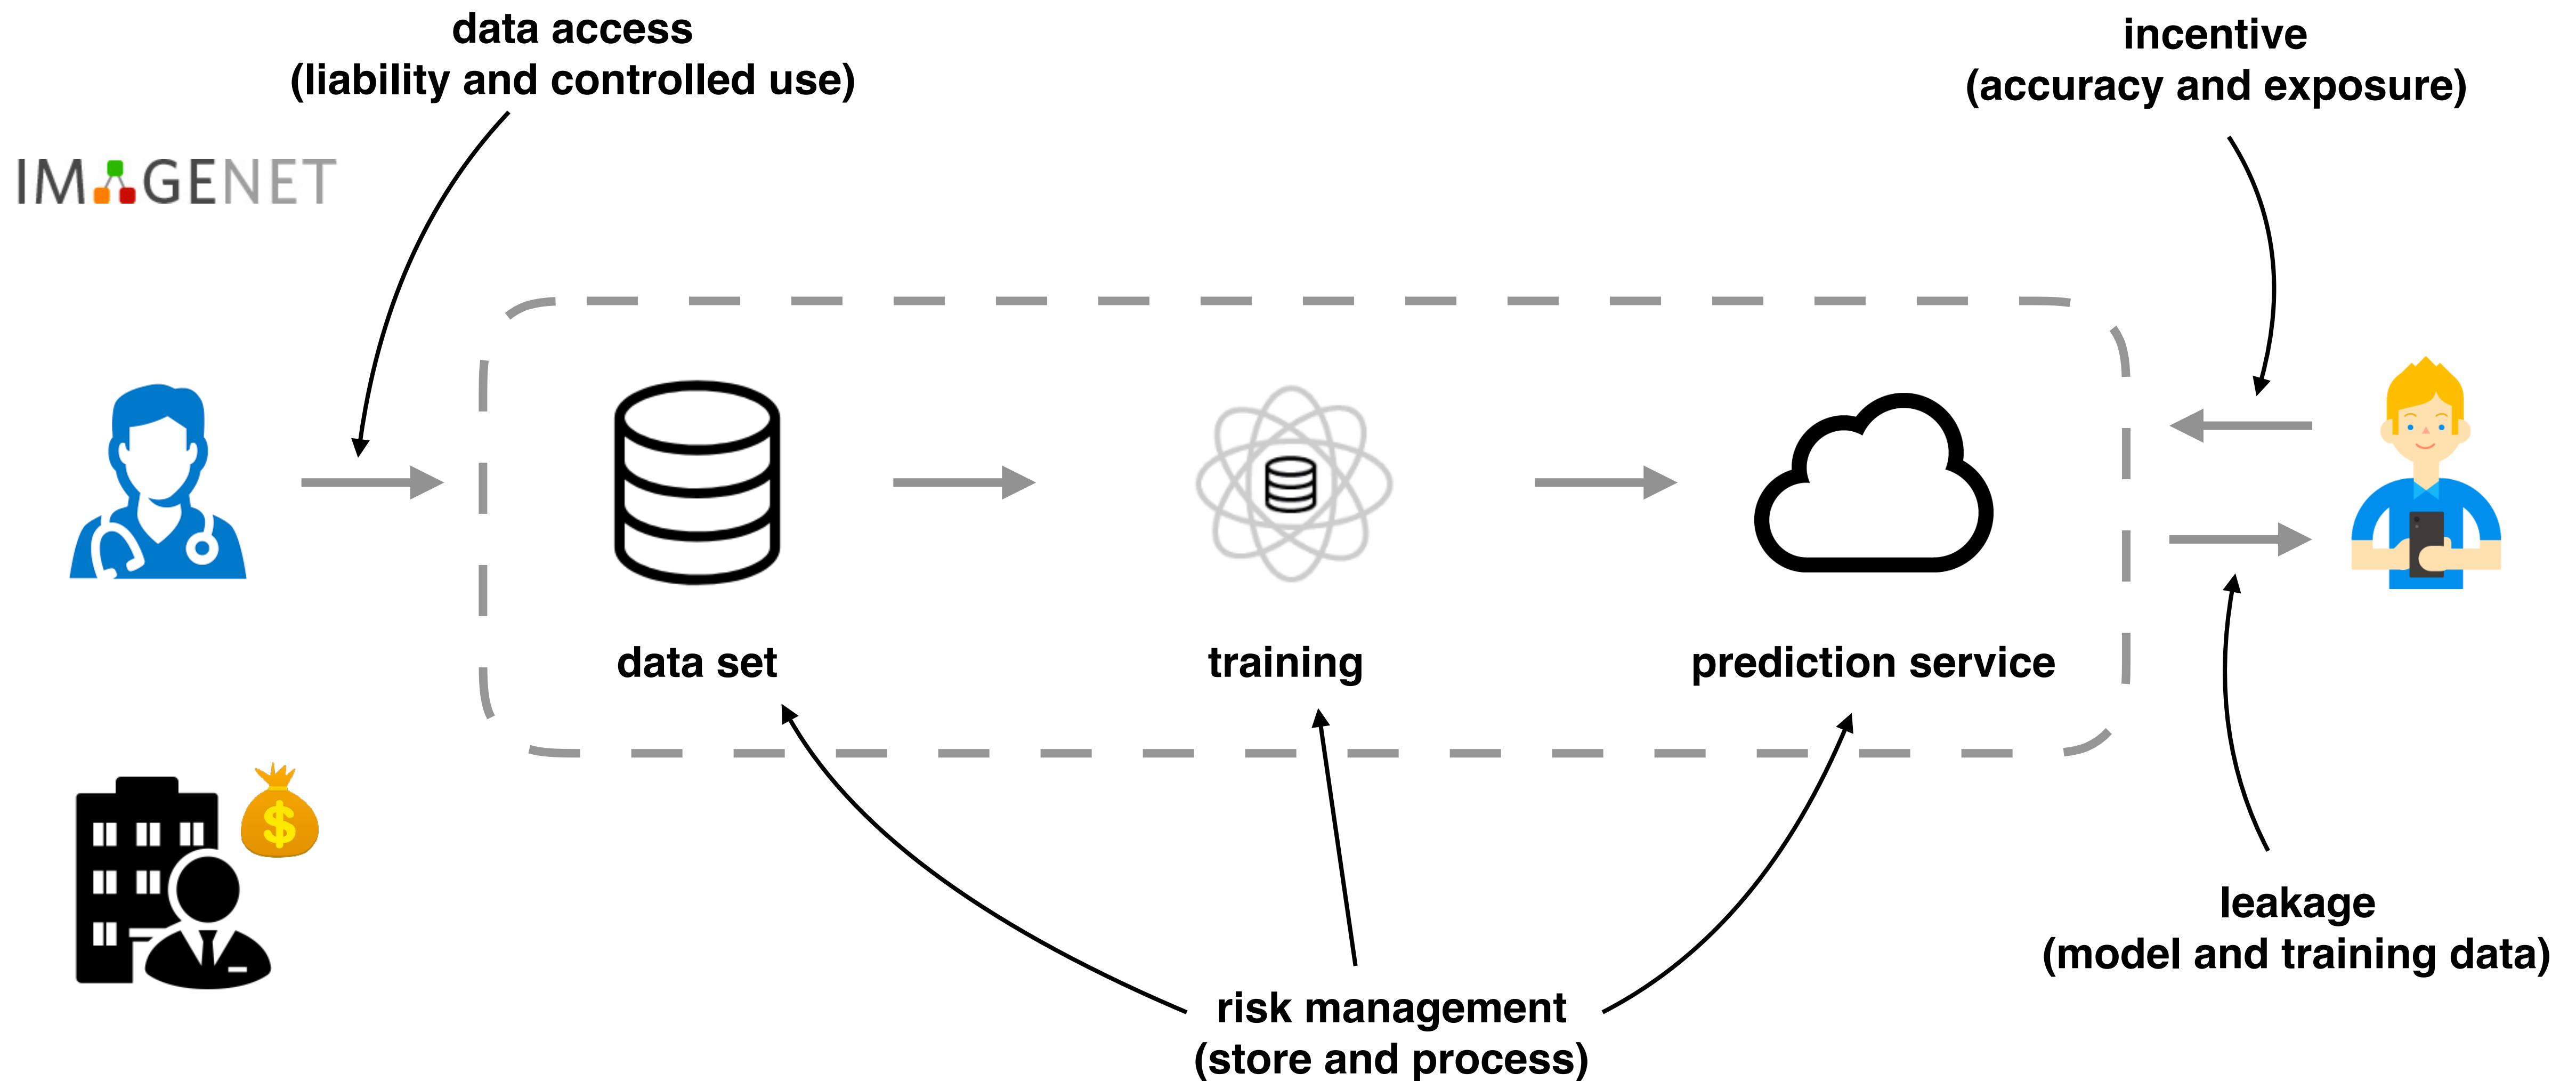
Brett Kuprel

12:30-12:40pm

Join Brett Kuprel, and see how TensorFlow was used by the artificial intelligence lab and medical school of Stanford to classify skin cancer images. He'll describe the project steps: from acquiring a dataset, training a deep network, and evaluating of the results. To wrap up, Brett will give his take on the future of skin cancer image classification.

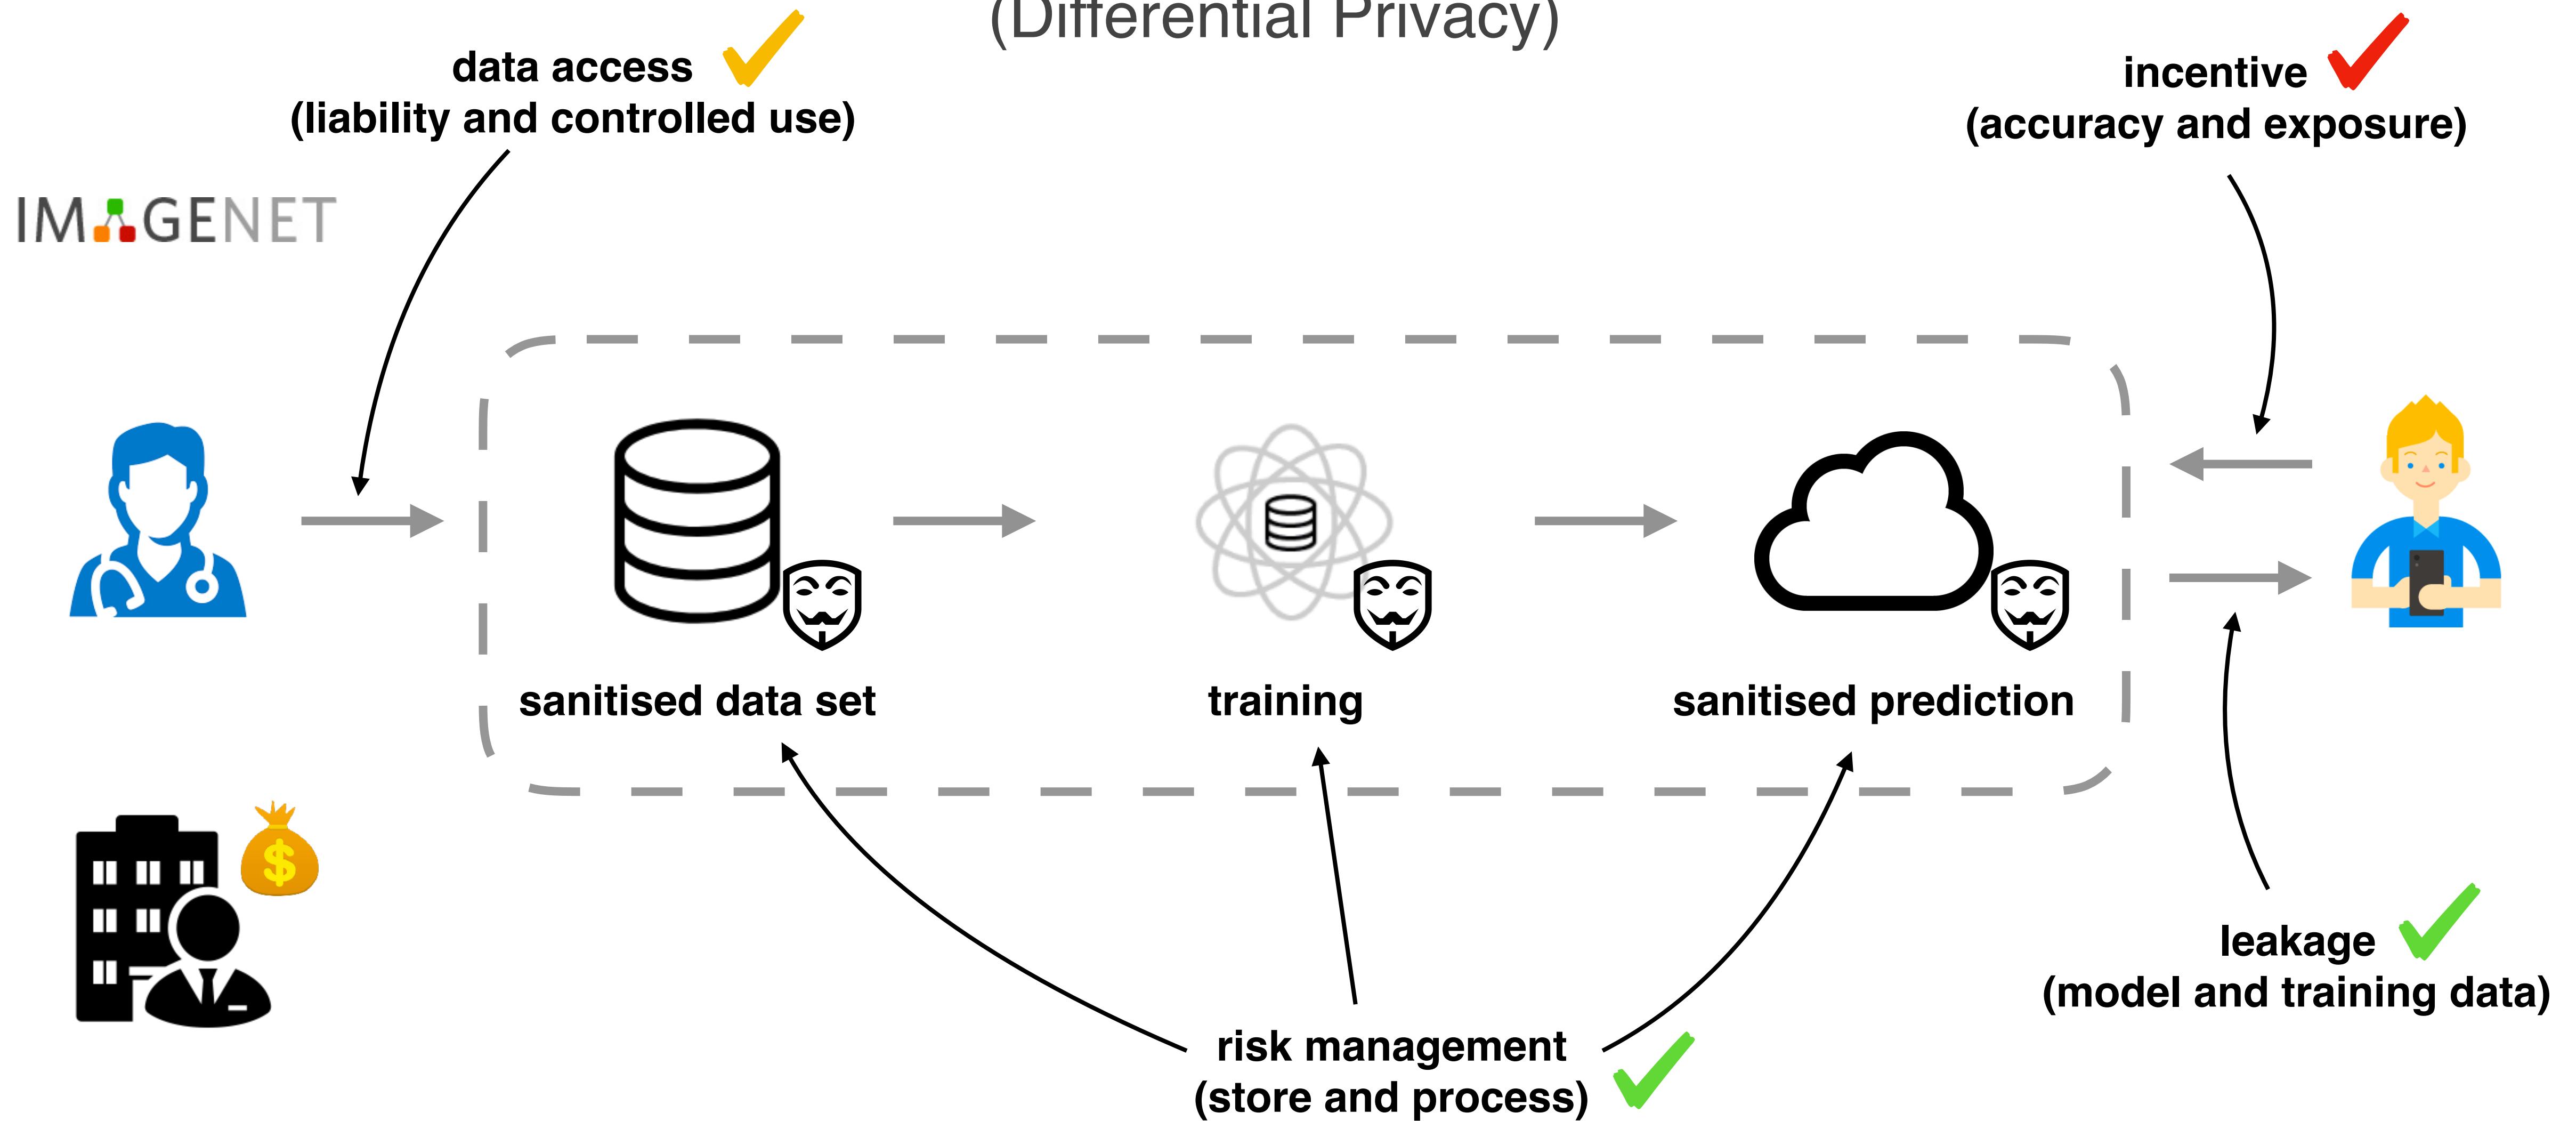


Potential Bottlenecks



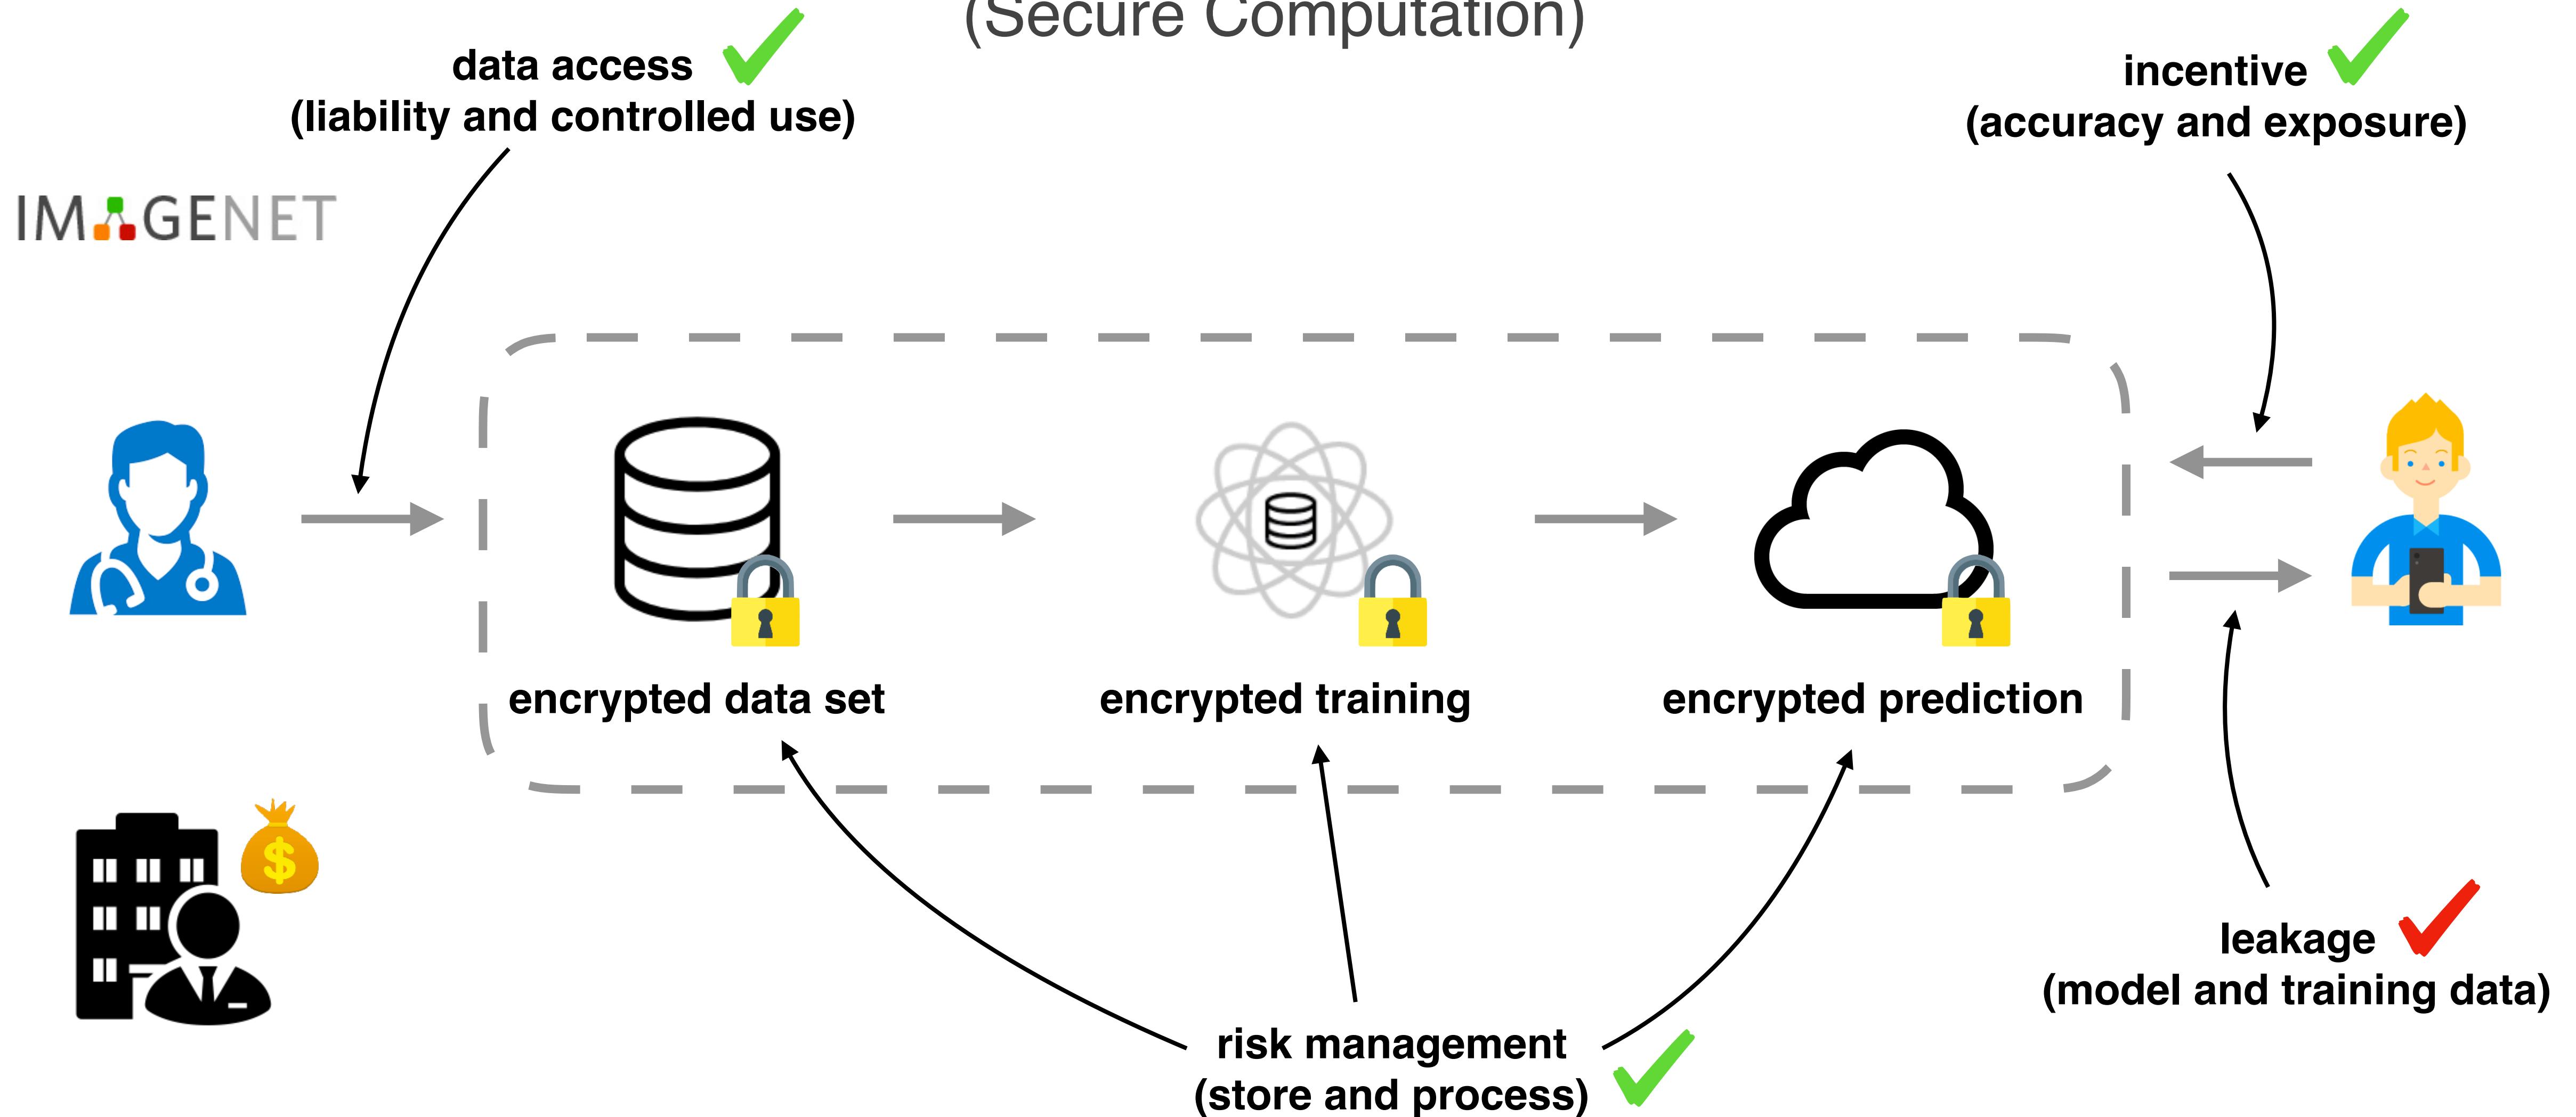
Sanitization

(Differential Privacy)

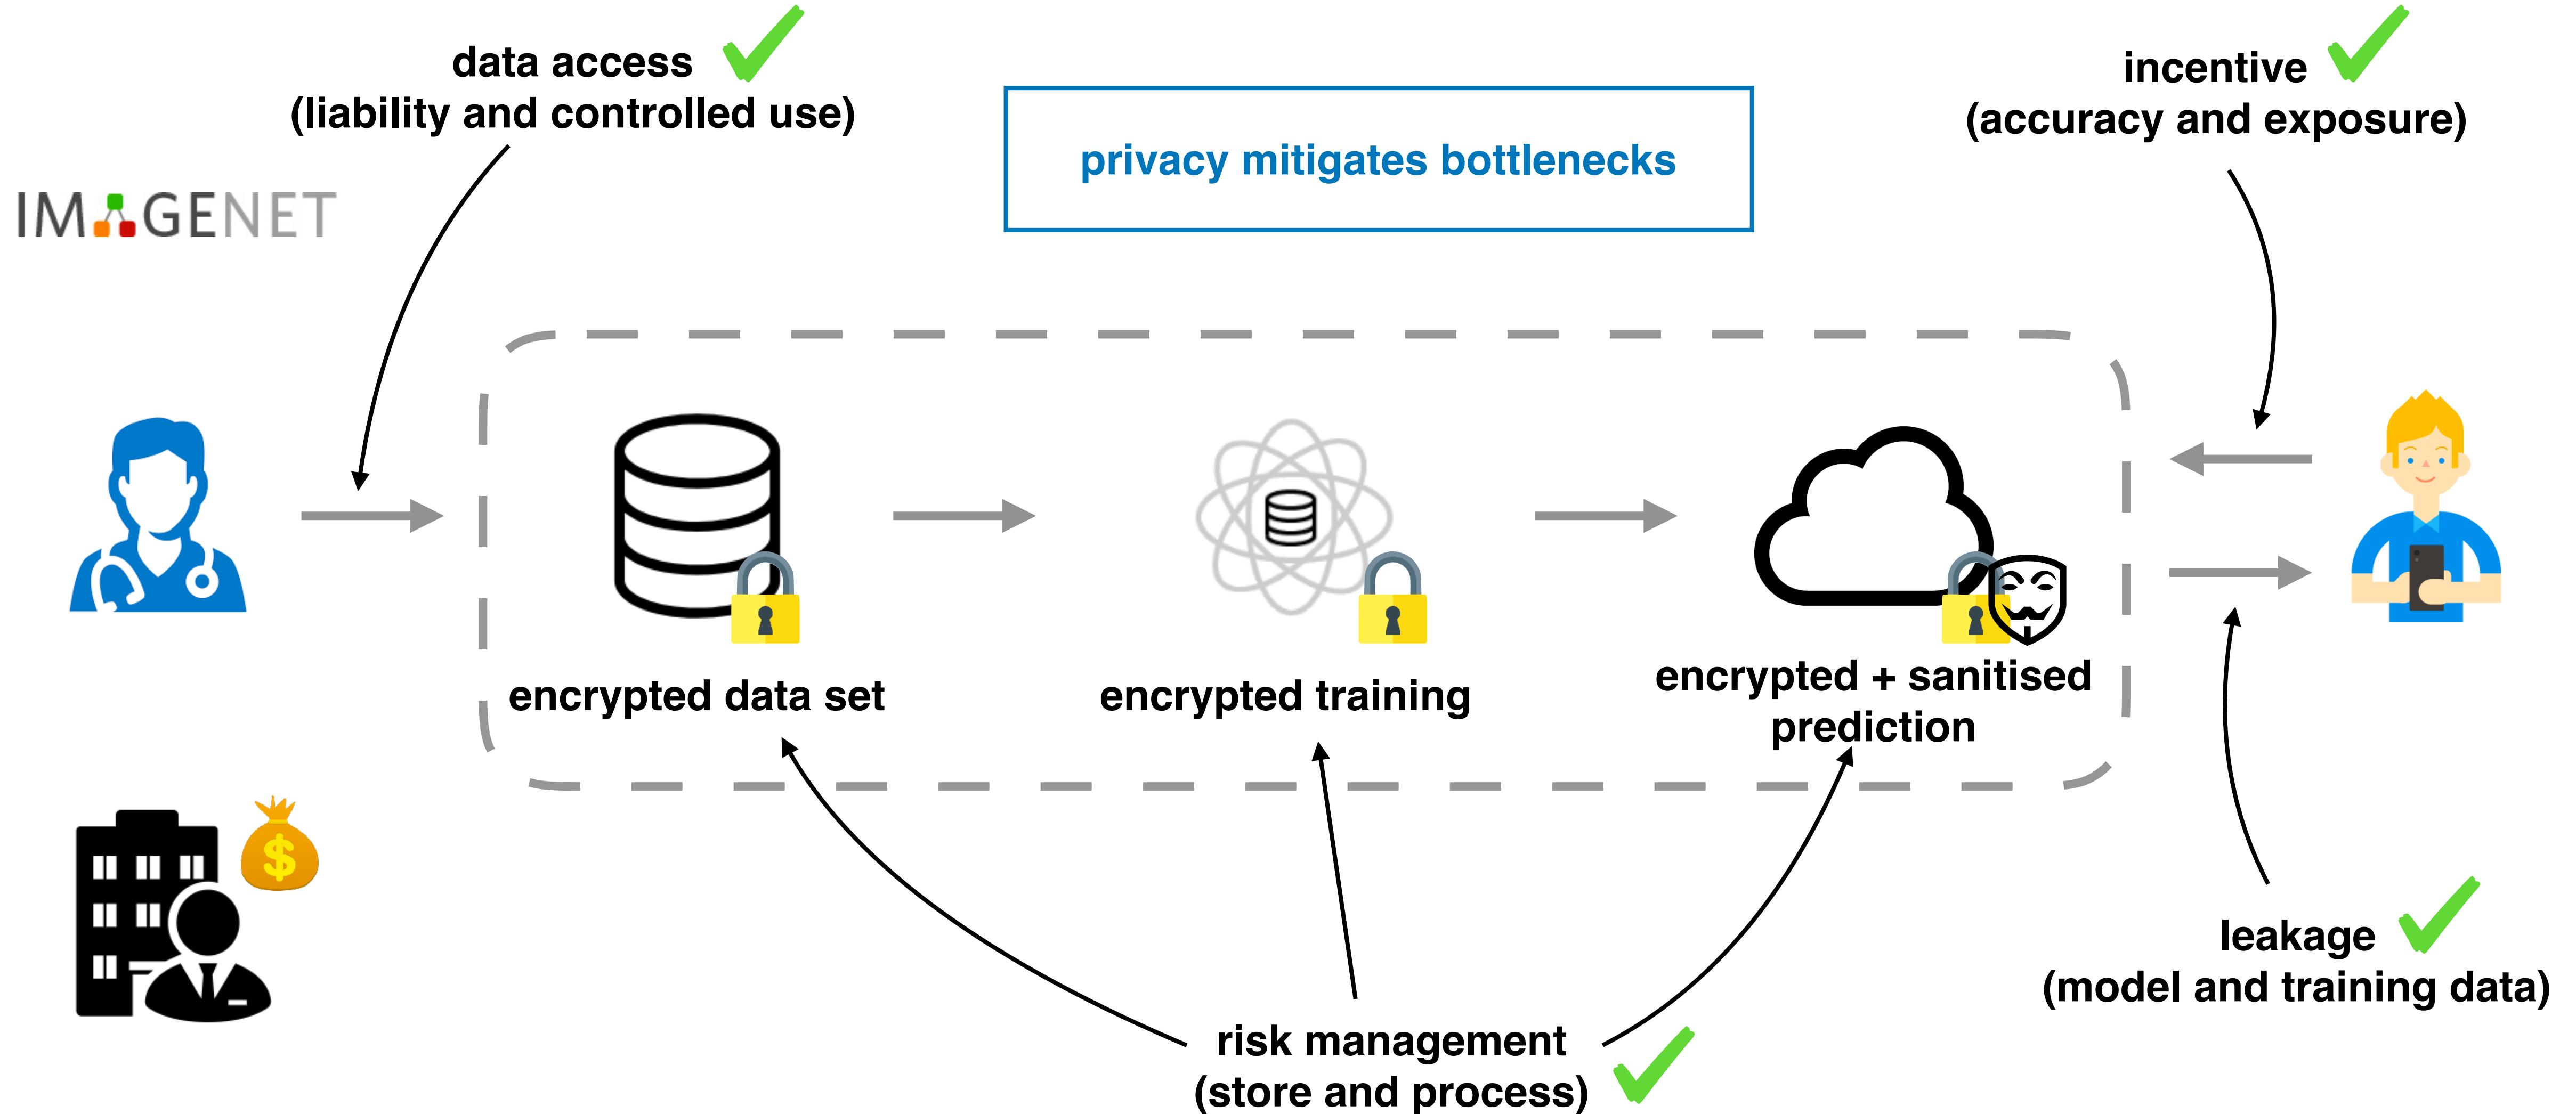


Encryption

(Secure Computation)



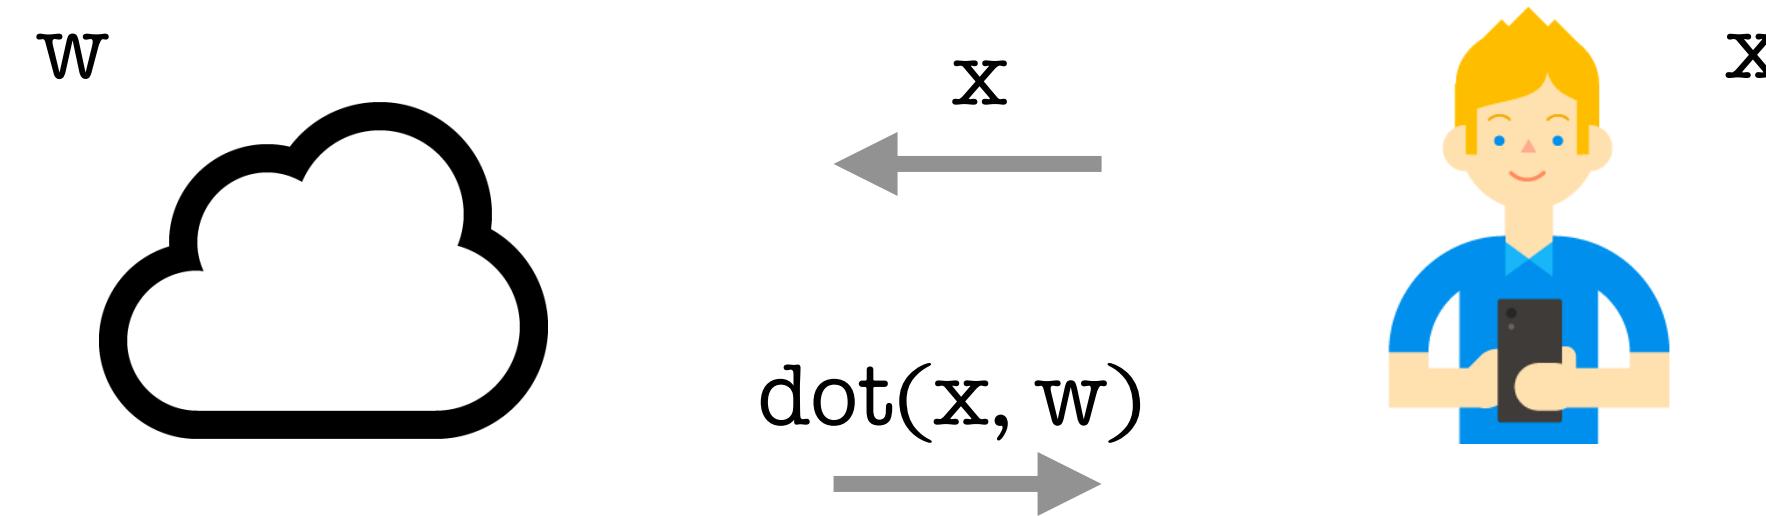
Hybrid



How?

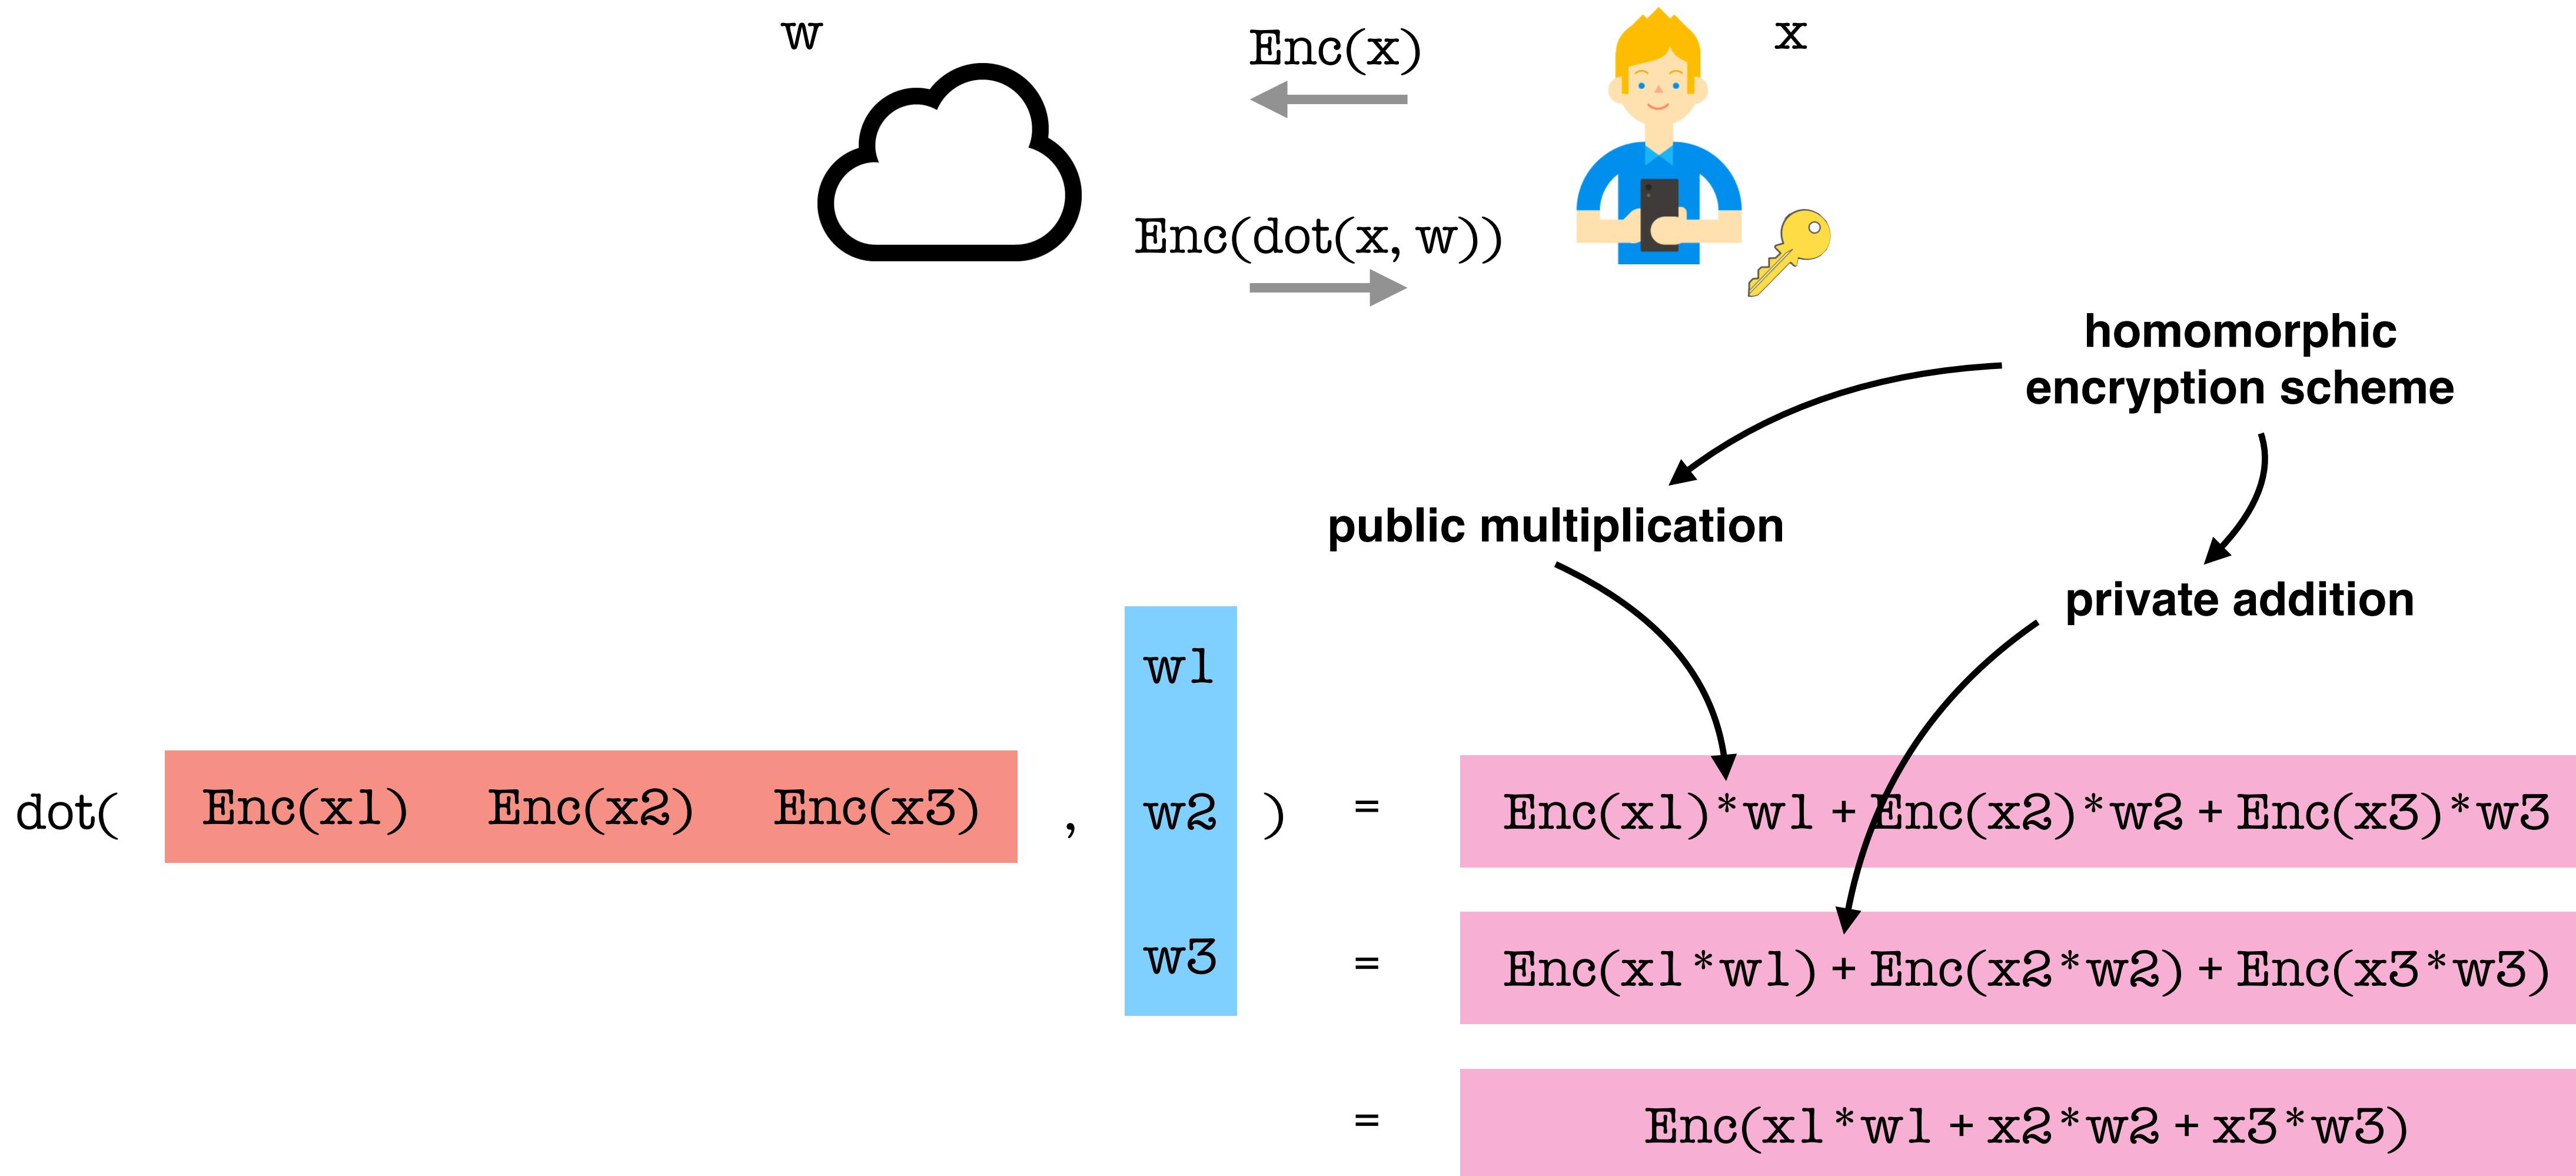
Computing on encrypted data

Prediction with Linear Model

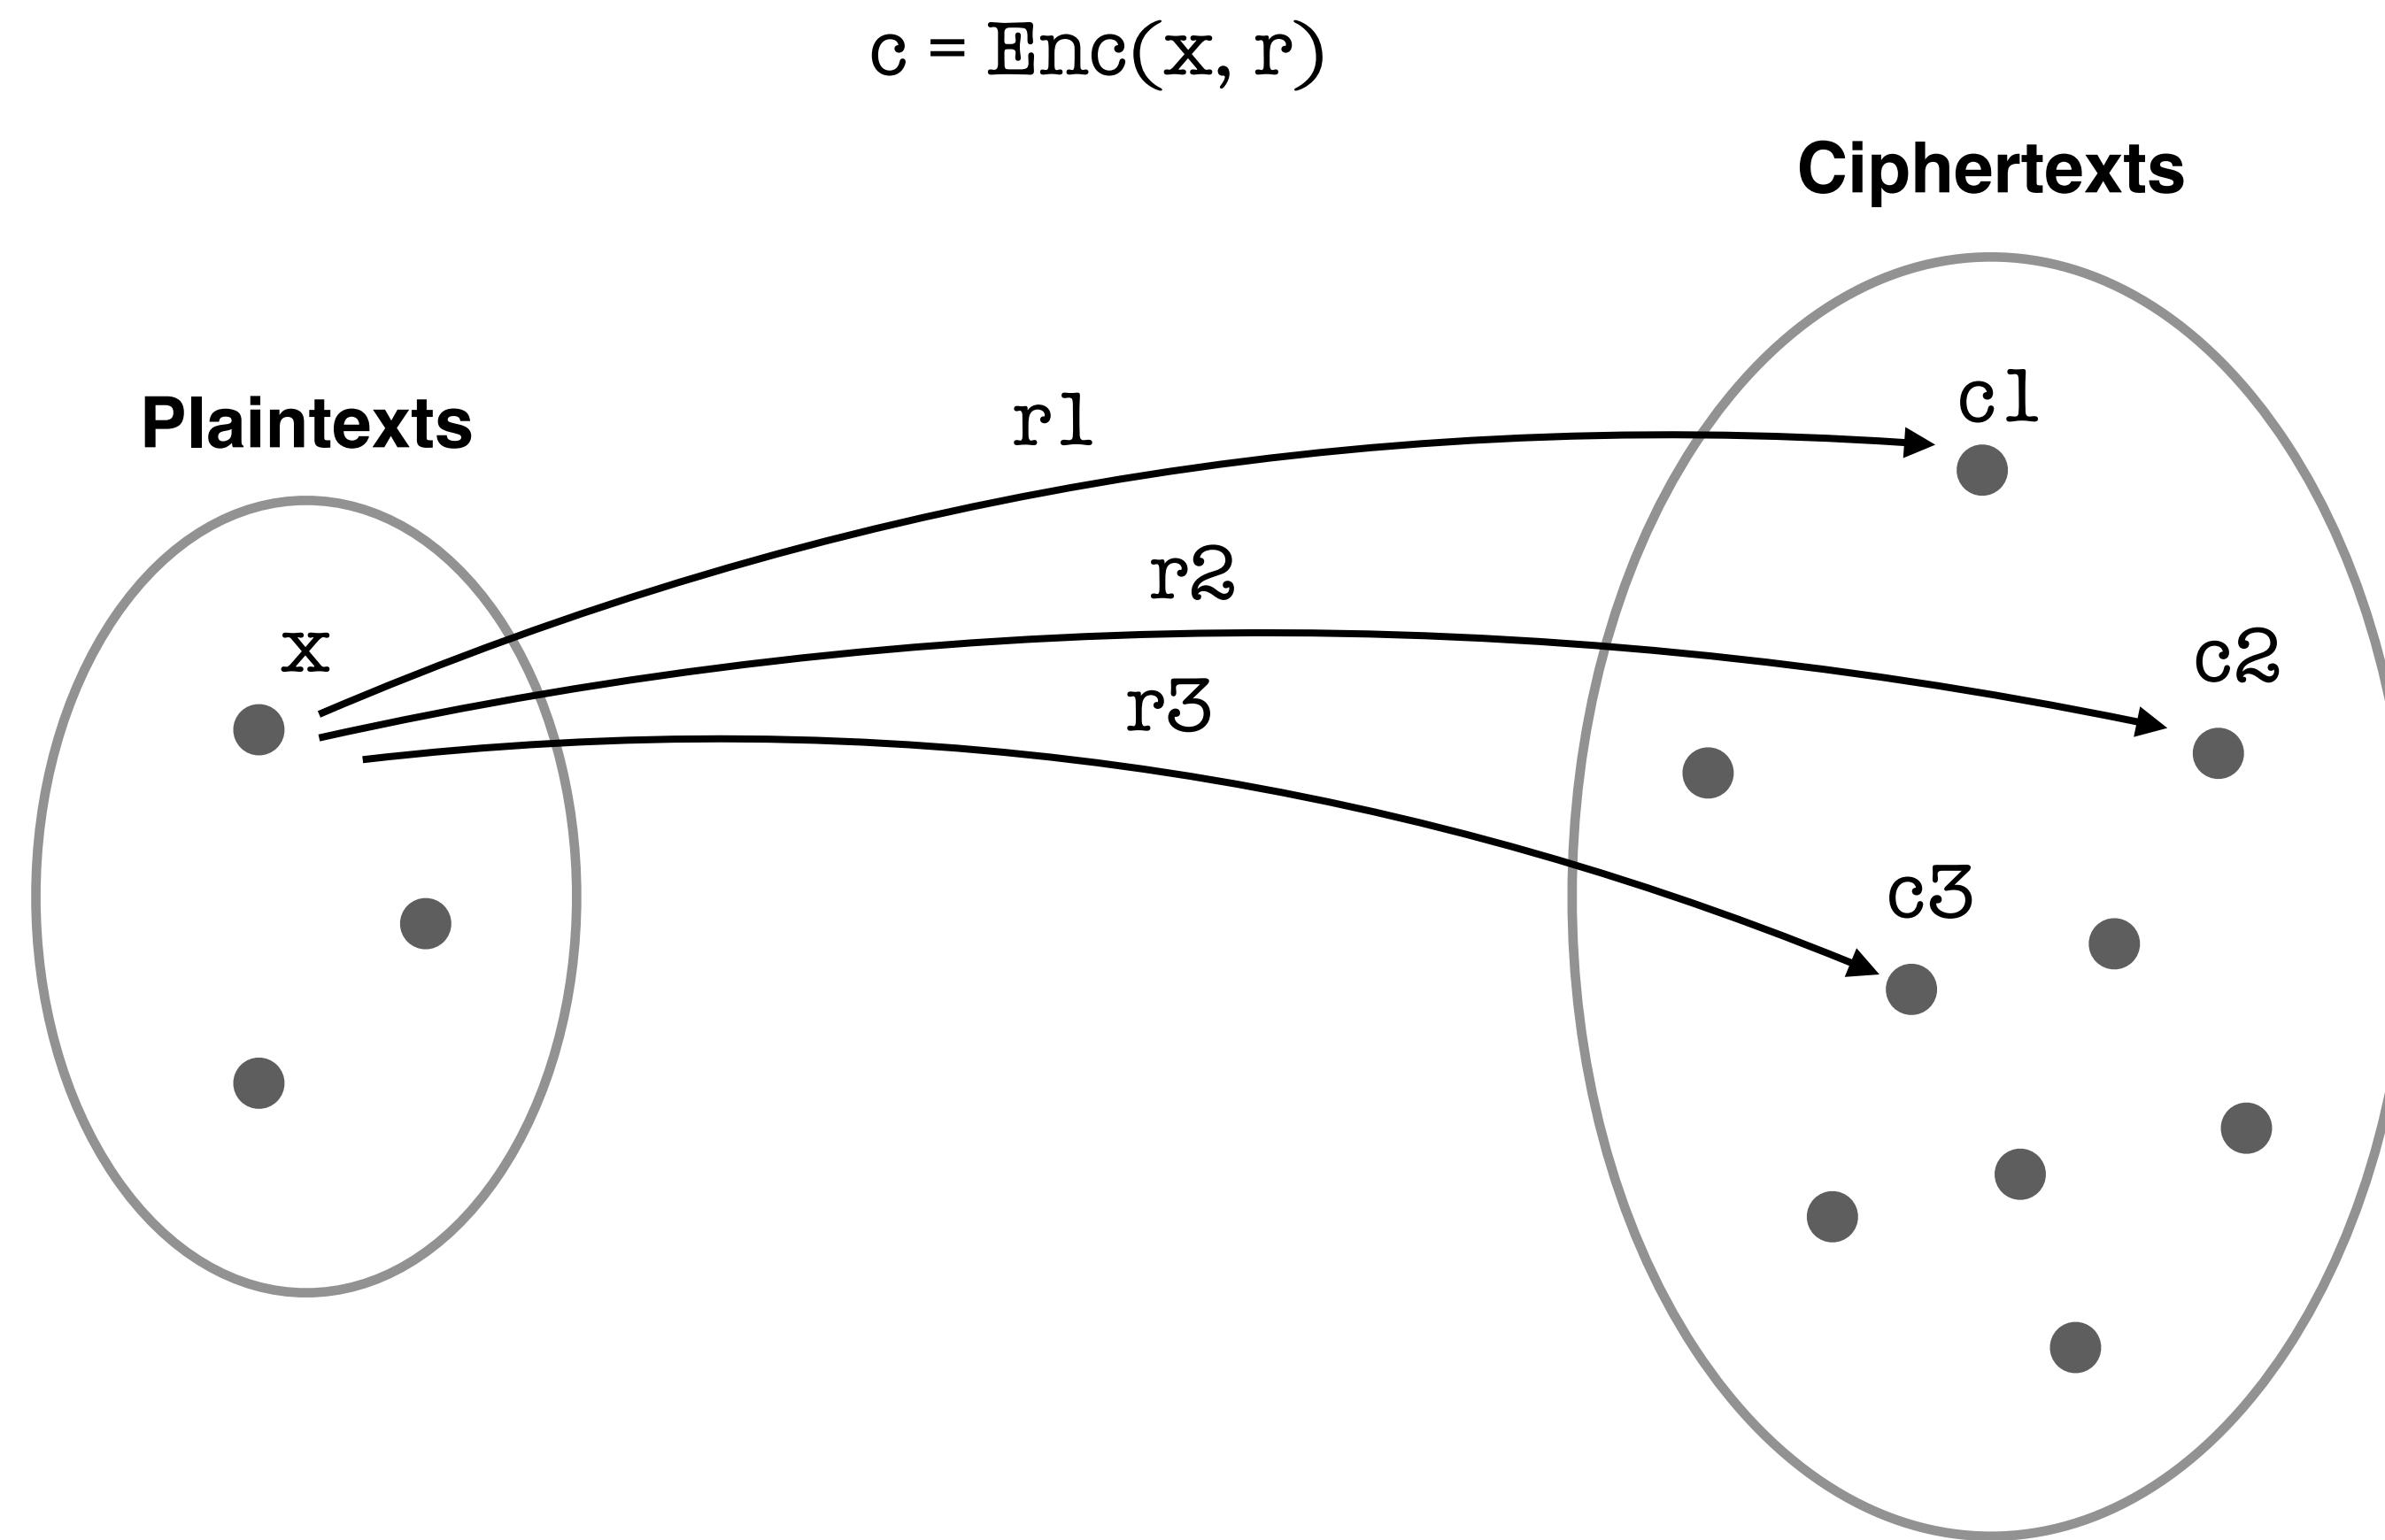


$$\text{dot}(\begin{matrix} x_1 & x_2 & x_3 \end{matrix}, \begin{matrix} w_1 \\ w_2 \\ w_3 \end{matrix}) = x_1 * w_1 + x_2 * w_2 + x_3 * w_3$$

... using Homomorphic Encryption



Probabilistic Encryption



Paillier Homomorphic Encryption

typically ~4000 bits:
computation is significantly
more expensive

$$c = \text{Enc}(x, r) = g^x * r^n \bmod n^2$$

public encryption key

public encryption key

$$\begin{aligned} g &= 36 \\ n &= 35 \\ n^2 &= 1225 \end{aligned}$$

$$\begin{aligned} \text{Enc}(5, 2) &= 36^5 * 2^{35} \bmod 1225 = 718 \\ \text{Enc}(5, 4) &= 36^5 * 4^{35} \bmod 1225 = 674 \end{aligned}$$

Private Addition in Paillier

$$\begin{aligned} & \text{Enc}(\mathbf{x}, \mathbf{r}) * \text{Enc}(\mathbf{y}, \mathbf{s}) \\ = & (\mathbf{g}^{\mathbf{x}} * \mathbf{r}^{\mathbf{n}} \bmod \mathbf{n}^2) * (\mathbf{g}^{\mathbf{y}} * \mathbf{s}^{\mathbf{n}} \bmod \mathbf{n}^2) \\ = & \mathbf{g}^{\mathbf{x+y}} * (\mathbf{r} * \mathbf{s})^{\mathbf{n}} \bmod \mathbf{n}^2 \\ = & \text{Enc}(\mathbf{x+y}, \mathbf{r*s}) \end{aligned}$$

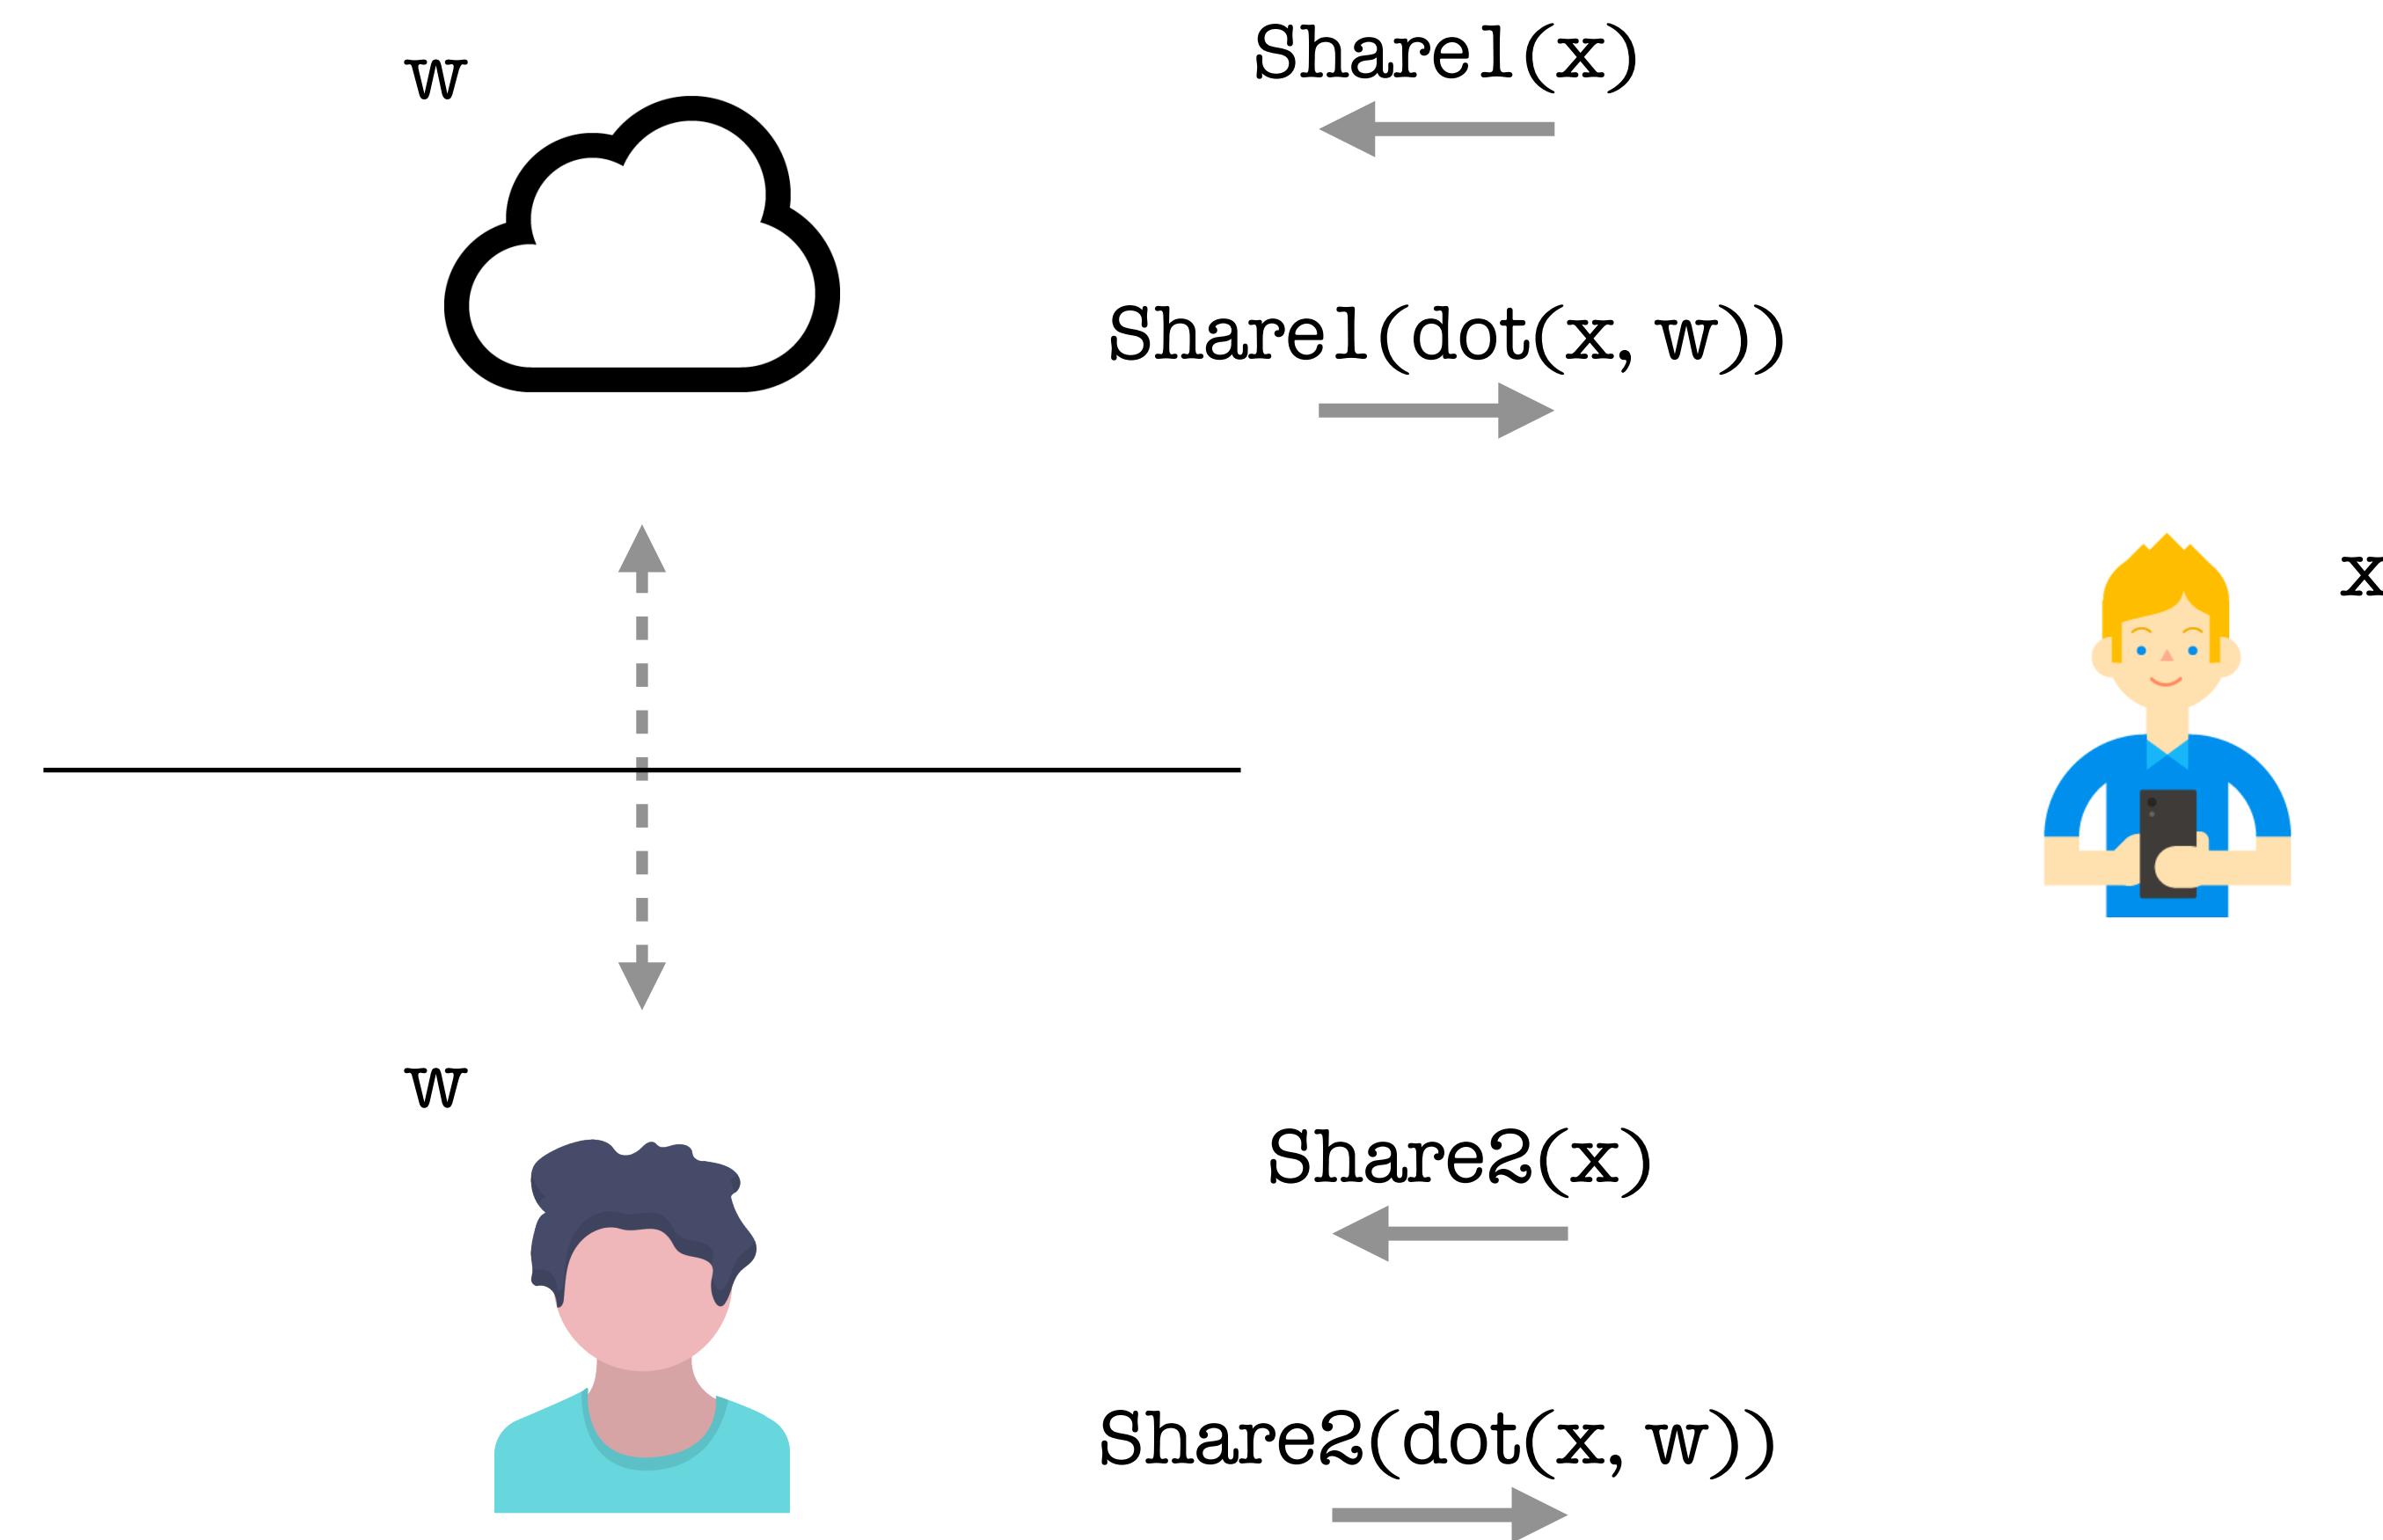
$$\begin{aligned} & \text{Enc}(5, 2) * \text{Enc}(5, 4) \\ = & 718 * 674 \\ = & 57 \\ = & 36^{10} * 8^{35} \\ = & \text{Enc}(10, 8) \end{aligned}$$

Public Multiplication in Paillier

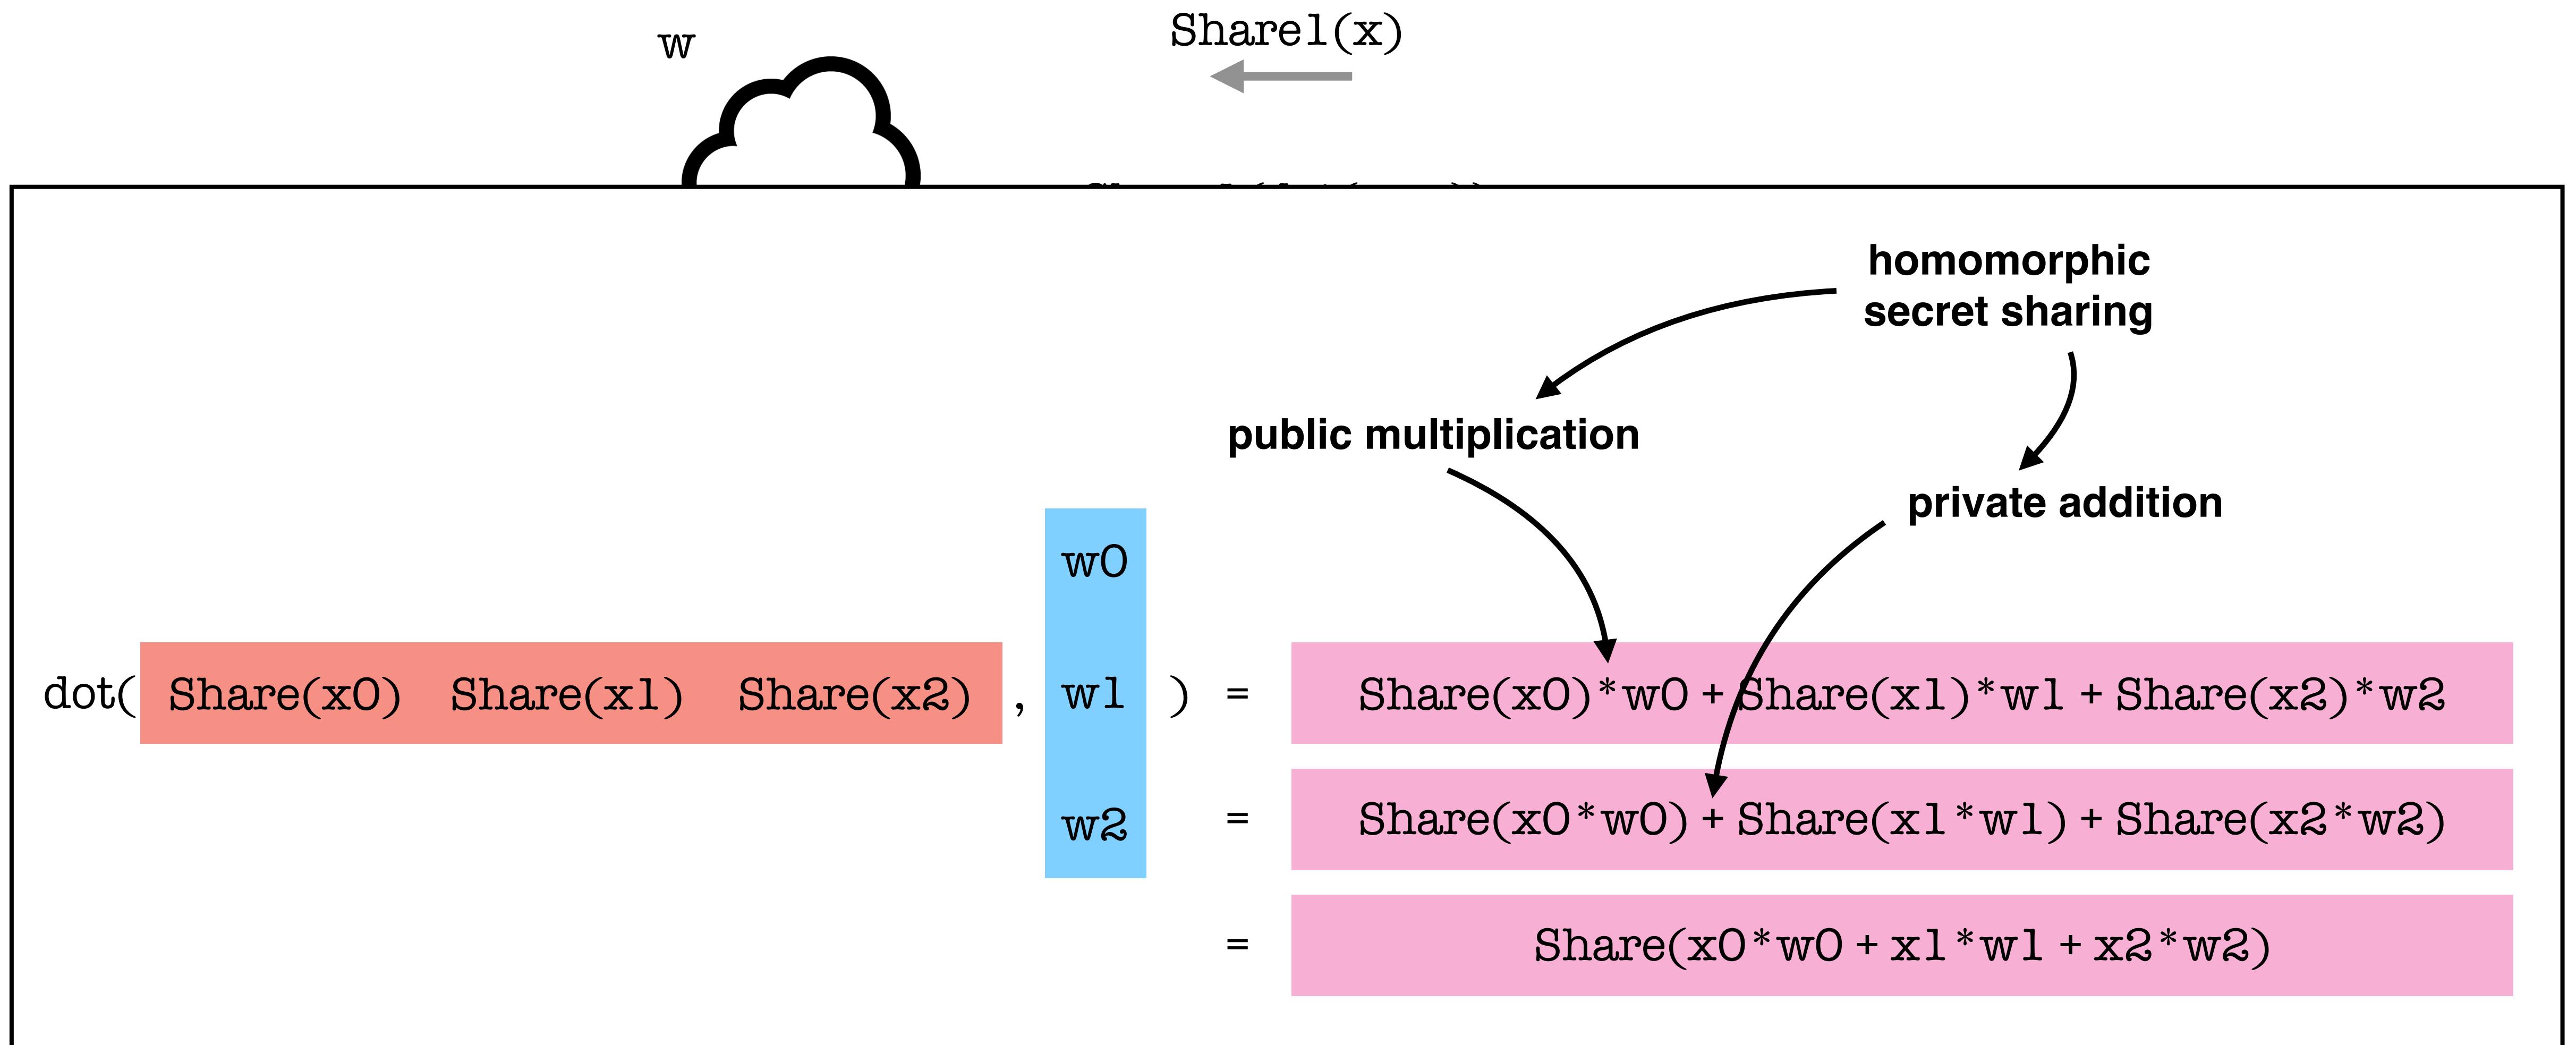
$$\begin{aligned} & \text{Enc}(\text{x}, \text{r})^{\wedge w} \\ &= (\text{g}^{\wedge x} * \text{r}^{\wedge n} \bmod n^2)^{\wedge w} \\ &= \text{g}^{\wedge(x * w)} * (\text{r}^{\wedge w})^{\wedge n} \bmod n^2 \\ &= \text{Enc}(\text{x} * w, \text{r}^{\wedge w}) \end{aligned}$$

$$\begin{aligned} & \text{Enc}(5, 2)^{\wedge 2} \\ &= 718 * 718 \\ &= 1024 \\ &= 36^{\wedge 10} * 4^{\wedge 35} \\ &= \text{Enc}(10, 4) \end{aligned}$$

... using Secret Sharing



... using Secret Sharing



Secret Sharing

$$\text{Share1}(x, r) = r \bmod m$$

$$\text{Share2}(x, r) = x - r \bmod m$$

$$x = \text{Share1}(x, r) + \text{Share2}(x, r) \bmod m$$

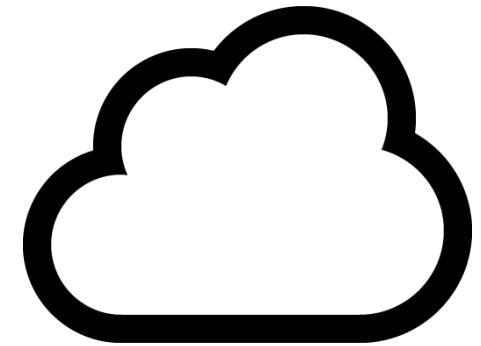
$$m = 10$$

$$\text{Share1}(5, 7) = 7 \bmod 10 = 7$$

$$\text{Share2}(5, 7) = 5 - 7 \bmod 10 = 8$$

$$7 + 8 = 15 = 5 \bmod 10$$

Private Addition with Secret Sharing



x₁

y₁

z₁ = x₁ + y₁



x₂

y₂

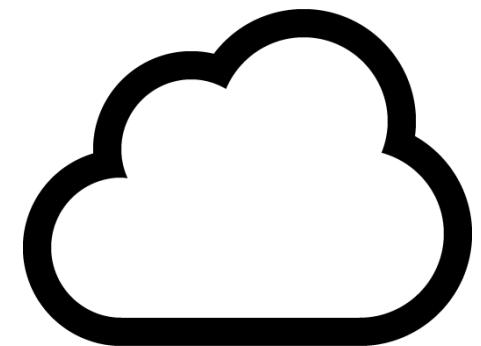
z₂ = x₂ + y₂

$$x = x_1 + x_2$$

$$y = y_1 + y_2$$

$$\begin{aligned}x + y &= (x_1 + x_2) + (y_1 + y_2) \\&= (x_1 + y_1) + (x_2 + y_2) \\&= z_1 + z_2\end{aligned}$$

Public Multiplication with Secret Sharing



x1

w

z1 = x1 * w



x2

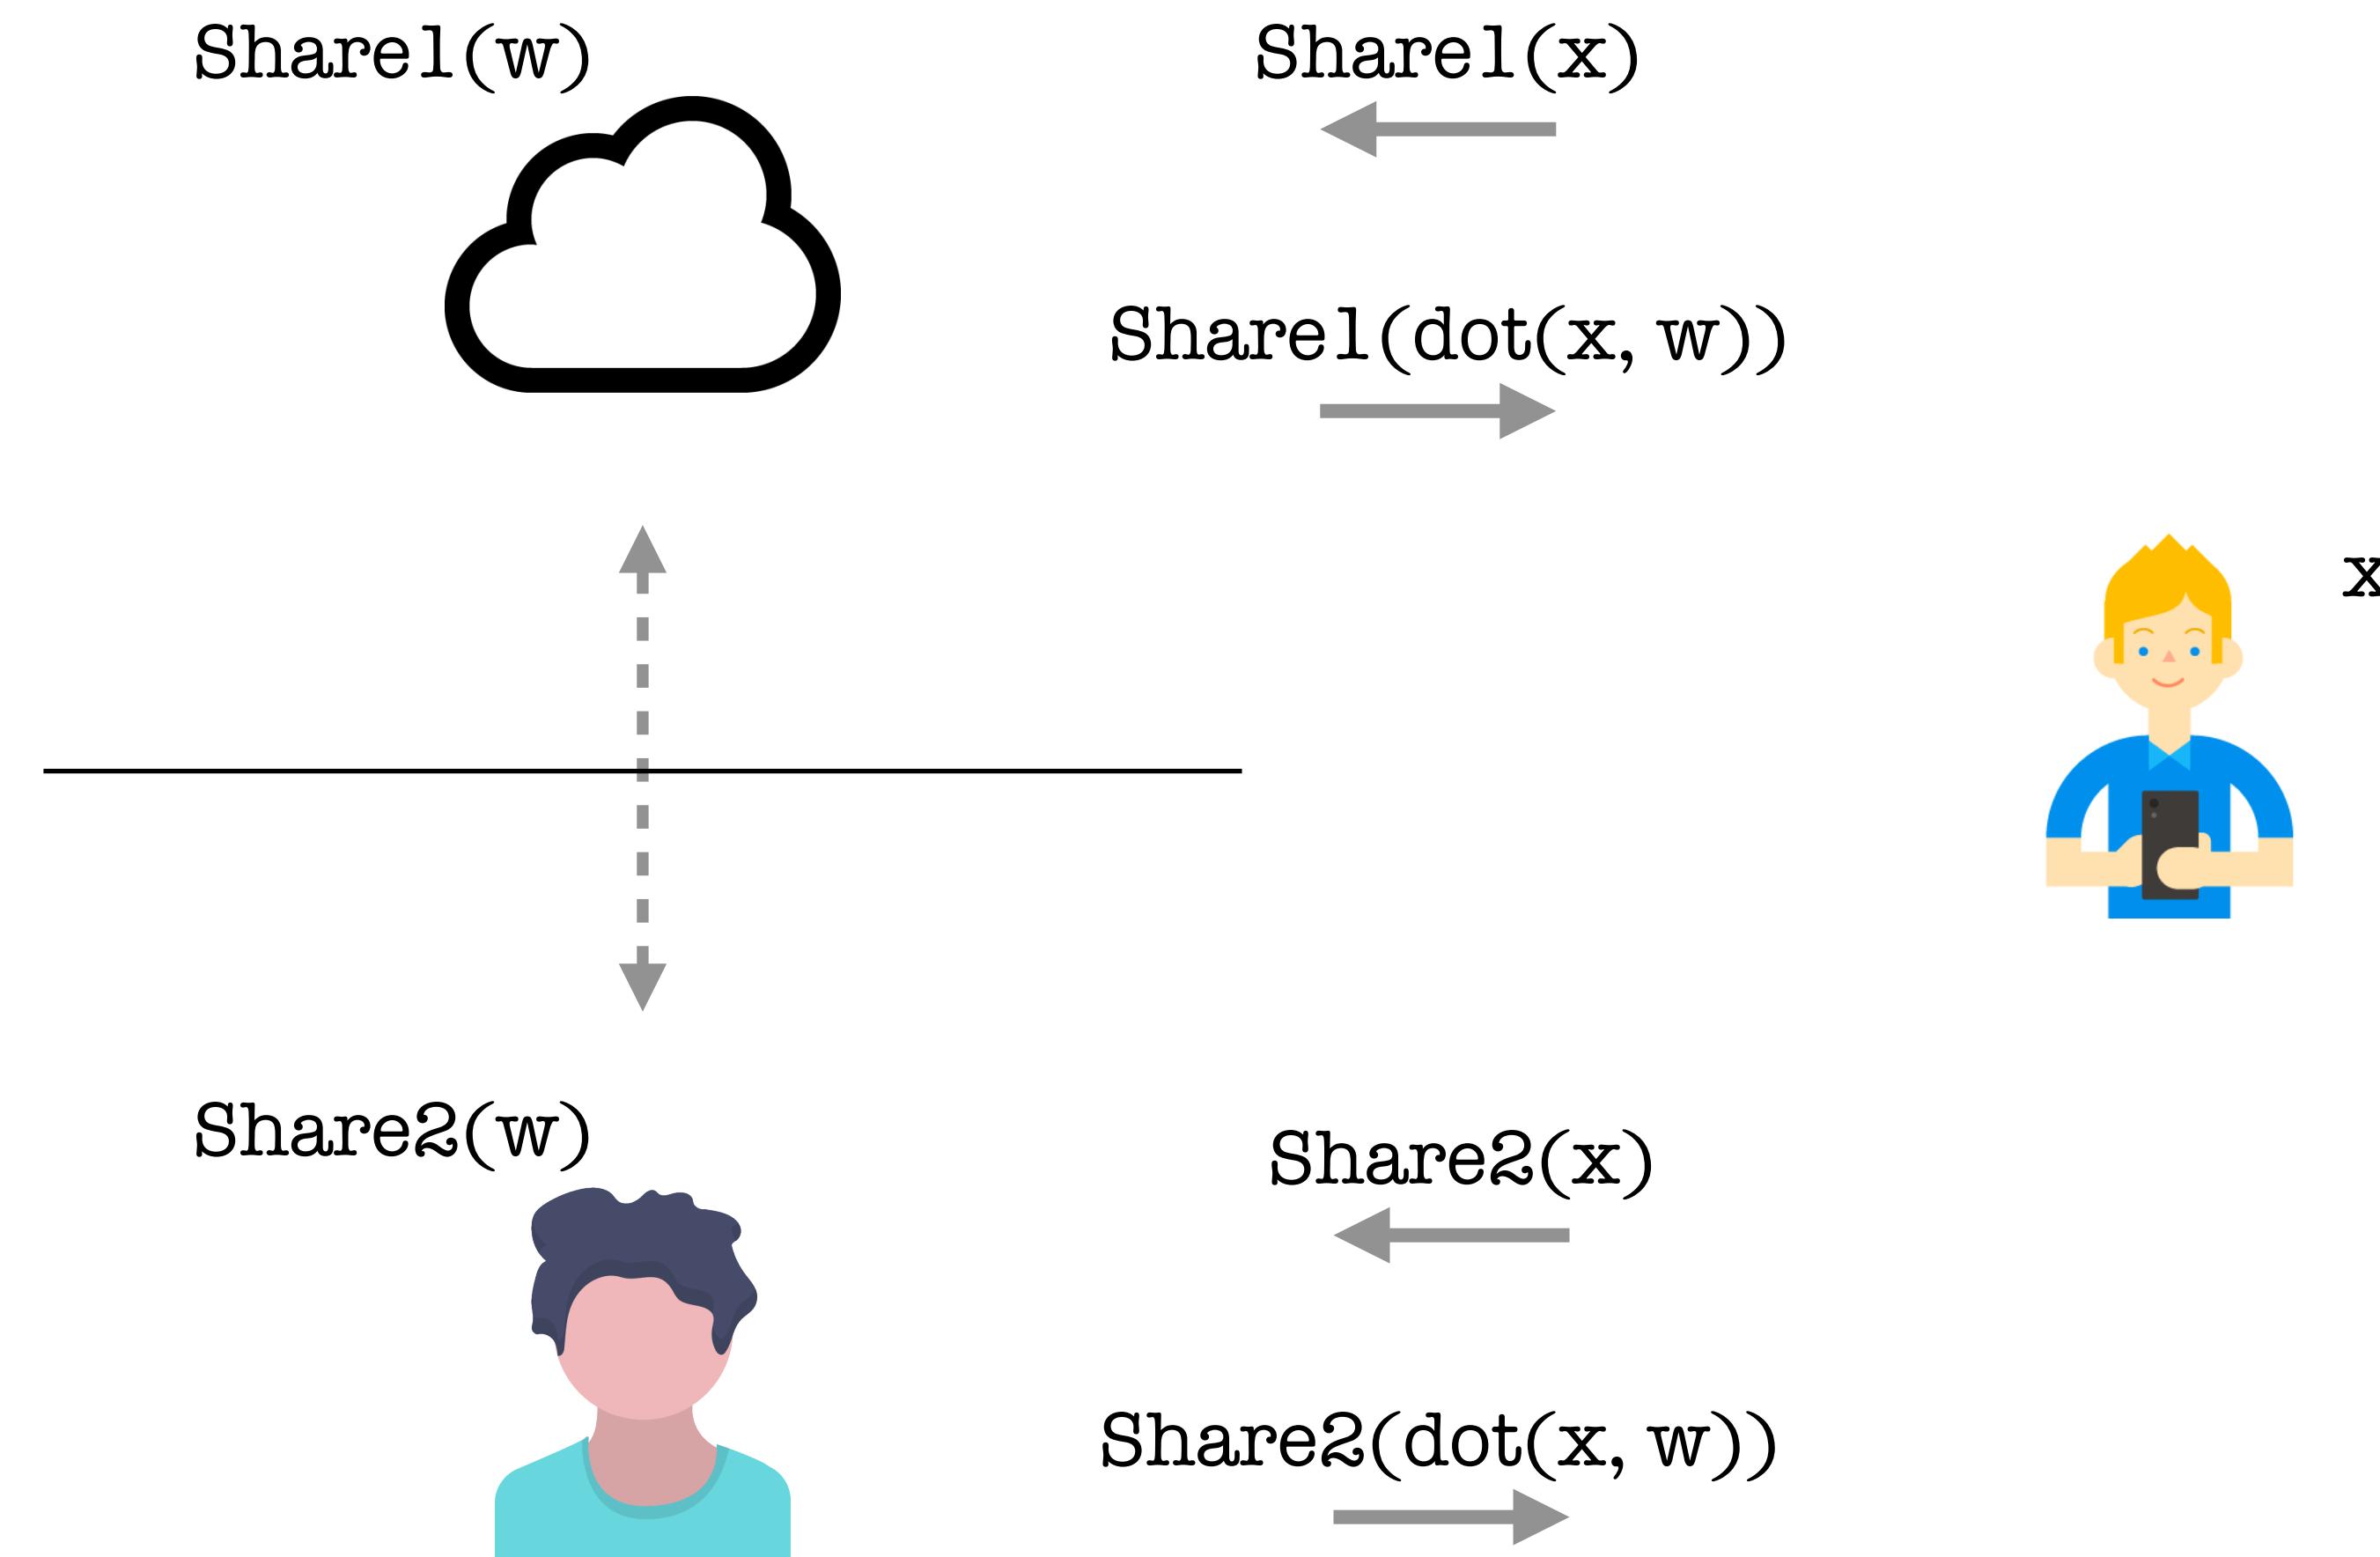
w

z2 = x2 * w

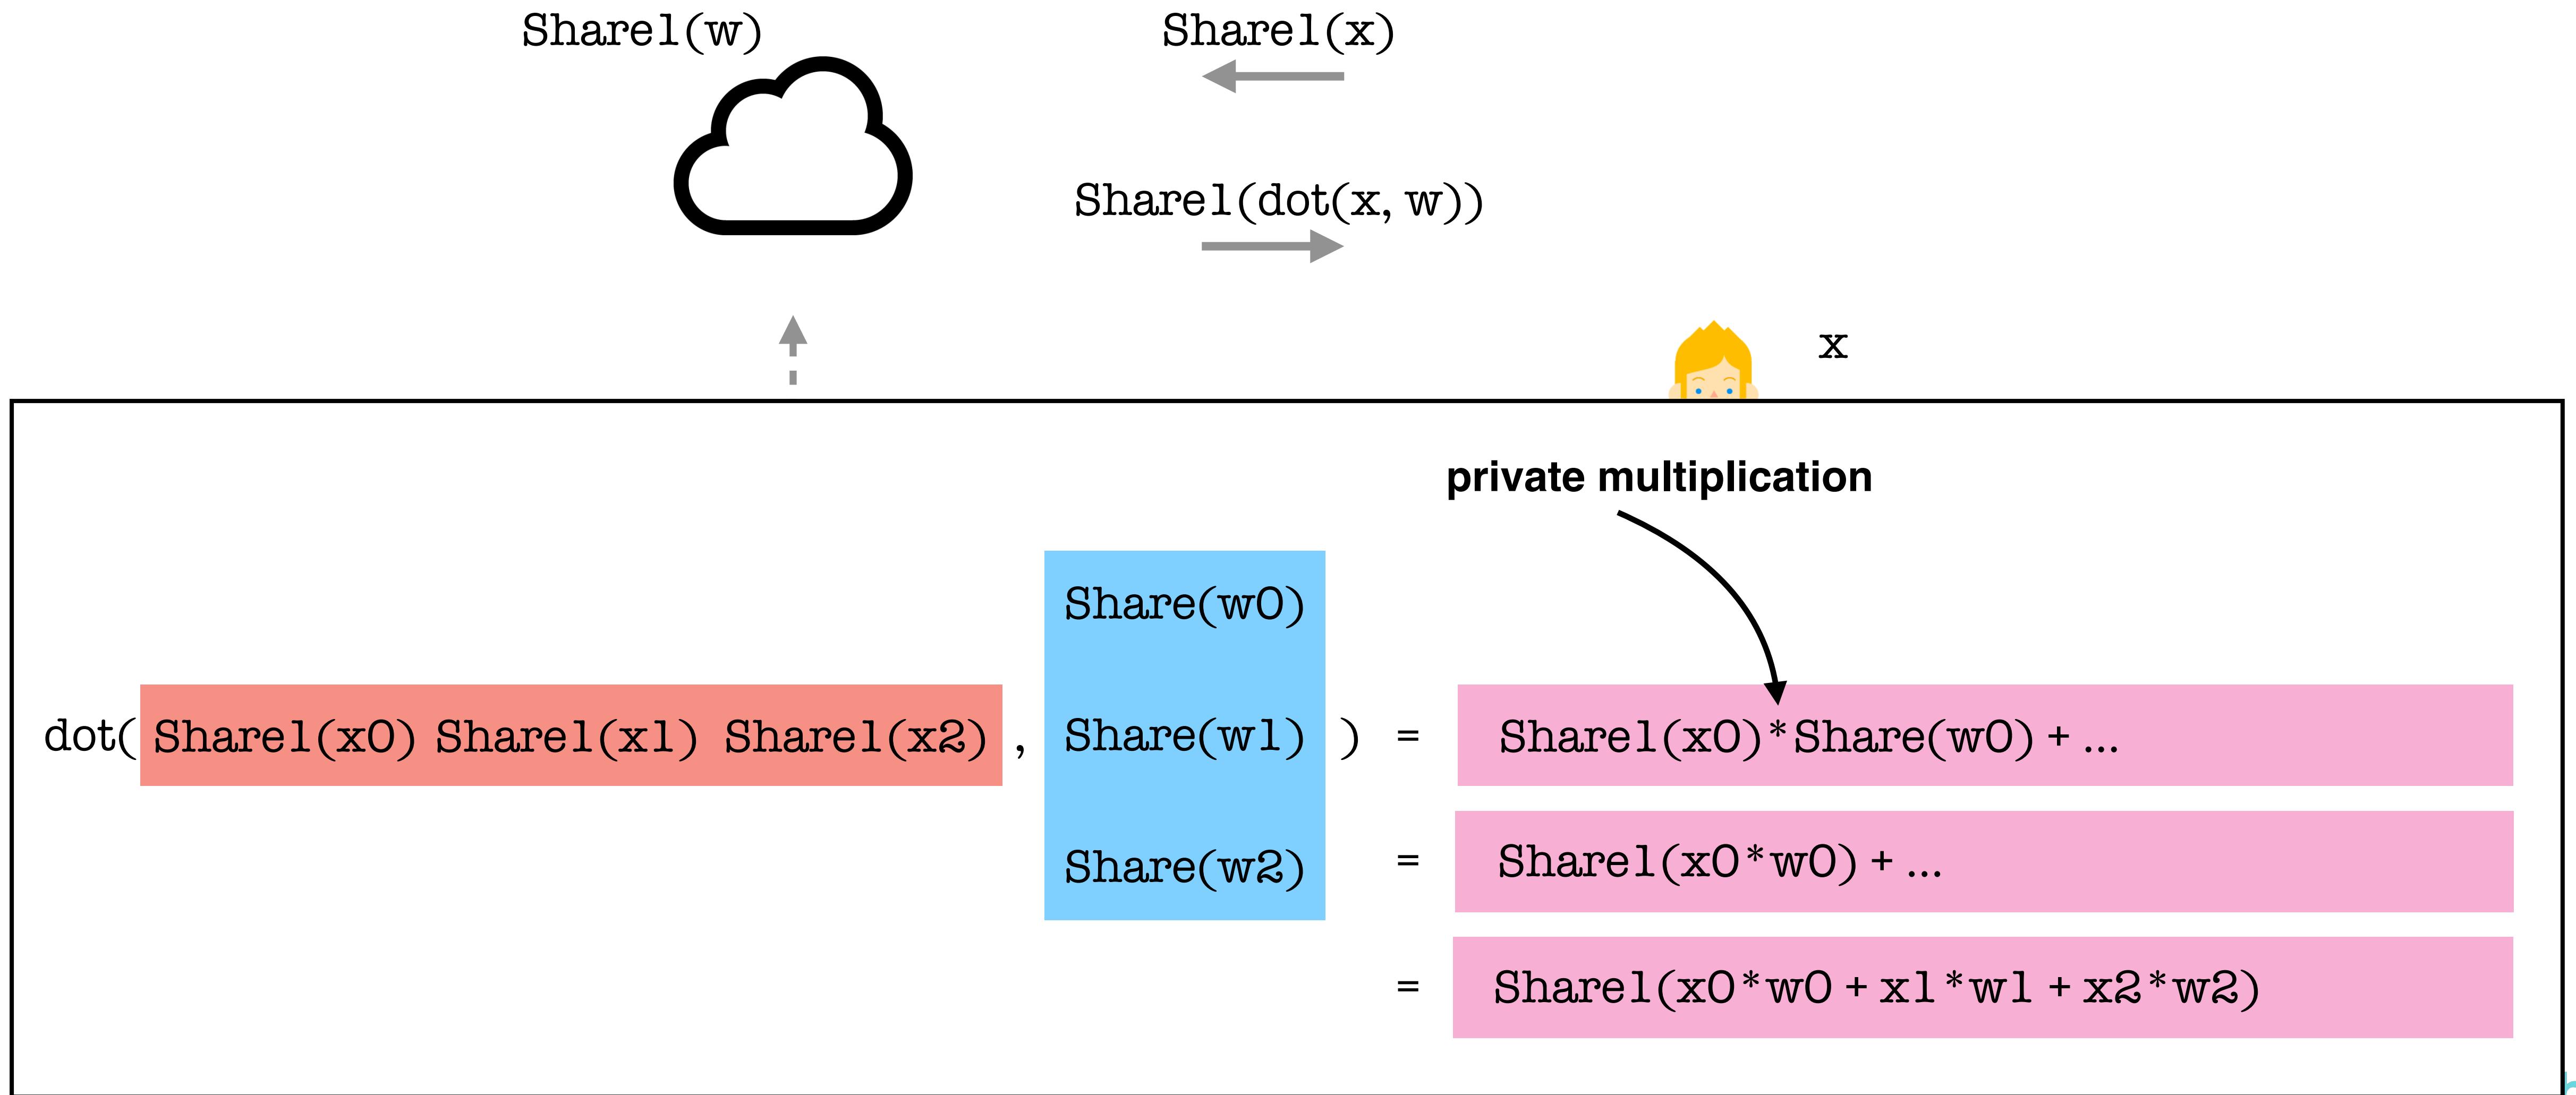
$$x = x1 + x2$$

$$\begin{aligned} x * w \\ &= (x1 + x2) * w \\ &= (x1 * w) + (x2 * w) \\ &= z1 + z2 \end{aligned}$$

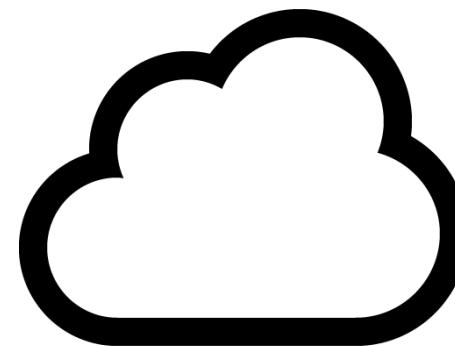
... using Secret Sharing, with Private Model



... using Secret Sharing, with Private Model



Private Multiplication with Secret Sharing



(a_1 , b_1 , c_1)

x_1

y_1

alpha

beta

$$\begin{aligned} z_1 &= \text{alpha} * \text{beta} \\ &+ \text{alpha} * b_1 \\ &+ \text{beta} * a_1 \\ &+ c_1 \end{aligned}$$



(a_2 , b_2 , c_2)

x_2

y_2

alpha

beta

$$\begin{aligned} z_2 &= \text{alpha} * b_2 \\ &+ \text{beta} * a_2 \\ &+ c_2 \end{aligned}$$

$$\begin{aligned} a &= a_1 + a_2 \\ b &= b_1 + b_2 \\ c &= a * b = c_1 + c_2 \end{aligned}$$

$$x = x_1 + x_2$$

$$y = y_1 + y_2$$

$$\text{alpha} = x - a$$

$$\text{beta} = y - b$$

$$\begin{aligned} x * y &= \dots \\ &= z_1 + z_2 \end{aligned}$$

Multidisciplinary Challenge

Data science
(use-cases, workflow, monitoring)

Cryptography
(techniques, protocols, trust)

Machine learning
(models, approx, precision)

Engineering
(distributed, multi-core, readability)

need common language

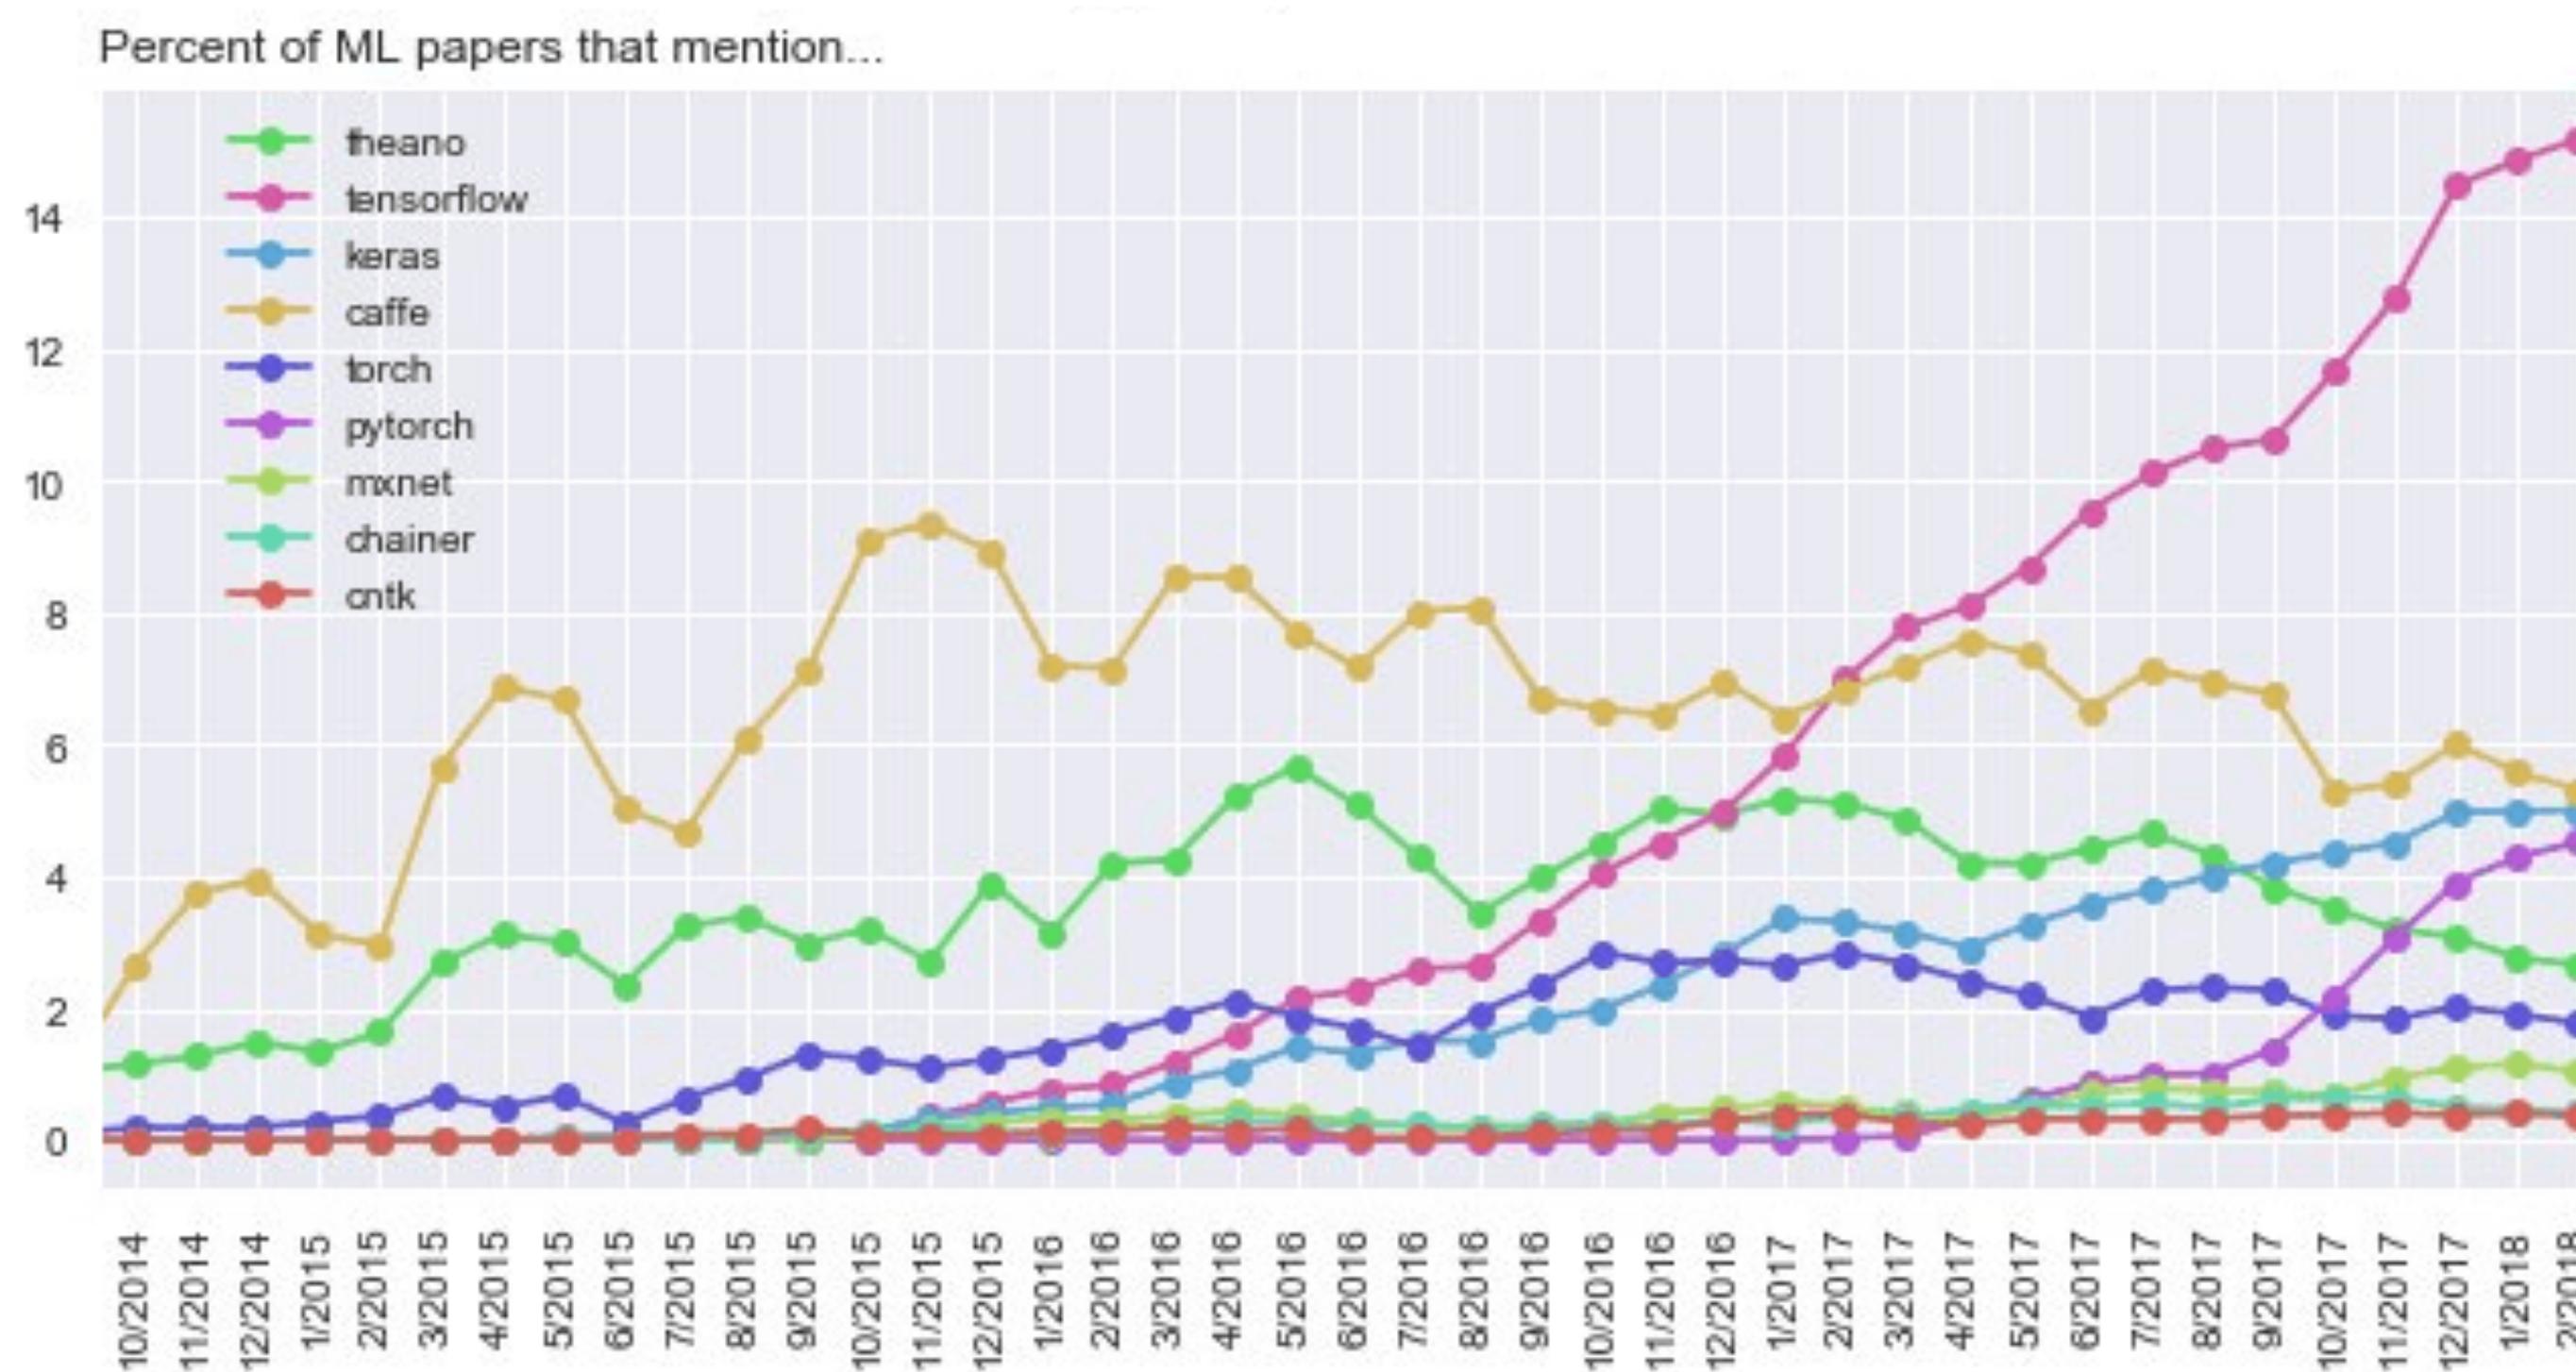
TF Encrypted

Making it accessible

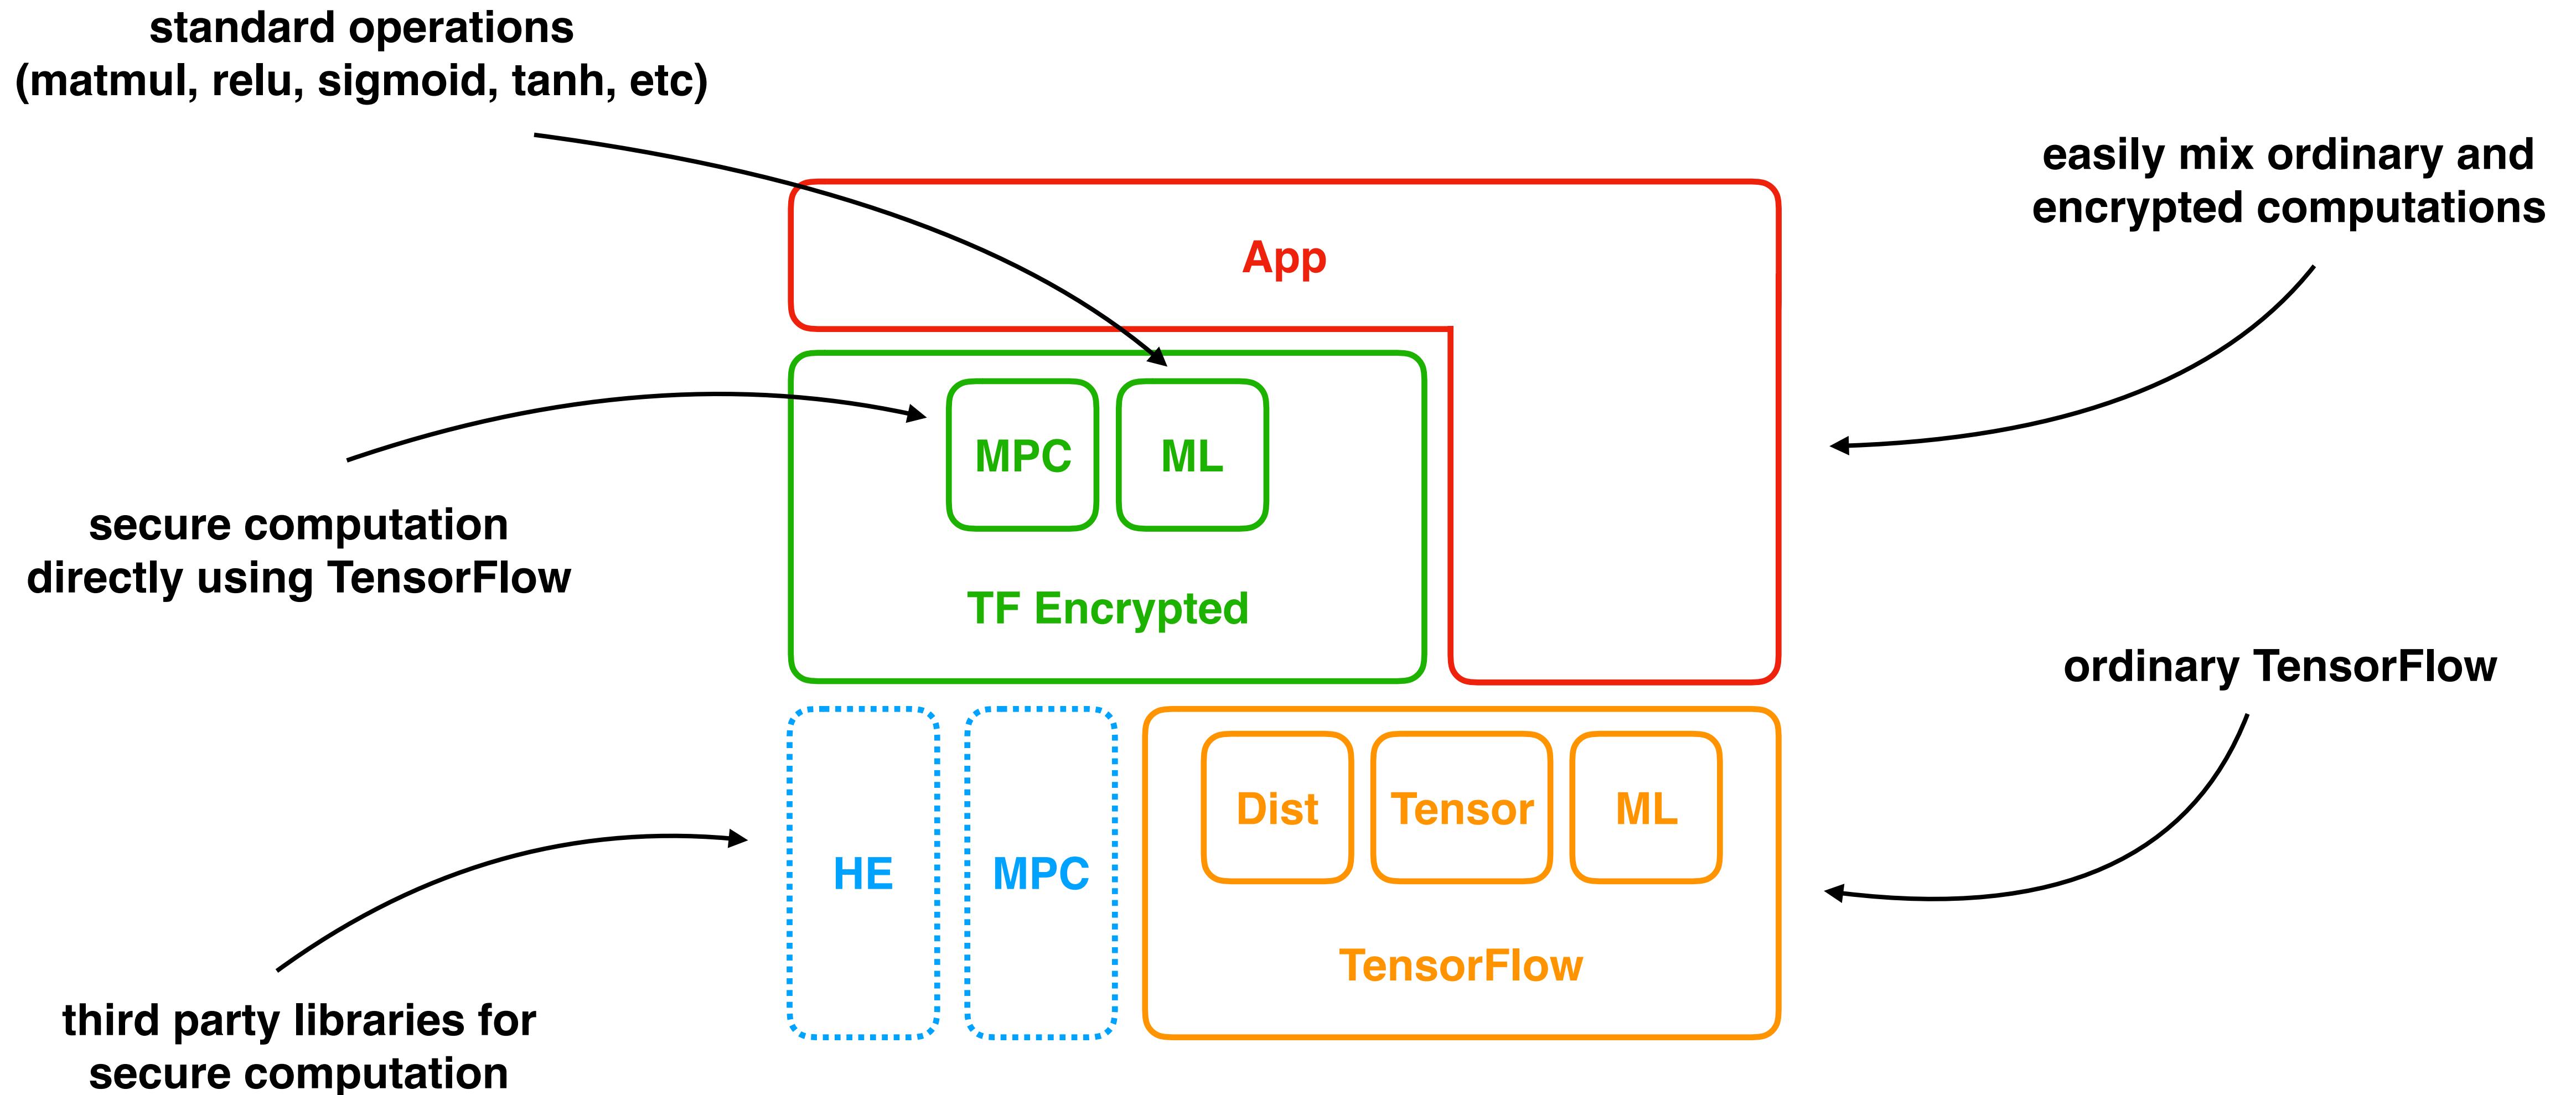
TensorFlow

platform for research and production-level training and deployment

popular and backed by Google



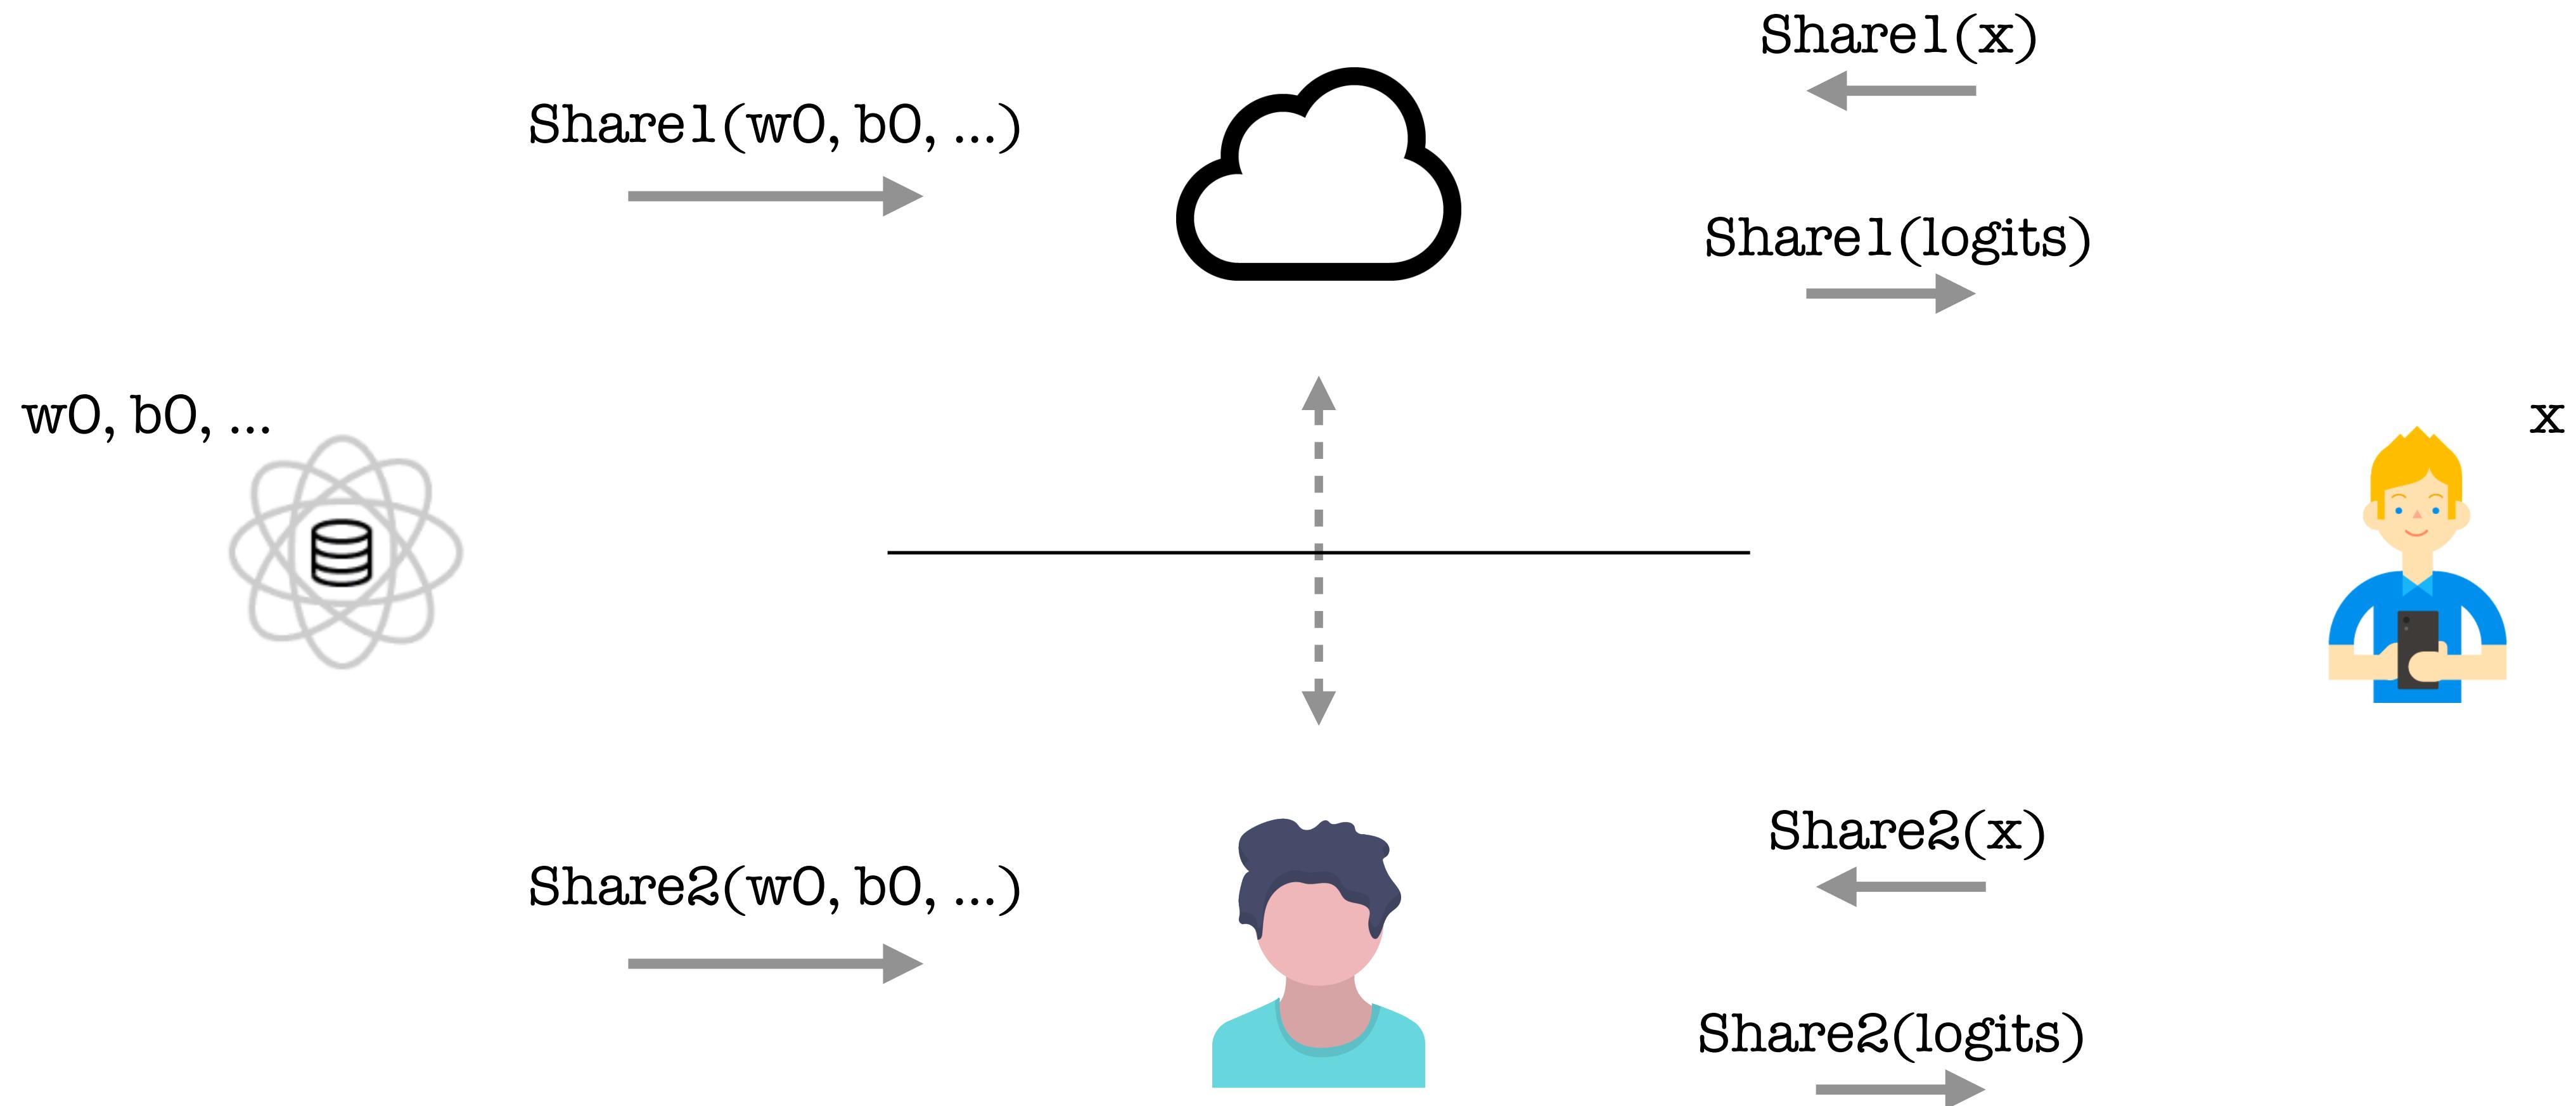
TF Encrypted Architecture



Prediction

Encouraging use

Participants

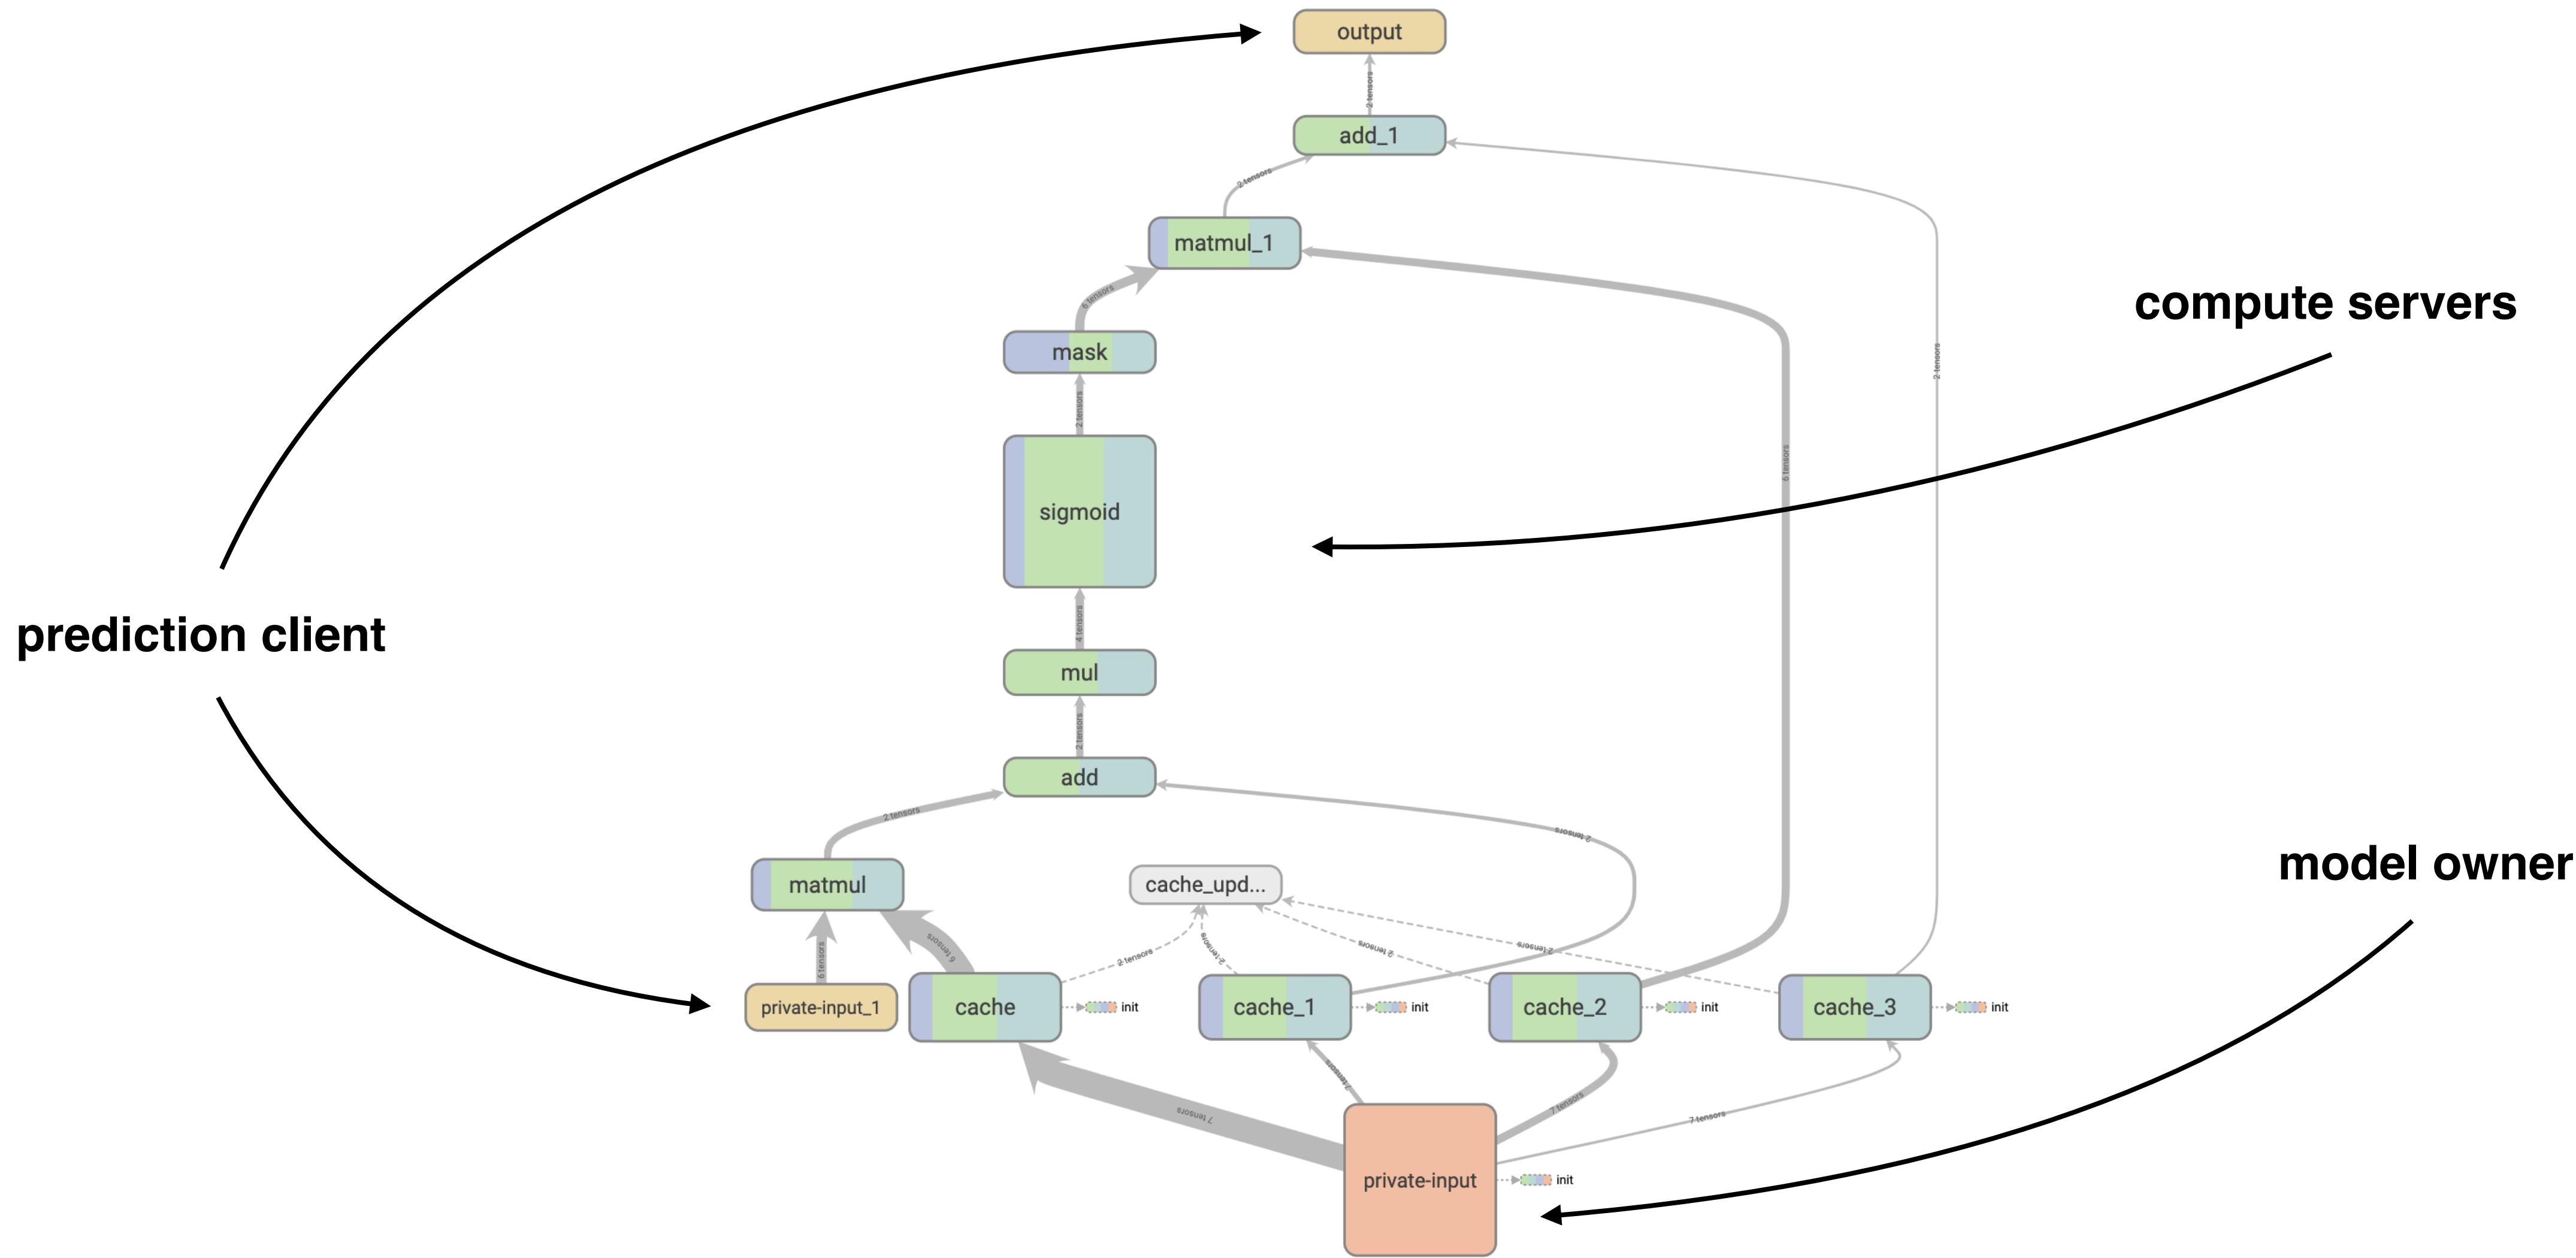


Private Prediction with TF Encrypted

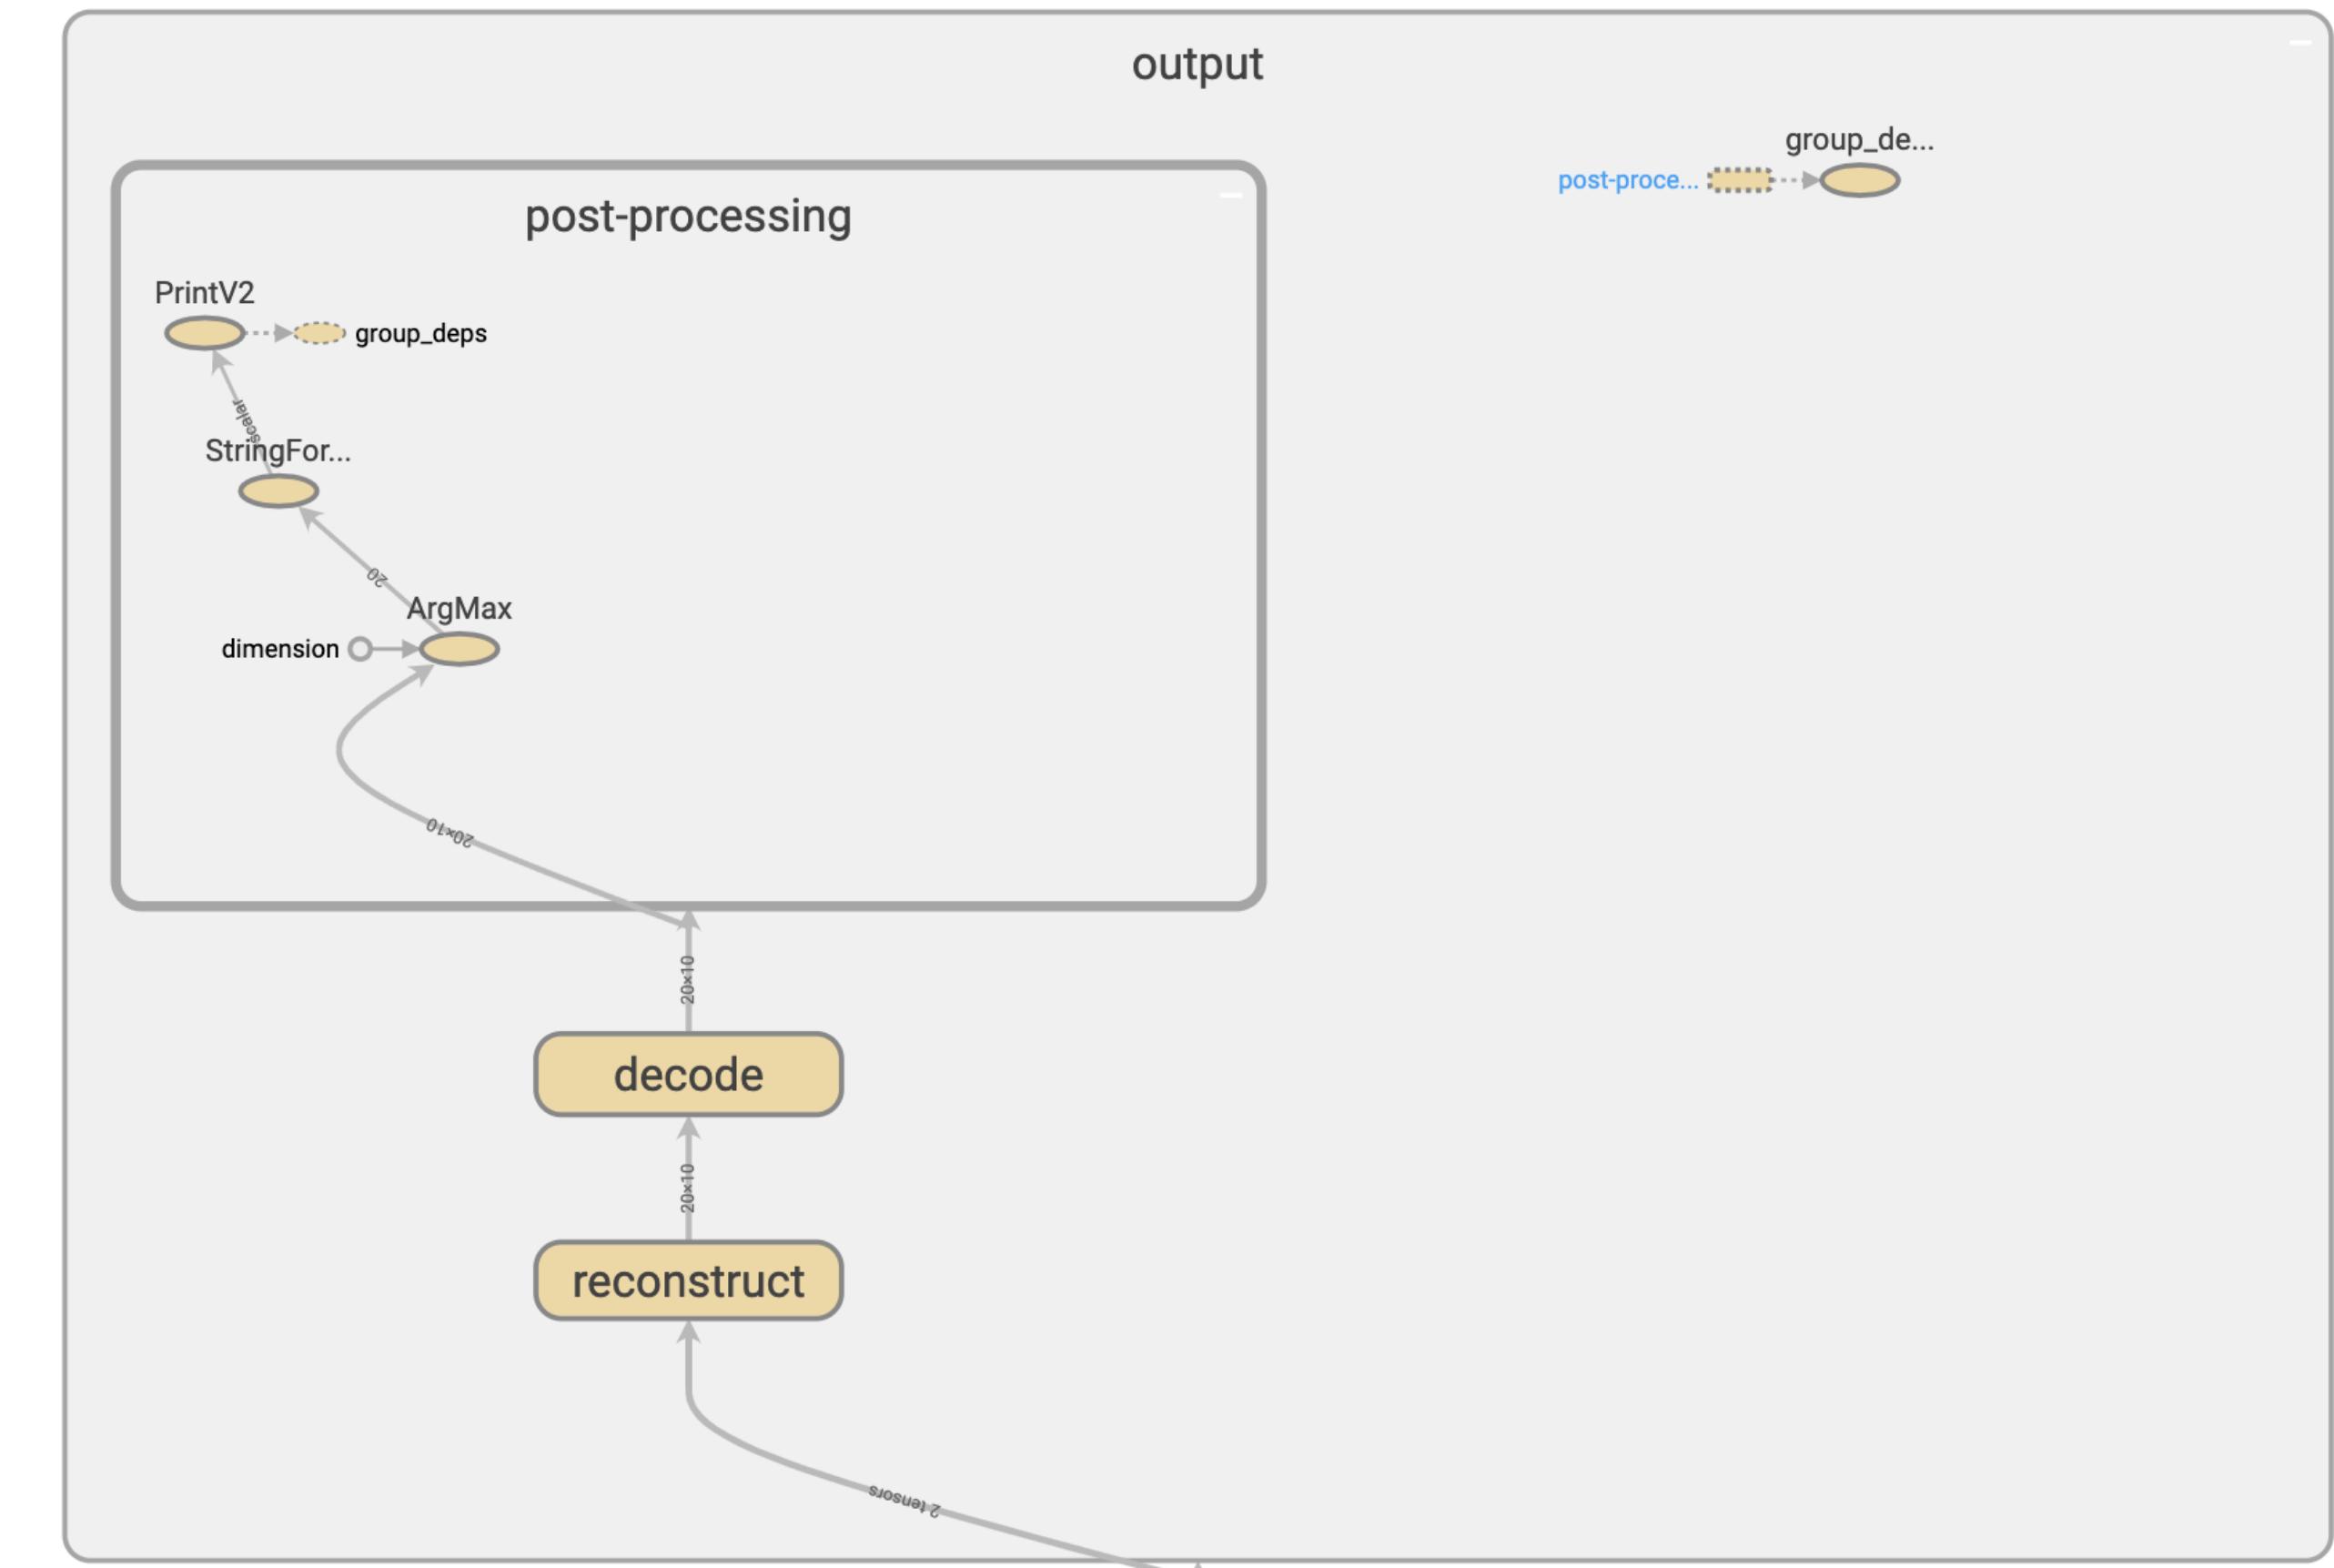
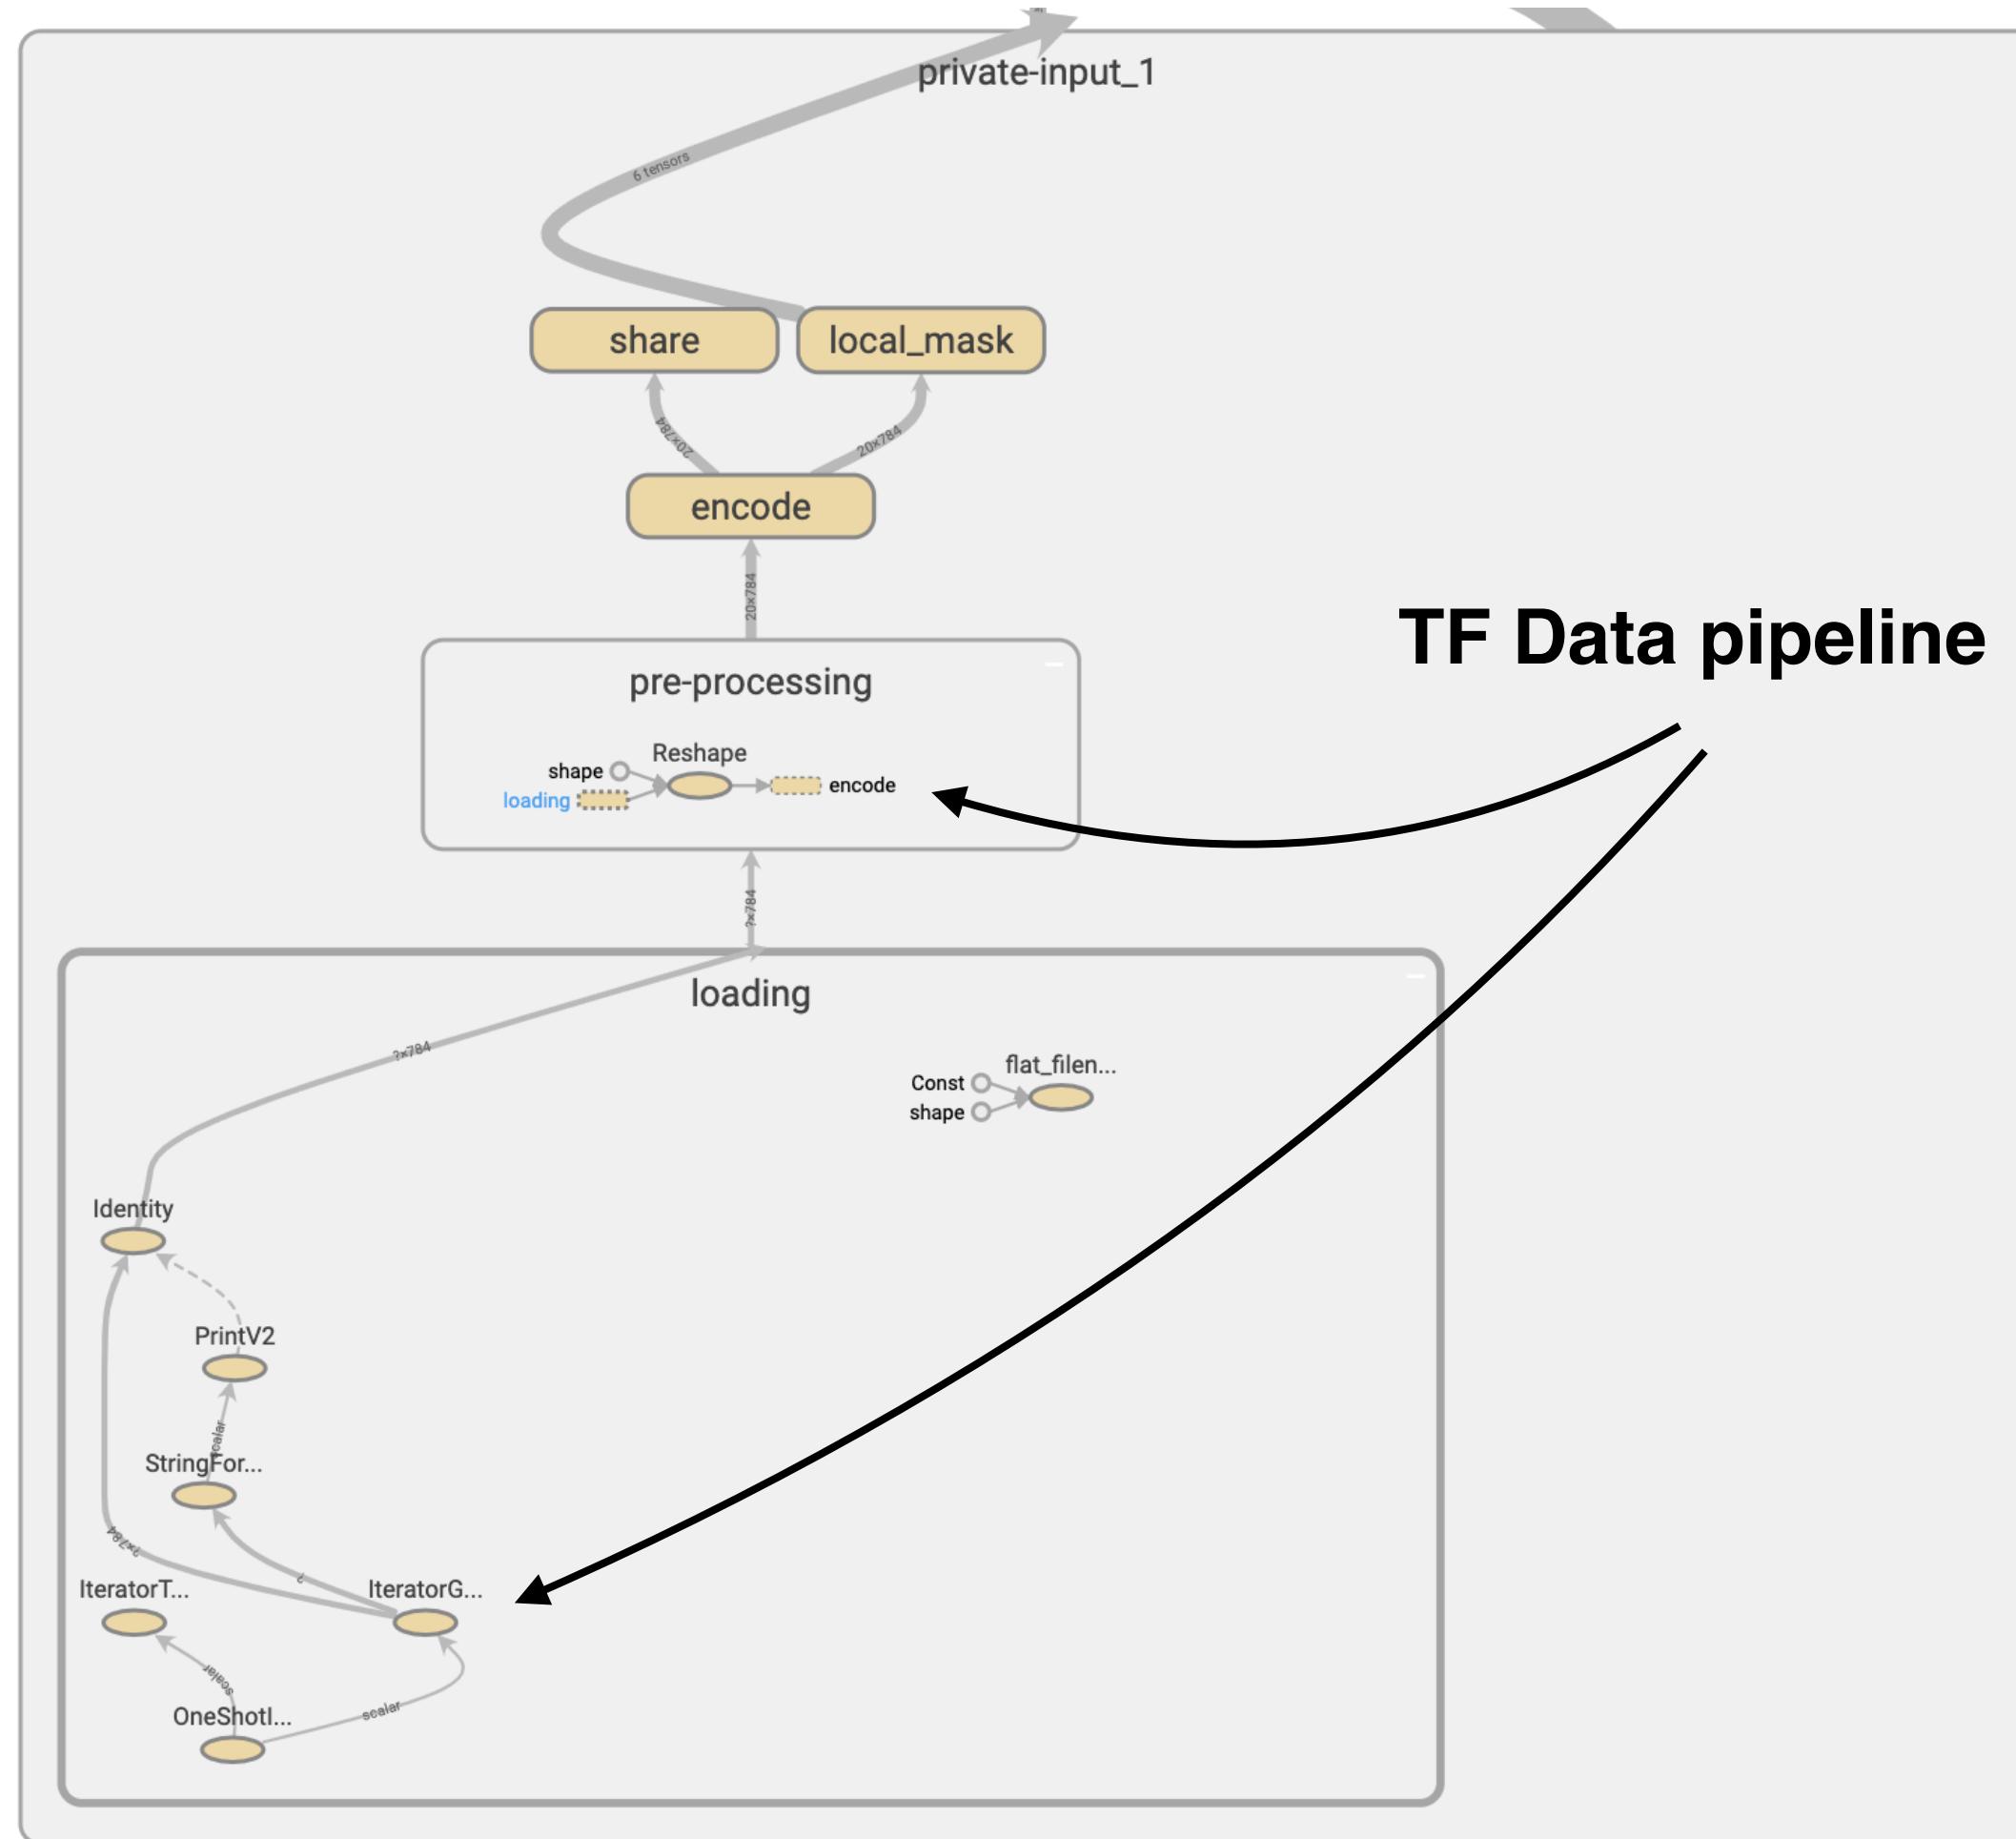
```
1 import tensorflow as tf
2
3 def provide_weights(): """ Load from disk """
4 def provide_input(): """ Pre-process """
5 def receive_output(logits): return tf.Print([], [tf.argmax(logits)])
6
7 # get model weights
8 w0, b0, w1, b1, w2, b2 = provide_weights()
9
10 # get prediction input
11 x = provide_input()
12
13 # compute prediction
14 layer0 = tf.nn.relu(tf.matmul(x, w0) + b0)
15 layer1 = tf.nn.relu(tf.matmul(layer0, w1) + b1)
16 logits = tf.matmul(layer2, w2) + b2
17
18 # process result of prediction
19 prediction_op = receive_output(logits)
20
21 # run graph execution in a tf.Session
22 with tf.Session() as sess:
23     sess.run(tf.global_variables_initializer())
24     sess.run(prediction_op)
```

```
1 import tensorflow as tf
2 import tf_encrypted as tfe
3
4 def provide_weights(): """ Load from disk """
5 def provide_input(): """ Pre-process """
6 def receive_output(logits): return tf.Print([], [tf.argmax(logits)])
7
8 # get model weights as private tensors from owner
9 w0, b0, w1, b1, w2, b2 = tfe.define_private_input("model-owner", provide_weights)
10
11 # get prediction input as private tensors from client
12 x = tfe.define_private_input("prediction-client", provide_input)
13
14 # compute private prediction on servers
15 layer0 = tfe.relu(tfe.matmul(x, w0) + b0)
16 layer1 = tfe.relu(tfe.matmul(layer0, w1) + b1)
17 logits = tfe.matmul(layer1, w2) + b2
18
19 # process result of prediction on client
20 prediction_op = tfe.define_output("prediction-client", logits, receive_output)
21
22 # run secure graph execution in a tf.Session
23 with tfe.Session() as sess:
24     sess.run(tf.global_variables_initializer())
25     sess.run(prediction_op)
```

Overall Computation



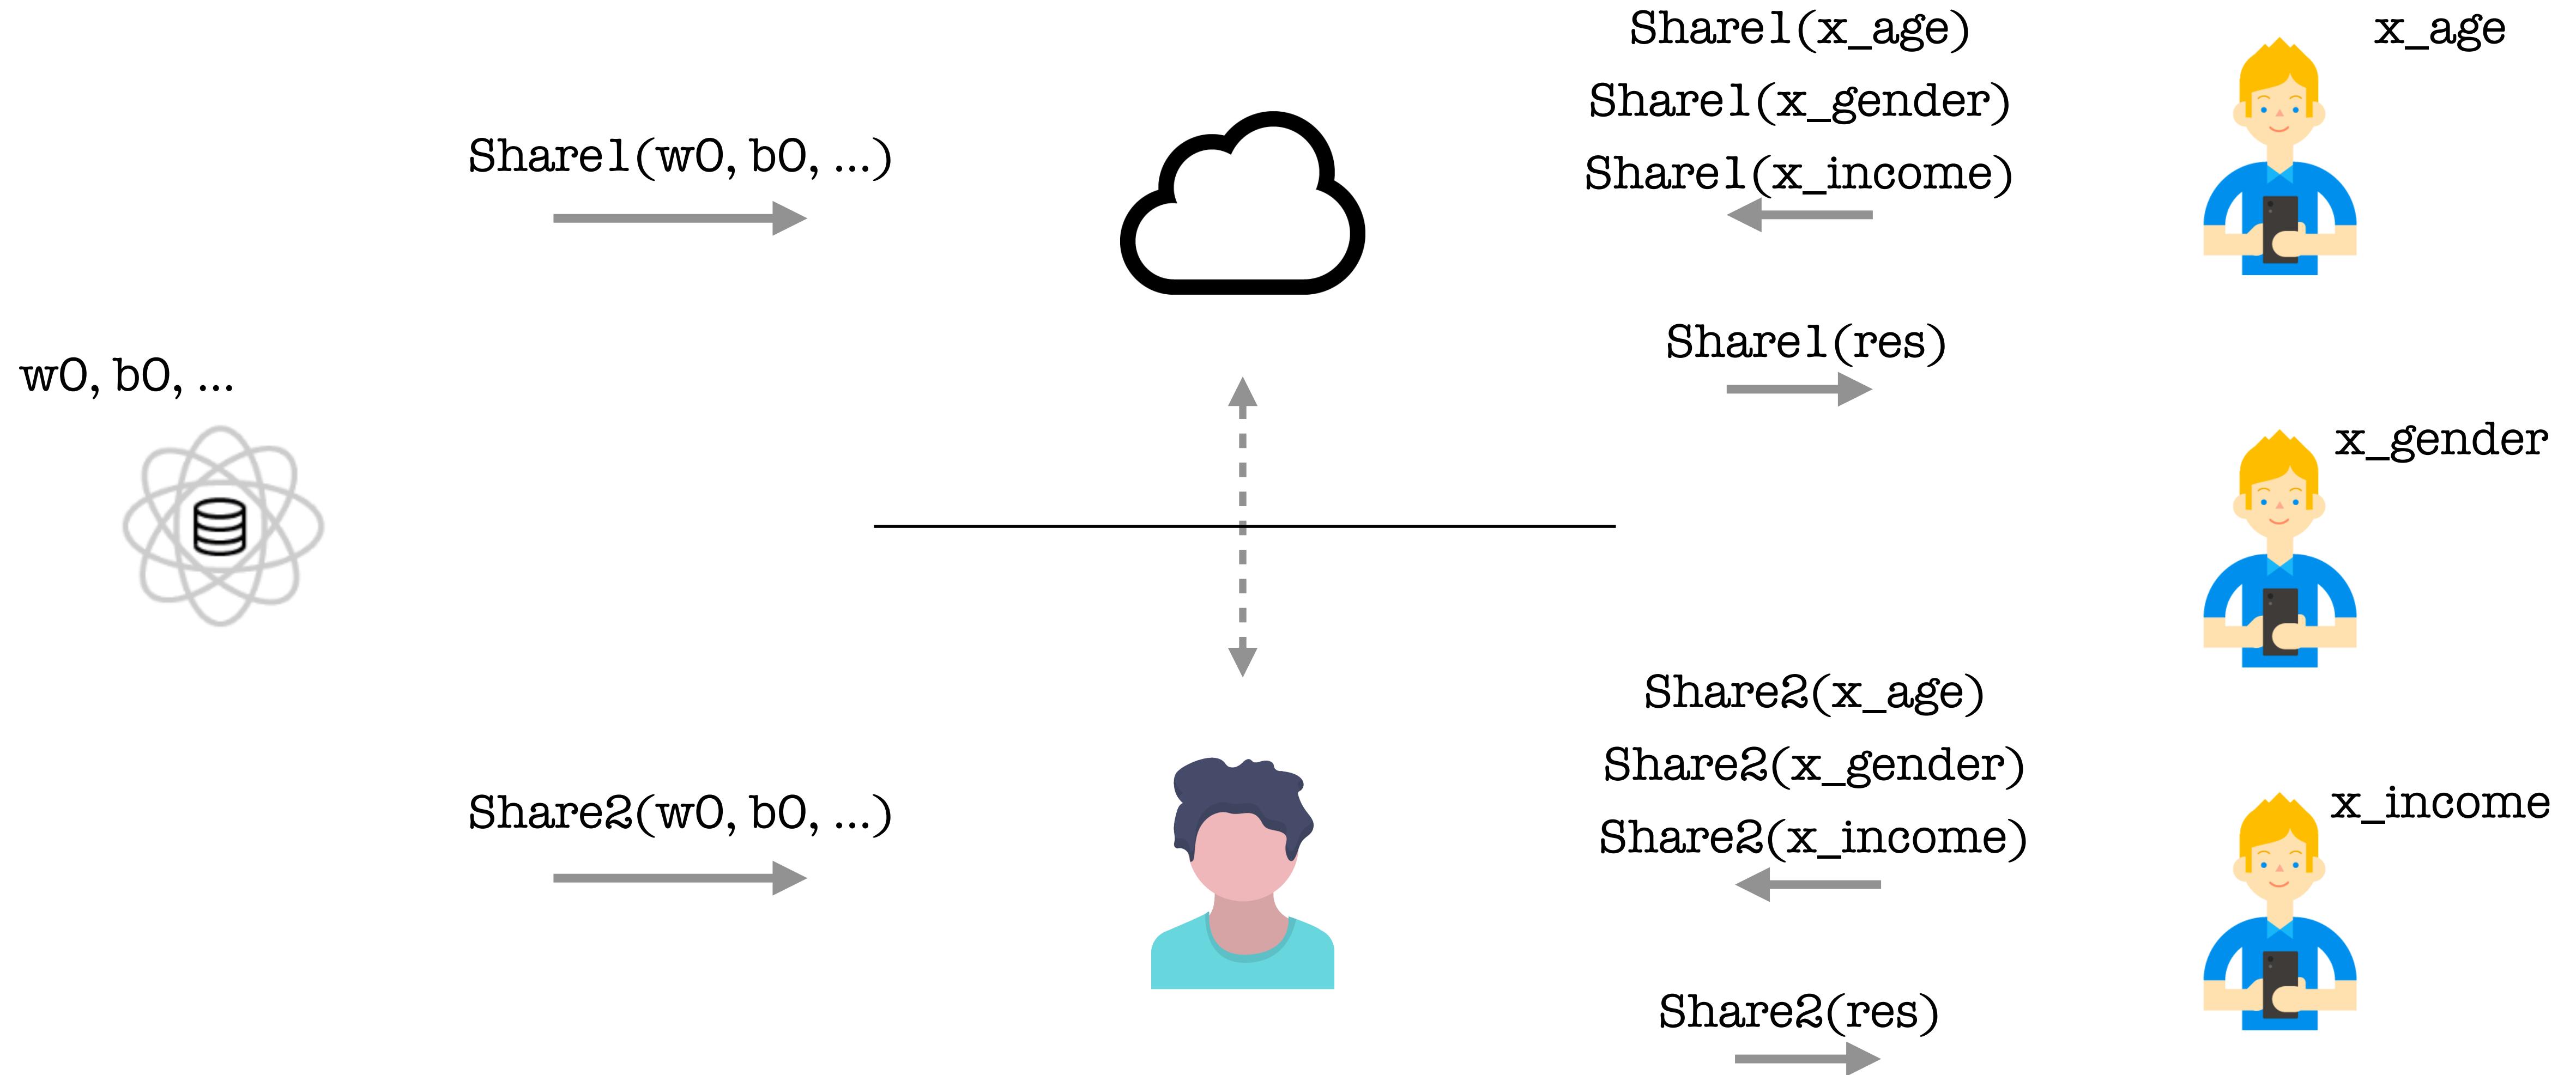
Local Processing



Joint Prediction

Combining knowledge for nuance

Participants



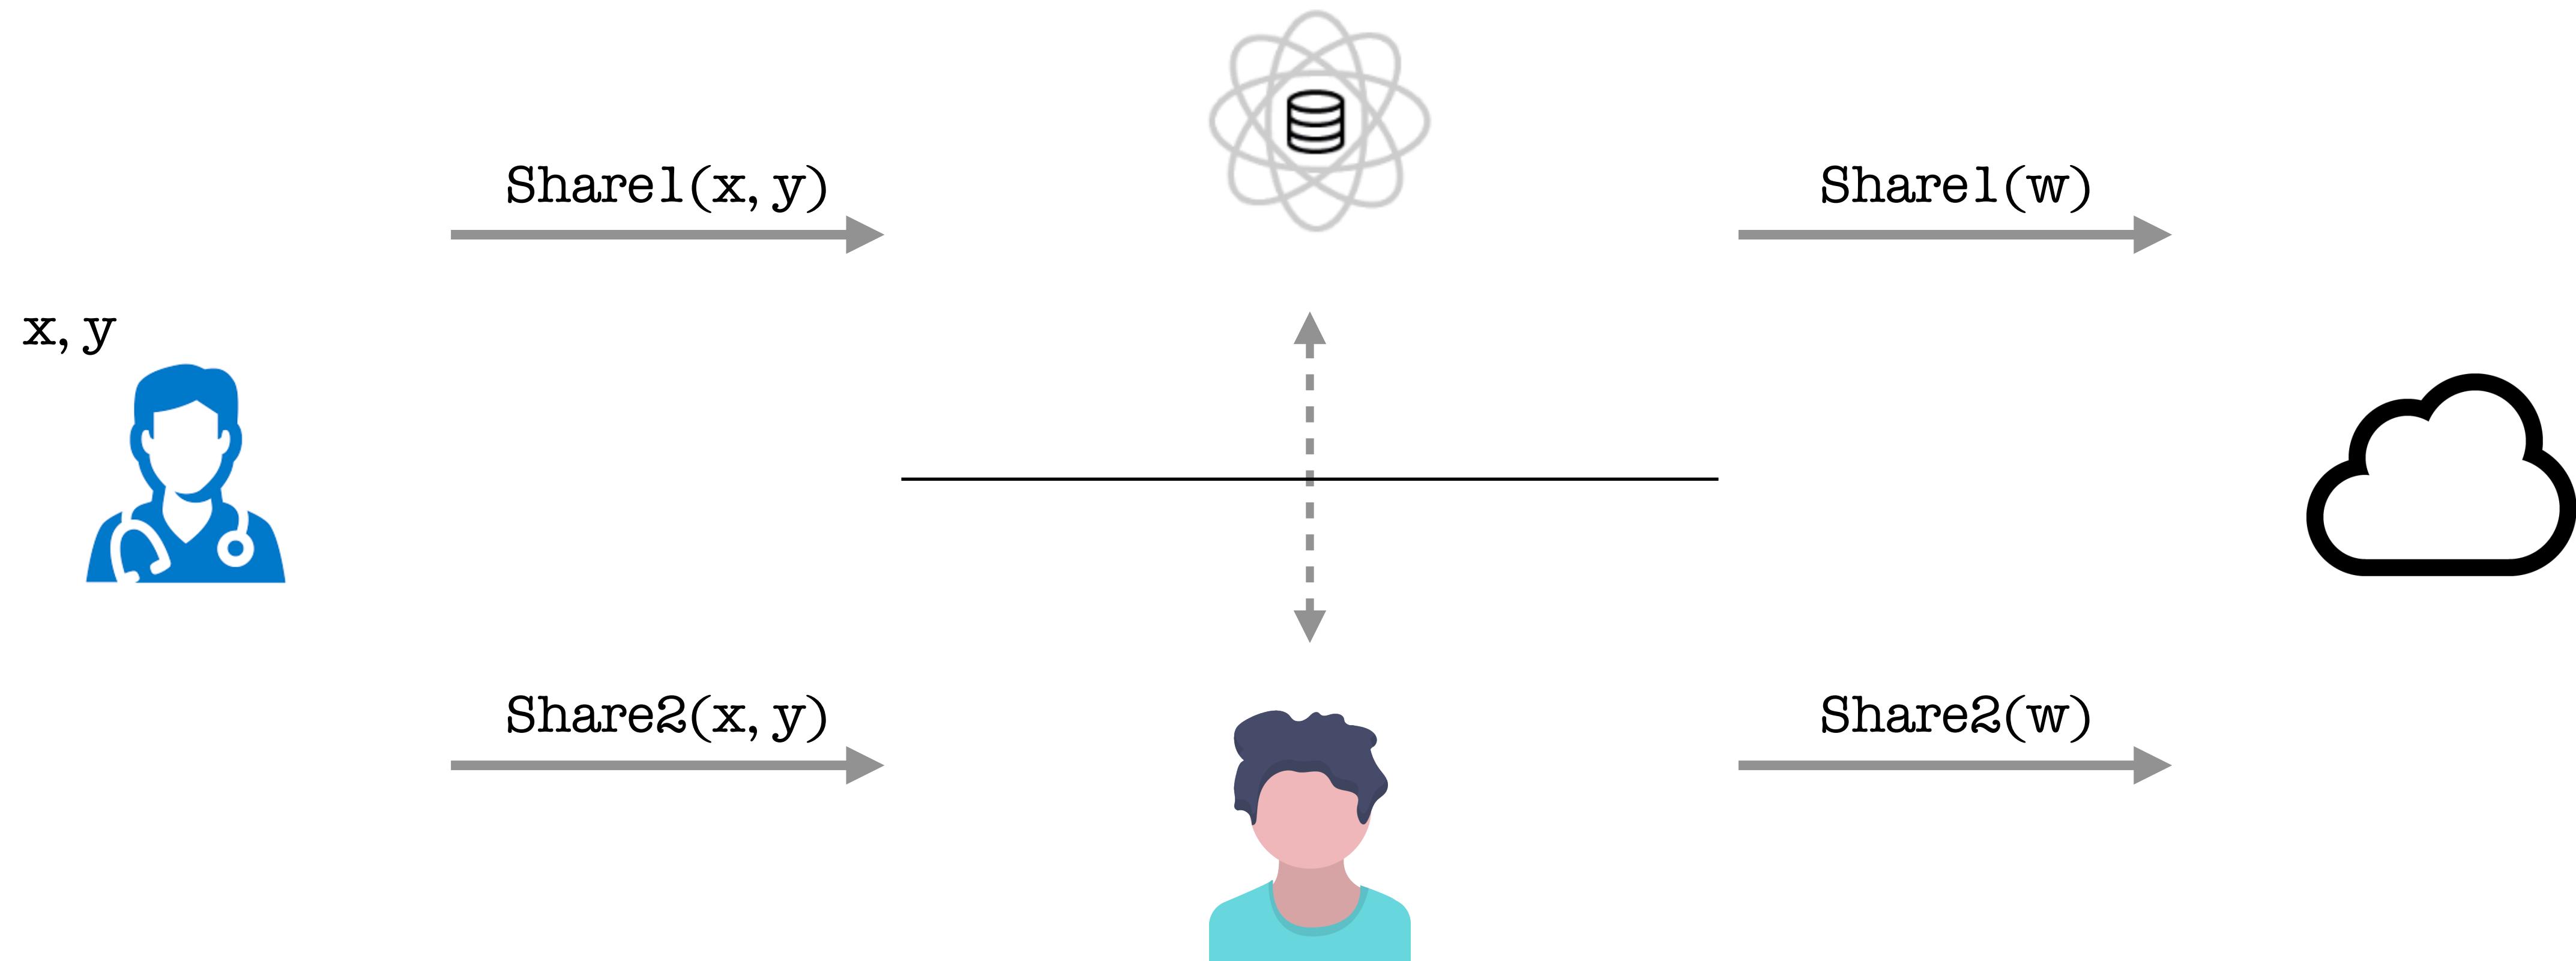
Private Joint Prediction with TF Encrypted

```
1 w, b = tfe.define_private_input("model-owner", provide_weights)
2
3 x_age = tfe.define_private_input("age-provider", provide_age)
4 x_gender = tfe.define_private_input("gender-provider", provide_gender)
5 x_income = tfe.define_private_input("income-provider", provide_income)
6
7 x = tfe.concat([x_age, x_gender, x_income], axis=1)
8
9 y = tfe.matmul(x, w) + b
10
11 prediction_op = tfe.define_output("result-receiver", y, receive_output)
12
13 with tfe.Session() as sess:
14     sess.run(prediction_op)
```

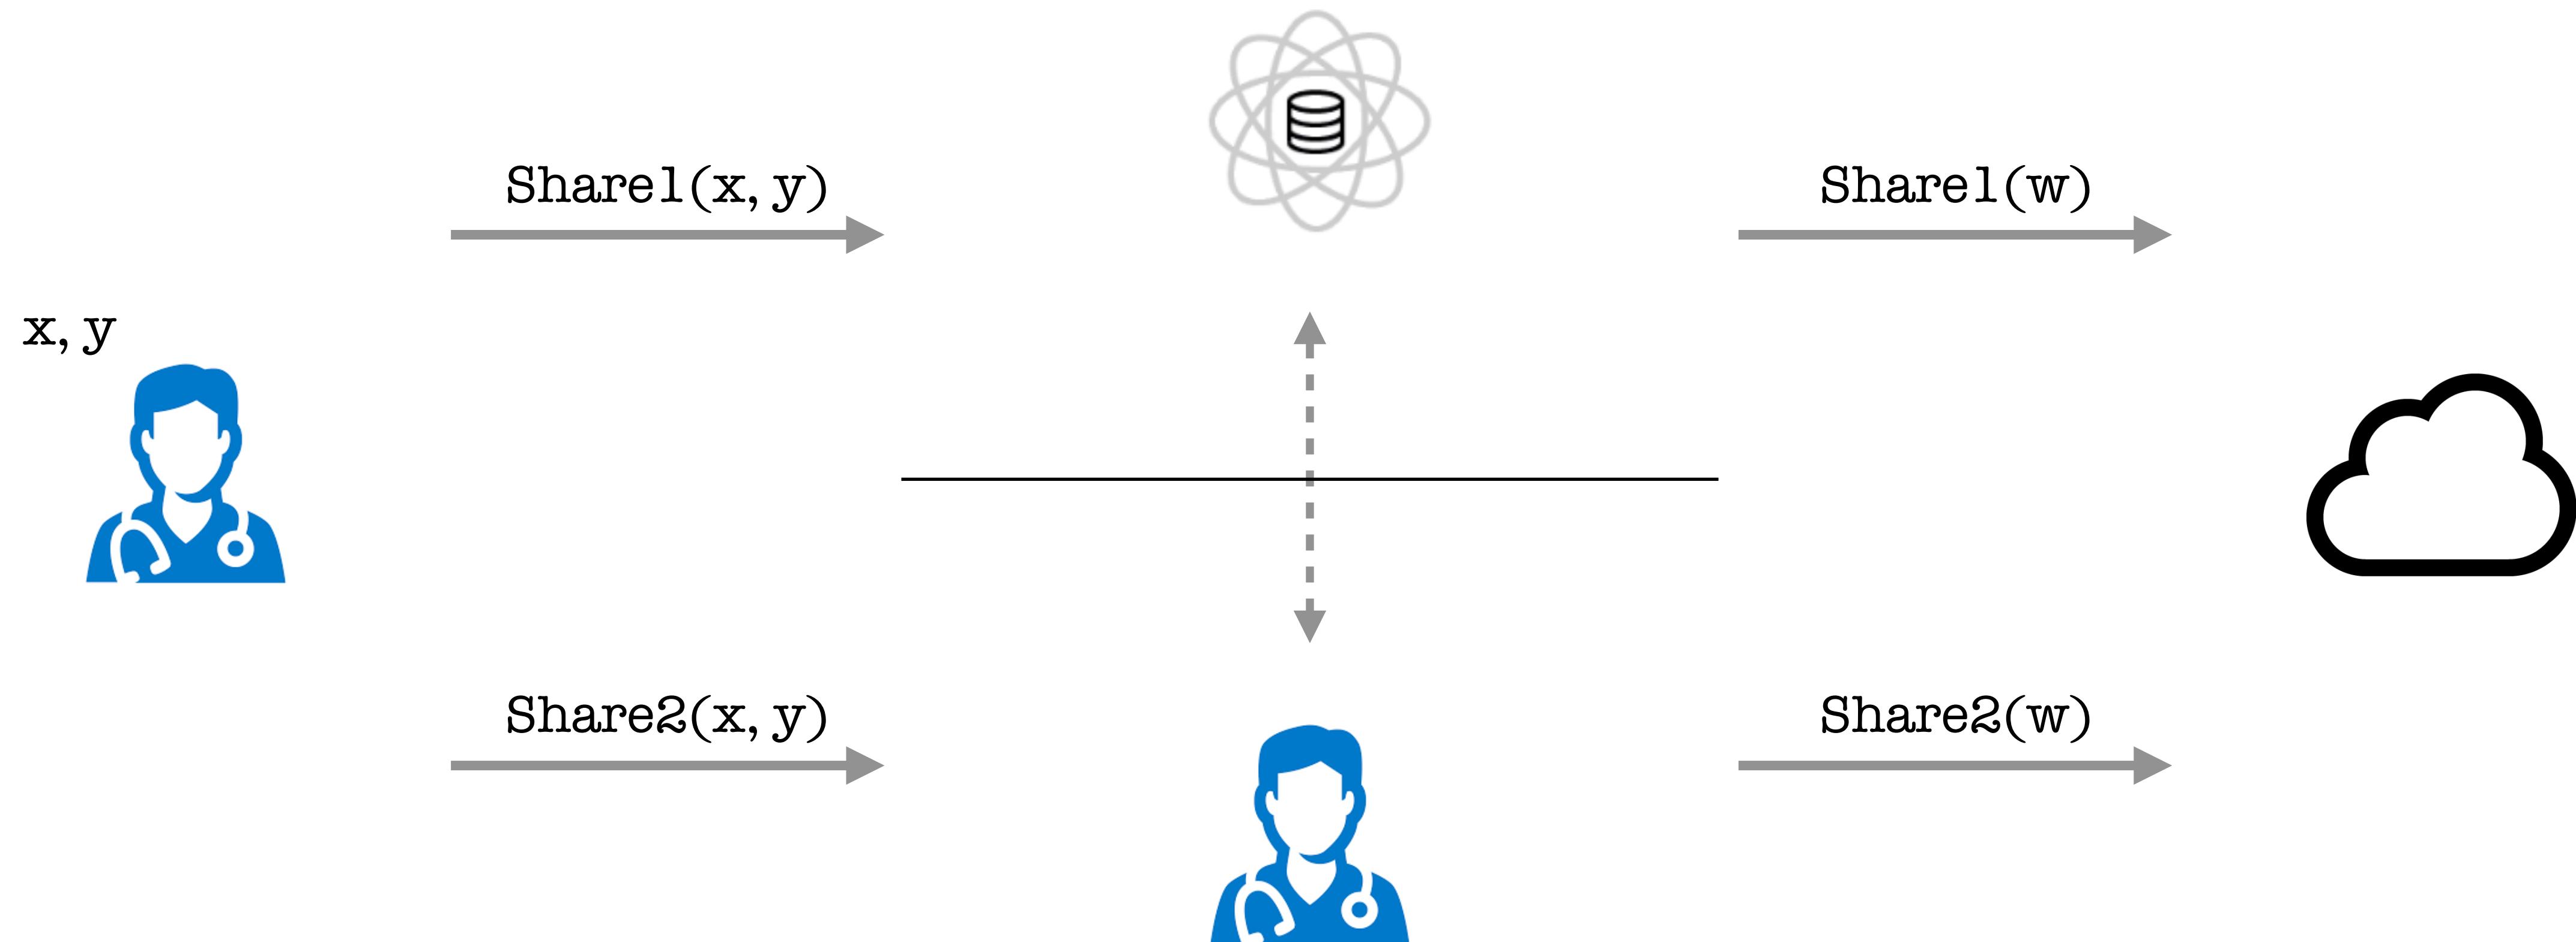
Training

Learning without seeing

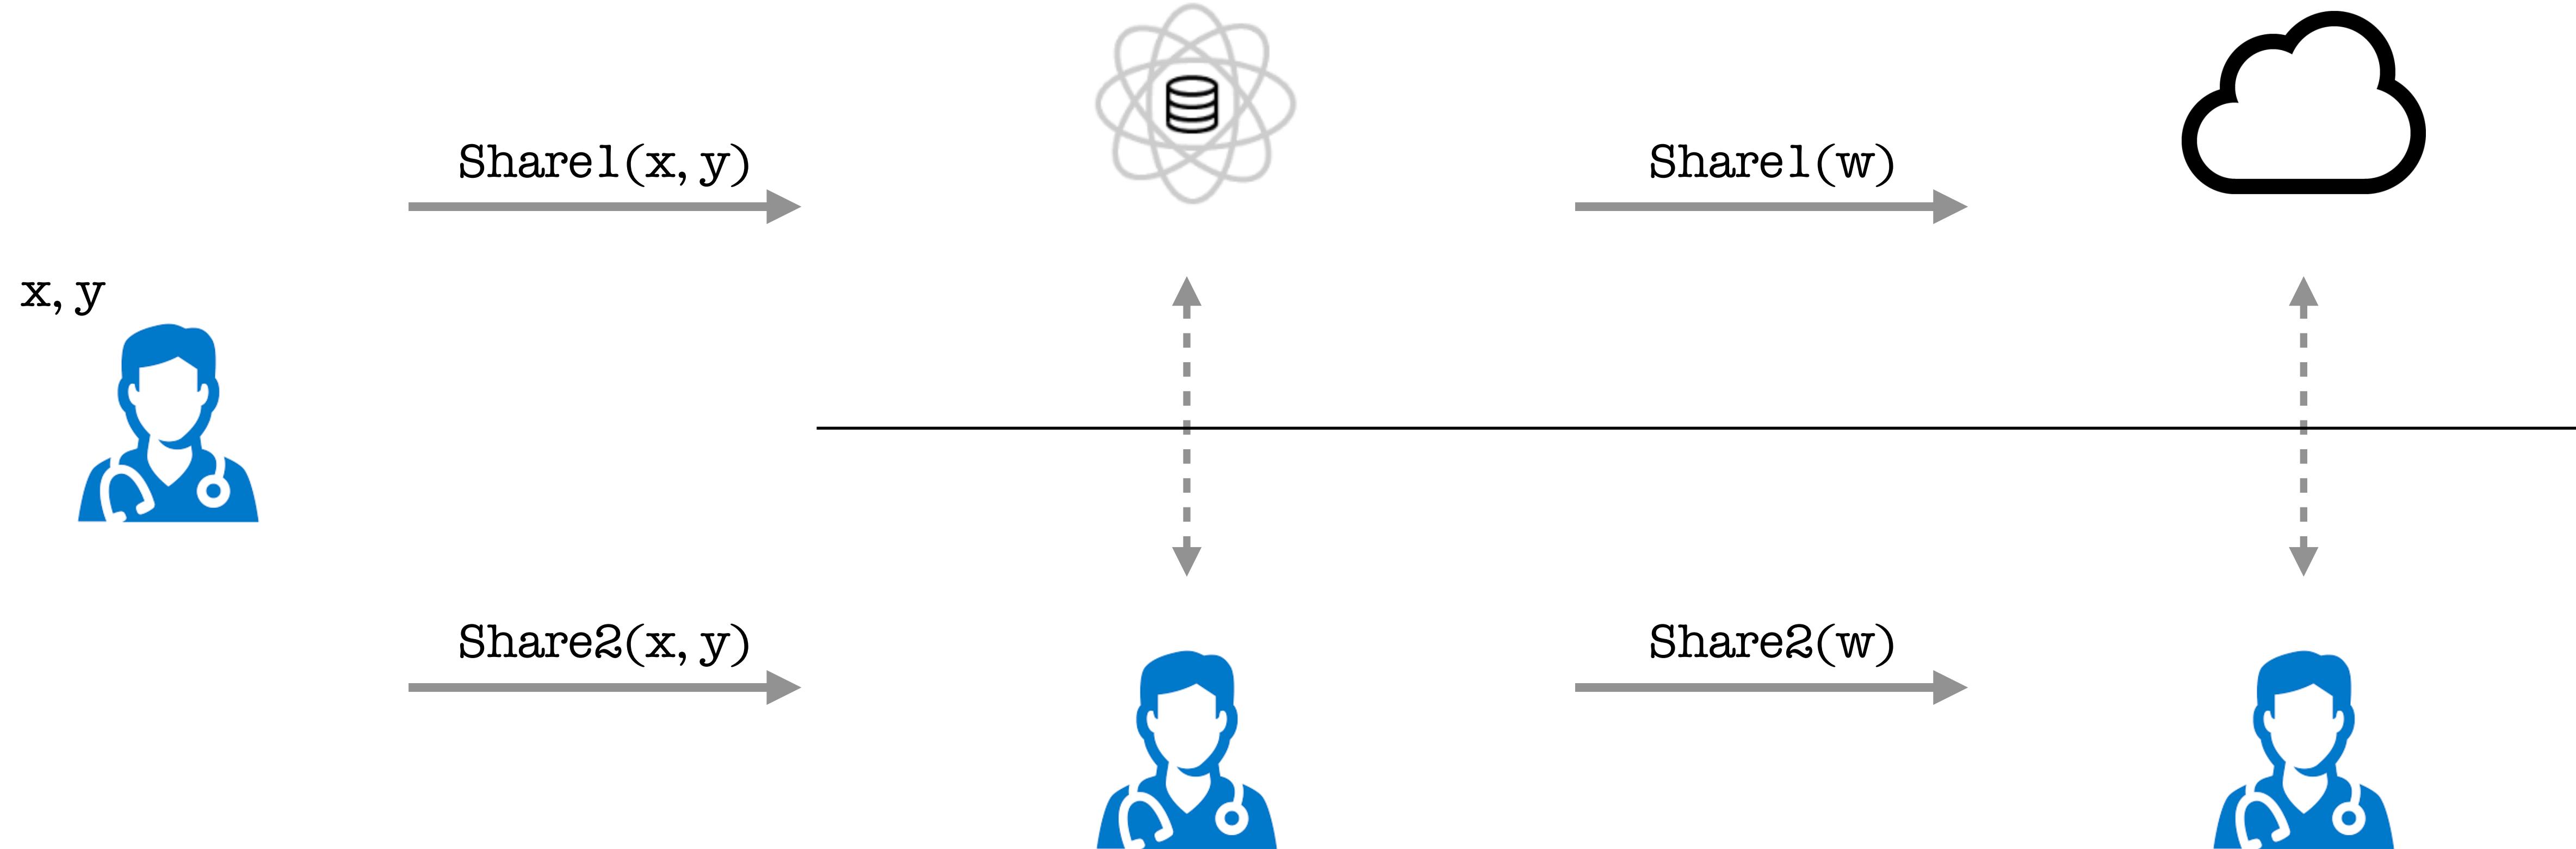
Participants



Participants



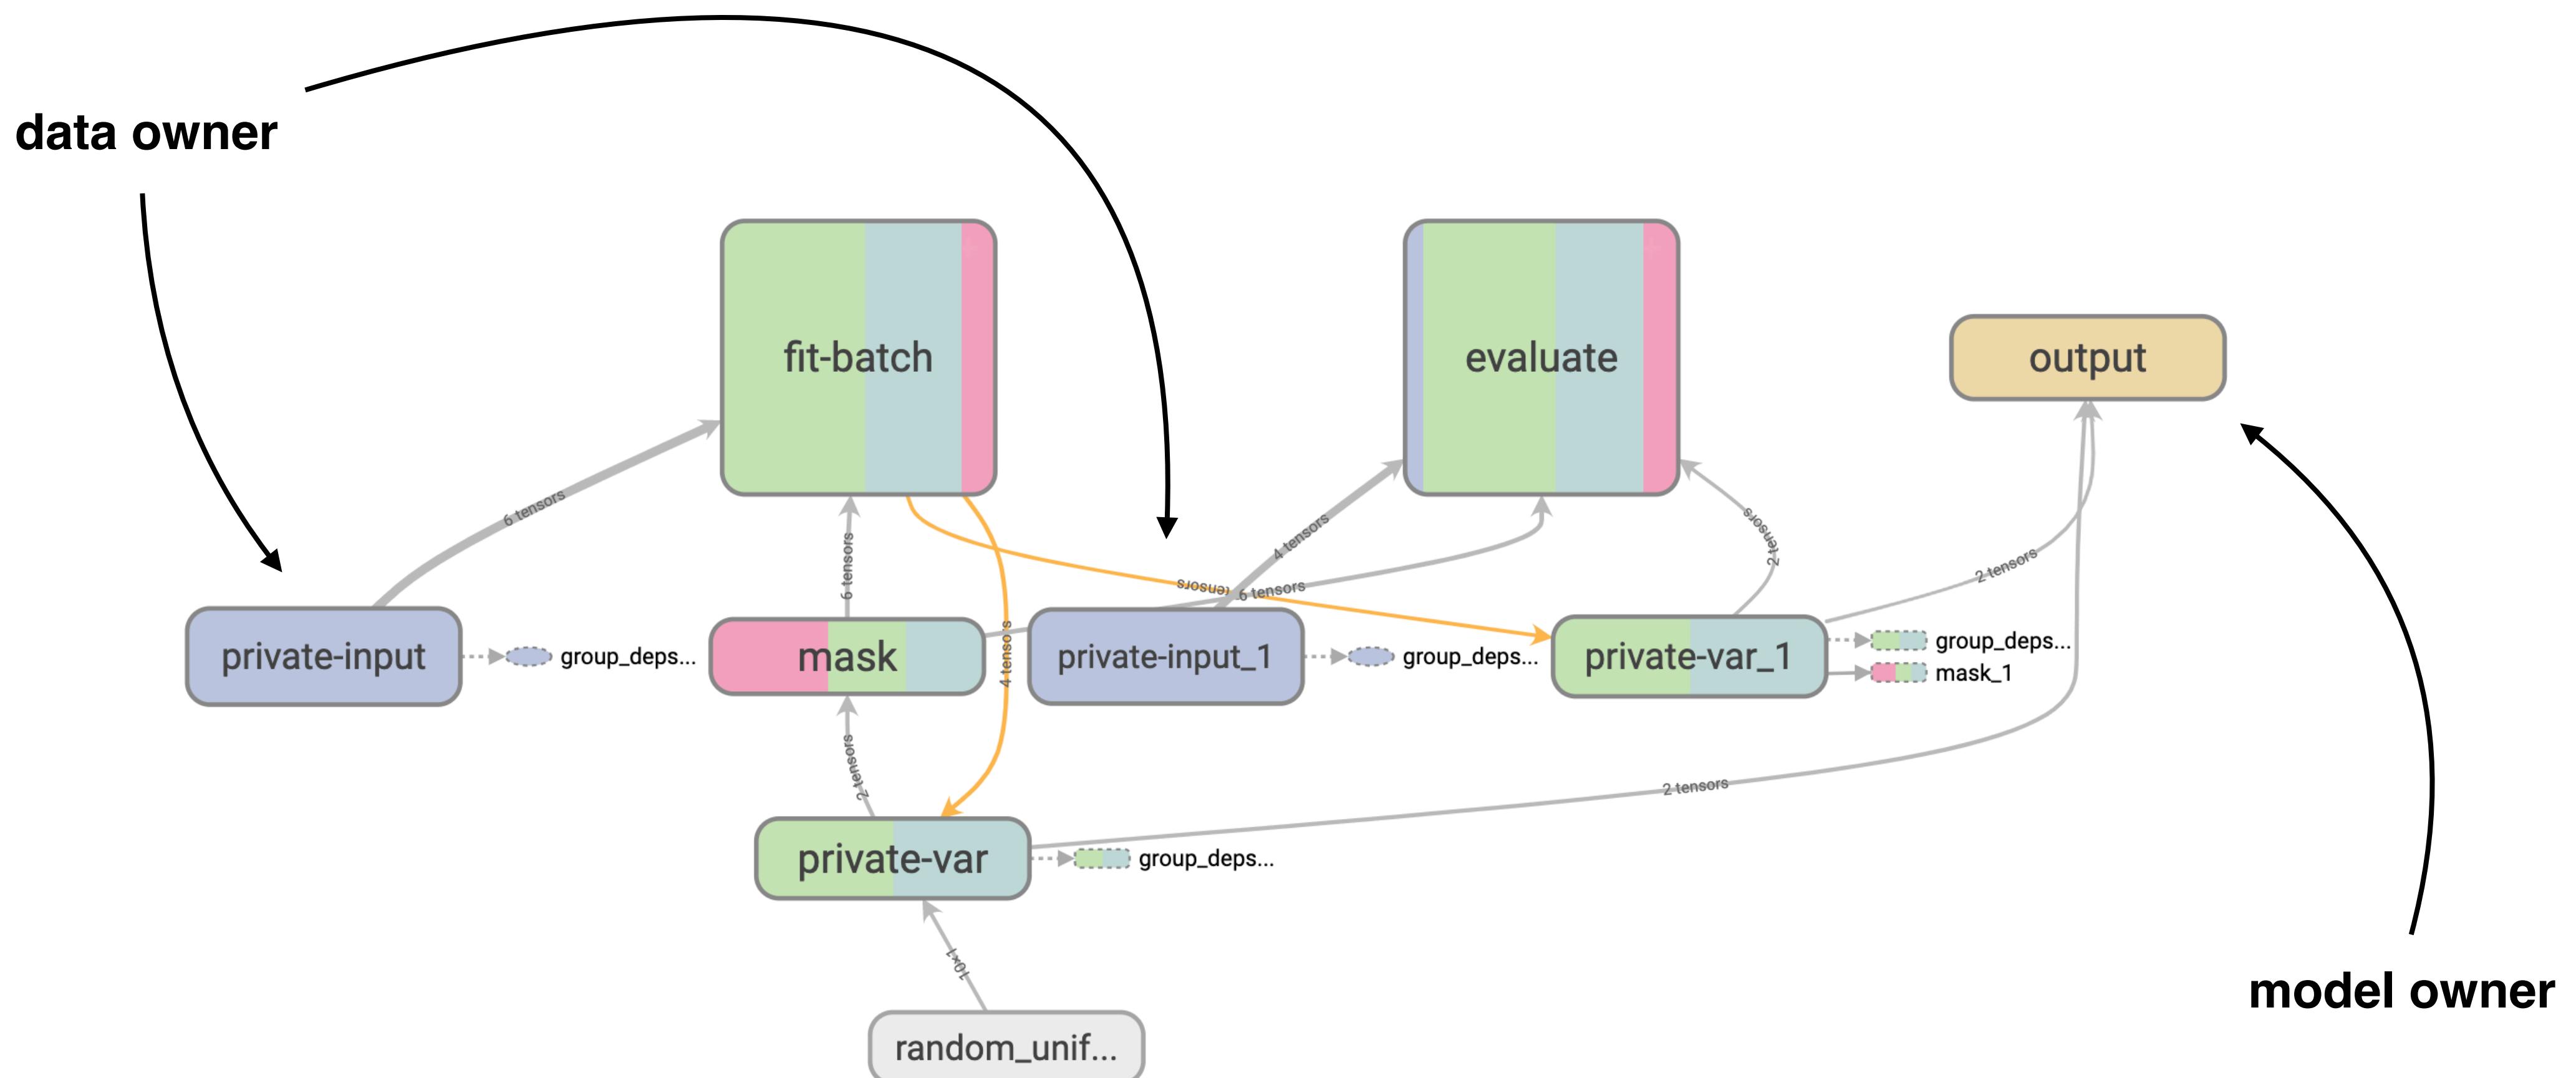
Participants



Private Training with TF Encrypted

```
1  class LogisticRegression(Model):
2
3      @property
4      def weights(self):
5          # ...
6
7      def forward(self, x):
8          # ...
9
10     def backward(self, x, dy, learning_rate=0.01):
11         # ...
12
13     def loss_grad(self, y, y_hat):
14         # ...
15
16 model = LogisticRegression()
17
18 data_owner = DataOwner('data-owner')
19 model_owner = ModelOwner('model-owner')
20
21 x_train, y_train = tfe.define_private_input(data_owner.player_name, data_owner.provide_training_data)
22 x_test, y_test = tfe.define_private_input(data_owner.player_name, data_owner.provide_testing_data)
23
24 reveal_weights_op = tfe.define_output(model_owner.player_name, model.weights, model_owner.receive_weights)
25
26 with tfe.Session() as sess:
27     sess.run([tfe.global_variables_initializer(), data_owner.initializer])
28
29     model.fit(sess, x_train, y_train, epochs=10)
30     model.evaluate(sess, x_test, y_test, data_owner)
31
32     sess.run(reveal_weights_op)
```

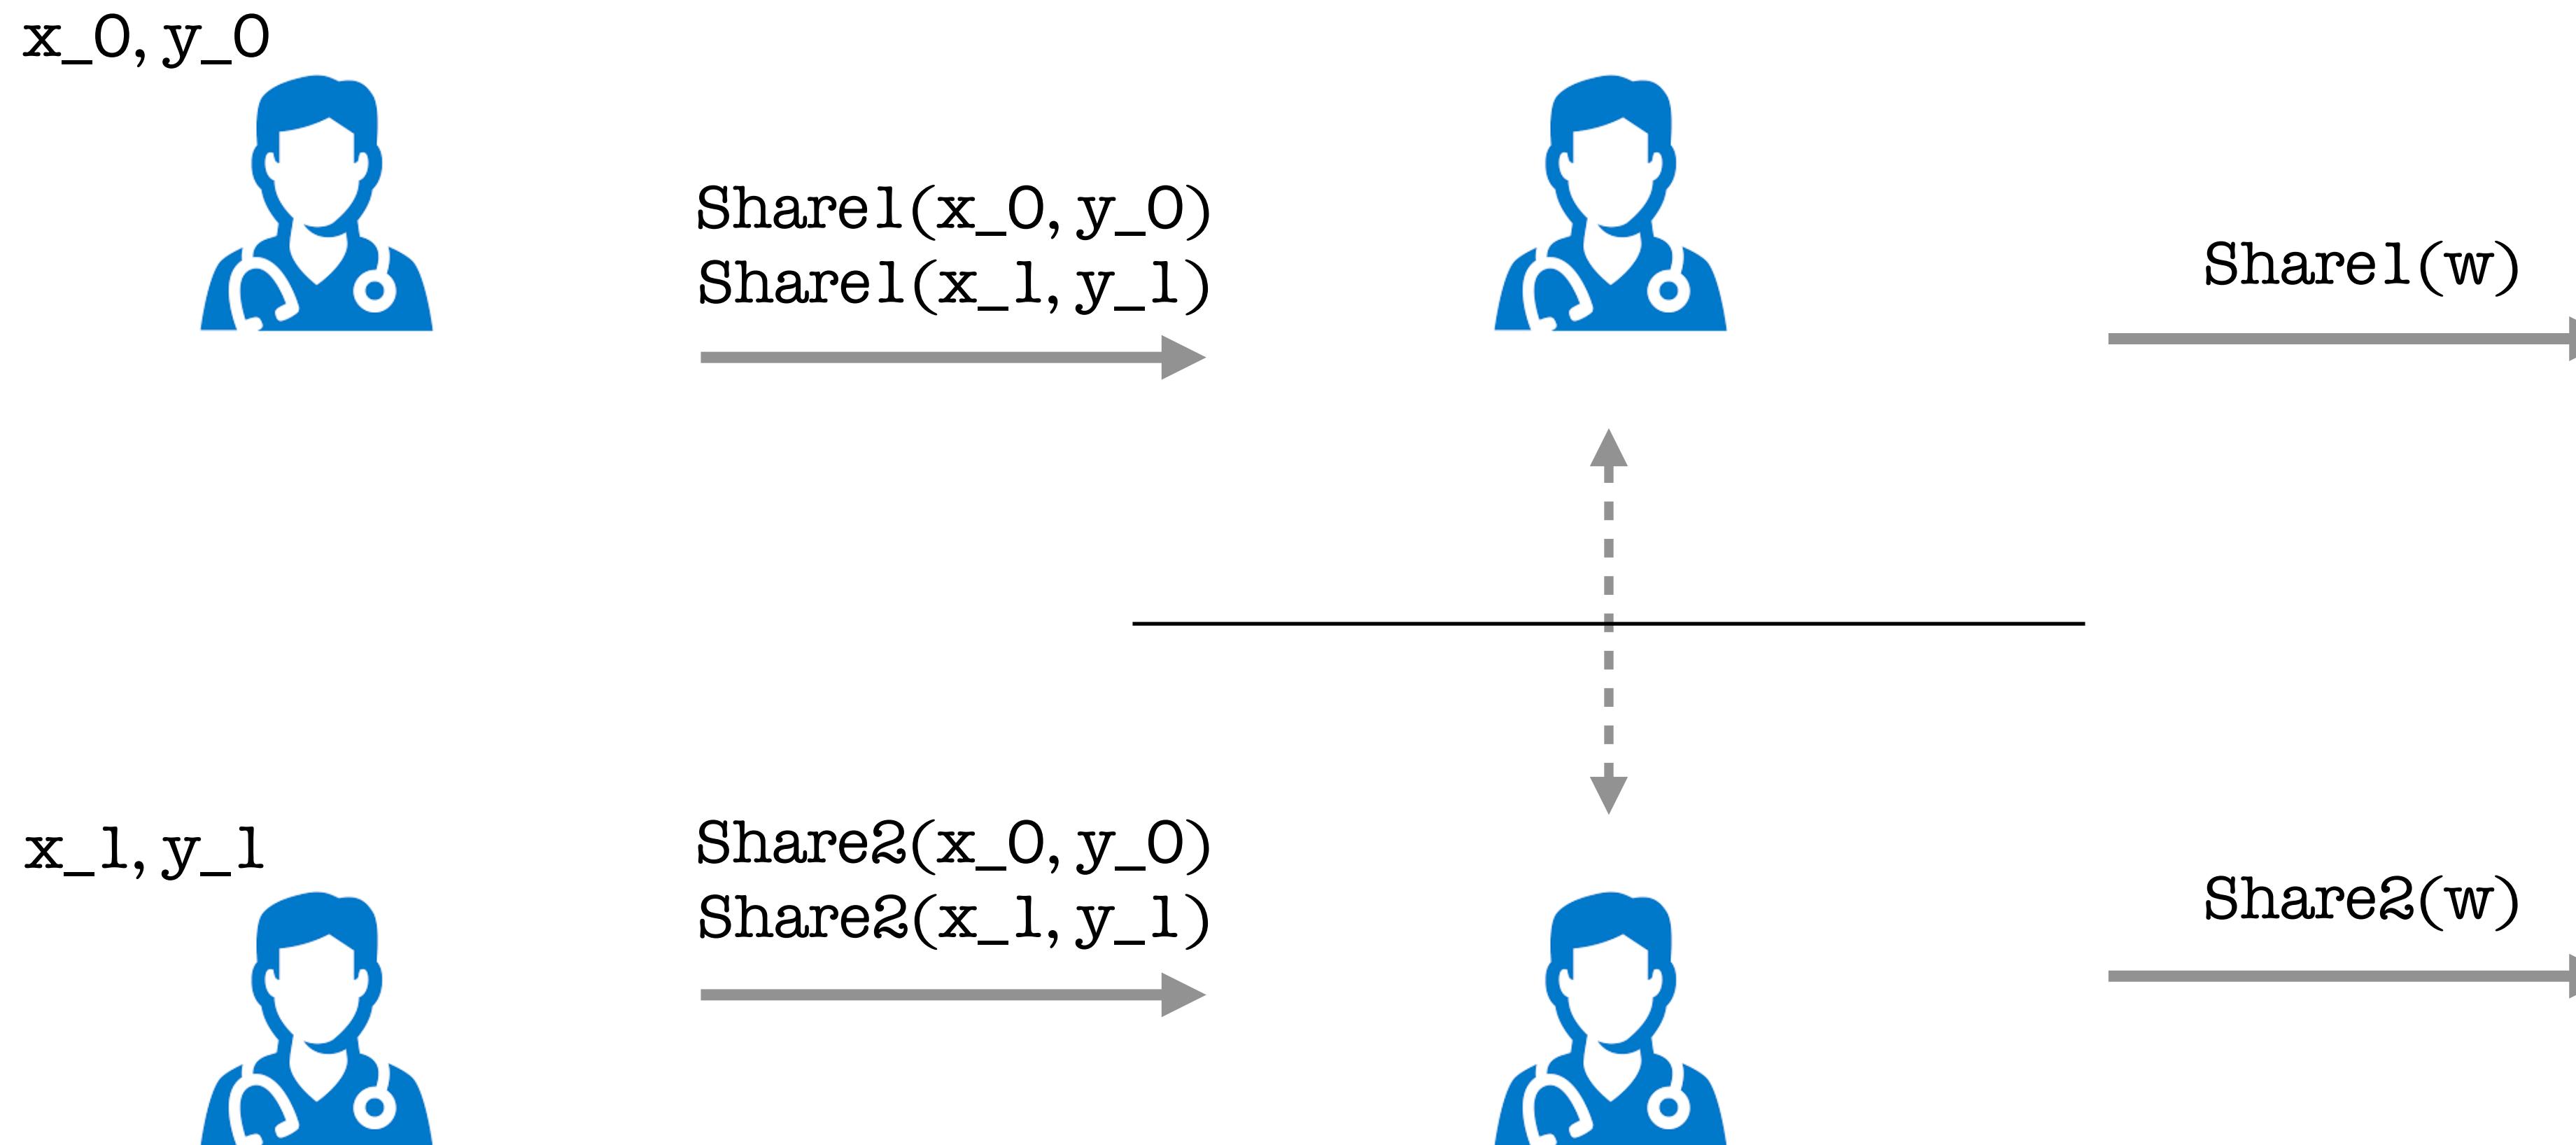
Overall Computation



Joint Training

Combining insights for better models

Participants



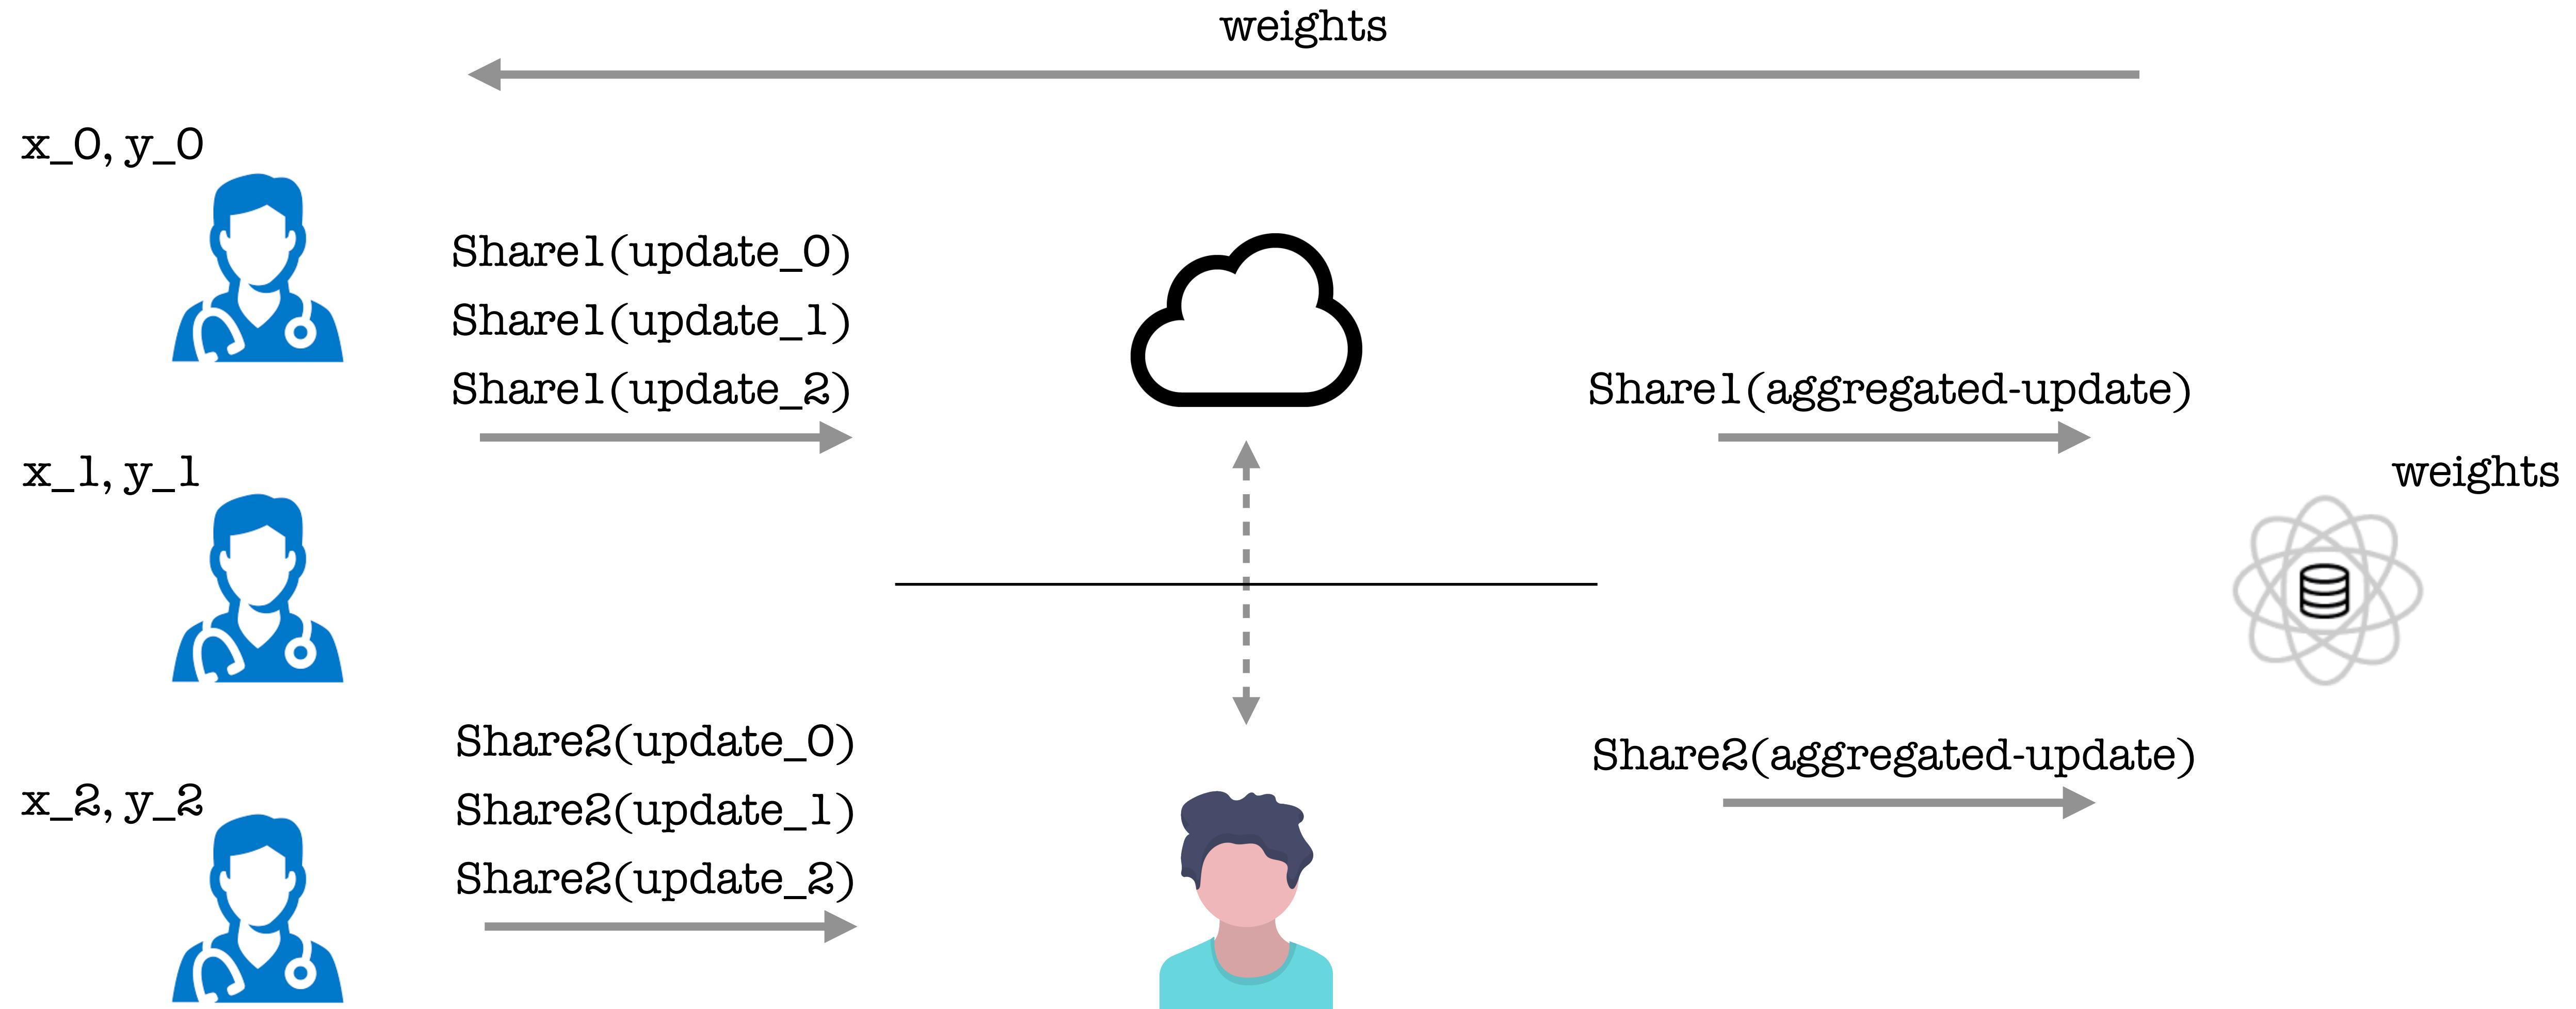
Private Joint Training with TF Encrypted

```
1 data_owner_0 = DataOwner('data-owner-0')
2 data_owner_1 = DataOwner('data-owner-1')
3
4 tfe.set_protocol(tfe.protocol.Pond(data_owner_0.player_name, data_owner_1.player_name))
5
6 x_train_0, y_train_0 = tfe.define_private_input(data_owner_0.player_name, data_owner_0.provide_training_data)
7 x_train_1, y_train_1 = tfe.define_private_input(data_owner_1.player_name, data_owner_1.provide_training_data)
8
9 x_train = tfe.concat([x_train_0, x_train_1], axis=0)
10 y_train = tfe.concat([y_train_0, y_train_1], axis=0)
```

Federated Learning

Keeping data decentralized

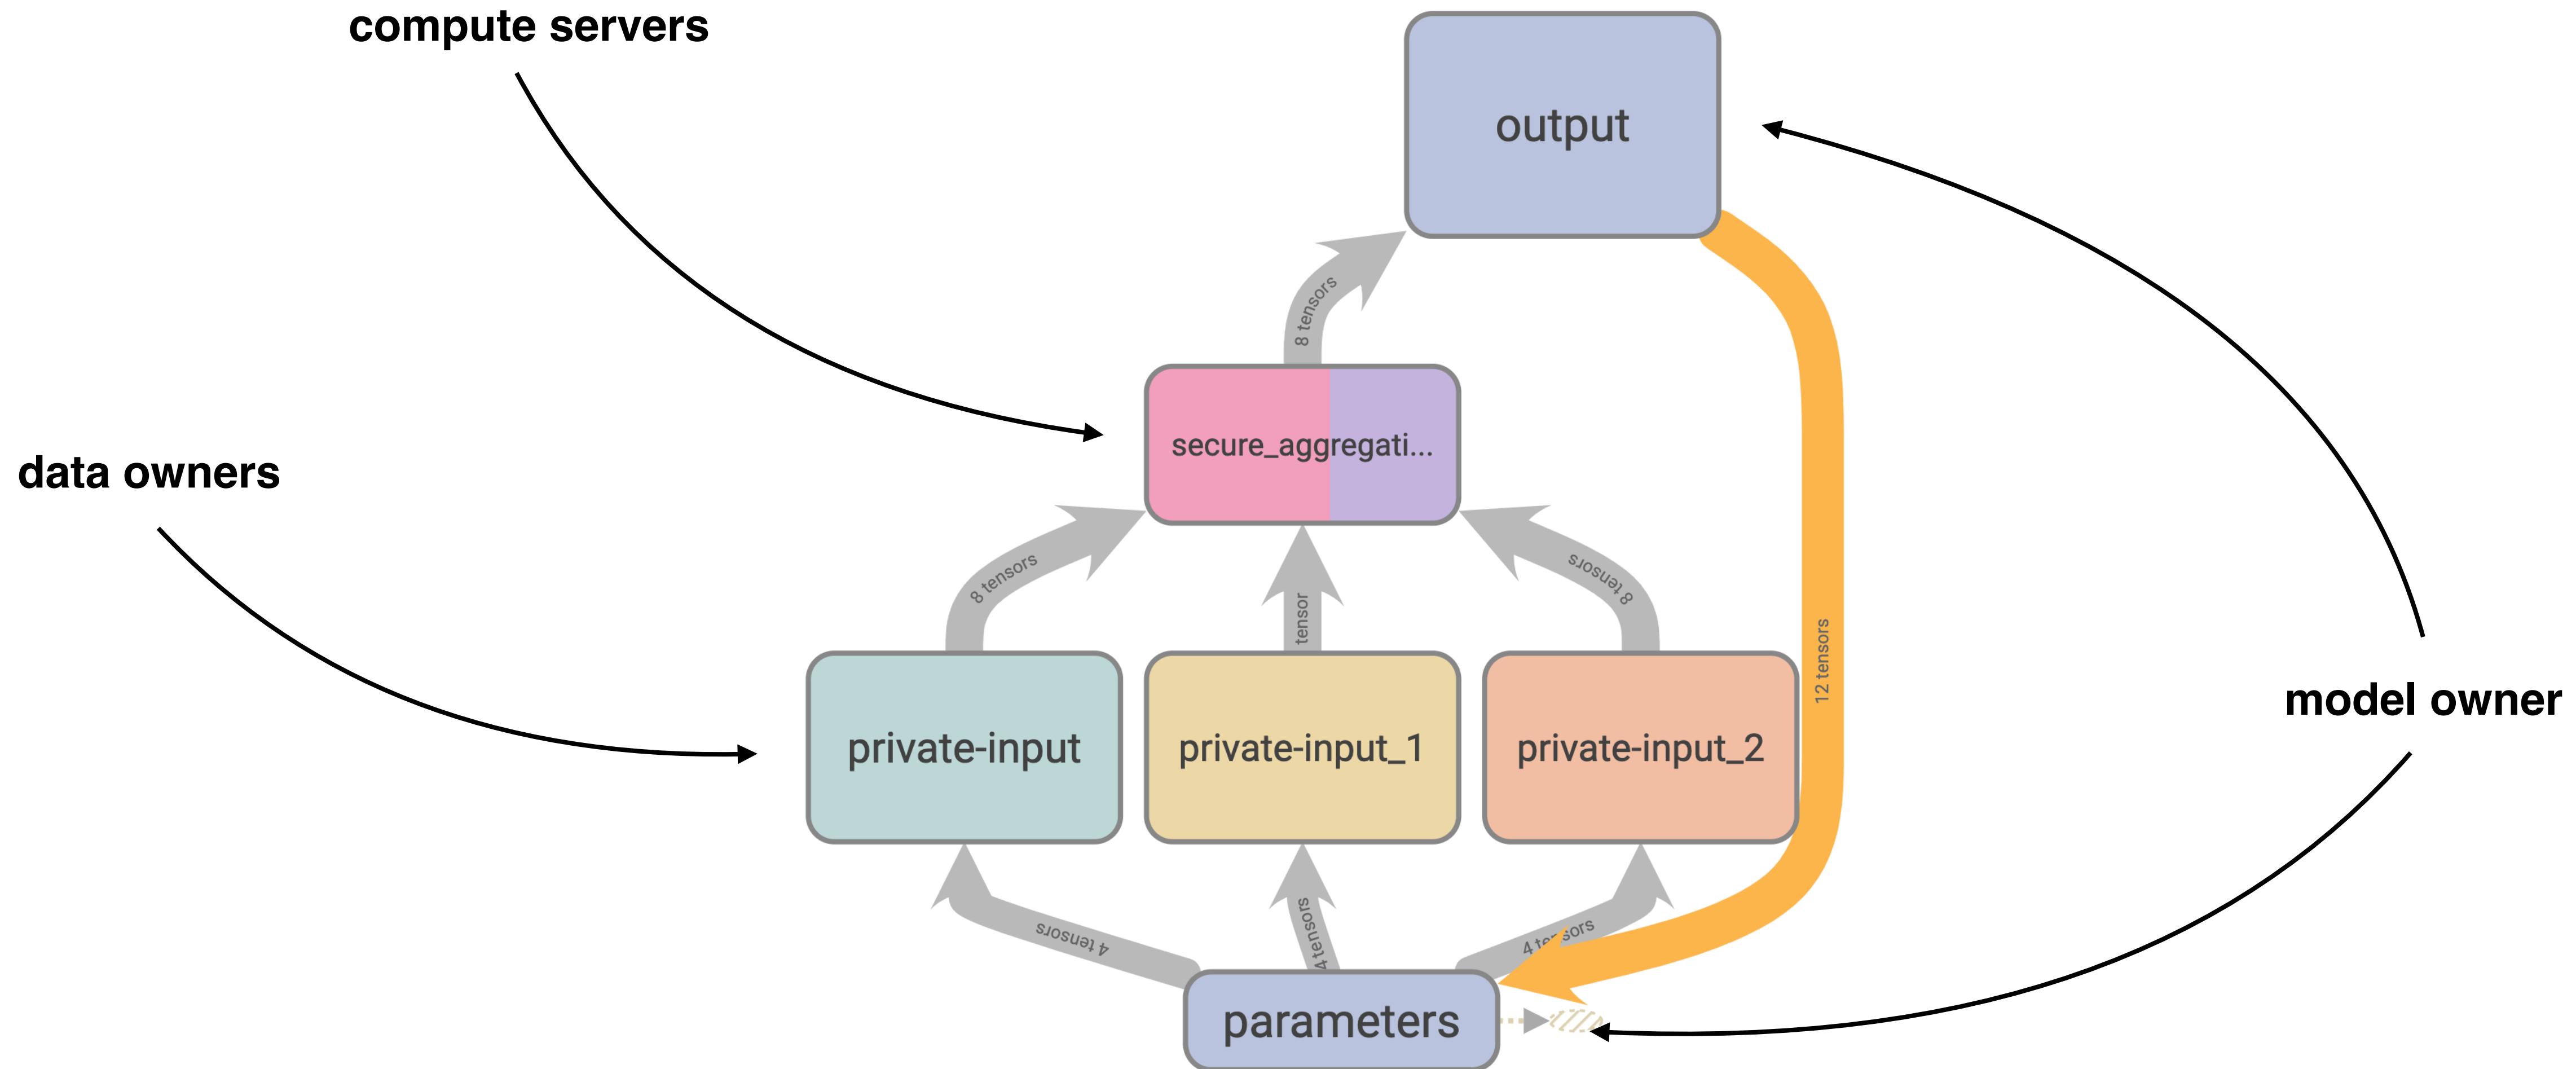
Participants



Secure Federated Learning in TF Encrypted

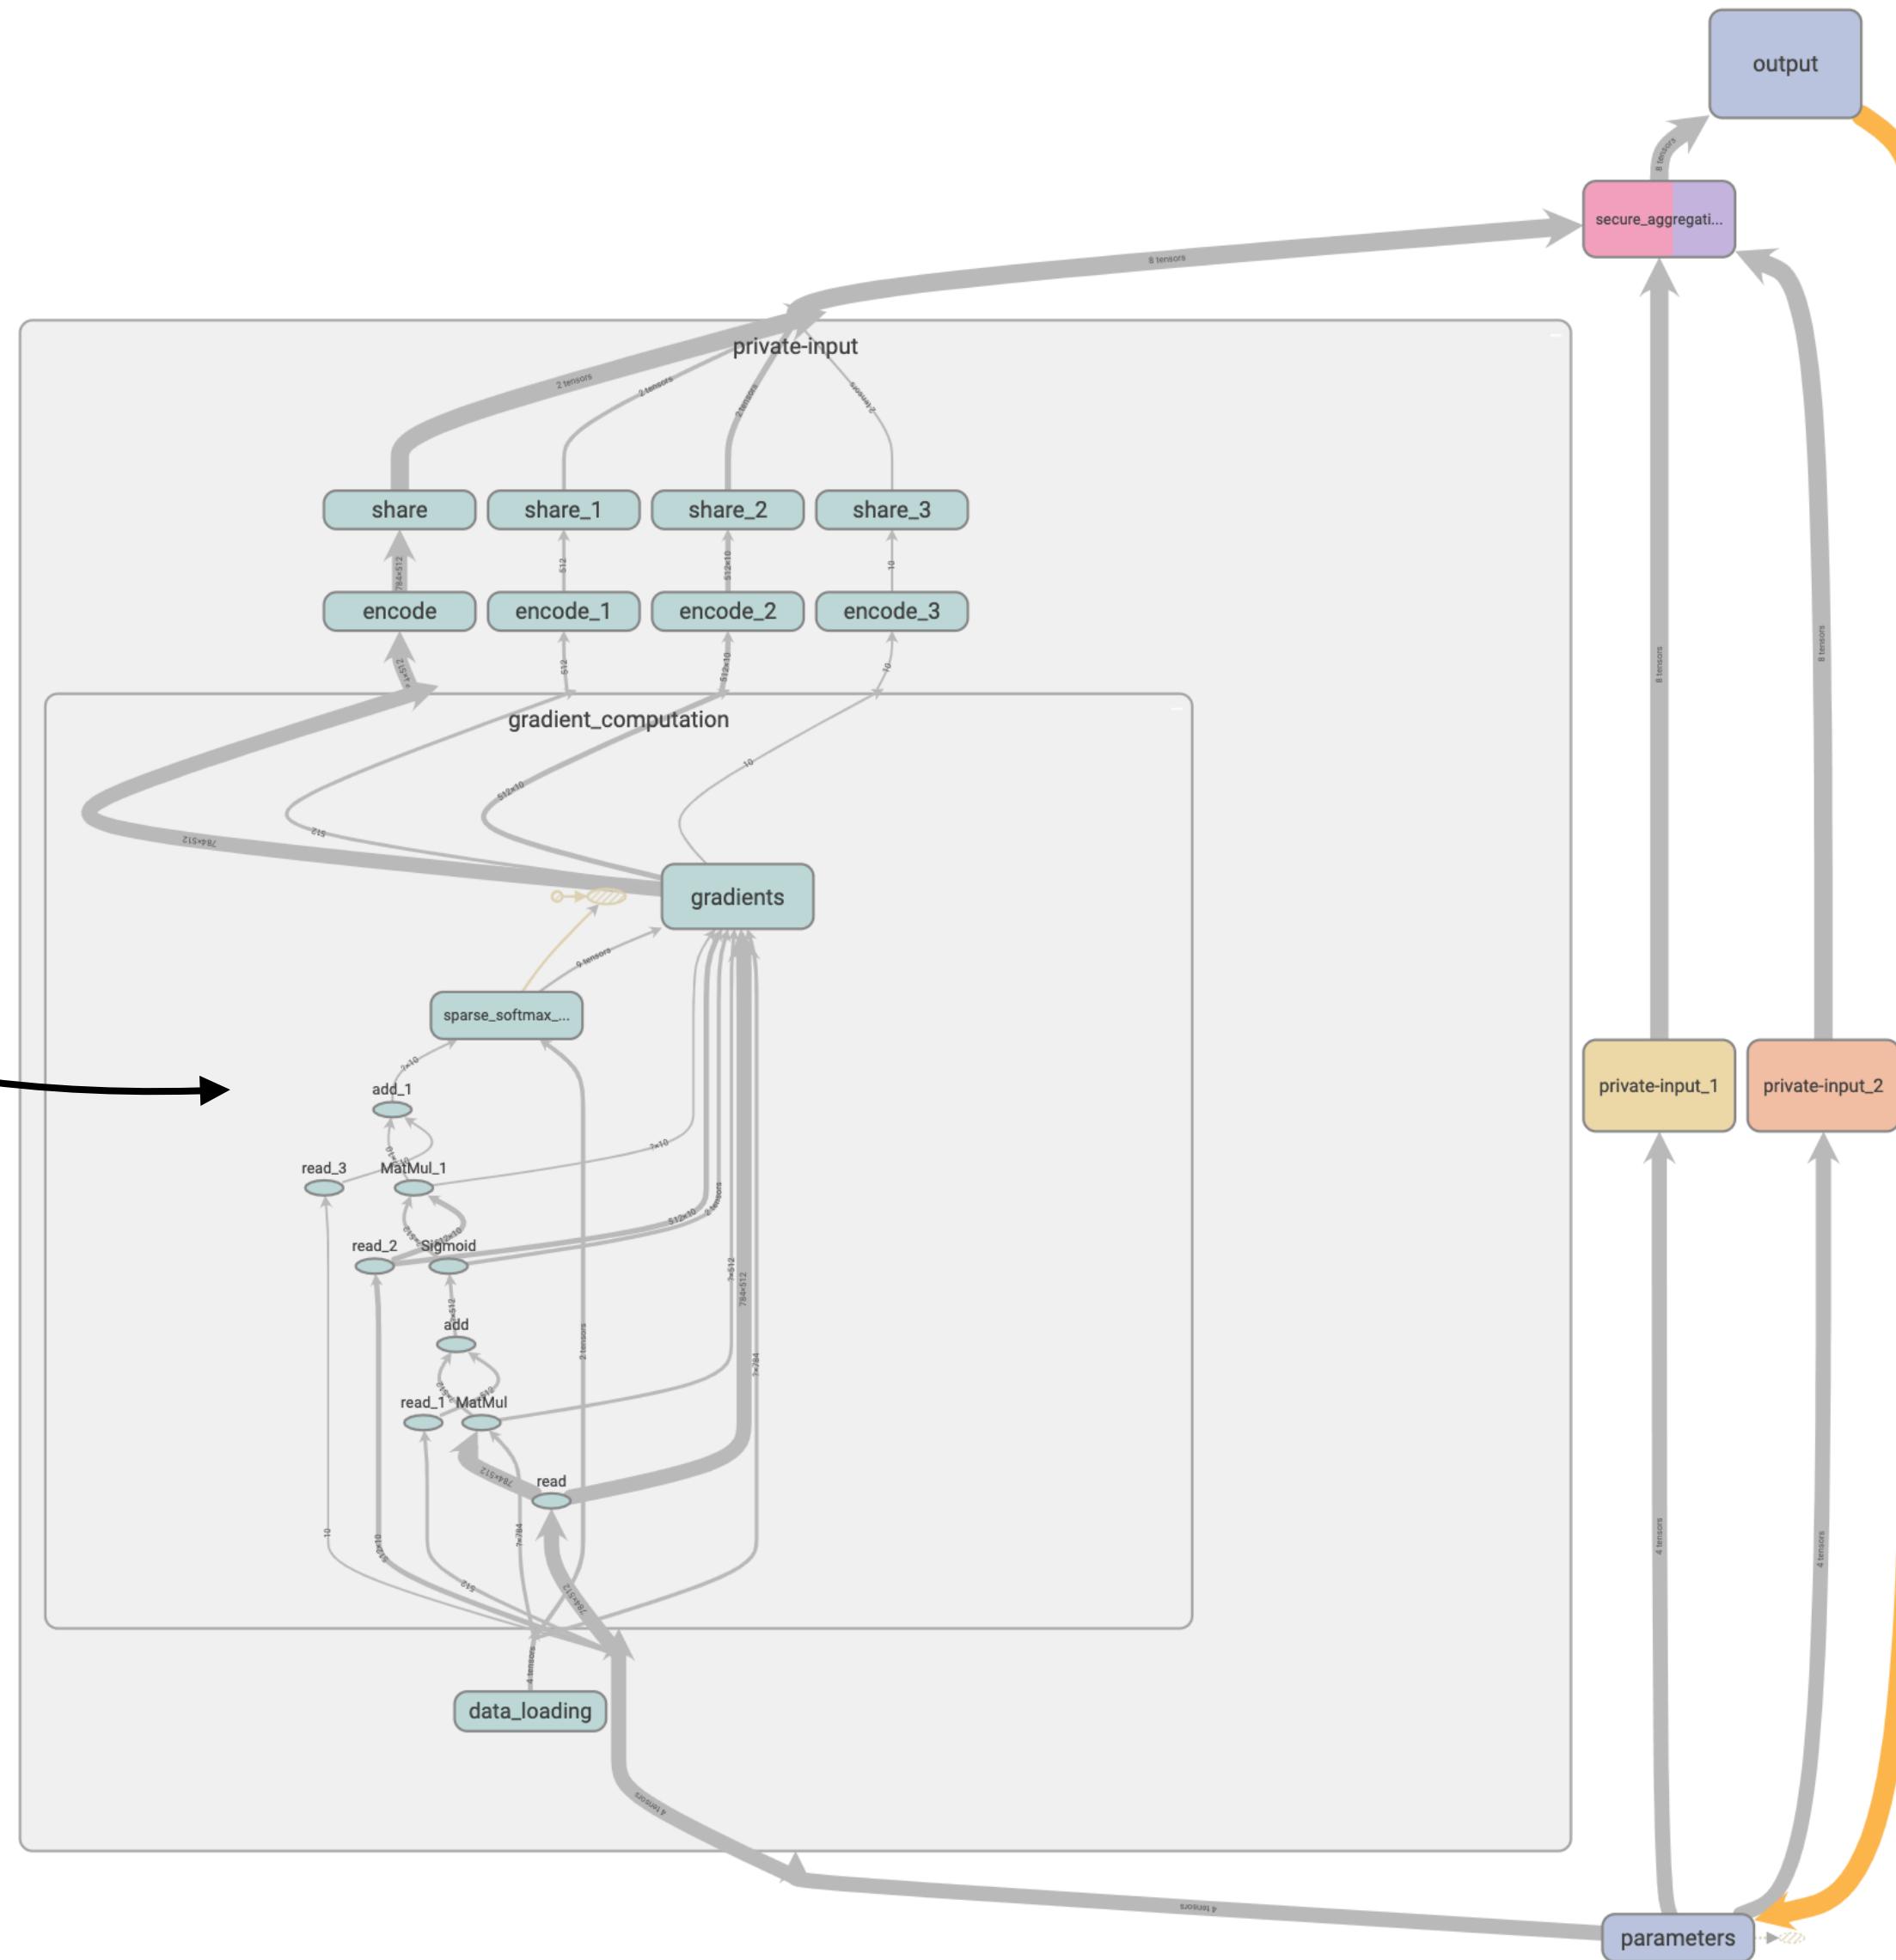
```
1 model_owner = ModelOwner('model-owner')
2 data_owners = [
3     DataOwner('data-owner-0', model_owner.build_training_model),
4     DataOwner('data-owner-1', model_owner.build_training_model),
5     DataOwner('data-owner-2', model_owner.build_training_model),
6 ]
7
8 model_grads = zip(*
9     tfe.define_private_input(data_owner.player_name, data_owner.compute_gradient)
10    for data_owner in data_owners
11 )
12
13 aggregated_model_grads = [
14     tfe.add_n(grads) / len(grads)
15     for grads in model_grads
16 ]
17
18 iteration_op = tfe.define_output(model_owner.player_name, aggregated_model_grads, model_owner.update_model)
19
20 with tfe.Session() as sess:
21     sess.run(tf.global_variables_initializer())
22
23     for i in range(model_owner.ITERATIONS):
24         sess.run(iteration_op)
```

Overall Computation



Local Optimization

TF optimization



Roadmap

High-level API (Private Keras, Pre-trained Models, Owned Data)

Tighter integration (TF Data, TF 2.0, TF Privacy, TF Federated)

Third-party cryptographic libraries (HE, MPC)

Improved performance

Wrap-Up

You can **compute on encrypted data**,
without the ability to decrypt

Privacy-preserving ML mitigate **bottlenecks** and
enable access to sensitive information

Secure computation **distributes trust and control**,
and is complementary to e.g. differential privacy

Privacy-preserving ML is a multidisciplinary field
benefitting from **adaptations** on both sides

TF Encrypted focuses on
usability and integration

Thank you!

github.com/tf-encrypted/

@mortendahlcs

@dropoutlabsai

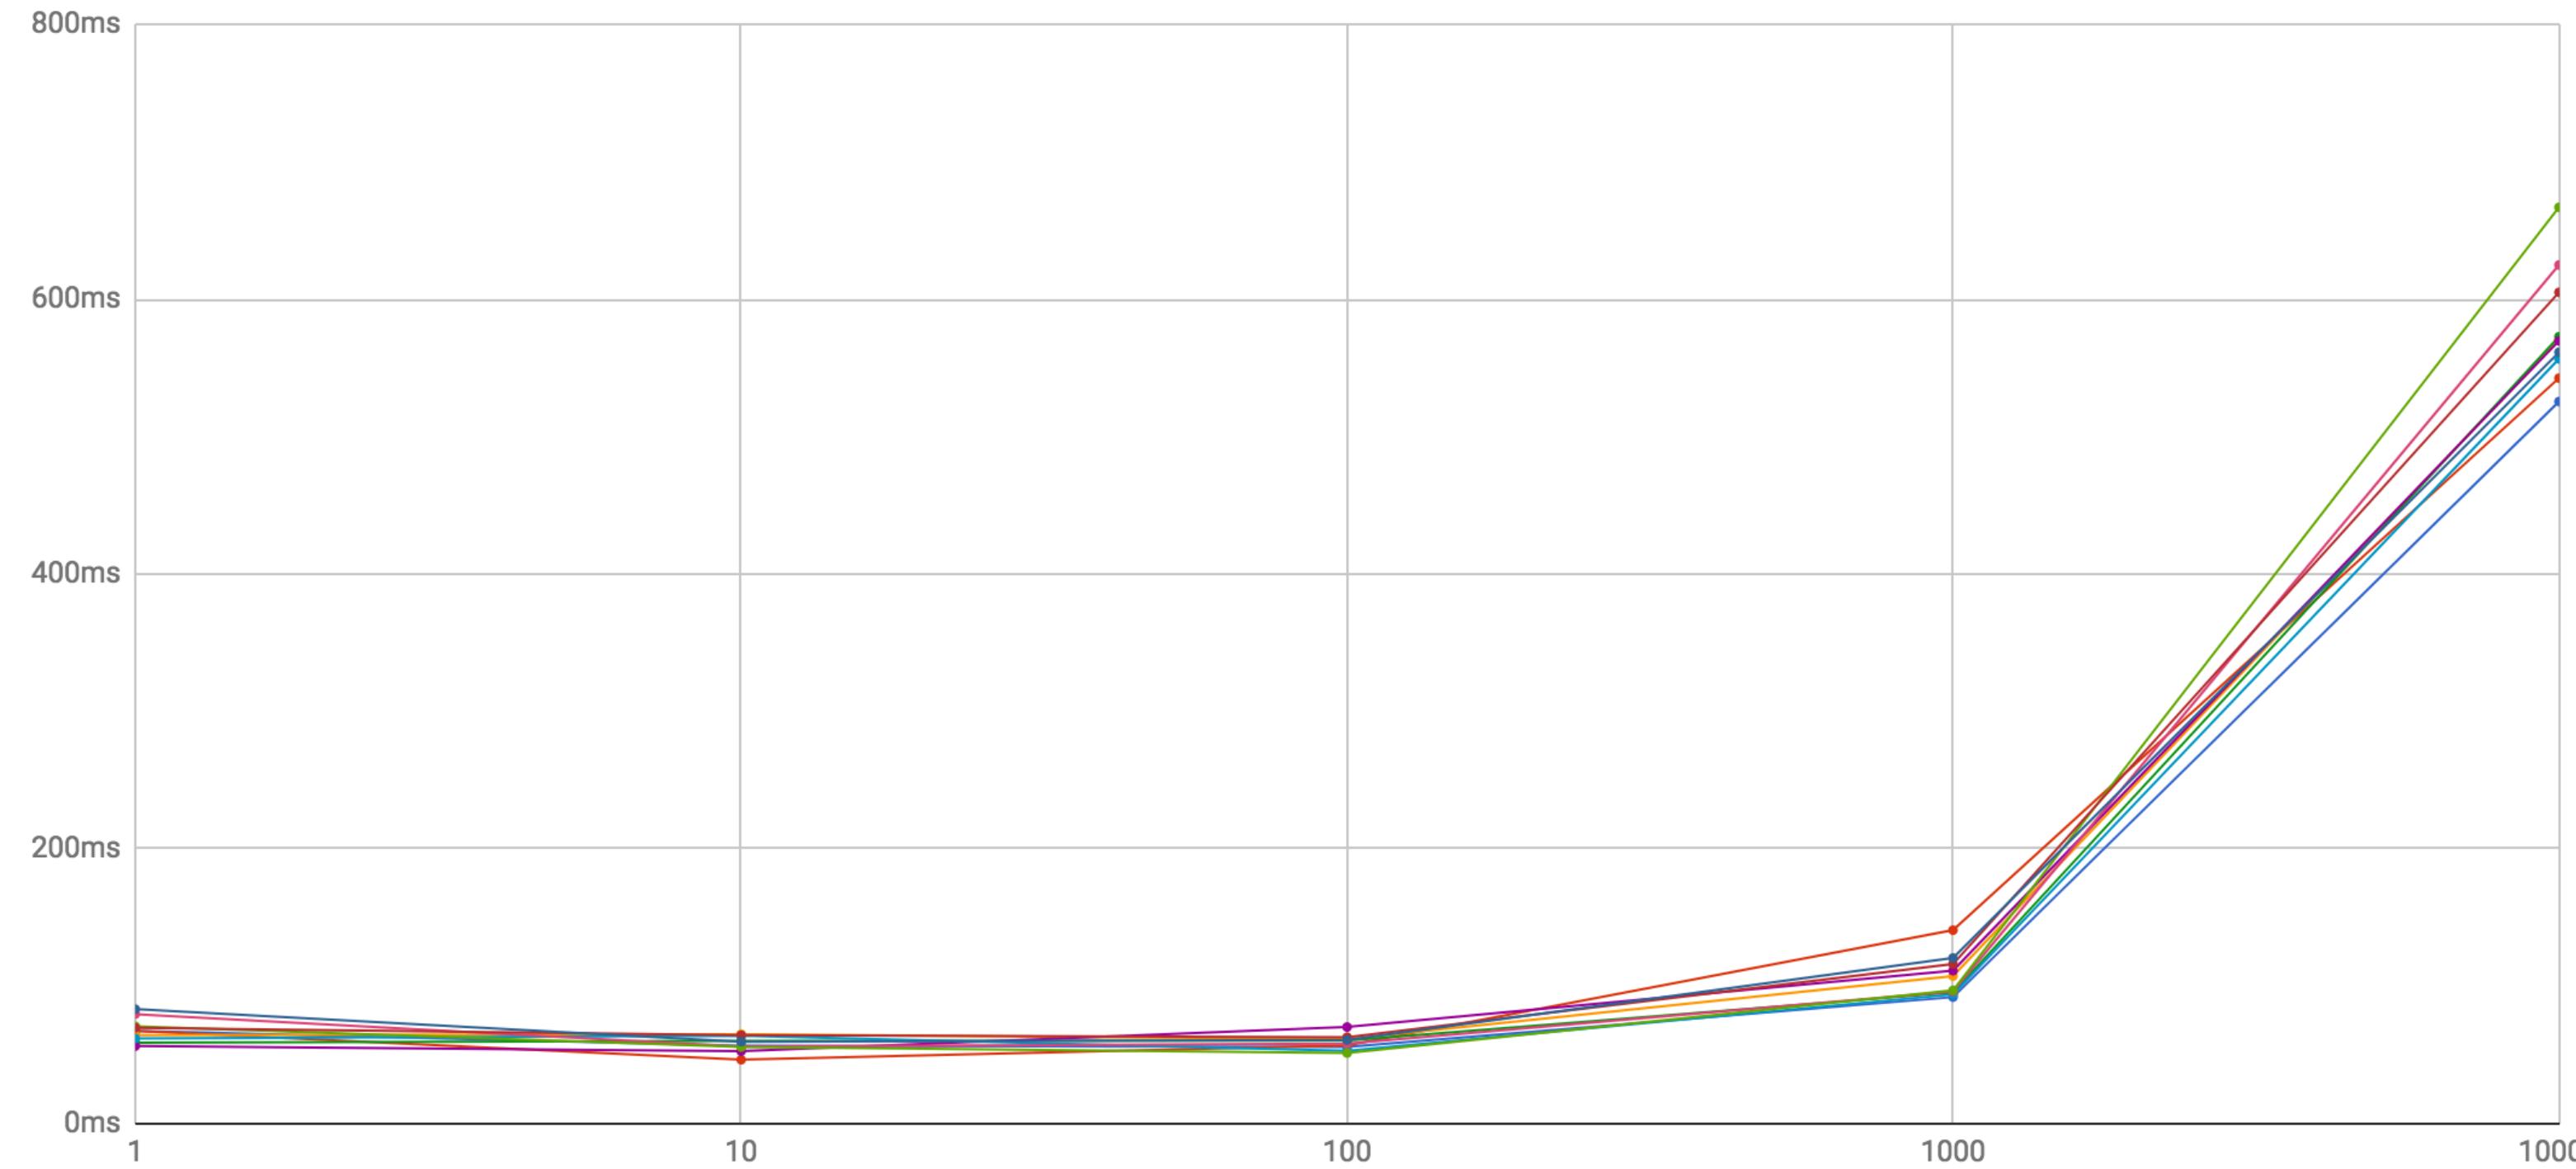
Appendix

Configuration

```
{  
    "server0": "10.0.0.1:4000",  
    "server1": "10.0.0.2:4000",  
    "model-owner": "10.0.0.3:4000",  
    "prediction-client": "10.0.0.4:4000",  
}
```

Performance

100 features, servers on Google cloud (2 vCPU, 10 GB)



More at <https://arxiv.org/abs/1810.08130>

Resources

Recent research papers using secure computation

CryptoNets: Applying Neural Networks to Encrypted Data with High Throughput and Accuracy, Dowlin et al.

SecureML: A System for Scalable Privacy-Preserving Machine Learning, Mohassel and Zhang

DeepSecure: Scalable Provably-Secure Deep Learning, Rouhani et al.

Gazelle: A Low Latency Framework for Secure Neural Network Inference, Juvekar et al.

ABY3: A Mixed Protocol Framework for Machine Learning, Mohassel and Rindal

SecureNN: Efficient and Private Neural Network Training, Wagh et al.

Blind Justice: Fairness with Encrypted Sensitive Attributes, Kilbertus et al.

(also great summary in <https://eprint.iacr.org/2017/1190>)

Specialised projects

TF Encrypted (<https://github.com/tf-encrypted/>)

PySyft (<https://github.com/OpenMined/PySyft>)

Secure computation frameworks

SCALE-MAMBA (<https://homes.esat.kuleuven.be/~nsmart/SCALE/>)

MP-SPDZ (<https://github.com/n1analytics/MP-SPDZ>)

ABY (<https://github.com/encryptogroup/ABY>)

OblivC (<http://oblivc.org/>)

(much more at <https://github.com/rdragos/awesome-mpc>)